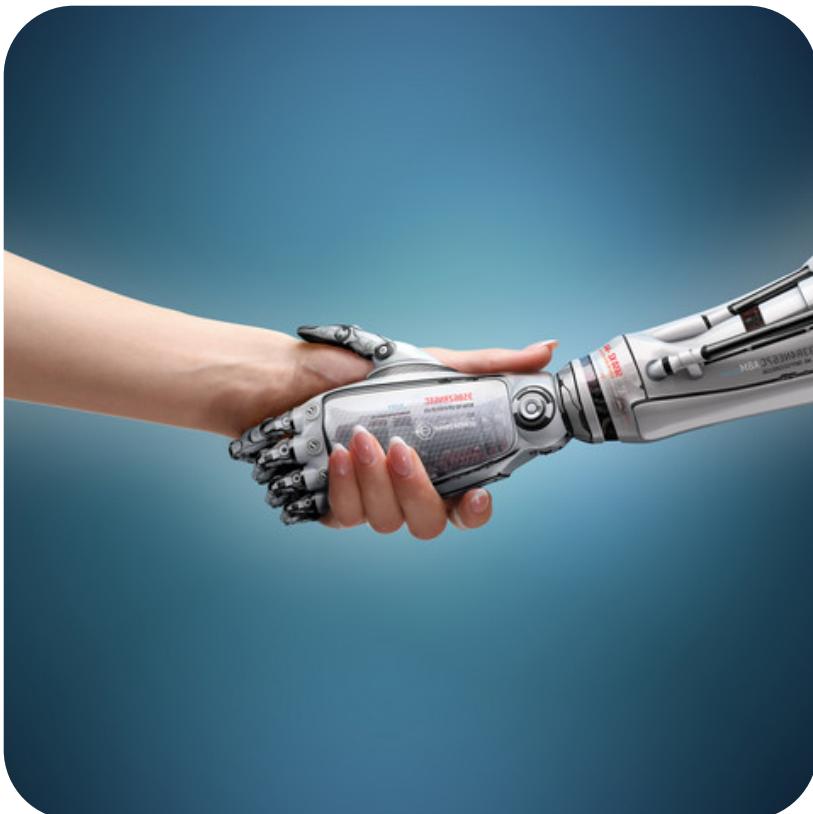


Artificial Intelligence

V02: Intelligent Agents

Agents, environments and rationality
Environments types and properties
Agent types



Based on material by Stuart Russell, UC Berkeley



Educational objectives

- **Remember** the definition of a rational agent and PEAS
- **Explain** why a rational agent might **neither** be omniscient, prophetic or **successful** and still be called «rational»
- **Argue** how expressiveness of an agent is a **mixed blessing**

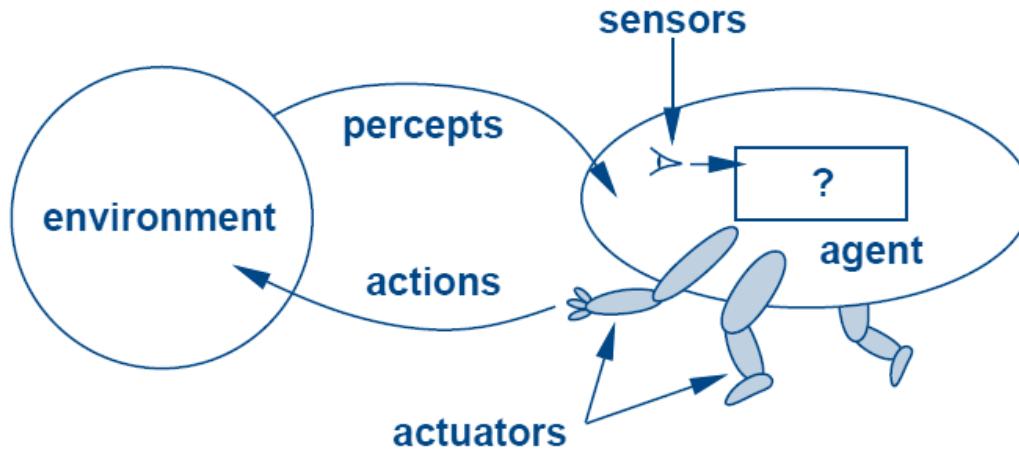


“In which we discuss the nature of agents, perfect or otherwise, the diversity of environments, and the resulting menagerie of agent types.”

➔ Reading: AIMA, ch. 2

1. AGENTS, ENVIRONMENTS AND RATIONALITY

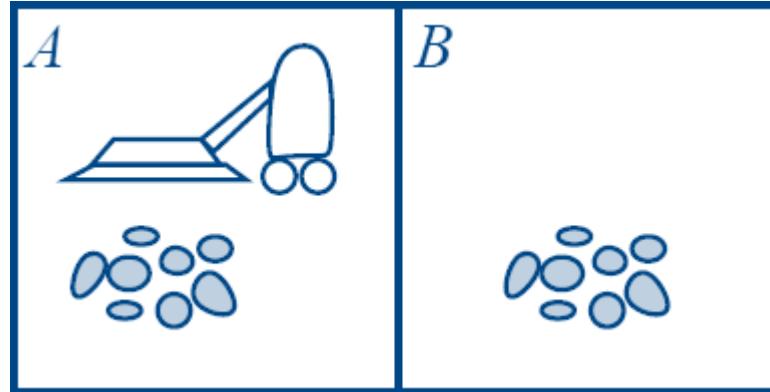
Agents and environments



Vocabulary

- Agents include humans, robots, softbots, thermostats, etc.
- The agent function maps from percept sequence (complete history) to actions:
$$f: P^* \rightarrow A$$
- The agent program runs on the physical architecture to produce f

Example: Vacuum-cleaner world



Environment

- Percepts: location and content, e.g., [A, Dirty]
- Actions: *Left*, *Right*, *Suck*, *NoOp*

A vacuum-cleaner agent

Agent function f : tabulated

Percept sequences	Corresponding actions
[A, Clean]	Right
[A, Dirty]	Suck
[B, Clean]	Left
[B, Dirty]	Suck
[A, Clean], [A, Clean]	Right
[A, Clean], [A, Dirty]	Suck
...	...

Agent program: implementation of abstract functional description

```
function Reflex-Vacuum-Agent([location, status]) returns an action
  if status = Dirty then return Suck
  else if location = A then return Right
  else if location = B then return Left
```

Questions

- What is the **right** function?
- Can it be implemented in a **small** agent **program**?

Rationality

Need of a **fixed performance** measure to evaluate **environment sequence**

- One point per square cleaned up in time T ?
 - One point per clean square per time step, minus one per move?
 - Penalize for $> k$ dirty squares?
- Is the sequence of results **desirable**?

A **rational agent** chooses whichever action maximizes the **expected value** of the performance measure **given the percept sequence to date**

- Rational ≠ omniscient
Percept may not supply all relevant information
 - Rational ≠ prophetic
Action outcome may not be as expected
- Rational ≠ successful
Rational = exploration, learning, autonomy



2. ENVIRONMENTS TYPES AND PROPERTIES

Environment specifications: PEAS

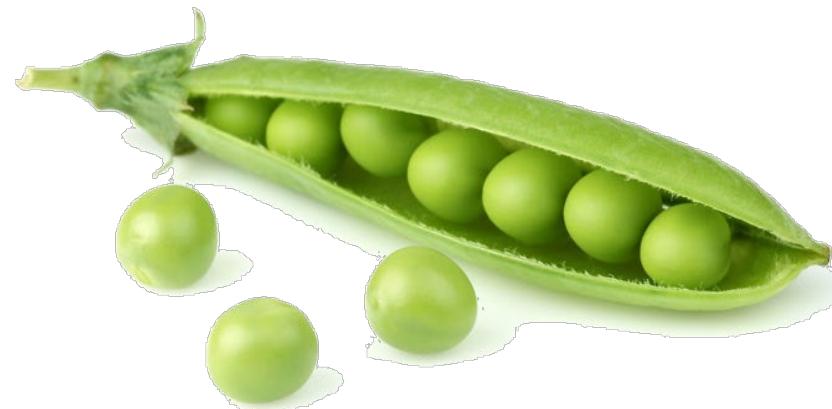
Performance measure, Environment, Actuators, Sensors

To design a rational agent, we must **specify** the **task environment**

- PEAS specifies environment **types**: how we as humans would perceive **external features**

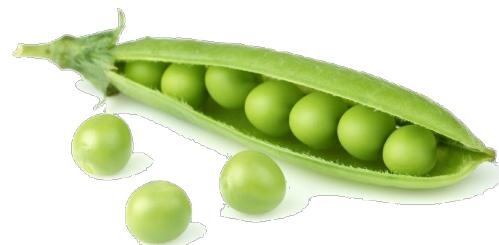
Example: the task of designing an automated taxi

- **Performance measure?** → safety, destination, profits, legality, comfort
- **Environment?** → streets/freeways, traffic, pedestrians, weather
- **Actuators?** → steering, accelerator, brake, horn, speaker/display
- **Sensors?** → video, accelerometers, gauges, engine sensors, keyboard, GPS



PEAS for an internet shopping agent

- **Performance measure?** → price, quality, appropriateness, efficiency
- **Environment?** → current and future WWW sites, vendors, shippers
- **Actuators?** → display to user, follow URL, fill in form
- **Sensors?** → HTML pages (text, graphics, scripts)



Environment properties

The *internal* features of an environment

- Fully **observable** vs. **partial observable** vs. unobservable
 - Do sensors give full access to the *relevant* state of the environment?
- **Single agent** vs. **multiple agents** → **competitive** vs. (partially) **cooperative**
 - Do others optimize a performance measure dependent on our agent?
- **Deterministic** vs. **stochastic** vs. nondeterministic
 - Do actions have certain consequences, or is the outcome probabilistic (other's actions don't count)?
- **Episodic** vs. **sequential**
 - Do current actions influence future decisions (probably not in classification settings)?
- **Static** vs. **dynamic**
 - Does the world keep turning while our agent decides what to do?
- **Discrete** vs. **continuous**
 - Regarding states, time, percepts and actions
- **Known** vs. **unknown**
 - Are the rules/laws governing the environment known to the agent (not strictly a property of the env.)?

Examples: Environments and their properties

	Solitaire	Poker	Image analysis	Internet shopping	Taxi
Observable?	(x)	(x)	x	-	-
Single-agent?	x	-	x	x (except auctions)	-
Deterministic?	x	- (stochastic)	x	(x)	-
Episodic?	-	-	x	-	-
Static?	x	x	(x)	(x)	-
Discrete?	x	x	-	x	-

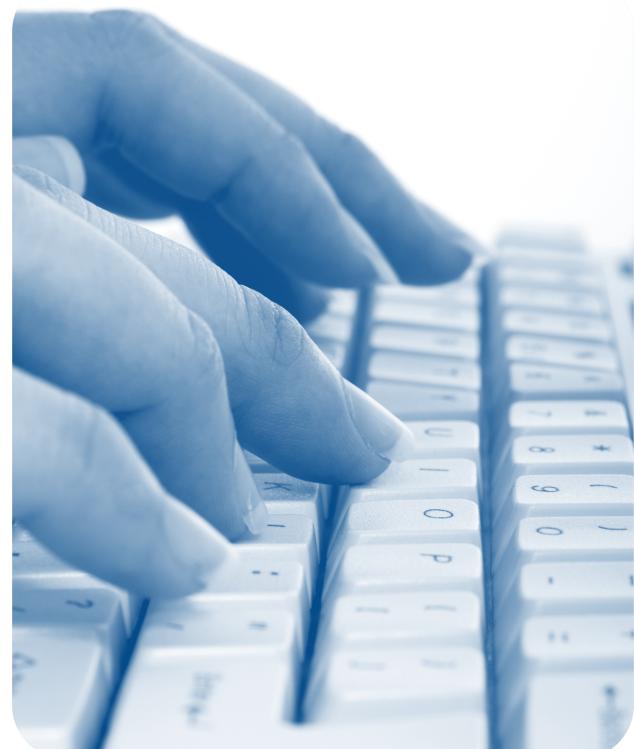
The **environment properties** largely **determine** the **agent design**

- The real world is (of course) partially observable, multi-agent, stochastic, sequential, dynamic, continuous

Exercise: Examine OpenAI Gym Retro

- Go to <https://blog.openai.com/retro-contest/> for a list of gaming environments assembled specifically to build AI agents
- Why did OpenAI create Gym Retro and deprecate their much broader «Universe» archive (see <https://github.com/openai/universe>)?
- Select one environment of your liking: How would you classify it according to the properties of the last slide?

OpenAI



3. AGENT TYPES

Four basic agent types

Agent := architecture + program

In order of increasing generality

- simple **reflex agents**: select action based on last percept
- **reflex agents with state**: regards history
- **goal-based agents**
- **utility-based agents**

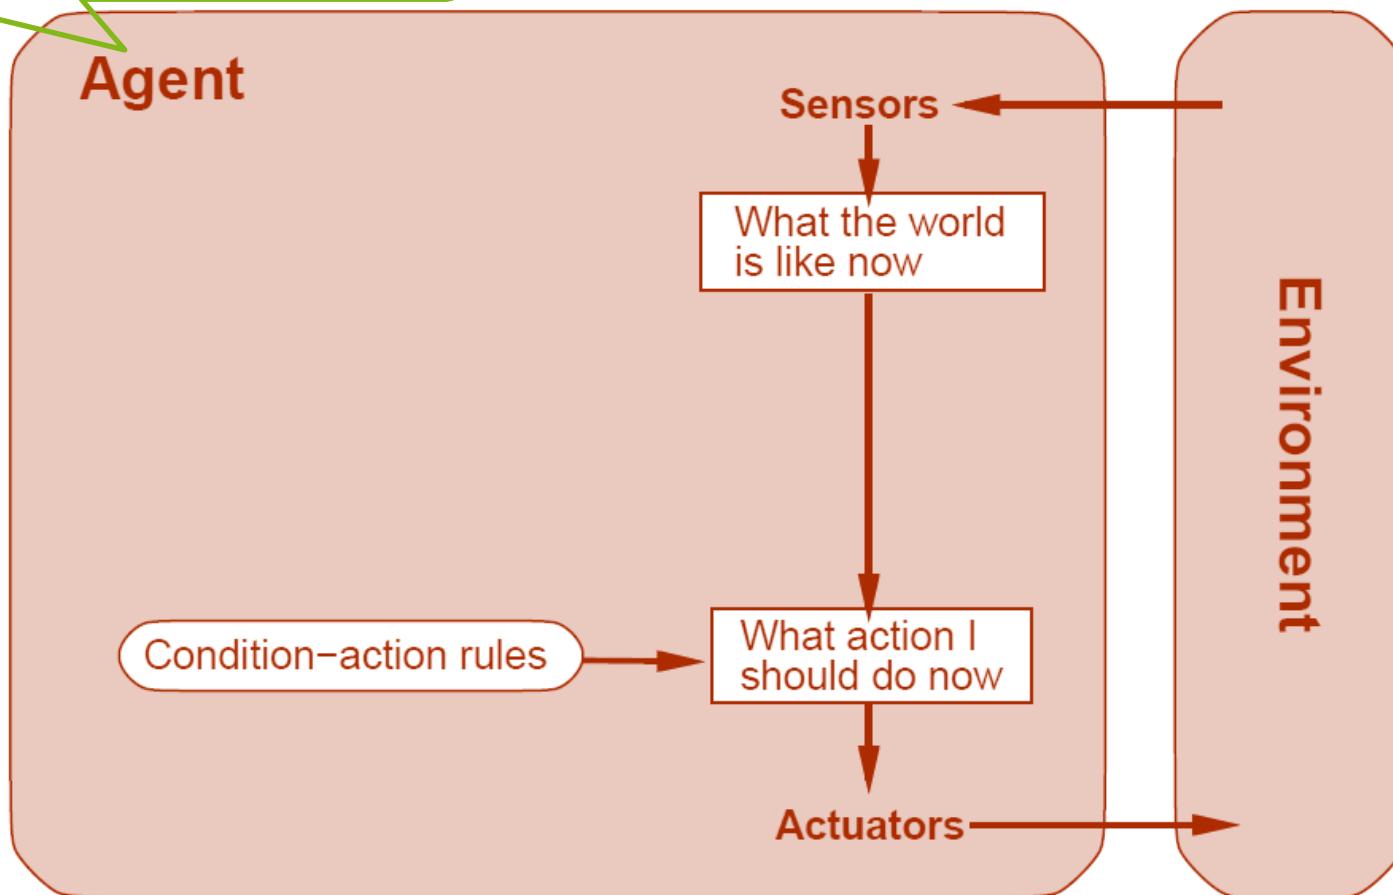
All these can be turned into **learning agents**

Goal of this section: to give an **abstract taxonomy** of the **different AI methods** encountered later in the course.



Simple reflex agents

No explicit goal: the agent **just maps from last percept to next action**, any implicit goal is reached eventually



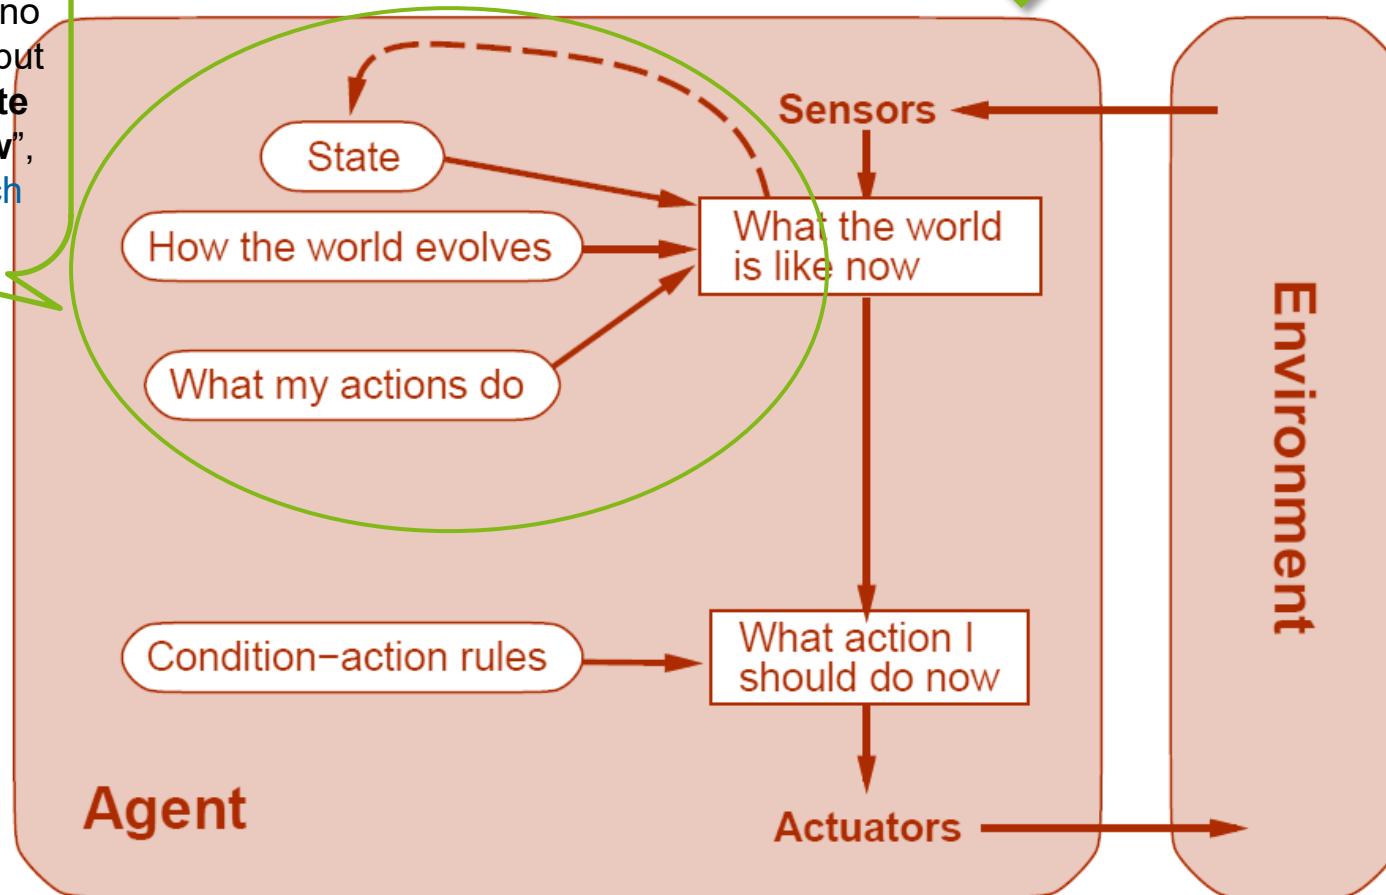
Example of simple reflex agent

```
function Reflex-Vacuum-Agent([location, status]) returns an action
  if status = Dirty then return Suck
  else if location = A then return Right
  else if location = B then return Left
```

Reflex agents with state

A.k.a. model-based reflex agents

Still no goal, still no planning ahead; but a **more complete view of the “now”**, e.g. [Local search](#)

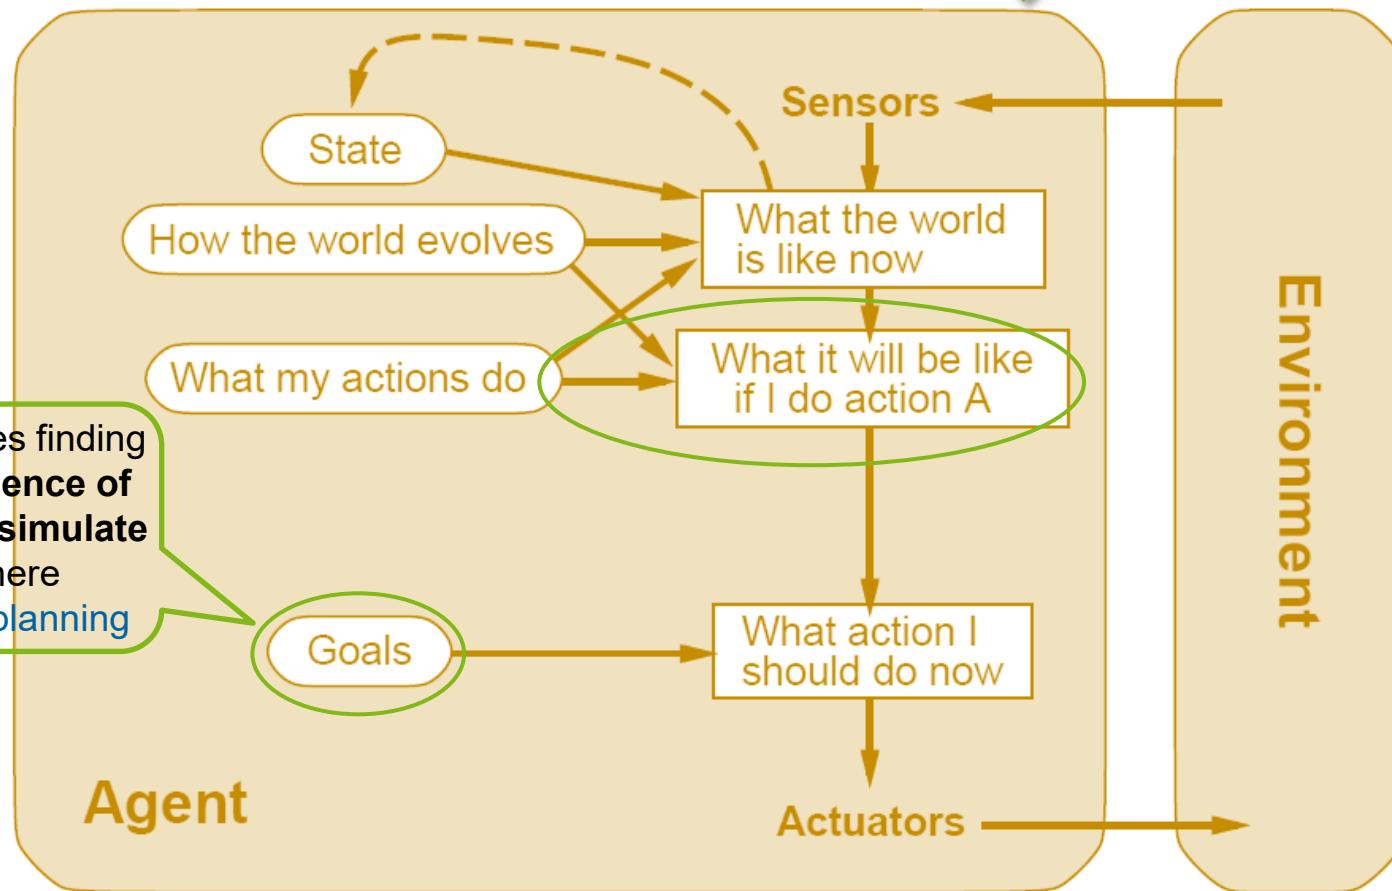
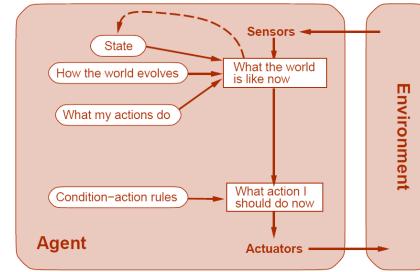


Example of stateful reflex agent

```
function Reflex-Vacuum-Agent([location, status]) returns an action
  static: last_A, last_B, numbers, initially INF
  if status = Dirty then ...
```

Goal-based agents

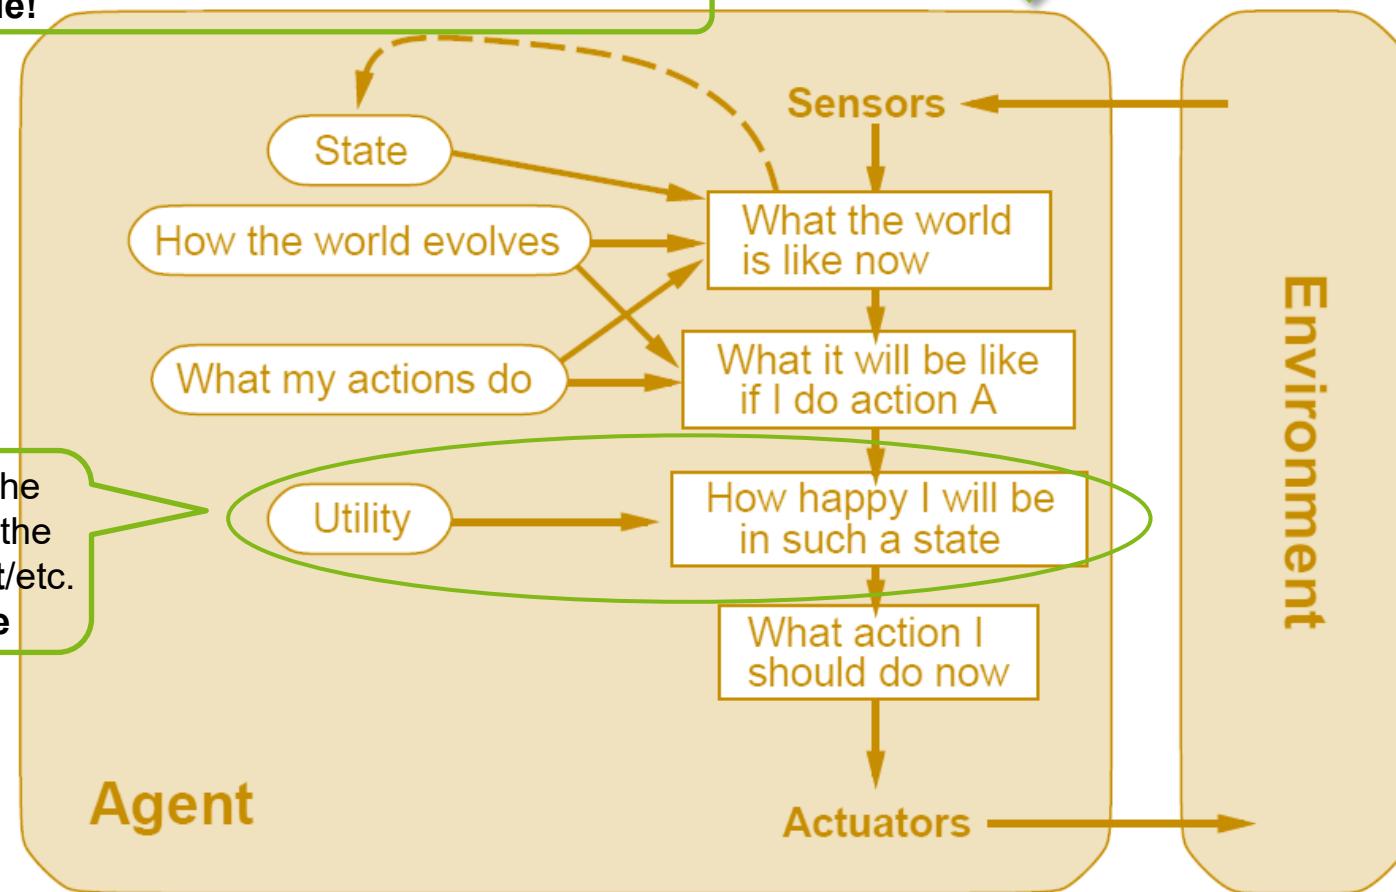
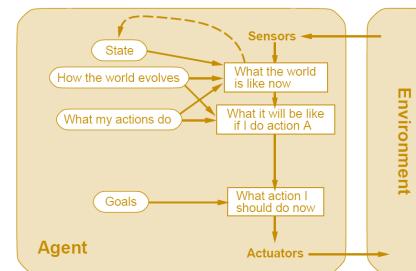
A.k.a. model-based agents



Utility-based agents

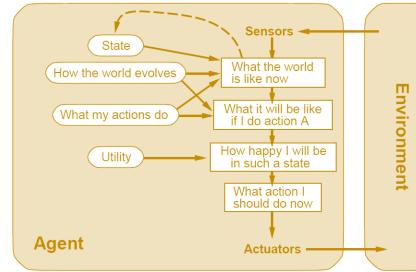
Finally reaching intelligence – or not?

Just maximizing (expected) utility will create rational agents; but it is not simple!



Learning agents

Learning is applicable to each agent type

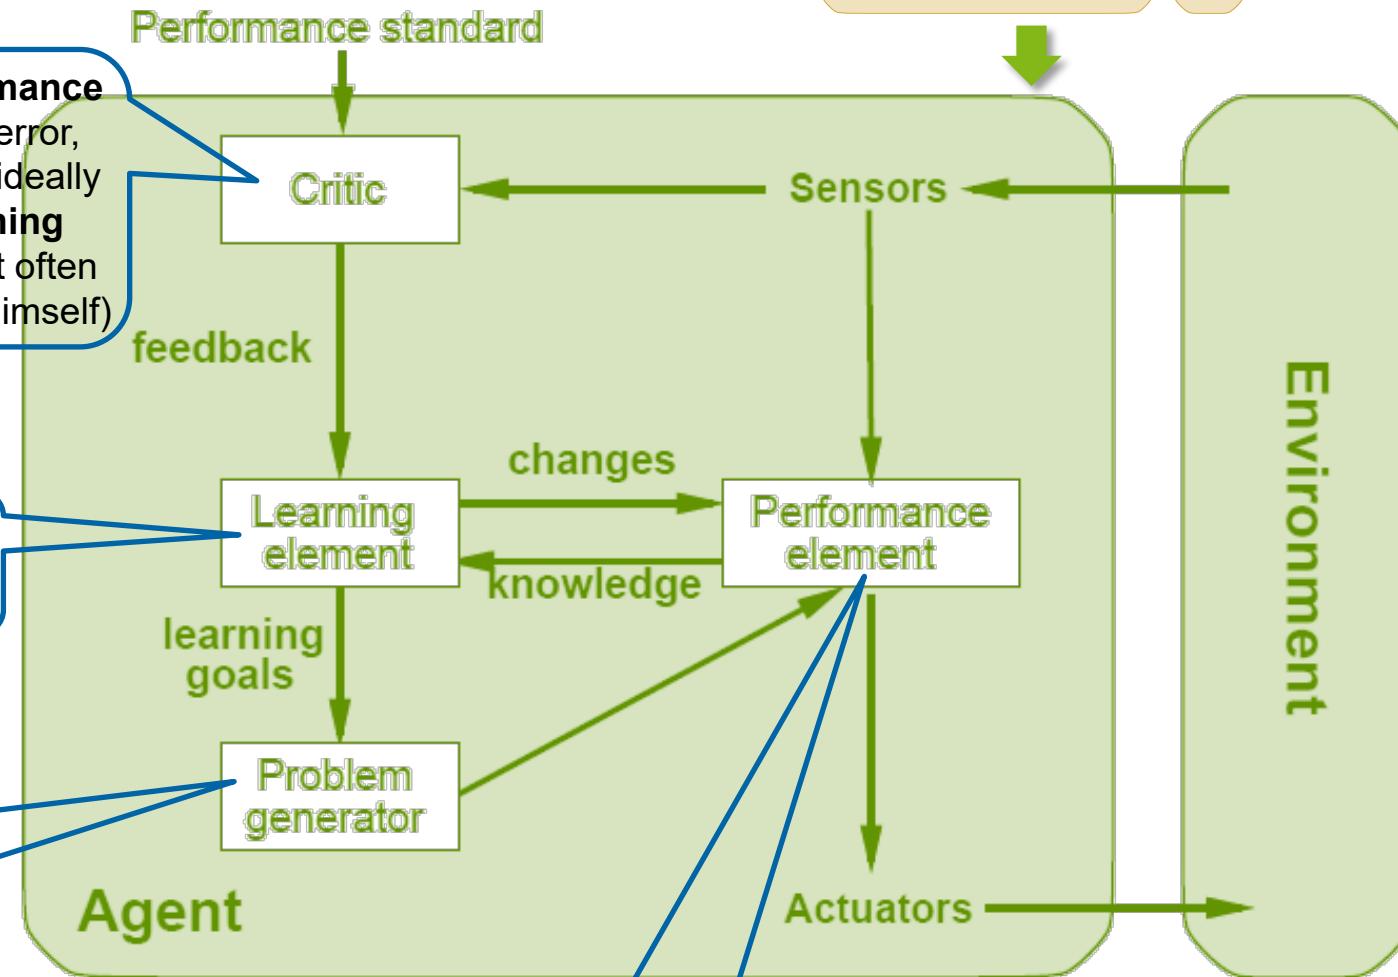


At least a **performance measure** (e.g., error, recall/precision); ideally an **active learning component** (most often the ML engineer himself)

The machine **learning algorithm** per se

Balancing **exploration / exploitation tradeoff**

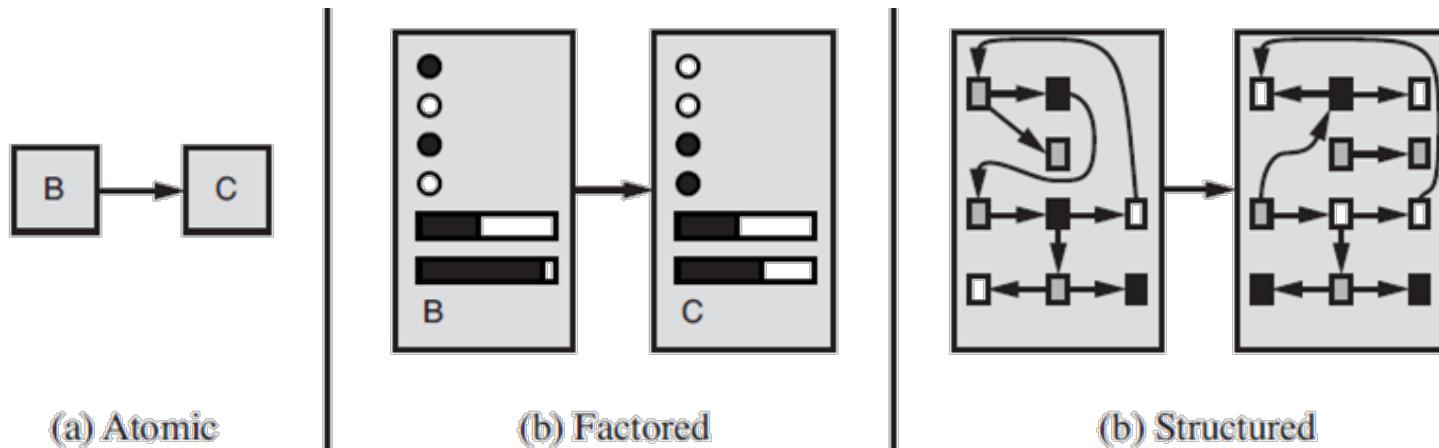
The model (of a certain representation)



A representation taxonomy

Consider the representation of any building block, e.g. «*What my actions do*»

- **Atomic**: states are “just different” from each other
 - search, game-playing, hidden Markov models, Markov decision processes
- **Factored**: states described by vectors (attributes), allowing for overlap and uncertainty
 - constrained satisfaction, propositional logic, planning, Bayesian networks, machine learning
- **Structured**: states as entities and their relationships with each other
 - first-order logic, first-order probability models, knowledge-based learning, NLP



Expressiveness revisited

Why a more capable agent is not always better

- Atomic – factored – structured is ordered by increasing expressiveness
 - A mixed blessing:
 - More expressive → **captures more**, often **much more concise**
 - More expressive → **learning/reasoning becomes much harder**
- ➔ Intelligent systems may need to operate at several points on the axis (task-dependent)



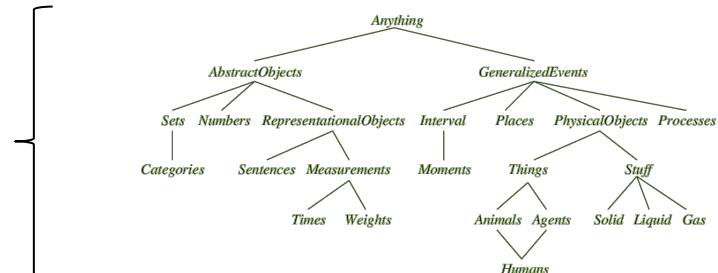
A model of practical AI

Inspired by E. Mogenet @ Zurich ML Meetup #31



AI Knowledge engineering (symbolic):

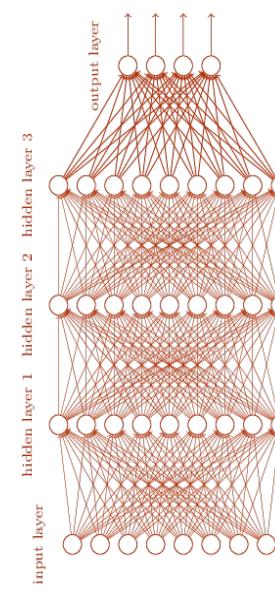
- ↓ Ontologies
- ↓ Logical inference



Gap to be filled by: common sense DB, NLP

Machine Learning (sub-symbolic):

- ↑ Hierarchical unsupervised learning
- ↑ Solid computer vision stack
- ↑ Images of the world



A model of more general future AI

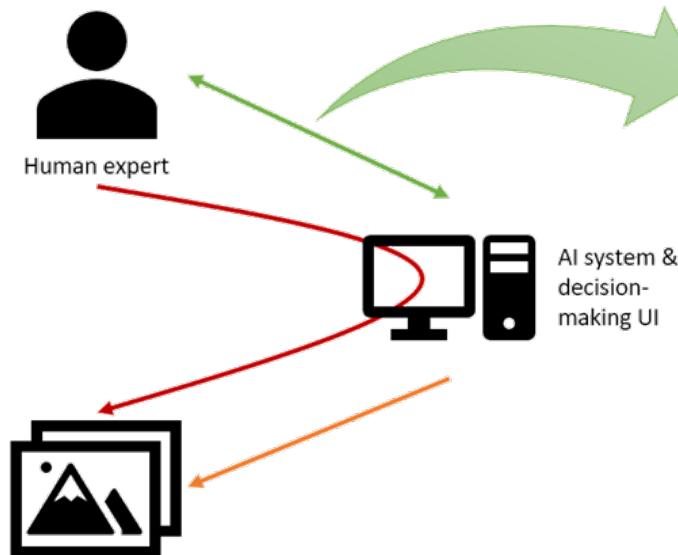
By T. Stadelmann with T. Wäfler, 2021

Successive stages of human-machine collaboration

1 (Mere tool): Human uses machine as mere UI

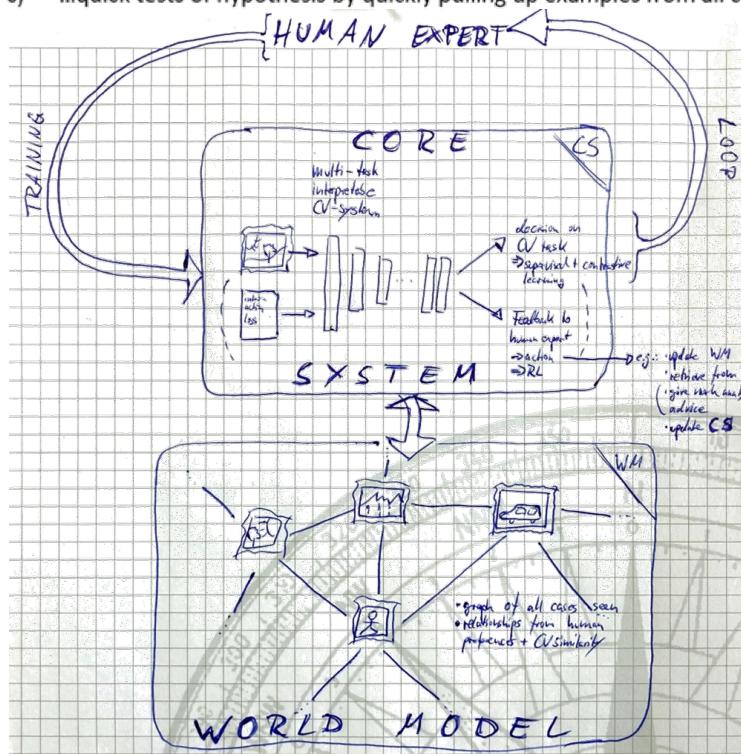
2 (Conventional computer vision): Trained CV system makes predictions that the human may consider (no learning interactions after training)

3 (Proposed co-learning): CV system continually learns from human preferences expressed in interactions; human learns from AI insights on own process



How to co-learn? High-level sources for mutual learning

1. AI system can adapt continually to human preferences using
 - a) ...explicit corrections made to AI decisions (feedback) >> continual learning (active / transfer learning)
 - b) ...implicit observations of human decision-making process >> contrastive learning, reinforcement learning
 - c) ...explicit hints to analogous situations not considered similar by AI so far >> deep case-based reasoning
2. Human can learn from AI as it provides
 - a) ...different features & local vision approach (complementary strengths) >> XAI
 - b) ...hints from analyzing the human interaction with the system via UI >> behavior analysis, recommender systems
 - c) ...quick tests of hypothesis by quickly pulling up examples from all cases seen so far >> graph-NN, image retrieval



Review

- **Agents** interact with **environments** through **actuators** and **sensors**
- The **agent function** describes what the agent does in all circumstances
- The **performance measure** evaluates the environment sequence
- A **perfectly rational agent** **maximizes expected performance**
- **Agent programs** implement (some) agent functions
- **PEAS** descriptions **define** (specify) task **environments**
- **Environments** are categorized along several **dimensions**
 - observable? single-agent? deterministic?
 - episodic? static? discrete?
- Several basic **agent architectures** exist (all also learnable)
 - reflex, reflex with state,
 - goal-based, utility-based

