

Artificial Intelligence


V03: Problem solving through search

Searching as a problem-solving strategy

- Uninformed search
- Heuristic (informed) search

Based on material by

- Stuart Russell, UC Berkeley
- Inês de Castro Dutra, Cooperating Intelligent Systems, U. Porto



Educational objectives

- **Know** classical search algorithms and selection criteria based on time and space complexity
- **Understand** how intelligent behavior evolves out of efficient algorithms
- **Know how** to inform search methods by heuristics
- **Be able to model** a real world problem to be solved by searching

“In which we see how an agent can look ahead to find a sequence of actions that will eventually achieve its goal.”

➔ Reading: ALMA, ch. 3



1. SEARCHING AS A PROBLEM-SOLVING STRATEGY

Example: On holiday in Romania

Task: Catch flight that leaves tomorrow from Bucharest

Initial state

- Currently in Arad

Formulate goal

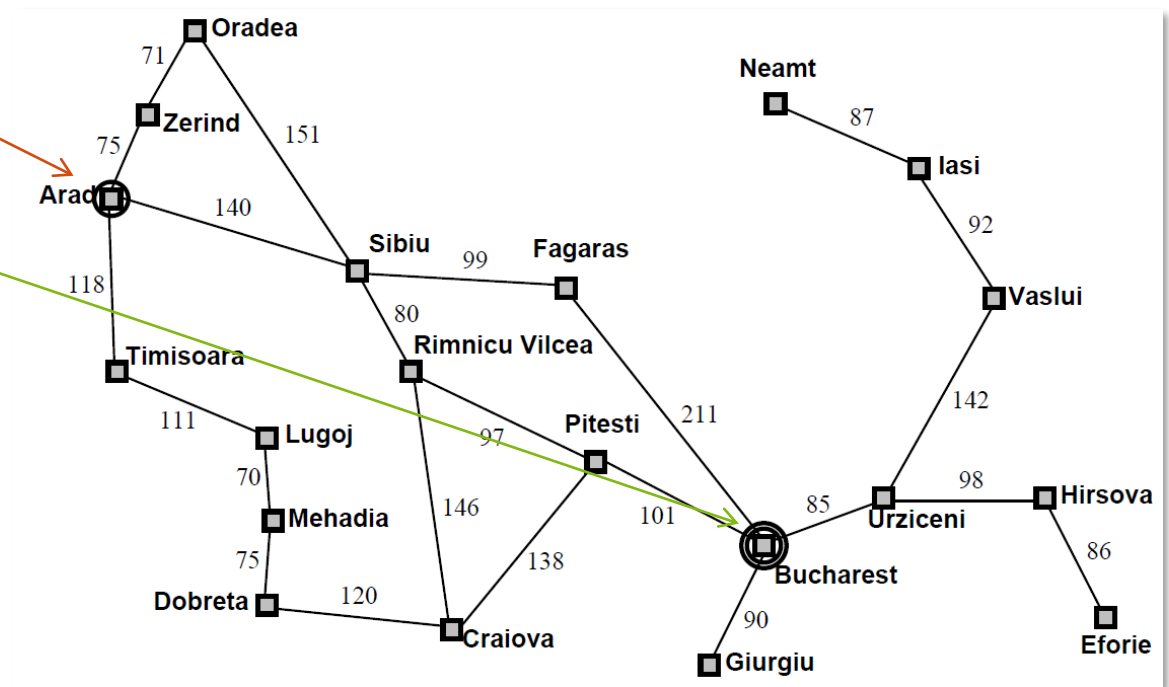
- be in Bucharest

Formulate problem

- **states**: various cities
- **actions**: drive between cities

Find solution

- sequence of cities
e.g., Arad→Sibiu→Fagaras→Bucharest

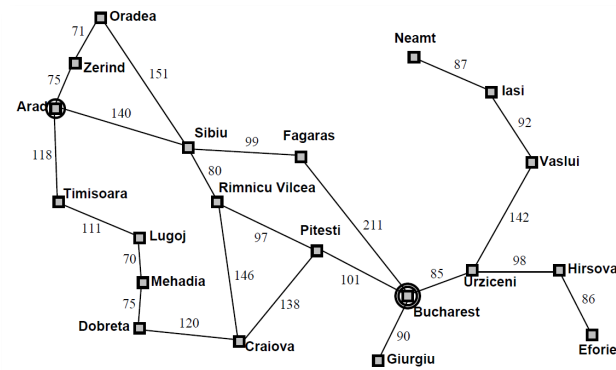


Problem formulation

For deterministic & fully observable environments

Problem is defined by four items

- **initial state**
e.g., $In(Arad)$
- **successor function** $S(x)$
set of action-state pairs, e.g.
 $S(Arad) = \{ \langle Arad \rightarrow Zerind; Zerind \rangle, \dots \}$
- **goal test**
explicit or implicit, e.g.
 $x = In(Bucharest)$ or $NoDirt(x)$
- **path cost** (additive)
e.g., sum of distances, number of actions, etc.
 $c(x, a, y) \geq 0$ is the step cost



Selecting a proper state space

- Real world is very complex
→ state & action space **must be abstracted**
- Abstract state: **set of** real states
- Abstract action: complex **combination of** real actions
e.g., $Arad \rightarrow Zerind$ represents a complex set of possible routes, detours, rest stops, etc.
- Abstract solution: **set of real paths** that are solutions in the real world
- For guaranteed realizability, **any** real state $In(Arad)$ must get to some real state $In(Zerind)$

→ See also appendix on **modeling**

→ Each abstract action **should be easier** than the original problem


Examples of problems solvable by searching

"Toy" problem: helps to identify strengths and weaknesses of different methods

8-puzzle

- States? integer locations of tiles (ignoring intermediate positions)
- Actions? move blank to left, right, up, down (ignoring unjamming etc.)
- Goal test? equals given goal state
- Path cost? 1 per move


Note: Optimal solution of n-Puzzle family is NP-hard (→ see appendix)



Real-world problem

Robotic assembly



- States? real-valued coordinates of robot joint angles; parts to be assembled
- Actions? continuous motions of robot joints
- Goal test? complete assembly
- Path cost? execution time



Other real-world problems




Protein design: find a sequence of amino acids that folds into a structure with certain properties



VLSI layout: place components and optimize wiring

All TSP-related problems of finding a shortest path



Diversity of search approaches

...solving increasingly complex problem types

→ this lecture

Uninformed (blind) search

- **All it can do:** generate successors of tree-nodes, distinguish goal- from non-goal states
- Suitable environments: **fully observable, deterministic, discrete** (episodic, static, single agent)

Extensions of today's methods exist to **non-deterministic** and **partially observable** as well as **(semi-)dynamic** environments (**online** search) (→ see AIMA, ch. 4.3-4.5)

Heuristic (informed) search

- **Knows whether** one non-goal state is “**more promising**” than another
- Suitable environments: as above, but **larger**

More informed search methods

Online search

- Environments are **dynamic** (i.e., not fully known from the beginning → percepts become important)

Local search

- Cares only to **find a goal state rather than** the optimal **path**
- Suitable environments: also **continuous** state/action spaces (hill climbing, simulated annealing)

Adversarial search

- Search in the face of an opponent (i.e., **dynamic multi-agent** environments; also **stochastic** and **partially observable** forms)

→ next lecture

2. UNINFORMED SEARCH

Uninformed search

Approach


- **Tree search**: iteratively **expand nodes** until a goal node is hit
- Different **strategies**: **order** of node expansion

Evaluation criteria for strategies

- **completeness**: does it **always find a solution** if one exists?
- **optimality**: does it always find a **least-cost solution**?
- **time complexity**: **number of nodes** generated/expanded
- **space complexity**: maximum number of **nodes in memory**

Time and space complexity are measured in terms of

- b : maximum **branching factor** of the search tree
- d : **depth of the least-cost solution**
- m : maximum **depth of the state space** (may be ∞)



Example

Growth of time and memory requirements

- Algorithm: breadth-first search (\rightarrow ADS: exponential time & space complexity $O(b^d)$)
Assumptions: $b = 10$, 1 mio nodes/sec, 1 kB/node
Question: what d is easily manageable?

Depth	Nodes	Time	Memory
2	110	.11 milliseconds	107 kilobytes
4	11,110	11 milliseconds	10.6 megabytes
6	10^6	1.1 seconds	1 gigabyte
8	10^8	2 minutes	103 gigabytes
10	10^{10}	3 hours	10 terabytes
12	10^{12}	13 days	1 petabyte
14	10^{14}	3.5 years	99 petabytes
16	10^{16}	350 years	10 exabytes

- \rightarrow Practical advice: **Exponential-complexity** search problems **cannot be solved by uninformed methods** for any but the smallest instances
- \rightarrow See appendix for some **recap** on **complexity theory**

Uninformed search strategies

→ Details: ADS or AIMA ch. 3.4

Expand the shallowest
unexpanded node

Expand node with
lowest path cost $g(n)$

Expand
deepest node

DFS only
up to level l

Try DLS with $l = 1, l = 2, \dots$ until
goal is reached

Criterion	Breadth-First	Uniform-Cost	Depth-First	Depth-Limited	Iterative Deepening	Bidirectional (if applicable)
Complete?	Yes ¹	Yes ^{1,2}	No	No	Yes ¹	Yes ^{1,4}
Optimal cost?	Yes ³	Yes	No	No	Yes ³	Yes ^{3,4}
Time	$O(b^d)$	$O(b^{1+\lceil C^*/\epsilon \rceil})$	$O(b^m)$	$O(b^\ell)$	$O(b^d)$	$O(b^{d/2})$
Space	$O(b^d)$	$O(b^{1+\lceil C^*/\epsilon \rceil})$	$O(bm)$	$O(b\ell)$	$O(bd)$	$O(b^{d/2})$

Practical advice

- **Depth-first tree search** is a **major work horse** for many AI tasks (due to linear space complexity)
- **Iterative deepening** is **not wasteful** (a tree with nearly the same b at each level has most nodes in the bottom level → generating higher-level states multiple times doesn't matter)
- **Iterative deepening** is **preferred uninformed search method** (for large search space and d is unknown)
- **Bi-directional search** can **help** a lot, but $O(b^{d/2})$ space complexity is major drawback

3. HEURISTIC (INFORMED) SEARCH

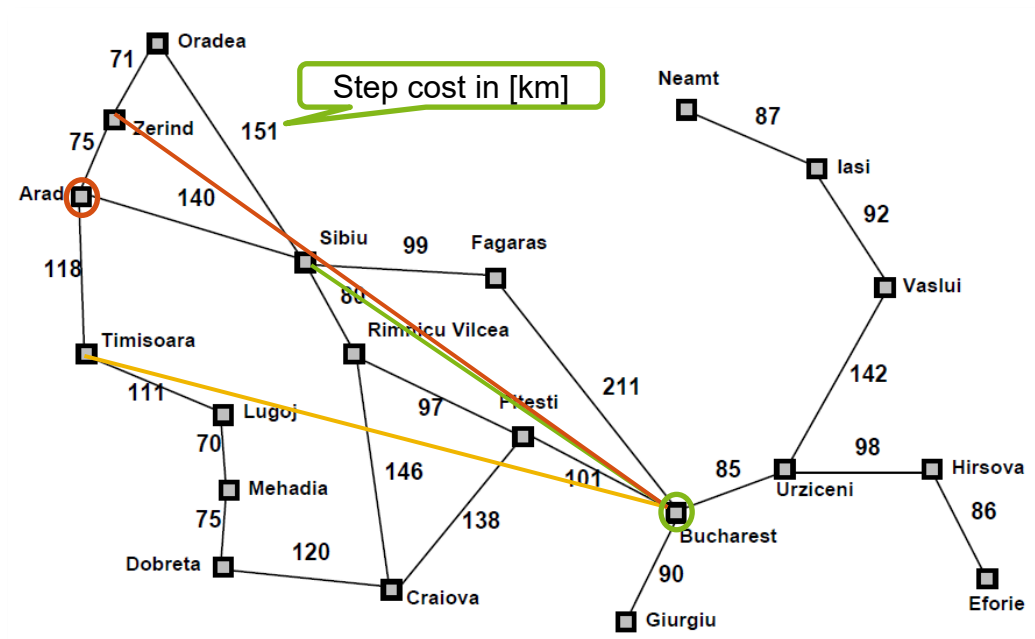
Tree-/graph search using additional knowledge

...beyond the definition of the problem

Best-first search

- **Select** the node to be expanded next **based on** some **evaluation function** $f(\text{node})$
- Typically, f is implemented by a **heuristic** $h(\text{node})$ (measure of “desirability”)
- $h(\text{node})$ facilitates **pruning** of the search tree: options are eliminated without examination

What could be a **good heuristic** for the distance to Bucharest (being in Arad)?



Arad	366	Mehadia	241
Bucharest	0	Neamt	234
Craiova	160	Oradea	380
Drobeta	242	Pitesti	100
Eforie	161	Rimnicu Vilcea	193
Fagaras	176	Sibiu	253
Giurgiu	77	Timisoara	329
Hirsova	151	Urziceni	80
Iasi	226	Vaslui	199
Lugoj	244	Zerind	374

Values of h_{SLD} —straight-line distances to Bucharest.

Typical implementations

Greedy search

- Expand node with lowest *subsequent* cost estimate according to some h , i.e. $f(n) = h(n)$
- n may only *appear* to be closest to the goal

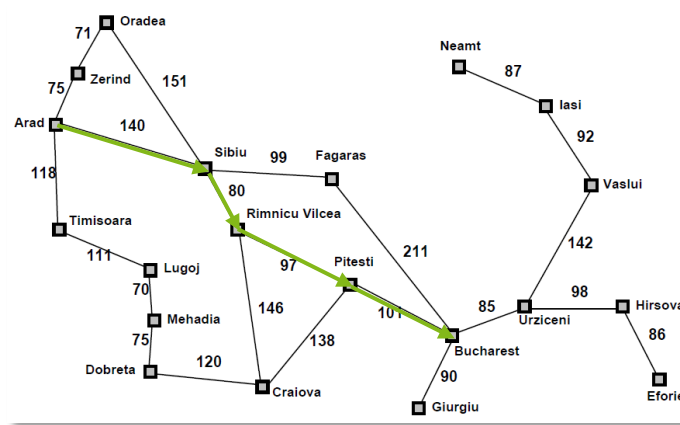
A*

- Obvious improvement: **consider full path cost**, i.e. $f(n) = g(n) + h(n)$
($g(n)$ cost so far to reach n , $h(n)$ estimated cost to goal from n , $f(n)$ estimated total path cost)
- $h(n)$ needs to be **admissible**: $\leq \text{true cost}$ and ≥ 0 (e.g., $h_{\text{straight line distance}}$)
- A* search is optimal, complete
- A* has time complexity $O(2^{(\text{error of } h) \cdot d})$ and **keeps all nodes in memory**

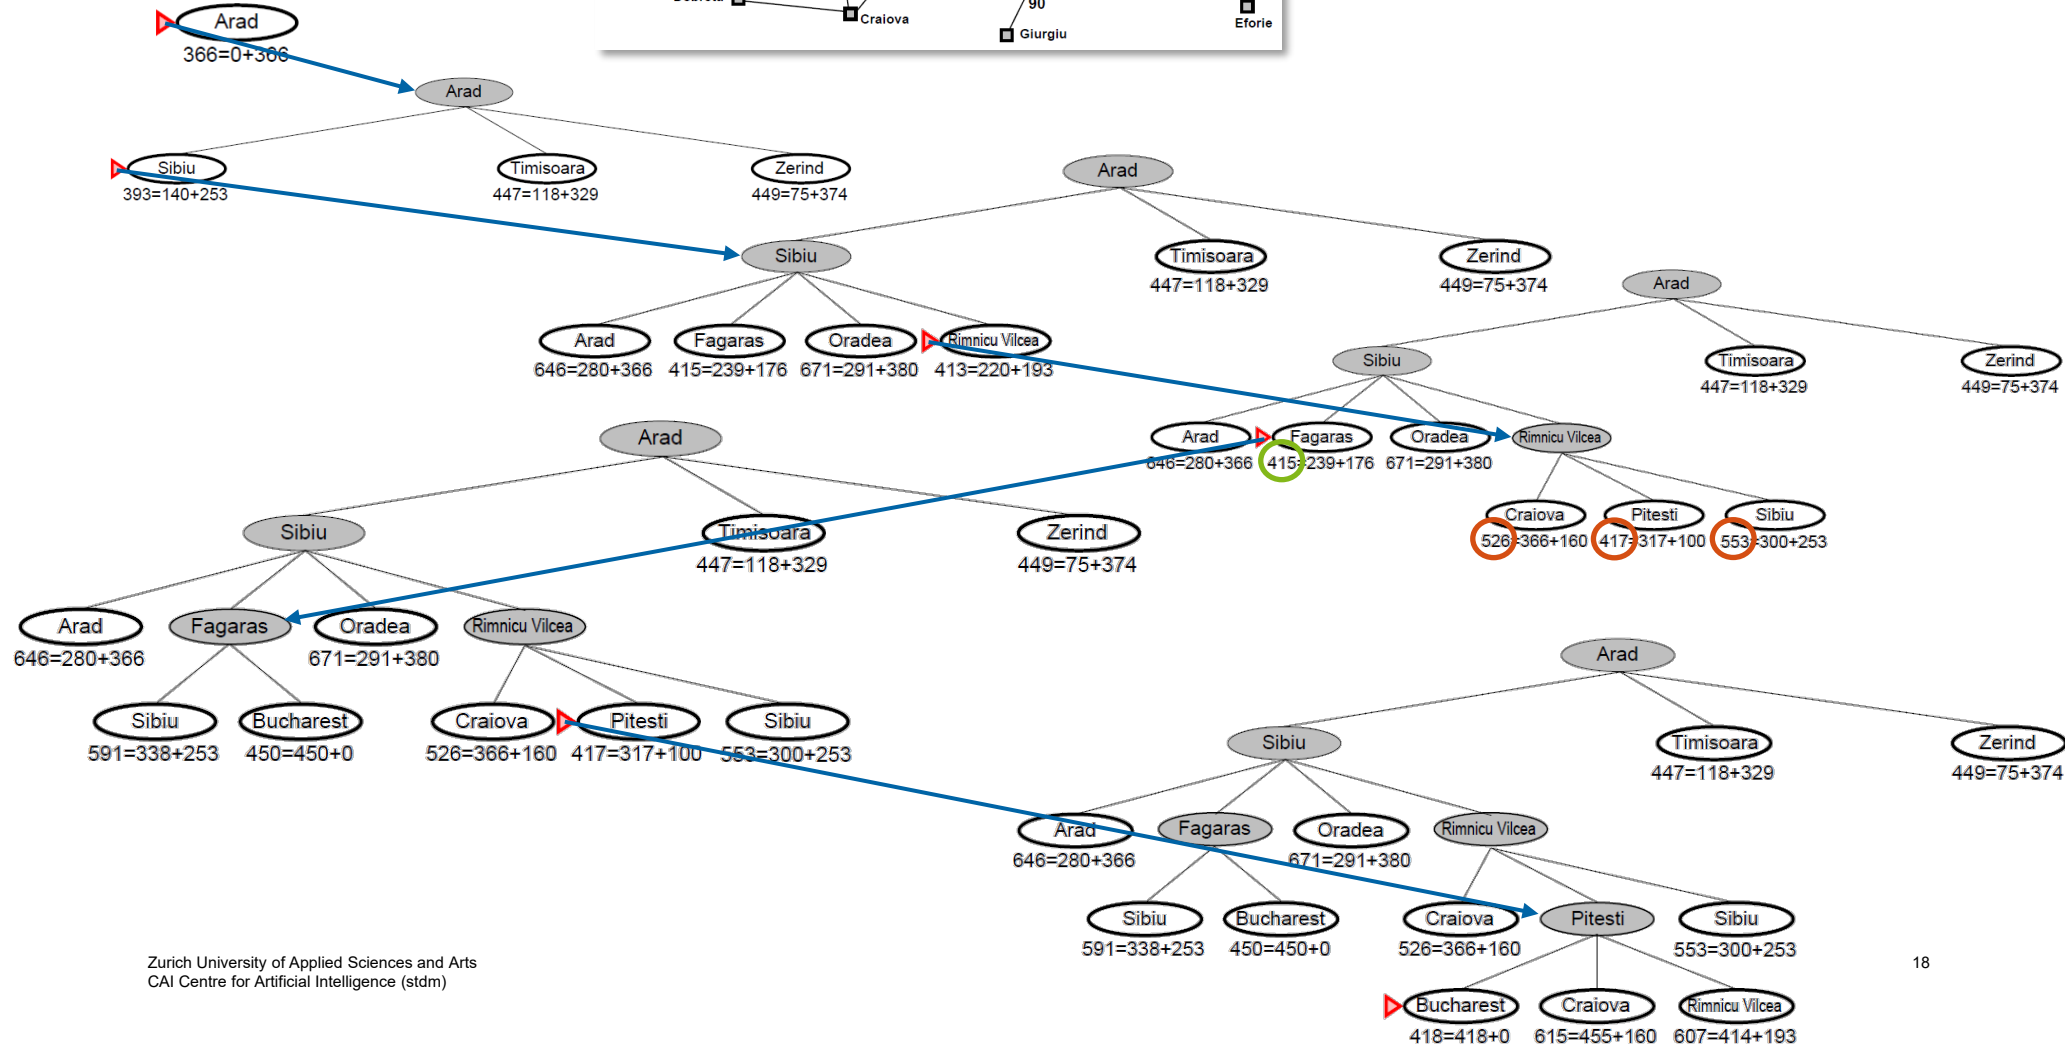
SMA* - simplified **memory-bounded** A*

- A* usually runs out of space first \rightarrow SMA* overcomes this by
- ...filling the memory up, then **starting to forget** the worst expanded nodes
- ...ancestors of forgotten **subtrees remember** the value of the **best path** within them
- ...thus, subtrees are only **regenerated if no better** solution exists

A* Example



Arad	366
Bucharest	0
Craiova	160
Drobeta	242
Eforie	161
Fagaras	176
Giurgiu	77
Hirsova	151
Iasi	226
Lugoj	244
Mehadia	241
Neamt	234
Oradea	380
Pitesti	100
Rimnicu Vilcea	193
Sibiu	253
Timisoara	329
Urziceni	80
Vaslui	199
Zerind	374



Succeeding with search

Learning to search

- **Learn a heuristic** function: use inductive supervised learning on features of a state
- **Alternative: construct a metalevel state space**, consisting of all internal states of search program
Example: For A* searching for a route in Romania, the search tree is its internal state
- Actions in metalevel space: computations that alter the metalevel state
In the example: Expanding a node
- Solution in metalevel space: a path as depicted on the last slide
→ can be input to machine learning algorithms to avoid unnecessary expansions

Practical advice


- **A* is impractical** for large scale problems
- Practical, **robust choice: SMA***
- **Have good heuristic functions!** A well-designed heuristic would have $b^* \approx 1$
(b^* is the effective branching factor)

A closer look on heuristic functions

Example: 8-puzzle

Two proposals – which is better?

- $h_1(n)$ = **number of misplaced tiles**
- $h_2(n)$ = total **Manhattan distance** (i.e., no. of horizontal/vertical squares from desired location of each tile)



$$h_1(S) = 6$$
$$h_2(S) = 4 + 0 + 3 + 3 + 1 + 0 + 2 + 1 = 14$$



Dominance

The 8-puzzle example continues

If $h_2(n) \geq h_1(n) \forall n \rightarrow h_2$ **dominates** h_1 and **is better for search**

Typical search costs

Algorithm	#nodes expanded with $d = 14$	#nodes expanded with $d = 24$
Iterative deepening	3'473'941	~54'000'000'000
$A^*(h_1)$	539	39'135
$A^*(h_2)$	113	1'641

Simple improvement

- Given any admissible heuristics h_a, h_b :
- $h(n) = \max(h_a(n), h_b(n))$ is also admissible and dominates h_a, h_b

Relaxed problems

Improving heuristics intelligently

Relaxation as a key

- Admissible **heuristics** can be **derived from** the **exact solution cost of a relaxed version** of the problem
- A relaxed problem has fewer constraints on the actions
- Relaxation can be automatized!
E.g., «Absolver» by (Prieditis, 1993) found best heuristic for 8-puzzle, first heuristic for Rubik's cube

Examples of relaxed 8-puzzle rules

- If **each tile can move anywhere** (in 1 step), then $h_1(n)$ **gives** the **shortest** solution
- If each tile can move **to any adjacent** square, then $h_2(n)$ **gives** the **shortest** solution

Intuition

- Removing constraints adds edges to the state graph
- Additional edges might provide „**short cuts**“
- The optimal solution cost of a relaxed problem (“short cut”) can be no greater than the optimal solution cost of the real problem

Where's the intelligence?

Man vs. machine

Uninformed search



- In the **abstraction** of the problem
- In the **choice of algorithm** that is optimal for the problem at hand
- In the **systematic exploration** of the state space graph

Heuristic search

- Additionally, in the **heuristic** function

Originally written in German during
his research stay at ETH

→ see also: Polya, «*How to solve it - a new aspect of mathematical method*», 1945



Exercise: Missionaries & cannibals

<https://aimacode.github.io/aima-exercises/search-exercises/>

Three missionaries and **3 cannibals** are on one side of a river, along with a **boat** that **can hold one or two** people. Find a way to **get everyone to the other side, without ever leaving** a group of **missionaries** in one place **outnumbered by the cannibals** in that place.

- Formulate the problem precisely:
Make only those distinctions necessary to ensure a valid solution. Draw a diagram of the complete state space.
- Implement and solve the problem optimally:
Use an appropriate search algorithm. Is it a good idea to check for repeated states?
- Why do you think people have a hard time solving this puzzle, given that the state space is so simple?



Review

- **Search** as an approach to AI **exists in its current form** more or less **since AI's inception**
- **Extensions** of search algorithms exist **to non-deterministic** and **partially observable** environments as well as **online** search
- **Problem formulation** usually **requires abstracting** away real-world details to define a state space that can feasibly be explored
- **Iterative deepening** search **uses only linear space** and not much more time than other uninformed algorithms
- **Graph search** can be **exponentially more efficient** than tree search
- **Good heuristics** can **dramatically reduce** search **cost**
- **A* search expands lowest $g + h$**
→ complete and optimal, also optimally efficient (up to tie-breaks, for forward search)
- **Admissible heuristics** can be **derived from** exact solution of **relaxed problems**



APPENDIX


Fun fact: implement depth-first search in a maze by keeping your left hand on the wall.



On modeling and abstraction

After *AIMA*, sec. 3.1.2

- A **model** [is] an abstract mathematical description [...] and not the real thing
 - The process of removing detail from a representation is called **abstraction**
 - The abstraction is **valid** if we can **expand** any abstract solution into a solution in the more detailed world
 - The abstraction is **useful** if carrying out each of the actions in the abstraction is **easier** than the original problem
 - The choice of a **good abstraction** thus involves **removing as much detail as possible while retaining validity** and **ensuring that the abstract actions are easy** to carry out
- Were it not for the ability to construct useful abstractions, intelligent agents would be completely swamped by the real world



Recap on complexity theory

Problems are classified to be part of (attention: only intuitive “definitions”)

- **P** – can be solved in polynomial time by a deterministic algorithm
→ deemed to be solvable «efficiently»
- **NP** – can only be solved efficiently (i.e., in polynomial time) by guessing the solution
(i.e., by a non-deterministic algorithm)

When people talk about **efficient computation**, this **always means** (at most) **polynomial time**: *efficient ~ polynomial time*.

More terminology

- **NP-hard** – a problem x is said to be NP-hard if **all problems in NP can be reduced** to (i.e., converted into / stated as) x (i.e., can be solved by an algorithm for x) **efficiently**
→ Example: **Traveling salesman problem** (i.e., any problem in NP is at most as hard as x)
 - **NP-complete** – a problem x is said to be NP-complete if it is NP-hard and in NP
→ Example: The **satisfiability problem (SAT)** – is there an assignment of truth values to make a given formula of **propositional logic** true? (→ see V06 and AIMA ch. 7.5)
- ...which is all good (i.e., we don't have to care for efficiency) if $P = NP$ (tremendously unlikely!)

Further reading

- AIMA appendix A.1 (< 3 pages!)
- J. Koehler's lecture slides on complexity and AI: <https://user.enterpriselab.ch/~takoehle/teaching/ai/ProblemComplexity.pdf>
- Some more intuition: <http://stackoverflow.com/questions/1857244/what-are-the-differences-between-np-np-complete-and-np-hard>

Pseudocode for general tree- and graph search

```
function Tree-Search(problem, frontier) returns a solution, or failure
  frontier ← Insert(Make-Node(Initial-State(problem)), frontier)
  loop do
    if frontier is empty then return failure
    node ← Remove-Front(frontier) #choice of picked node defined by strategy
    if Goal-Test(problem) applied to State(node) succeeds return node
    frontier ← InsertAll(Expand(node, problem), frontier)

function Graph-Search(problem, frontier) returns a solution, or failure
  frontier ← Insert(Make-Node(Initial-State(problem)), frontier)
  explored ← empty
  loop do
    if frontier is empty then return failure
    node ← Remove-Front(frontier) #choice of picked node defined by strategy
    explored ← Insert(node, explored)
    if Goal-Test(problem) applied to State(node) succeeds return node
    frontier ← InsertAll(Expand(node, problem), frontier) only if not in frontier or explored set
```

→ ***Bold italic*** font shows the additions that handle repeated states in graph search


Missionaries & cannibals (contd.)

States

- $\theta = (M, C, B)$ signifies the number of missionaries, cannibals, and boats on the left bank
- The **start state is (3,3,1)** and the **goal state is (0,0,0)**

Actions (successor function)


- 10 possible, but only 5 available each move due to boat
- One cannibal/missionary crossing $L \rightarrow R$: subtract (0,1,1) or (1,0,1)
- Two cannibals/missionaries crossing $L \rightarrow R$: subtract (0,2,1) or (2,0,1)
- One cannibal/missionary crossing $R \rightarrow L$: add (1,0,1) or (0,1,1)
- Two cannibals/missionaries crossing $R \rightarrow L$: add (2,0,1) or (0,2,1)
- One cannibal and one missionary crossing: add/subtract (1,1,1)



Source: <http://www.cse.msu.edu/~michmer3/440/Lab1/cannibal.html>

Missionaries & cannibals states

- Assumes that passengers have to get out of the boat after the trip
- Red states** = missionaries get eaten



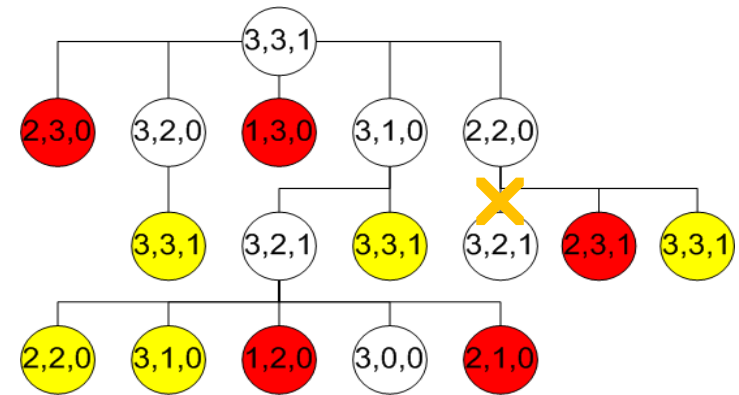
Breadth-first search (4 iterations) on missionaries & cannibals

States are generated by applying

- +/- (1,0,1)
- +/- (0,1,1)
- +/- (2,0,1)
- +/- (0,2,1)
- +/- (1,1,1)

Red states = missionaries get eaten

Yellow states = repeated states



Breadth-first search (final state) on missionaries & cannibals

- Breadth first search expanded 48 nodes
- This is an optimal solution (minimum number of crossings)
- Depth-first search expanded 30 nodes
- ...if repeated states are checked, otherwise we end up in an endless loop

