


Artificial Intelligence

V04: Local and adversarial search

From hill climbing search to genetic algorithms
Game playing
Resource limits and other difficulties

Based on material by Stuart Russell, UC Berkeley



2048 leaderboard link

<https://goo.gl/meh3Ro>

Educational objectives

- **Re-tell** the story of improving **local search** from hill-climbing to genetic algorithms
- **Remember** the **minimax**, α - β and **expectiminimax** algorithms
- **Implement** an **AI agent** for a given simple game

“In which we relax the simplifying assumptions of the previous lecture, thereby getting closer to the real world; including the problems that arise when we try to plan ahead in a world where other agents are planning against us.”

➔ Reading: ALMA, [ch. 4.1-4.2 (local search)]; ch. 6 (games)




1. FROM HILL CLIMBING SEARCH TO GENETIC ALGORITHMS

Local search

Example: n -queens problem

Task

- Put n queens on a $n \times n$ board with no two queens on the same row, column, or diagonal



Possible solution

- Initialize one queen per column
- Move **one queen up/down at a time** to reduce number of conflicts using heuristic h
- Almost always solves n -queens problems almost instantaneously (#states: n^n)
 - works for very large n , e.g., $n = 1'000'000$

Iterative improvement algorithms

Local search: search for optimal states instead of path's

- In many optimization problems, **path is irrelevant**; the **goal state itself is the solution**
 - State space: set of "complete" configurations;
 - Goal: find **optimal** configuration (or a configuration satisfying constraints)
- Examples: TSP, timetable

Iterative improvement

- In such cases: use **iterative improvement** algorithms
 - **Keep a single "current" state**, try to improve it
 - Constant space, suitable for online as well as offline search

Possible implementations

- **Hill climbing**
- **Simulated annealing**
- **Genetic algorithms**

Hill climbing search


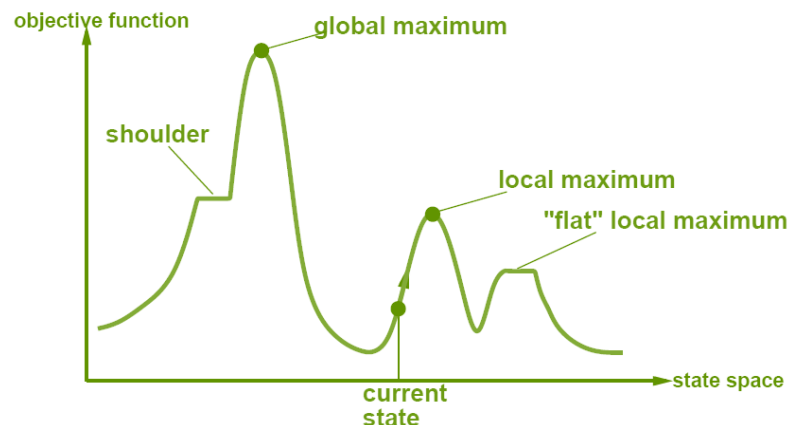
(a.k.a. gradient ascent/descent)

Systematic search for an optimum

- Analogy: «*Like climbing Everest in thick fog with amnesia*»
- Result: finds a state that is a **local maximum**
...by selecting only the highest-valued successor for expansion **iff** its value is better

The state space landscape

- Practical problems typically have an **exponential number of local maxima** to get stuck in
- **Random-restart** hill climbing overcomes local maxima → trivially complete
- **Random sideways** moves escape from shoulders (good), loop on flat maxima (bad)



All previously discussed search algorithms only work in discrete state and action spaces (otherwise the branching factor is infinite)

- 8

Simulated annealing

Towards optimizing hill climbing search

Idea (by [Metropolis et al., 1953] for physical process modelling)

- Escape local maxima by allowing some “bad” moves
- ...but **gradually decrease** their **size and frequency**

Application

- For “**good**” **schedule** of decreasing the temperature (→ see appendix), it **always reaches the best state**
- **Widely applied** for, e.g., VLSI layout, airline scheduling

Modern variants

- «[momentum](#)», «[Adam](#)» and other adaptation strategies for a «[learning rate](#)»
→ first link on the last slide



Local beam search

...and still optimizing hill climbing search

Idea

- **Keep k states** instead of 1; **choose top k** of all their **successors**
(not the same as k searches run in parallel! → Why?)
- **Searches** that find good states **recruit other** searches **to join** them


Not: each

Problem


- Quite often, all k states end up on same local hill

Idea contd.

- **Choose k successors randomly, biased towards good ones**
→ Observe the close analogy to natural selection!



Compare Don Knuth on
“the advantages of
unbiased sampling as a
way to gain insight into a
complicated subject” (e.g.,
ch. 2 in the above book)

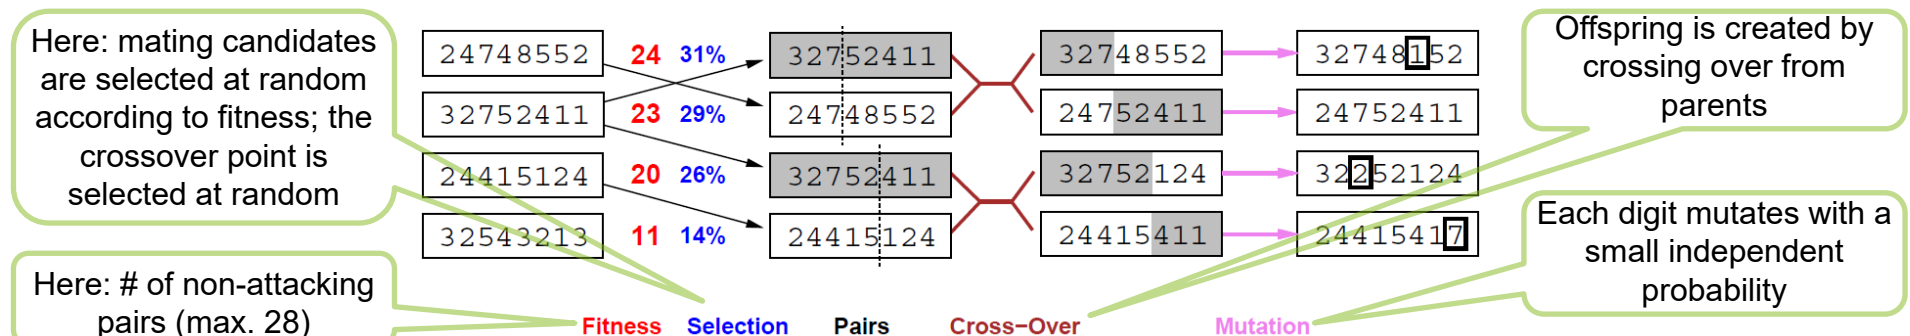


Genetic algorithms (GA)

...improving on the idea of local beam search

Idea

- Combine stochastic local beam search + generating successors from *pairs* of states
 → **uphill tendency + random exploration + exchange of information** among searches



Example: 8-queens states encoded as digit strings. The original population (left) is ranked by a fitness function, resulting in pairs for mating. The offspring is subject to mutation.

Application

- GAs **require states** encoded as strings
- Crossover helps **iif substrings are meaningful** components
- GAs \neq evolution

2. GAME PLAYING

Adversarial search

Games vs. search problems


“Unpredictable” opponent

→ solution is a **strategy** (specifying a move for every possible opponent reply)

Time limits

→ unlikely to find goal, must **approximate**

Which movie?



Early history:

- Computer considers possible lines of play (computer chess: Babbage, 1846)
- Algorithm for perfect play (minimax: Zermelo, 1912; game theory: von Neumann, 1944)
- Finite horizon, approximate evaluation (depth cut-off: Zuse, 1945; Wiener, 1948; Shannon, 1950)
- First chess program (Turing, 1951)
- Machine learning to improve evaluation accuracy (RL for checkers: Samuel, 1952-57)
- Pruning to allow deeper search (α - β search: McCarthy, 1956)

Types of games

	deterministic	stochastic
perfect information	chess , checkers («Dame»), go, othello («Reversi»)	backgammon, monopoly
only partial observability	battleship , kriegspiel (chess without seeing enemy pieces)	bridge (~ «Jass», «Skat»), poker , scrabble, <i>global thermonuclear war</i>

Which movie?



Deterministic (turn-based, 2-player) games

The search tree, e.g. of tic-tac-toe

„Max“

Player's name


- 1st player (moves first)
- Wants to **maximize utility** of terminal states
- Tree shows Max's perspective

„Min“

- 2nd player
- Wants to **minimize** (Max's) **utility**

Utility

- Numeric value („payoff“) of terminal state
- „zero-sum game“ iif total payoff (to all players) is constant over all game instances



Minimax: depth-first exploration of game tree


Perfect play for deterministic, fully observable games

Idea

- Choose move to position with highest **minimax** value
- Minimax value: **highest value among options minimized** by adversary
→ **best** achievable payoff **against best play**

Example

- Any 2-ply game tree (i.e., each player moves once)
- Max's best move at root: a_1 (leading to highest minimax value of 3)
- ...because Min's best reply will be b_1 (leading to lowest minimax value / utility of 3)



Minimax (contd.)

Algorithm and properties

```

function Minimax-Decision(state) returns an action
  inputs: state, current state in game
  return the a in Actions(state) maximizing Min-Value(Result(a, state))

function Max-Value(state) returns a utility value
  if Terminal-Test(state) then return Utility(state)
  v ← -∞
  for a, s in Successors(state) do v ← Max(v, Min-Value(s))
  return v

function Min-Value(state) returns a utility value
  if Terminal-Test(state) then return Utility(state)
  v ← ∞
  for a, s in Successors(state) do v ← Min(v, Max-Value(s))
  return v

```

$$\text{MINIMAX}(s) = \begin{cases} \text{UTILITY}(s) & \text{if } \text{TERMINAL-TEST}(s) \\ \max_{a \in \text{Actions}(s)} \text{MINIMAX}(\text{RESULT}(s, a)) & \text{if } \text{PLAYER}(s) = \text{MAX} \\ \min_{a \in \text{Actions}(s)} \text{MINIMAX}(\text{RESULT}(s, a)) & \text{if } \text{PLAYER}(s) = \text{MIN} \end{cases}$$

Diagram illustrating the Minimax algorithm structure with callouts:





- Min-Value** (pointing to the \min operation in the MIN player's case)
- Max-Value** (pointing to the \max operation in the MAX player's case)

Properties


- Complete? **Yes**, if tree is finite (e.g., chess has specific rules for this)
- Optimal? **Yes**, against an optimal opponent (Otherwise?)
- Time complexity? $O(b^m)$
- Space complexity? $O(bm)$ (for depth-first exploration)
- For chess, $b = 35$, $m = 100$ for “reasonable” games
 - ➔ exact solution completely **infeasible**; but do we need to explore every path?

α - β pruning example

Overcoming exponential (b^m) number of states to be explored



Recursion: A
 $\alpha = 3$
 $\beta = 3$



Successively tightening bounds on minimax values

- α is the **best value** (to Max) found so far in current subtree of a Max node ($A \rightarrow \alpha$)
- If any node v is worse than α , Max will not choose it \rightarrow **prune** that **branch**
- Similarly: β is best score Min is assured of in current subtree of a Min node ($B, C, D \rightarrow \beta$)

α - β pruning (contd.)

Algorithm and properties (changes to minimax in *bold-italic*)

```
function Alpha-Beta-Search(state) returns an action
  v  $\leftarrow$  Max-Value(state,  $-\infty$ ,  $\infty$ )
  return the a in Actions(state) with value v
```

```
function Max-Value(state,  $\alpha$ ,  $\beta$ ) returns a utility value
  if Terminal-Test(state) then return Utility(state)
  v  $\leftarrow$   $-\infty$ 
  for a in Actions(state) do
    v  $\leftarrow$  Max(v, Min-Value(Result(state, a),  $\alpha$ ,  $\beta$ ))
    if v  $\geq$   $\beta$  then return v
     $\alpha \leftarrow$  Max( $\alpha$ , v)
  return v
```

```
function Min-Value(state,  $\alpha$ ,  $\beta$ ) returns a utility value
  if Terminal-Test(state) then return Utility(state)
  v  $\leftarrow$   $\infty$ 
  for a in Actions(state) do
    v  $\leftarrow$  Min(v, Max-Value(Result(state, a),  $\alpha$ ,  $\beta$ ))
    if v  $\leq$   $\alpha$  then return v
     $\beta \leftarrow$  Min( $\beta$ , v)
  return v
```

Properties

- Pruning **does not affect final result**
- Good **move ordering improves** effectiveness of pruning
- With “perfect ordering”, time complexity = $O(b^{m/2})$
 - ➔ **doubles solvable depth**
 - ➔ a simple example of the value of reasoning about which computations are relevant (**metareasoning**)
 - ➔ unfortunately, $35^{100/2}$ (for chess) is still impossible!

3. RESOURCE LIMITS AND OTHER DIFFICULTIES

Adversarial search (contd.)

Resource limits

Towards real-world conditions

Standard approach

- Use **Cutoff-Test** instead of Terminal-Test
e.g., depth limit (perhaps add **quiescence search**: only cut off search at positions that don't drastically change their value in the near future, e.g. captures in chess; otherwise continue search)
- Use **Eval** instead of Utility
i.e., evaluation function that estimates desirability of position
- **Lookup** of start/end games


Example

- Suppose we have 100 seconds, explore 10^4 nodes/second
→ 10^6 nodes per move $\approx 35^{8/2}$
- α - β reaches depth 8
→ pretty good chess program




Eval(uation) functions

Designing or learning effective cutoff tests



(a) White to move



(b) White to move

Example: The two chess positions **differ only in the position of the rook** (“Turm”) at lower right.
 In (a), **Black has an advantage** of a knight (“Springer”) and two pawns (“Bauern”) → should be enough to win.
 In (b), **White will capture the queen** (“Dame”) → should be strong enough to win.

For chess, typically **linear weighted sum of features**


- $Eval(s) = w_1 f_1(s) + w_2 f_2(s) + \dots + w_n f_n(s)$
- Example:
 $w_1 = 9$ and $f_1(s) = (\text{number of white queens}) - (\text{number of black queens})$
- **Can be learned** with machine learning techniques

Nondeterministic (stochastic) games

Chance is introduced by e.g. dice-rolling or card-shuffling



Simplified example

- A game with coin-flipping
- Nondeterminism is **handled by an additional level in the tree**, consisting of **chance nodes**



Real-world example

- **2048**: numbers appear with probability $P(2) = \frac{9}{10}$ and $P(4) = \frac{1}{10}$ at random free board positions
- **Backgammon**: Before each move, dice-rolls determine the legal moves



Expectiminimax – maximizing expected value

Algorithm and properties

```

function ExpectiMinimax-Decision(state) returns an action
  inputs: state, current state in game
  return a in Actions(state) maximizing
    ExpectiMinimax-Value(Result(a, state))

function ExpectiMinimax-Value(state) returns a utility value
  if Terminal-Test(state) then
    return Utility(state)
  if state is a Max node then
    return highest ExpectiMinimax-Value of Successors(state)
  if state is a Min node then
    return lowest ExpectiMinimax-Value of Successors(state)
  if state is a chance node then
    return average of ExpectiMinimax-Value of Successors(state)

```

$$\text{EXPECTIMINIMAX}(s) = \begin{cases} \text{UTILITY}(s) & \text{if } \text{TERMINAL-TEST}(s) \\ \max_a \text{EXPECTIMINIMAX}(\text{RESULT}(s, a)) & \text{if } \text{PLAYER}(s) = \text{MAX} \\ \min_a \text{EXPECTIMINIMAX}(\text{RESULT}(s, a)) & \text{if } \text{PLAYER}(s) = \text{MIN} \\ \sum_r P(r) \text{EXPECTIMINIMAX}(\text{RESULT}(s, r)) & \text{if } \text{PLAYER}(s) = \text{CHANCE} \end{cases}$$

Properties

- Algorithm **works** just like **Minimax** – except chance-nodes are also handled
- Expectiminimax gives **perfect play**
- In case of only **1 player**, Expectiminimax becomes **Expectimax**
- Time complexity: $O(b^m n^m)$ (where n is the number of distinct random events, e.g. dice rolls)
 - ➔ **Possibilities** are **multiplied enormously** in games of chance
 - ➔ Simultaneously, **no** likely sequences exist to do effective **α - β pruning**


Nondeterministic games in practice

Example Backgammon

- Dice rolls increase b (21 possible rolls with 2 dice)
- Ca. 20 legal moves (can be 6,000 with 1-1 roll)
 - at depth 4: $20 \times (21 \times 20)^3 \approx 1.2 \times 10^9$ nodes
- As depth increases, probability of reaching a given node shrinks
 - value of lookahead is diminished (see appendix for consequences)

But

- «TDGammon» (Tesauro, 1992) uses depth-2 search \approx world-champion level
 - uses neural network and reinforcement learning (RL) to train Eval function via games against itself



- → see also <http://webdocs.cs.ualberta.ca/~sutton/book/ebook/node108.html>

Where's the intelligence?

Man vs. machine

- Searching over the right state spaces leads to **emergent “clever” behavior** in games (information gathering, alliances, bluffs)
- Many local and adversarial search methods can be **enhanced by learning** (e.g., learn good Eval functions by playing games against oneself)
- The brain might perform local (hill climbing) search as well
→ see <https://www.cs.toronto.edu/~hinton/backpropincortex2014.pdf>




How to do backpropagation in a brain

Geoffrey Hinton

Canadian Institute for Advanced Research
&
University of Toronto
&
Google Inc.

Google Inc.
&
University of Toronto

- Hill-climbing, minimax, α - β , etc. are still pure computation
- Intelligent behavior emerges from their composition on **suitable** data structures



Review

- **Local search** algorithms **evolve** a **small number of states** towards better utility (typical: 1 state)
 - In **continuous space**, local search by **linear programming** or **convex optimization** is extremely efficient in practice (polynomial time complexity!)
 - In non-deterministic environments, **keeping track of one's belief state** is paramount
 - Games are fun to work on – and dangerous
 - They illustrate several important points about AI
 - **perfection** is **unattainable** → must **approximate**
 - good idea to **think about what to think** about (metareasoning → pruning)
 - **uncertainty** constrains the assignment of values to states
 - optimal decisions depend on **information state** (belief state), not real state
- Games are to AI as grand prix racing is to automobile design






APPENDIX

Pseudocode for hill climbing search and simulated annealing

```
function Hill-Climbing(problem) returns a state that is a local maximum
  inputs: problem, a problem
  local variables: current, a node
                   neighbour, a node
  current  $\leftarrow$  Make-Node(Initial-State[problem])
  loop do
    neighbour  $\leftarrow$  a highest-valued successor of current
    if Value[neighbour]  $\leq$  Value[current] then return State[current]
    current  $\leftarrow$  neighbour
  end
```

```
function Simulated-Annealing(problem, schedule) returns a solution state
  inputs: problem, a problem
          schedule, a mapping from time to "temperature"
  local variables: current, a node
                   next, a node
                   T, a "temperature" controlling the probability of downward steps
  current  $\leftarrow$  Make-Node(Initial-State[problem])
  for t  $\leftarrow$  1 to  $\infty$  do
    T  $\leftarrow$  schedule[t]
    if T = 0 then return current
    next  $\leftarrow$  a randomly selected successor of current
     $\Delta E \leftarrow$  Value[next] - Value[current]
    if  $\Delta E > 0$  then current  $\leftarrow$  next
    else current  $\leftarrow$  next only with probability  $e^{\frac{\Delta E}{T}}$ 
```


Digression: Exact Eval values do matter ...for nondeterministic games



- Behavior is **preserved only by positive linear transformation** of Eval
- Hence **Eval** should be **proportional to the expected payoff**

→ Exact values **don't matter for deterministic** games → see appendix

Digression: Exact Eval values don't matter ...for deterministic games



- Behavior is preserved under any **monotonic** transformation of Eval
- Only the **order matters**: payoff in deterministic games acts as an **ordinal utility** function

Deterministic games in practice

Checkers

- «*Chinook*» ended 40-year-reign of human world champion Marion Tinsley in 1994
- Used an endgame database defining perfect play for all positions involving 8 or fewer pieces on the board (total: 443,748,401,247)



Chess

- «*Deep Blue*» defeated human world champion Gary Kasparov in a six-game match in 1997
- Searches 200 million positions per second, uses very sophisticated evaluation, and undisclosed methods for extending some lines of search up to 40 ply



Othello

- Human champions refuse to compete against computers, who are **too good**



Go

- 2010: human champions refuse to compete against computers, who are **too bad**
- 2016: Google DeepMind's «*AlphaGo*» unexpectedly defeats world champion Lee Sedol
- $b > 300 \rightarrow$ most programs use pattern knowledge bases to suggest plausible moves



Expectations & expected values

adapted from U Washington's CSE473, lecture 8

The expectation operator $E()$

- We can define a function $f(X)$ of a **random variable** X
- The **expected value** of a function is its **average** value **under** the **probability distribution** **over** the function's **inputs**:

$$E(f(X)) = \sum_x f(X = x)P(X = x)$$

Example

- How long to drive to the airport?
- Driving time D (in mins) as a function of traffic T :
 $D(T = none) = 20, D(T = light) = 30, D(T = heavy) = 60$
- What is your expected driving time?
Let probability $P(T) = \{none: 0.25, light: 0.5, heavy: 0.25\}$
 $\rightarrow E(D(T)) = D(none)P(none) + D(light)P(light) + D(heavy)P(heavy)$
 $\rightarrow E(D(T)) = 20 * 0.25 + 30 * 0.5 + 60 * 0.25 = 35 \text{ mins}$

Why averaging over clairvoyance is wrong

A common sense example of a journey

Day 1

- Road A leads to a small heap of gold pieces
 - Road B leads to a fork:
 - the left fork **leads to** a bigger heap of gold;
 - take the right fork and you'll be run over by a bus.
- „B“ (and then „left“) is the optimal choice

Day 2

- Road A leads to a small heap of gold pieces
 - Road B leads to a fork:
 - **take** the left fork and you'll be run over by a bus;
 - the right fork leads to a bigger heap of gold.
- „B“ (and then „right“) is the optimal choice

Day 3

- Road A leads to a small heap of gold pieces
 - Road B leads to a fork:
 - **guess** correctly and you'll find a bigger heap of gold;
 - guess incorrectly and you'll be run over by a bus.
- „B“ still seems optimal; but this **ignores** the resulting **belief state** (that includes ignorance & possibility of death!)



Averaging over clairvoyance **will never** select actions to **gather information** because it assumes future states to be of perfect knowledge after the initial deal.