


# Artificial Intelligence

## V05: Constraint satisfaction problems

Introduction to CSPs  
CSP solving  
Solving CSPs in practice

Based on material by Stuart Russell, UC Berkeley



# Recap V01 – V04

**AI:** a set of tools to solve (different kinds of) complex problems

**Search:** one compartment in the toolbox, suitable for specific problems (i.e., finding some optimal sequence of actions)

**Super power:** efficient algorithms plus human-made heuristics

**Fundamental limitation:** atomic states → no concept of similarity-degrees among states, hence no stronger reasoning than trial & error (goal test) → gigantic runtime for any but the simplest problems



# Educational objectives

- **Remember** what makes **CSP** solving **more powerful** than pure search techniques
- **Explain** how CSPs are solved **on the algorithmic level** by **backtracking** using the **MRV / degree- / least constraining value** heuristics and **forward checking / constraint propagation**
- **Formulate** a suitable problem as a **CSP**

*“In which we see how treating states as more than just little black boxes leads to new search methods and a deeper understanding of problem structure.”*

➔ Reading: ALMA, ch. 5



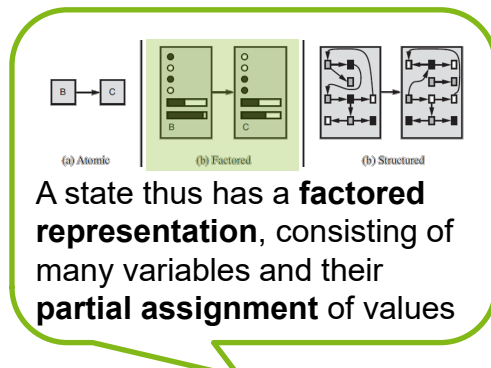


# 1. INTRODUCTION TO CSPS

# Constraint satisfaction problems (CSPs)

## Standard search problem

- State is a “**black box**” – any data structure that supports Goal Test, Eval, Successor



5	3			7			
6			1	9	5		
	9	8				6	
8				6			3
4			8		3		1
7				2			6
	6					2	8
			4	1	9		5
				8		7	9

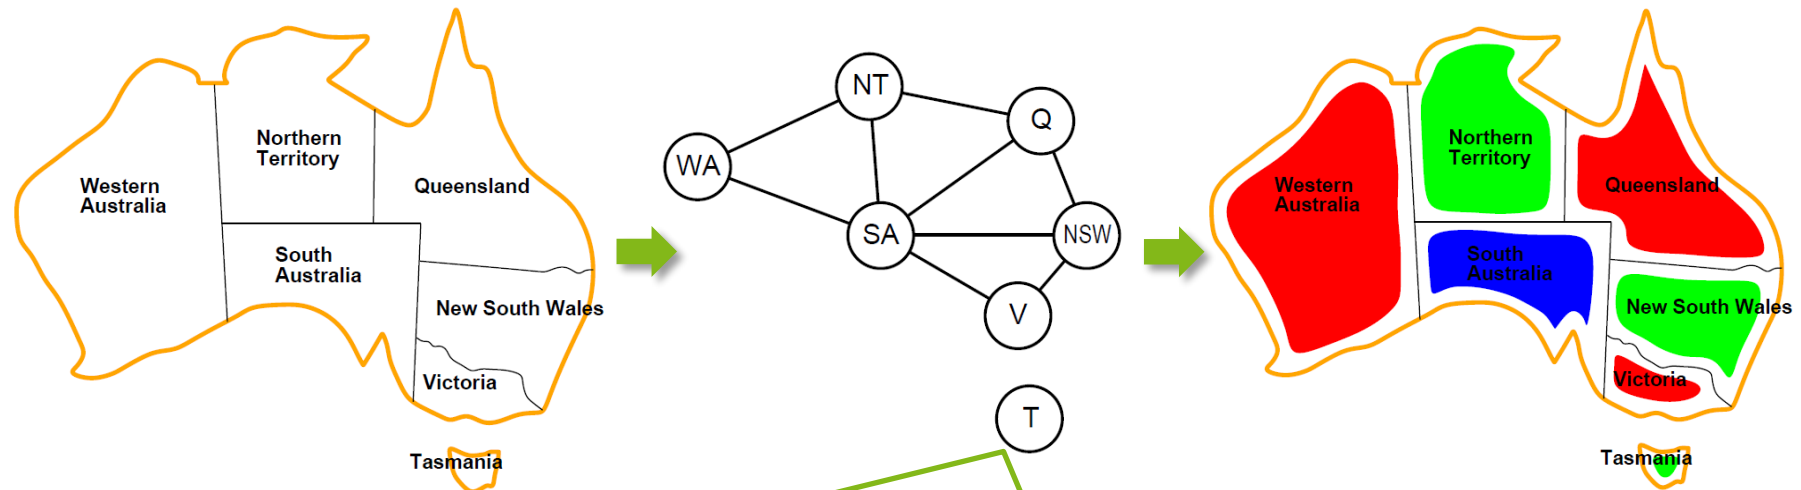
## CSP

- State** is defined by **variables**  $X_i$  with **values** from **domain**  $D_i$
- Goal Test** is a set of **constraints**: allowable combinations of values for subsets of variables

➔ Simple example of a **formal representation language**

➔ Allows useful **general-purpose algorithms** with **more power** than standard search

# Example: Map-coloring



Binary CSPs (each constraint relates at most two variables) have a **constraint graph**. General-purpose CSP algorithms use the graph structure to **speed up search**: E.g.,  $T$  is an independent subproblem!

Variables:  $WA, NT, Q, NSW, V, SA, T$

Domains:  $D_i = \{red, green, blue\}$

Constraints: adjacent regions must have different colors

- e.g.,  $WA \neq NT$  (if language allows this; otherwise  $(WA, NT) \in \{(red, green), (red, blue), (green, red), (green, blue), \dots\}$ )

Solutions: assignments satisfying all constraints

- e.g.,  $\{WA = red, NT = green, Q = red, NSW = green, V = red, SA = blue, T = green\}$

# Varieties of CSPs

## Discrete variables

- Finite domains of **size**  $d \rightarrow O(d^n)$  complete assignments ( $n$  is number of variables)
- Other finite domains (integers, strings, etc.)
  - e.g., job scheduling: variables are days (or integer-minutes) for each job
  - need a **constraint language**, e.g.,  $StartJob_1 + 5 \leq StartJob_3$
  - linear constraints solvable, nonlinear undecidable



## Continuous variables

- e.g., precise start/end times for Hubble Telescope observations
- linear constraints solvable in polynomial time by linear programming methods

## Varieties of constraints

- **Unary** constraints: involve a single variable, e.g.,  $SA \neq green$
- **Binary** constraints involve variable pairs, e.g.,  $SA \neq WA$  (**all constraints can be made binary**)
- **Higher-order** constraints involve 3 or more variables, e.g., column constraints in Sudoku
- **Preferences** (soft) constraints, e.g.,  $red IS\_BETTER\_THAN green$ 
  - often representable by a cost for each assignment: **constrained optimization problems (COP)**



# Examples

## Car assembly

(job scheduling, simplified)

- Variables:  $Axle_F$ ,  $Axle_B$ ,  $Wheel_{RF}$ ,  $Wheel_{LF}$ ,  $Wheel_{RB}$ ,  $Wheel_{LB}$ ,  $Nuts_{RF}$ ,  $Nuts_{LF}$ ,  $Nuts_{RB}$ ,  $Nuts_{LB}$ ,  $Cap_{RF}$ ,  $Cap_{LF}$ ,  $Cap_{RB}$ ,  $Cap_{LB}$ ,  $Inspect$  (tasks to be completed)
- Domains:  $D_i = \{1, 2, 3, \dots, 27\}$  (start time of tasks as integer, due to an overall runtime of 30 minutes)
- Constraints: (precedence constraints among tasks)
  - $Axle_F + 10 \leq Wheel_{RF}$ ;  $Axle_F + 10 \leq Wheel_{LF}$
  - $Axle_B + 10 \leq Wheel_{RB}$ ;  $Axle_B + 10 \leq Wheel_{LB}$
  - $Axle_F + 10 \leq Axle_B$  **or**  $Axle_B + 10 \leq Axle_F$
  - ...

Installing an axle takes 10 minutes and must be prior to wheel assembly

Only one shared tool for axle installing, so can't be simultaneous



## Cryptarithmic


(which letter represents which digit?)

- Variables:  $F, T, U, W, R, O, C_1, C_2, C_3$
- Domains:  $\{0, 1, 2, 3, 4, 5, 6, 7, 8, 9\}$
- Constraints:
  - $alldiff(F, T, U, W, R, O)$
  - $O + O = R + 10C_1$
  - $C_1 + W + W = U + 10C_2$
  - $C_2 + T + T = O + 10C_3$
  - $C_3 = F$

$$\begin{array}{r} T \ W \ O \\ + \ T \ W \ O \\ \hline F \ O \ U \ R \end{array}$$

$C_1, C_2, C_3$ : auxiliary variables for carryover

A so-called **global constraint** involves an **arbitrary number** of variables

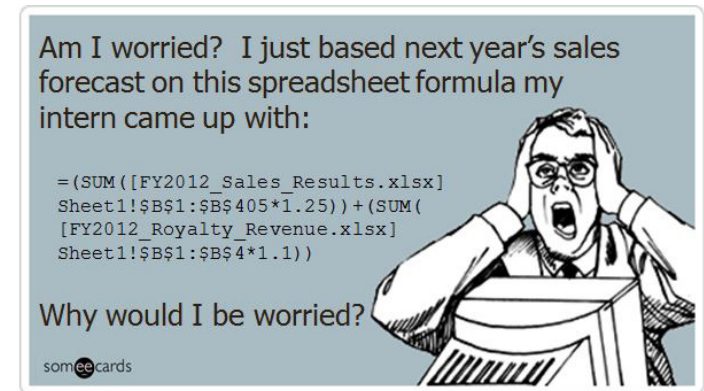


Constraint **hypergraphs** have square (hyper-)nodes for  $n$ -ary constraints



# Real-world CSPs

- **Assignment** problems  
e.g., who teaches what class
- **Timetabling** problems  
e.g., which class is offered when and where?
- **Optimization** with spreadsheets  
e.g., debugging (Abreu, Ribeiro & Wotawa, 2012)
- Other **scheduling** tasks  
e.g., in transportation or factory workflow
- Other **layout** tasks  
e.g., floor planning or hardware configuration



➔ Notice that many real-world problems involve real-valued variables

# Exercise: Formulating Sudoku as a CSP

→ see also P03

Sudoku puzzles are played on a 9x9 board and enjoyed by millions of people daily. The goal is to fill in each cell with a single digit, subject to several constraints:

- Each digit must be present in each row exactly once
- Each digit must be present in each column exactly once
- Each digit must be present in each box exactly once  
(the 9x9 board consists of 9 non-overlapping 3x3 boxes  
→ see thicker lines below)
- Each digit must be consistent with any digit already placed on the original board by the riddle issuer

→ Formulate the Sudoku riddle below as a CSP **using pen & paper** (i.e., decide on variables, domains and constraints)

5	3			7				
6			1	9	5			
	9	8					6	
8				6				3
4			8		3			1
7				2				6
	6					2	8	
			4	1	9			5
				8			7	9



## 2. CSP SOLVING

# Standard search formulation

## Seriously flawed, thus incremental

Let's start with the straightforward, dumb approach, then fix it

- **States** are defined by the **values assigned so far**
  - Initial state: the empty assignment  $\{\}$
  - Successor function: assign a value to an unassigned variable without conflict with current assignment  
→ fail if no legal assignment (not fixable!)
  - Goal test: the current assignment is complete
- CSPs all have a common structure  
→ This is the same for all CSPs, **no domain-specific adaptations** (transition models etc.) needed! 😊
- Every solution appears at depth  $n$  (for  $n$  variables)  
→ use **depth-first search**
- Path is irrelevant, so can also use **complete-state formulation** (as with local search)  
→ i.e., **evolve one state** instead of creating new ones
- Branching factor  $b = (n - l)d$  at depth  $l$   
→ hence  **$n! d^n$  leaves!** 😞 😞 😞




# Backtracking search

## First improvement

- Variable assignments are **commutative**  
e.g.  $[WA = red, then NT = green]$  same as  $[NT = green, then WA = red]$ 
  - Only need to consider assignments to a single variable at each node
  - $b = d$ , thus there are  $d^n$  leaves


## Backtracking search

- Using depth-first search with single-variable assignments for CSPs is called backtracking search
- It is the basic uninformed algorithm for CSPs
  - Can solve  $n$ -queens for  $n = 25$



Remember V04: simple heuristic solves 1'000'000-queens...

Backtracking search considers one variable per level



# Backtracking search

## Algorithm & suggested improvements

```
function Backtracking-Search(csp) returns solution/failure
    return Backtrack({}, csp)

function Backtrack(assignment, csp) returns solution/failure
    if assignment is complete then return assignment
    var ← Select-Unassigned-Variable(csp)
    for each value in Order-Domain-Values(var, assignment, csp) do
        if value is consistent with assignment then
            add {var = value} to assignment
            inferences ← Inference(csp, var, value)           #optional
            if inferences ≠ failure then                       #optional
                add inferences to assignment                  #optional
            result ← Backtrack(assignment, csp)
            if result ≠ failure then return result
        else remove {var = value} from assignment
    return failure
```

General-purpose methods can give huge gains in speed:

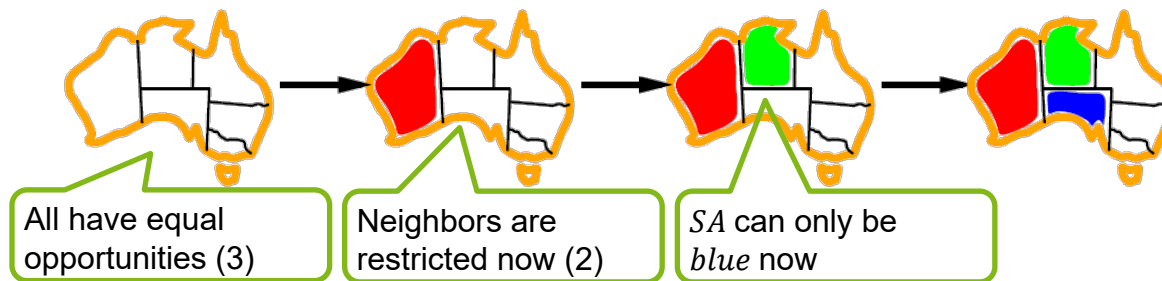
- **Which variable** should be assigned next?
  - In what **order** should its **values** be tried?
  - Can we **detect** inevitable **failure early**?
  - Can we **take advantage** of **problem structure**?
- ➔ can be achieved by implementing the *bold/italic* functions above

# Which variable should be assigned next?

## Ideas for *Select-Unassigned-Variable* (*csp*)


### Minimum remaining values (MRV):

- **Choose** the variable with the **fewest legal values**  
→ **failing fast** prunes large portions of the tree
- Can work up to 1'000 times better than picking just the next (or a random) unassigned variable (very problem dependent)



### Degree heuristic

- **Choose** the variable that adds **most constraints on remaining** variables  
→ In practice: Used as **tie-breaker** among MRV variables






# In what order should its values be tried?

Ideas for *Order-Domain-Values* (*var*, *assignment*, *csp*)

## Least constraining value

- Given *var*, **choose** the value that **rules out the fewest values** in the remaining *vars*  
→ Combining this with the previous 2 heuristics makes 1'000-queens feasible (instead 25)

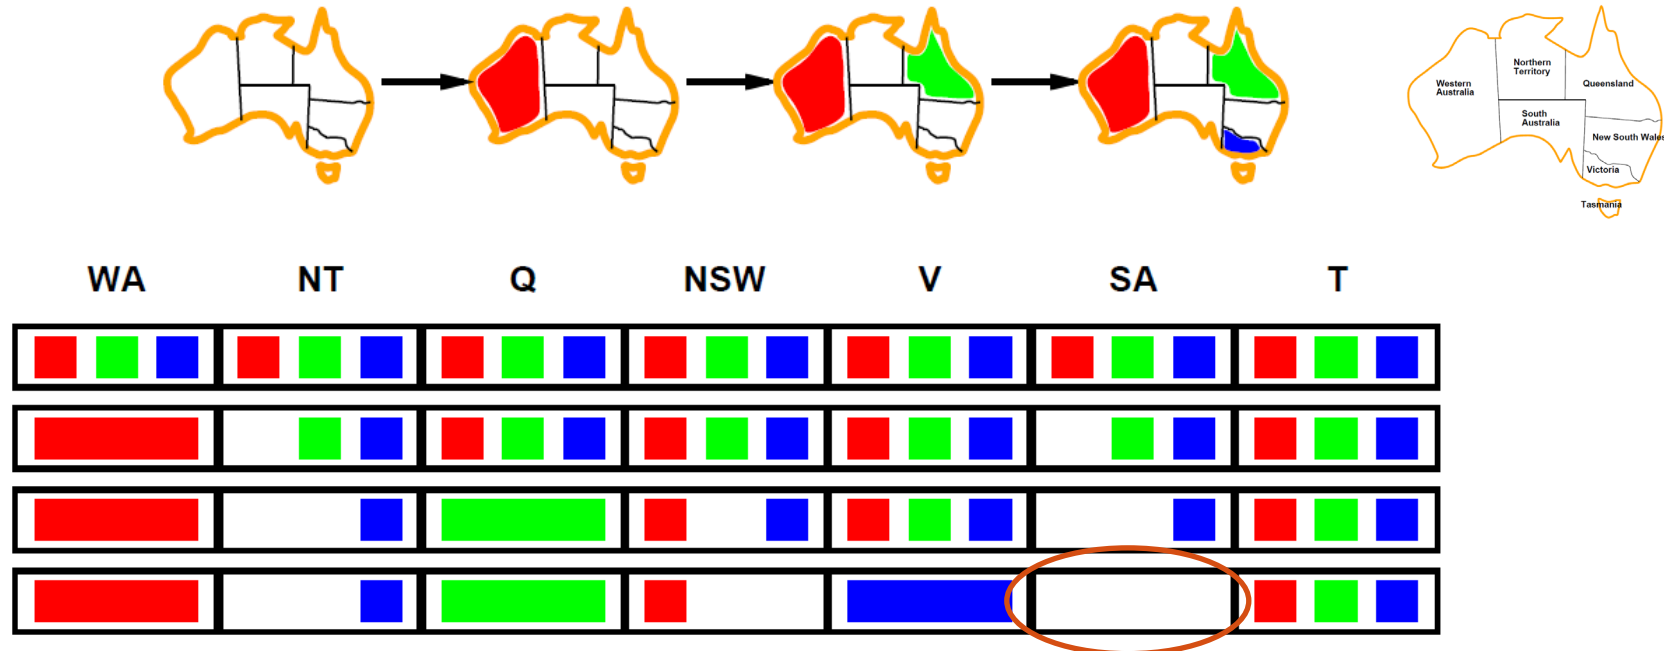


# Can we detect inevitable failure early?

Ideas for *Inference* (*csp*, *var*, *value*)

## Forward checking

- Idea: Keep **track** of **remaining legal values** for unassigned variables  
 → Terminate search when any variable has no legal values


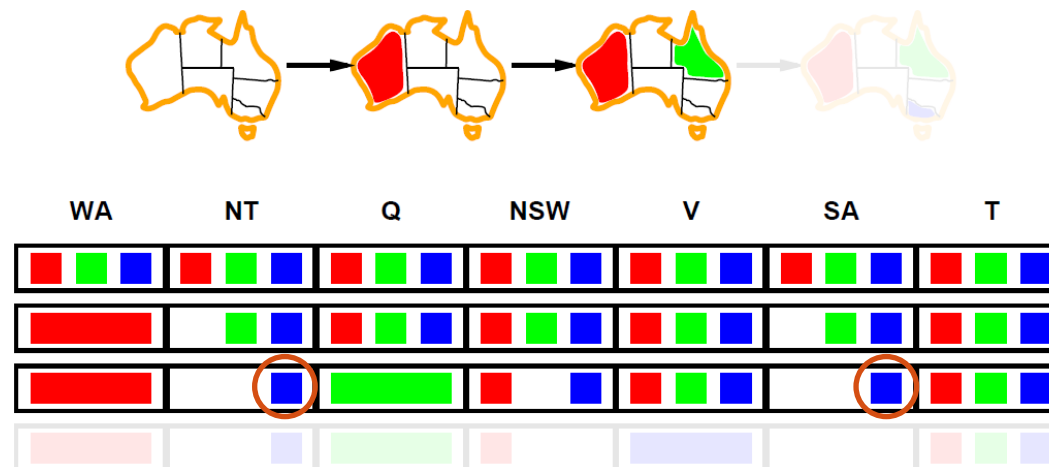


# Can we detect inevitable failure early? (contd.)

Ideas for *Inference* (*csp*, *var*, *value*)

## Constraint propagation

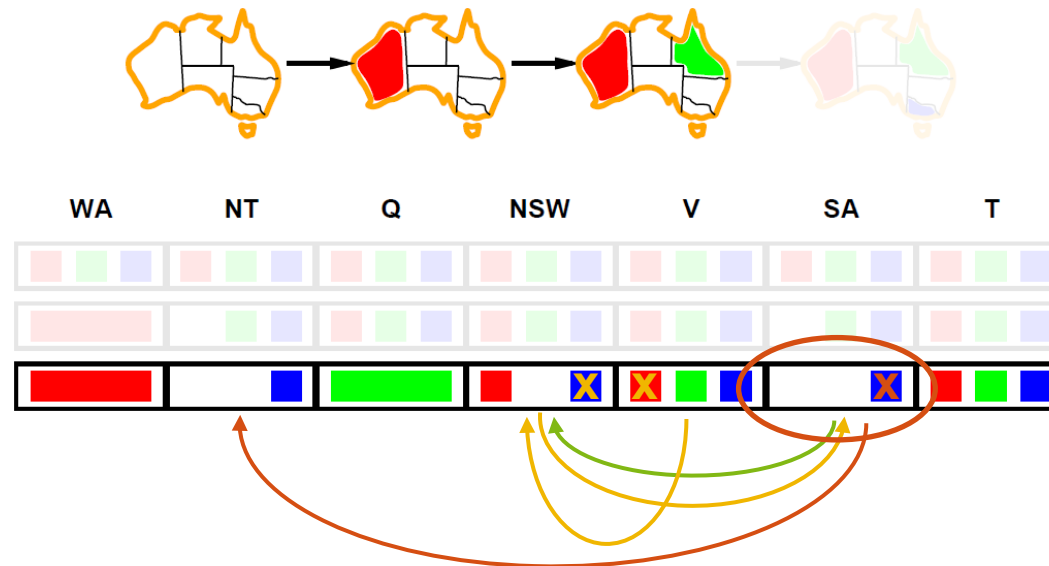
- **Forward checking** propagates information from assigned variables **only to immediate neighbours** (i.e., fails to do so recursively after a change in some domain)  
→ e.g., *NT* and *SA* cannot both be *blue*!



→ Constraint propagation would repeatedly **enforce constraints locally**

## Ideas for *Inference*(*csp*, *var*, *value*)

- $X \rightarrow Y$  is consistent *iff* for every value  $x$  of  $X$  there is some allowed  $y$  for  $Y$
- Arc consistency detects failure earlier than forward checking (by making **every** arc consistent)  
 ➔ Can be run as a preprocessor or after each assignment



-

# Backtracking search

## Revisiting suggested improvements

```
function Backtracking-Search(csp) returns solution/failure
  return Backtrack({}, csp)

function Backtrack(assignment, csp) returns solution/failure
  if assignment is complete then return assignment
  var ← Select-Unassigned-Variable(csp)
  for each value in Order-Domain-Values(var, assignment, csp) do
    if value is consistent with assignment then
      add {var = value} to assignment
      inferences ← Inference(csp, var, value)           #optional
      if inferences ≠ failure then                       #optional
        add inferences to assignment                   #optional
        result ← Backtrack(assignment, csp)
        if result ≠ failure then return result
      else remove {var = value} from assignment
  return failure
```

General-purpose methods can give huge gains in speed:


- **Which variable** next? MRV (fewest legal values), degree heuristic (most constraints on rest) on tie
- What **value first**? Least constraining value
- How **detect failure early**? Constraint propagation via arc consistency
- Can we **take advantage of problem structure**? → next

# Can we take advantage of problem structure?

## Exploiting structure in the constraint graph

### Example

- Tasmania and mainland are independent **subproblems**, identifiable as **connected components** of constraint graph  
→ can be solved individually, and solution combined
- Suppose each subproblem has  $c$  variables (out of  $n$  total)  
→ Worst-case solution cost is  $n/c \cdot d^c$  (linear in  $n$ )
- This is a **dramatic improvement!**
  - E.g.,  $n = 80$ ,  $d = 2$ ,  $c = 20$ :
    - $2^{80} = 4$  billion years (at 10 million nodes/second)
    - $4 \cdot 2^{20} = 0.4$  seconds (at 10 million nodes/second)

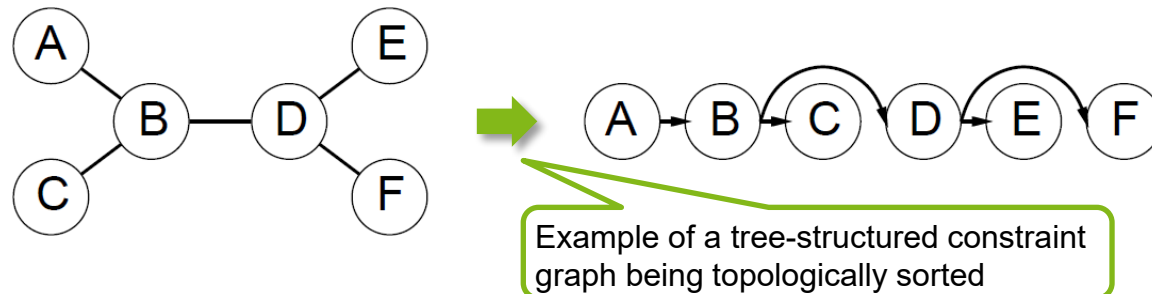


# Can we take advantage of problem structure?

## Exploiting structure in the constraint graph (contd.)

### Tree-structured CSPs

- A (constraint) graph is a **tree** if any **2 variables** are **connected by only 1 path** (i.e., no loops)
- **Theorem:** If the constraint graph has **no loops**, the CSP can be solved in  $O(nd^2)$  time
  - Compare to **general CSPs**, where **worst-case** time is  $O(d^n)$
  - Also applies to logical and probabilistic reasoning
  - Important example of the relation between **syntactic restrictions** and the **complexity of reasoning**



### Algorithm for tree-structured CSPs

- Do a **topological sort**: **Choose** a variable as **root**, then **order variables** from root to leaves such that every node's parent precedes it in the ordering
- Create **directed arc-consistency** by: For  $j$  from  $n$  down to 2, make  $(Parent(X_j), X_j)$  arc consistent
- For  $j$  from 1 to  $n$ , **assign**  $X_j$  consistently with  $Parent(X_j)$



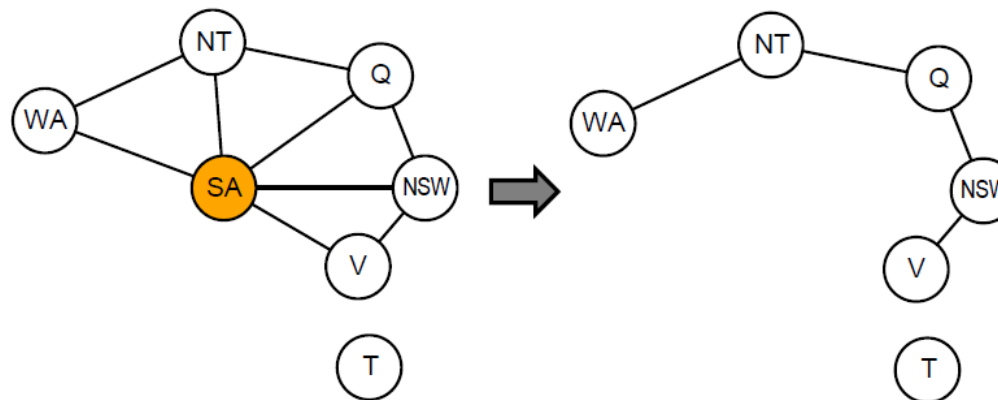


### 3. SOLVING CSPS IN PRACTICE

# Exploiting non-optimal structure

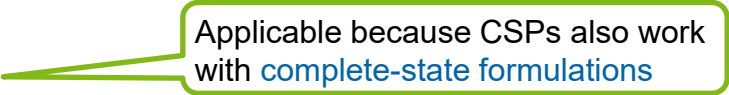
## Nearly tree-structured CSPs

- Many real-world CSPs can be converted to tree-structured problems  
→ then solved by **divide & conquer**
- ...by choosing a **cycle cutset**: a set of variables that if removed make the graph a tree



- ...and subsequent **cutset conditioning**: instantiate (in all ways) the variables in the cutset, then prune choices from remaining variables in the tree  
→ Very fast for small cutset size  $c$ : Runtime is  $O(d^c \cdot (n - c)d^2)$  (linear in  $n$ )


# Other advice

- **Exploiting structure in the values by breaking symmetry** reduces search space up to  $d!$  (e.g., we must give  $WA$ ,  $NT$ ,  $SA$  3 different colors, but have  $3!$  options to do so)  
→ can be reduced by adding a symmetry-breaking constraint like  $NT < SA < WA$
- **Local search** (→ see V04) is **very effective** for CSPs 
  - **Min-conflicts** heuristic very useful: start with random full assignment, subsequently change the variable that minimizes remaining conflicts
  - E.g., hill climbing search with min-conflicts solves  $n$ -queens in constant time with high probability (even for  $n = 10'000'000$ )
- **Constraint learning** (→ see appendix) is one of the **most important techniques** in modern CSP solvers  
(together with backtracking search, the MRV / degree- / least constraining value heuristics, and forward checking / arc consistency)
- **Trade-off** between the **cost of enforcing consistency** and the **reduction in search time**  
(some researchers favor pure forward checking, some full arc consistency after each assignment)  
→ full arc consistency pays off for harder CSPs)
- Comparing CSP algorithms is done empirically (**no algorithm dominates** on all CSPs)

# Where's the intelligence?

## Man vs. machine

- If classical **search is brute force**...
- ...**CSP** solving **enhances** it using the following powerful ingredients:
  - **General-purpose** heuristics  
(MRV etc. → not problem- or domain specific!)
  - **Inference** over constraints  
(constraint propagation → allows e.g., for intelligent **backjumping**)
  - **Exploiting structure** in the problem definition to vastly prune the search space  
(e.g., symmetric values, tree-like constraint graph → implements a **general divide & conquer** approach)
- CSP solving thus can reduce the **time complexity** of some problems **from exponential to linear**, by **acting more “clever”**
- **Human** intelligence goes into **stating the task** as a CSP



# Review

- CSPs are a special kind of problem:
  - **states** defined by values of a **fixed set of variables**
  - **goal test** defined by **constraints on variable values**
- **Backtracking = depth-first search** with one variable assigned per node
  - **Variable ordering** and **value selection heuristics** help significantly
  - **Forward checking prevents** assignments that guarantee **later failure**
  - **Constraint propagation** (e.g., arc consistency) does **additional** work to constrain values and detect **inconsistencies**
- The CSP representation allows **analysis of problem structure**
  - **Tree-structured** CSPs can be solved in **linear time**
  - **Iterative min-conflicts** is usually effective in **practice**
- **Methods** can **handle** problems with **up to 100'000 variables**, and up to **1'000'000 constraints** in practice





## APPENDIX

# Arc consistency

## AC-3 Algorithm

```
function AC-3(csp) returns the CSP, possibly with reduced domains
  inputs: csp, a binary CSP with variables  $(X, D, C)$ 
  local variables: queue, a queue of arcs, initially all the arcs in csp
  while queue is not empty do
     $(X_i, X_j) \leftarrow \text{Remove-First}(\text{queue})$ 
    if Revise(csp,  $X_i, X_j$ ) then
      if size of  $D_i = 0$  then return false
      for each  $X_k$  in  $X_i.\text{Neighbors} - \{X_j\}$  do
        add( $X_k, X_i$ ) to queue
  return true

function Revise(csp,  $X_i, X_j$ ) returns true iff we revise the domain of  $X_i$ 
  revised  $\leftarrow$  false
  for each  $x$  in  $D_i$  do
    if no value  $y$  in  $D_j$  allows  $(x, y)$  to satisfy the constraint  $X_i$  and  $X_j$  then
      delete  $x$  from  $D_i$ 
      revised  $\leftarrow$  true
  return revised
```

- After applying AC-3, either every arc is consistent or some variable has an empty domain  
→ CSP not solvable
- Time complexity:  $O(n^2 d^3)$  (can be reduced to  $O(n^2 d^2)$ , but detecting all is NP-hard)
- Trivia: Name stems from this algorithm being the third one in the paper (Mackworth, 1977)



# Can we detect inevitable failure early? (contd.)

Ideas for *Inference*(*csp*, *var*, *value*)

## Constraint learning

- If `Backtrack()` fails on  $X_i$ , it **backs up to the last variable** and tries another value  
→ would be more intelligent to track back to one of the variables that caused  $D_i = \{\}$
- Forward checking etc. already has this information  
→ can be stored in a **conflict set**
- Constraint learning **adds new constraints** on the fly for sets of assignments (so-called **no-goods**) that **repeatedly** caused `Backtrack()` to fail