

DVOP1 - Cloud Native Application Composition

Prof. Dr. Thomas M. Bohnert
Christof Marti

Content

- Architecture Recap
- Decomposition
- Composition
 - Centralised v.s. Decentralised
 - Declarative v.s. Imperative
 - Standards & Defacto
- Technologies
 - docker-compose
 - Kubernetes Helm

Recap: Architecture

SOA Principles:

- Standardized protocols (e.g., SOAP, REST)
- Abstraction (from service implementation)
- **Loose coupling**
- Reusability
- **Composability**
- Stateless services
- Discoverable services

Microservices architecture is a SOA architectural style to develop applications

- as a **suite of “small” services**,
- each **running in its own process**
- and communicating with **lightweight** mechanisms (REST APIs or Messaging).
- They are **built around business capabilities** following the “do one thing well” principle.

Services are highly **decoupled** (yet ***composed!***)
and focus on doing a small task.


Decomposition

- Key enabler of SOA and Microservice Architectures
- **Reduce** the application into a set of independent functional services
 - A service is a unit of software that is independently replaceable and upgradeable
 - A service encapsulates functionality and enforces an API
 - A service is independently deployable
 - A service lends itself to a continuous delivery software development process
 - Complex applications are **composed** of small, independent services (processes) and communicate with each other using APIs

Decomposition & Domain Driven Design (DDD)

Overview

- A software development methodology used for complex systems
- Supports evolution of the design
- Primary focus of designs are domains
- Requires collaboration between domain (business) and technical expert

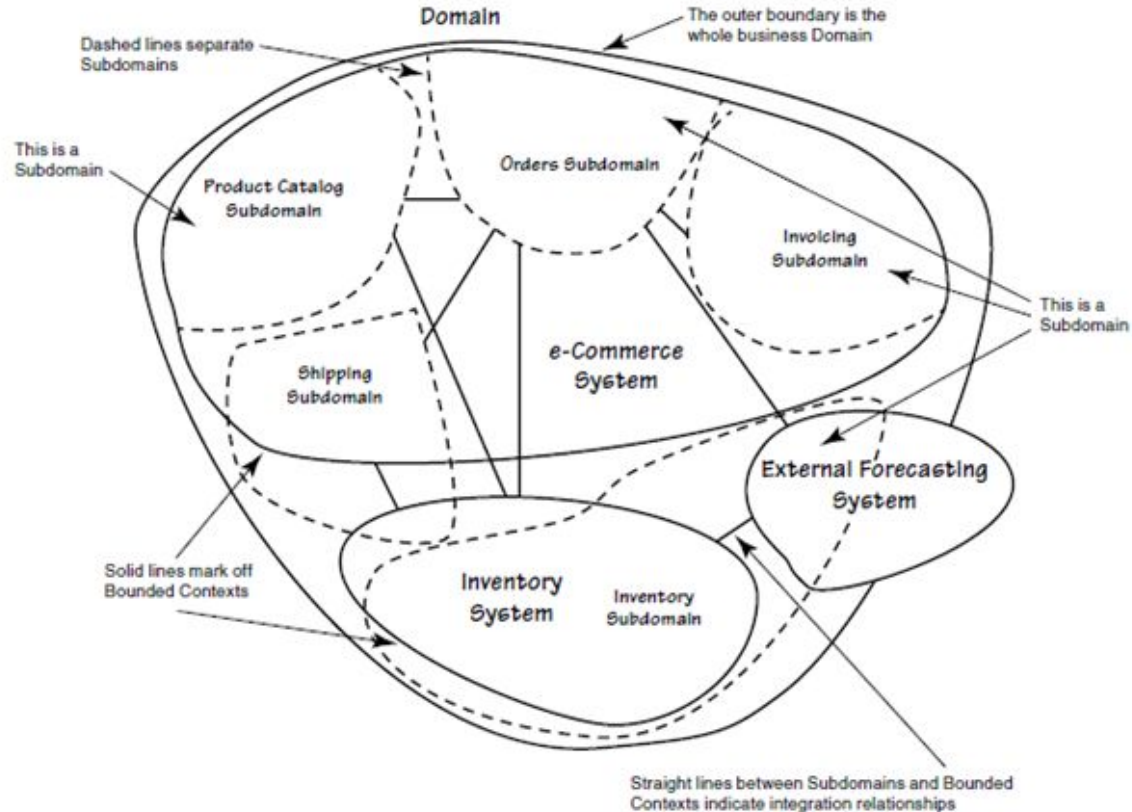


Core Concepts

- **Context:** the setting in which a concept appears determines its meaning
- **Domain:** specific area of knowledge or activity. Composed of **sub-domains**.
- **Model:** system of abstractions that describes selected aspects of a domain
- **Ubiquitous Language:** language structured around the domain model
- **Bounded Context:** Explicit definition of the context where a model applies (organisation, usage boundaries).

<https://www.pearson.com/us/higher-education/program/Evans-Domain-Driven-Design-Tackling-Complexity-in-the-Heart-of-Software/PGM168436.html>

Decomposition & Domain Driven Design (DDD)



How to Decompose: by sub-domain

Define services corresponding to Domain-Driven Design (DDD) sub-domains

- Refers to the application's problem space - the business - as the domain
- A Domain consists of multiple sub-domains

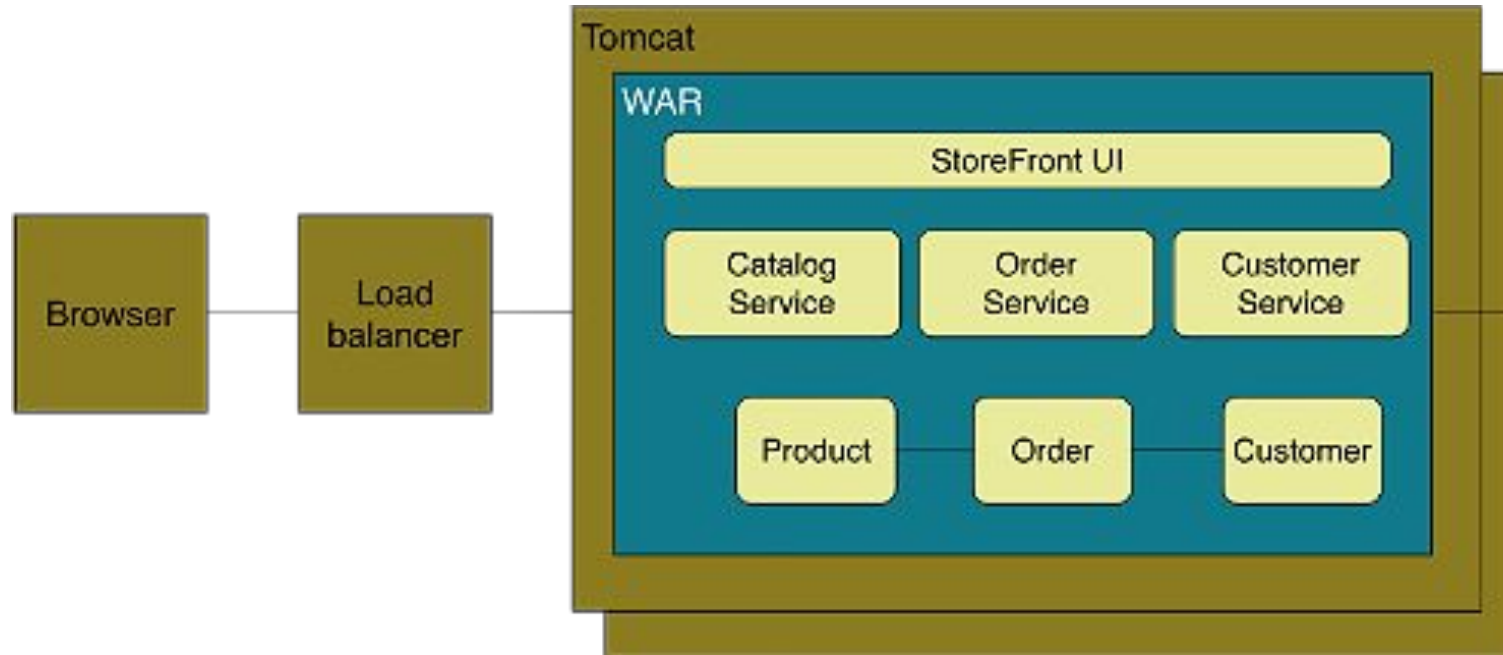
Each subdomain corresponds to a different part of the business. Subdomains can be classified as follows:

- **Core** - key differentiator for the business and the most valuable part of the application
- **Supporting** - related to what the business does but not a differentiator (implemented in-house or outsourced)
- **Generic** - not specific to the business and are ideally implemented using off-the-shelf software

How to identify the sub-domains?

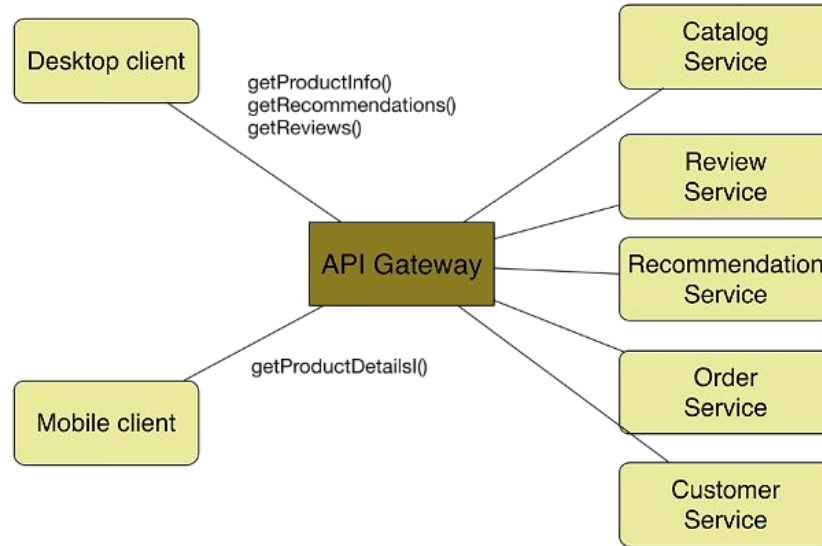
- Requires an understanding of the business.
- Like business capabilities, subdomains are identified by analyzing the business and its organizational structure and identifying the different areas of expertise using an iterative process.
- Good starting points for identifying subdomains are:
 - organization structure - different groups within an organization might correspond to subdomains
 - high-level domain model - subdomains often have a key domain object

Decomposing: An Example



Monolithic Application Design

Decomposing: An Example - decomposed



- How to assemble the application for deployment?

Application composition...

Application Composition

- Bringing these services together is known as composition.
- Composition's goal is to automatically bring together (deploy, provision) many components to deliver a functional system (e.g. replicated database system) or application (e.g. a 3-tier web application with API) that operates reliably.
- These services may be configured independently, or may depend on each other.
- Composition does not manage the lifecycle of an application
- **Overarching process' goal:**
Present the complete application to the end-user

Centralised Application Composition

Centralised is a common (enterprise) approach


- centralised global model (representation of system)
- logically single controller

used by technologies:

- TOSCA
- Docker Compose

Centralised:

- Great for management (global control, oversight)
- Poor for scalability (the internet is not centralised!)



Decentralised Application Composition

Decentralised is **less common** in enterprise:


- decentralized local model (independent)
- logically multiple controllers with local state only, managed independently

Examples:

- Habitat.sh: closest approach today
- DNS, BGP, P2P


Decentralised:

- Great for scalability
- Difficult to implement correctly and control




Composition Interactions

Orchestrated Compositions:
centralised process
one participant
global configuration



Choreographed Compositions:
decentralised process
multiple participants
local configurations



Composition Model Specification

Declarative

Example: docker-compose

```
1 version: '2.1'
2
3 services:
4   messaging:
5     image: redis
6     network:
7       - mynet
8
9 networks:
10  mynet:
11    driver: "bridge"
```

- Runtime modifications difficult
- Debugging is difficult
- Independent of a language
- Requires design of document structure
 - States the system as it should be delivered
- Typical of today's approaches
 - Docker compose, K8S, CF manifest

Imperative

Example: [kelda.io](https://keldai.io)

```
1 const app = new Container({
2   name: 'appName',
3   image: 'anyDockerImage',
4 });
5
6
7 const Redis = require('@kelda/redis');
8 const redis = new Redis(3, 'AUTH_PASSWORD');
9 allowTraffic(publicInternet, app, 80);
10
```

- Runtime modification easier (active, flexible)
- Could be debugged
- Uses existing languages, or specific DSL
 - less learning for dev
- Not very common today
 - More research-based
 - another example → Push2Cloud

Application Composition Standards

Industry standards are defined and driven out of the product and interest of the industry. Most of these standards are developed in closed groups and require a consensus mechanism (slow process).

- **TOSCA, CAMP, WS-BPEL**

De facto standards are those defined by a (single) entity and are adapted by the community so quickly/widely that the gain is huge popularity.

- **Docker Compose, Helm (Kubernetes)**

De jure standards are those driven from government related entities. Community driven, industry or de facto standards can become de jure standards if a government officially embraces one (e.g. TCP/IP protocol, ASCII encoding).

Topology and Orchestration Specification for Cloud Applications (**TOSCA**)



- YAML template document.
- Definition of building blocks for a cloud application.
- Modeling of components and the relationships between them.
- Standard (declarative description model and format) that enables portable applications and services.

TOSCA Entities

Nodes

- Represent **Components** of an application or service and their **Properties**. Example nodes include:
 - **Infrastructure**: Compute, Network, Storage, etc.
 - **Platform**: OS, VM, DB, Web Server, etc.
 - **Granular**: functional Libraries, Modules, etc.
- Include **Operations** which are the management functions for the node
 - e.g. *deploy()*, *start()*, *stop()*, *connect()*, etc.
- Export their dependencies on other nodes as **Requirement** and **Capabilities**

Relationships


- Represent the logical **Relationships** between nodes
 - e.g. "*hostedOn*", "*connectsTo*", etc.
- Describes the valid **Source** and **Target nodes** they are designed to couple
 - e.g. source "*web application*" node is designed to "*connectTo*" a target "*database*" node
- Have their own **Properties** and **Constraints**

Artifacts

- Describe **Installables** and **Executables** required to instantiate and manage a service. Currently, they include:
- **Implementation Artifacts**:
 - Executables or **Plans** that implement a Node's or Relationship's Operations (e.g. a Bash script)
- **Deployment Artifacts**:
 - Installables of the components (e.g. a TAR file)

Service Templates

- **Group** the nodes and relationships that make up a service's topology
 - Allowing modeling of sub-topologies
- Service Templates "*look like nodes*" enabling:
 - **Composition** of applications from one or more service templates
 - **Substitution** of abstract Node types with available service templates of the same type



TOSCA Service Template Instance

General definition of a “service template instance” (STI)

- A solution stack is a set of software subsystems or components needed to create a complete platform
- Example:
 - LAMP
 - Wordpress

Basic Definition of a STI applied to a Resource Orchestrator

- Collection of resources that will be managed by Resource Orchestrator (e.g. K8s or docker swarm).
- Might include instances (VMs, containers), networks, subnets, routers, ports, router interfaces, security groups, security group rules, autoscaling rules, etc.

TOSCA Template

YAML format (once XML)

Consists of a header (meta data)

- Version
- Description

and 3 content sections

- Node templates
 - Definition of resources
- Inputs - parameters
- Outputs - results of creation

```
tosca_definitions_version: tosca_simple_yaml_1_0
description: Monitoring Service Template

imports:
  - custom_types/tosca_compute.yaml
  - custom_types/tosca_floating_ip.yaml
  - custom_types/tosca_security_group.yaml
  - custom_types/tosca_router.yaml

topology_template:
  inputs:
    compute_image: ...
    compute_flavor: ...
    key_name: ...
    public_net: ...
    private_net_name: ...

  node_templates:
    my_server: ...
    private_net: ...
    router: ...
    floating_ip: ...
    port: ...
    server_security_group: ...

  outputs:
    it.hurtle.mon.dashboard: ...
```

Example:

- Monitoring as a Service
 - an IaaS application
 - one VM
 - attached to a network
 - has an external IP address
- Accessed under URL returned by
it.hurtle.mon.dashboard

Complete TOSCA Example

```
tosca_definitions_version: tosca_simple_yaml_1_0
description: Monitoring Service Template
```

```
imports:
- custom_types/tosca_compute.yaml
- custom_types/tosca_floating_ip.yaml
- custom_types/tosca_security_group.yaml
- custom_types/tosca_router.yaml
```

```
topology_template:
```

```
  inputs:
    compute_image:
      type: string
      description: Compute instance image
      default: cactiSnapshot2
    compute_flavor:
      type: string
      description: Compute instance flavor
      default: m1.small
    key_name:
      type: string
      description: Key name
      default: mauiskey
    public_net:
      type: string
      description: Name or ID of the public
network
  default:
77e659dd-f1b4-430c-ac6f-d92ec0137c85
  private_net_name:
    type: string
    description: Name of private network to be
created
  default: fr-tosca-mon-net
```

```
node_templates:
```

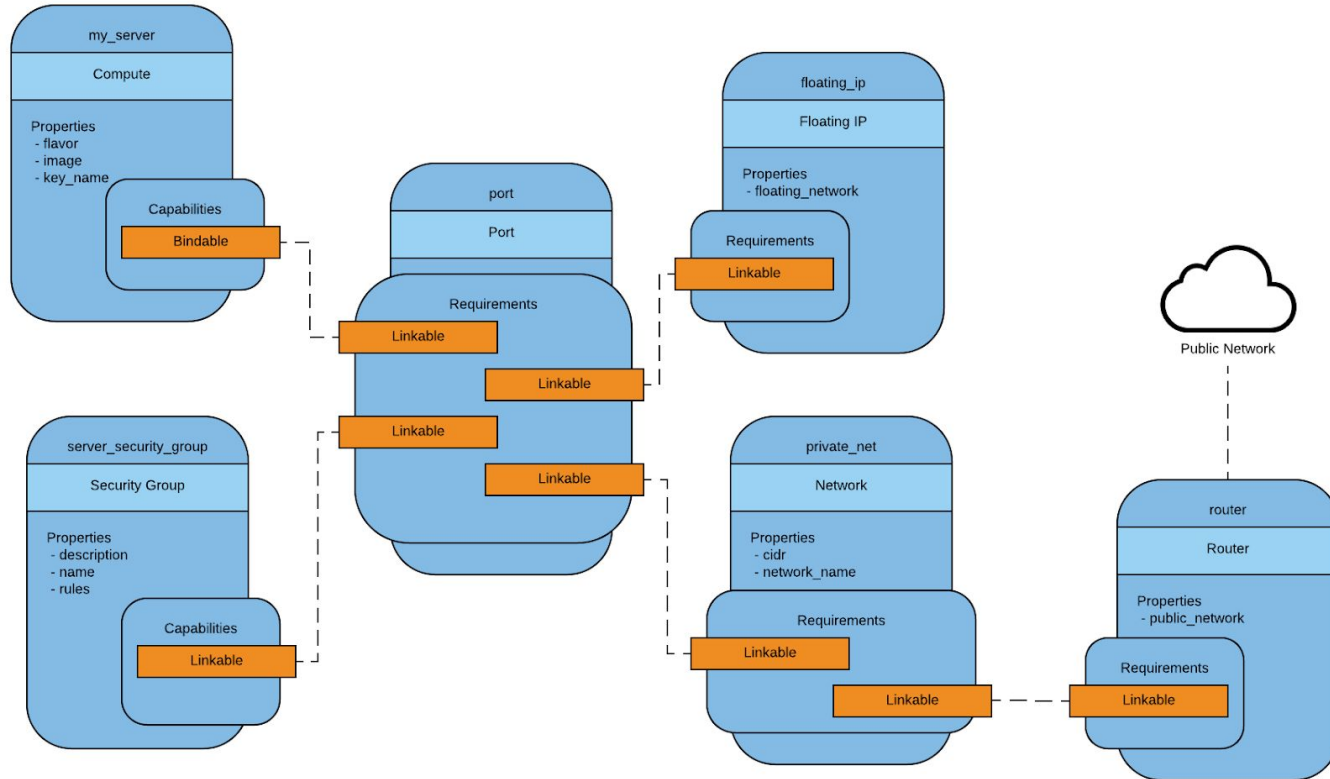
```
  my_server:
    type: hot.nodes.Compute
    properties:
      flavor: { get_input: compute_flavor }
      image: { get_input: compute_image }
      key_name: { get_input: key_name }
  private_net:
    type: tosca.nodes.network.Network
    properties:
      cidr: '192.168.1.0/24'
      network_name: { get_input: private_net_name }
  router:
    type: tosca.nodes.network.Router
    properties:
      public_network: { get_input: public_net }
    requirements:
      - link: private_net
  floating_ip:
    type: tosca.nodes.network.FloatingIP
    properties:
      floating_network: { get_input: public_net }
    requirements:
      - link: port
  port:
    type: tosca.nodes.network.Port
    properties:
      order: 0
    requirements:
      - link: private_net
      - link: server_security_group
      - binding: my_server
```

```
  server_security_group:
    type: tosca.nodes.network.SecurityGroup
    properties:
      description: Test group to demonstrate Neutron
security group functionality with Heat.
      name: test-security-group
    rules:
      - protocol: tcp
        port_range_min: 22
        port_range_max: 22
      - protocol: tcp
        port_range_min: 80
        port_range_max: 80
```

```
outputs:
```

```
  it.hurtle.mon.dashboard:
    description: Monitoring dashboard
    value:
      str_replace:
        template: http://host/cacti/
        params:
          host: { get_attribute: [my_server,
public_address ] }
```

TOSCA Template Logical Diagram



TOSCA Execution Environment

Service Template Instance (STI) is created by a Resource Orchestrator (RO)

- ROs manages instances (VMs, containers), networks, subnets, routers, ports, router interfaces, security groups, security group rules, autoscaling rules, etc.
- ROs supports multiple cloud environments (e.g. AWS, OpenStack, K8S, Docker Swarm)

Native RO (executes TOSCA)

- Cloudify
- OpenBaton
- Alien4Cloud

Adapted RO (Translator-based)

- OpenStack TOSCA to Heat Translator
- OpenStack Tacker

Application Composition: (2) Docker Compose [de facto]

- Docker allows a user to deploy many containers, but requires many commands to do so!
- Link a new web container to a previously created db container (see Container Linking)
 - `docker run -d -P --name web --link db:db training/webapp python app.py`
- Each new container to be linked needs to run a similar command (e.g. to have a new instance of the web frontend connected to the same database)
 - No automation! Manual, error prone process.
- **Docker Compose** aims to solve the process of **deploying applications comprised of multiple containers (services)**.
- Docker Compose is a local software which communicated with a locally accessible docker engine (daemon).
 - Can be a single docker engine or a clustered docker swarm
- **Define a multi-container application in a single file**, then deploy an application in a single command: `docker-compose up`.

Docker Compose Application Model

Compose is based on YAML template files describing an application.
It has a simple basic structure of:

- **version**
 - Defines featureset of Docker Compose
 - Higher versions features are not backward compatible
 - 3 major versions: 1, 2 and 3 (latest).
- **services**
 - Definition of the containers to be used in the application composition
 - Each service can be connected to multiple networks and volumes
- **volumes**
 - Definition of 1 or more storage resources to be used
- **networks**
 - Definition of 1 or more network resources to be used



```
version: '2.1'
services:
  myapp:
    image: myapp:1.0
    environment:
      ADMIN_PASSWORD: shush
```


Docker Compose – Service Definition

api:	⇒ service name
build:	⇒ builds a container if it doesn't exist
- context: ./	⇒ path to Dockerfile
image: me/my_image	⇒ names the container to use
depends_on: db	⇒ dependency on other service(s)
environment:	⇒ environment variables to parameterise the service
MY_VAR: value	
ports:	⇒ expose port 5000 of the container to port 5000 on the host
- "5000:5000"	
links:	⇒ links this service to another service (db)
- db	
volumes:	⇒ data volume is mounted in /info directory in the container
- data:/info	
networks	⇒ the api service is connected to the my-net docker network
- my-net	

Docker Compose – Extending Services and Service Health Checks

api:

extends:

```
file: ./base.yml
service: simple-api
...
```

⇒ takes a service “simple-api” from an existing definition

In the file base.yml

⇒ now can override or add new elements to the existing definition

api:

healthcheck:

test: "curl -f http://localhost:5000" ⇒ this is the command for the check - binary in the container

interval: 10s

⇒ how often to execute the test

timeout: 10s

⇒ how long to wait before trying again

retries: 5

⇒ how many times to try before it's a failed check

execute `docker-compose ps` to see the current status of health checks

Name	Command	State
ccp2dvop1lab_dbext_1	docker-entrypoint.sh mysqld	Up
ccp2dvop1lab_horizon_1	/bin/sh -c /usr/sbin/apach ...	Up (unhealthy)
ccp2dvop1lab_keystone_1	/etc/bootstrap.sh	Up (healthy)

failed
successful

Docker Compose – Volumes and Network

volumes:

⇒ Declares volumes for use in the application

`data:`

⇒ One volume named `data`

`driver: "local"`

⇒ uses the default `local` storage driver

- Others include: `overlay(2)`, `aufs`, `devicemapper`, `btrfs`, `zfs`

Execute `docker volume list` to see existing volumes

`docker volume --help` shows all management options

networks:

`my-net:`

⇒ Declares networks for use in the application

`driver: "bridge"`

⇒ One network named `my-net`

⇒ uses the default `bridge` storage driver

- Others include: `overlay` and `macvlan`

Execute `docker network list` to see existing networks

`docker network --help` shows all management options

Docker Compose – Creating the Application

- **docker-compose -f <MANIFEST_FILE> up**

Deploys an application from the manifest (default `docker-compose.yml` in the current directory).

Following steps happen with the manifest on the previous pages:

Building api... ⇒ Build web image from Dockerfile


Starting api_1... ⇒ Start api container first

api_1 | * **Running on http://0.0.0.0:5000/**

- API starts first, but as it does not exist in the local registry or docker hub, it is built and then executed.
 - **Note:** a container can only be linked to **after** its creation.
- **docker-compose up --scale api=2 --no-recreate**
Specify how many instances of a container type should be active.
 - e.g. `docker-compose up --scale api=2 --no-recreate` ensures that two containers (no more, no less) of the type `api` are active, it will create one more if there is only one, or stop any additional one. The flag `--no-recreate` ensures that the service is not restarted.
- To delete the application: **docker-compose -f <MANIFEST_FILE> down**
- More commands: **docker-compose --help**

Application Composition: (1) Helm [de facto]

- A “package manager” for Kubernetes (K8s)
- Bundles related k8s **manifests** together in a chart
 - Manifest is a k8s application description
 - Set of manifests used to deliver an application is known as a **chart**
- Enables reuse and composition through dependencies
- Provides templating for k8s manifests
 - allows for template variable overriding (from sub-charts, value files, command line options)
- Installing a chart creates a release
 - Maintains versions, allows for rolling updates & roll-back



Helm – Key Concepts

- **Chart**
 - a package; bundle of Kubernetes resources (pods, services, ...)
- **Release**
 - a chart instance is loaded into Kubernetes
 - same chart can be installed several times into the same cluster (separate namespace); each will have its own Release
- **Repository**
 - a repository of published Charts (e.g. like docker hub but for k8s)
 - Public repository here: <https://artifacthub.io/> (Private can be hosted too)
- **Template**
 - a K8s configuration file mixed with Go/Sprig* template



Helm-Chart package structure

Package structure of a Helm-Chart:

see: <https://helm.sh/docs/topics/charts/>

mypackage/

Chart.yaml	# A YAML file containing information about the chart
LICENSE	# OPTIONAL: A plain text file containing the license for the chart
README.md	# OPTIONAL: A human-readable README file
values.yaml	# The default configuration values for this chart
values.schema.json	# OPTIONAL: A JSON Schema for imposing a structure on the values.yaml file
charts/	# A directory containing any charts upon which this chart depends.
crds/	# Custom Resource Definitions
templates/	# A directory of templates that, when combined with values, # will generate valid Kubernetes manifest files.
templates/NOTES.txt	# OPTIONAL: A plain text file containing short usage notes

Chart.yaml structure

apiVersion: The chart API version (required)
name: The name of the chart (required)
version: A SemVer 2 version (required)
kubeVersion: A SemVer range of compatible Kubernetes versions (optional)
description: A single-sentence description of this project (optional)
type: The type of the chart (optional)
keywords:
- A list of keywords about this project (optional)
home: The URL of this projects home page (optional)
sources:
- A list of URLs to source code for this project (optional)
dependencies: *# A list of the chart requirements (optional)*
- **name:** The name of the chart (nginx)
 version: The version of the chart ("1.2.3")
 repository: (optional) The repository URL ("https://example.com/charts") or alias ("@repo-name")
 tags: *# (optional)*
 - Tags can be used to group charts for enabling/disabling together
 import-values: *# (optional)*
 - ImportValues holds the mapping of source values to parent key to be imported. Each item can be a string or pair of child/parent sublist items.
 alias: (optional) Alias to be used for the chart. Useful when you have to add the same chart multiple times
maintainers: *# (optional)*
- **name:** The maintainers name (required for each maintainer)
 email: The maintainers email (optional for each maintainer)
 url: A URL for the maintainer (optional for each maintainer)
icon: A URL to an SVG or PNG image to be used as an icon (optional).
appVersion: The version of the app that this contains (optional). Needn't be SemVer. Quotes recommended.
deprecated: Whether this chart is deprecated (optional, boolean)
annotations:
 example: A list of annotations keyed by name (optional).

see: <https://helm.sh/docs/topics/charts/>

Example Chart.yaml

```
apiVersion: v2
name: helm-os
description: A Helm chart for the osi application
type: application
version: 0.1.0
appVersion: "1.0"
dependencies:
  - name: apache
    version: 1.2.3
    repository: https://example.com/charts
  - name: postgresql
    version: 0.1.0
  - name: microservice
    version: 0.1.0
    alias: microservice-order
  - name: microservice
    version: 0.1.0
    alias: microservice-shipping
  - name: microservice
    version: 0.1.0
    alias: microservice-invoicing
```

Version of the chart (semantic version number)

Version of the app (Text)

Subcharts this charts depends on

Download from repo to **chart** folder

If no repository is given, it has to be put
to the **chart** folder manually

name of the chart in the chart folder

alias can be used to install multiple similar
components using the same chart

Example: Deployment Template

```
apiVersion: apps/v1
kind: Deployment
metadata:
  labels:
    app: {{ quote .Values.name }}
    version: {{ quote .Values.deploymentVersion }}
  name: {{ quote .Values.name }}
spec:
  replicas: {{ .Values.replicaCount }}
  selector:
    matchLabels:
      app: {{ quote .Values.name }}
  strategy: {}
  template:
    metadata:
      labels:
        app: {{ quote .Values.name }}
    spec:
      containers:
        - name: {{ quote .Values.name }}
          image: "{{ .Values.image.repository }}:{{ .Values.image.tag | default .Chart.AppVersion }}"
          imagePullPolicy: {{ .Values.image.pullPolicy }}
          ports:
            - containerPort: {{ .Values.port }}
          resources: {}
```

see: https://helm.sh/docs/chart_template_guide/

{{ Placeholder }}

values will be taken from the --set arguments, the provided values file or the default values.yaml of the chart.

Template Functions allow more complex operations (quote, print, b64enc, b64dec, include, ...).
see https://helm.sh/docs/chart_template_guide/function_list/

If no value is found an alternative default value can be specified.

Also chart (.Chart) or release (.Release) specific values can be used

Example: values.yaml

```
name: apache
deploymentVersion: 1.0
replicaCount: 1
port: 80
host: web.160.85.253.<X>.nip.io
```

```
image:
  repository: registry.localhost:5000/ccp2-apache
  tag: latest
  pullPolicy: Always
```

```
service:
  type: ClusterIP
  port: 80
```

see: https://helm.sh/docs/chart_template_guide/

Helm commands

Looking up charts

see: <https://helm.sh/docs/helm/>

- `helm search hub postgresql` in artifacthub.io
- `helm search repo postgresql` in the repo list

Managing repos

- `helm repo list`
- `helm repo add bitnami https://charts.bitnami.com/bitnami`
- `helm repo delete bitnami`

Manage applications

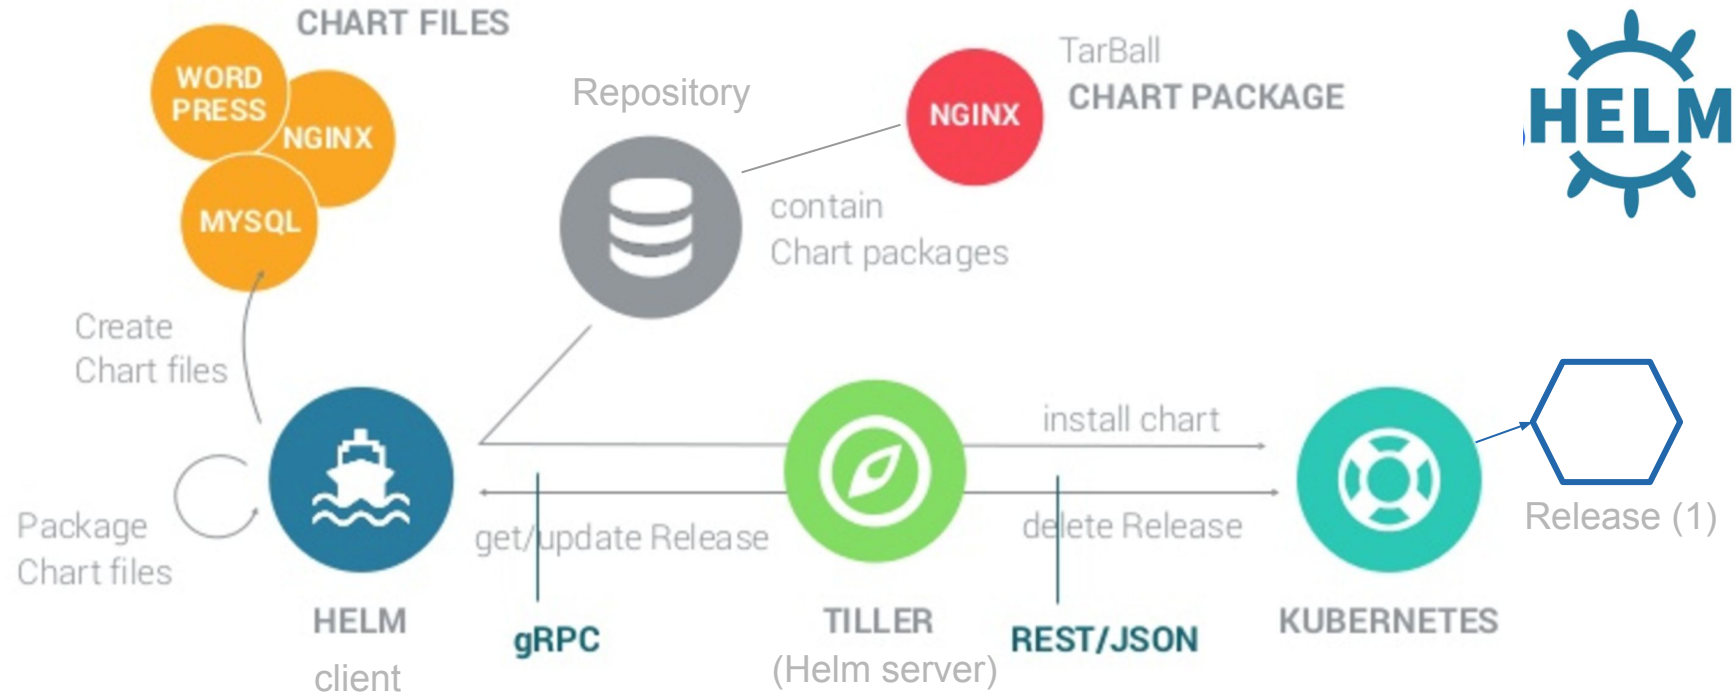
- `helm install -f myvalues.yaml myDB bitnami/postgresql`
- `helm install --set auth.username=dbuser myDB bitnami/postgresql`
- `helm list`
- `helm uninstall myDB`

Manage packages

- `helm pull bitnami/postgresql` Downloads the chart package (tgz)
- `helm package ./ois-package/` Create a chart package from a path

Appendix

Helm – Architecture v1 (deprecated)



Helm – Getting Started

Prerequisites

- A Kubernetes Cluster
 - Many ways to install from small to large
 - See minikube (small local install) or kops (large AWS EC2 install)
- The helm binary



Install & Upgrade

- `helm init` \Rightarrow sets environment and installs Tiller on the k8s cluster
- `helm init --upgrade`
- `helm --help` \Rightarrow all you can do with Helm

TOSCA Custom Type

 **tosca_router.yaml** 378 Bytes 

```
1  tosca_definitions_version: tosca_simple_yaml_1_0
2  node_types:
3    tosca.nodes.network.Router:
4      derived_from: tosca.nodes.Root
5      properties:
6        public_network:
7          type: string
8      requirements:
9        - link:
10          capability: tosca.capabilities.network.Linkable
11          relationship: tosca.relationships.network.LinkTo
12          node: tosca.nodes.network.Network
```

Note:

- Linkable
- LinksTo

TOSCA Concepts

TOSCA Concepts in OOP Terms:

- **Service Template** = class
- **Service Template Instance** = Instance of a template (Object)
- **Inputs** = Constructor Arguments
- **Node Templates** = Variables/Content of the class
- **Outputs** = Getter Methods/Attributes

