

# DVOP2 - DevOps, Continuous Deployment

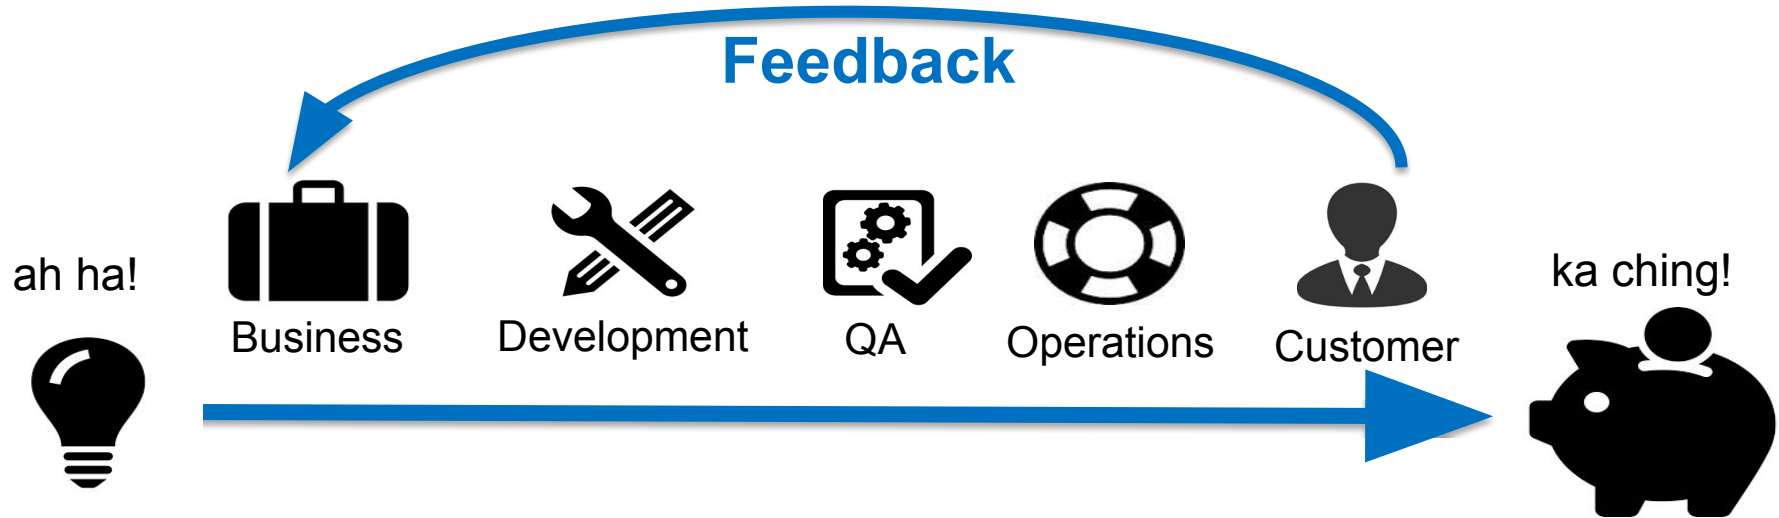
Prof. Dr. Thomas M. Bohnert  
Christof Marti

# Content

---

- Continuous Delivery / Deployment Intro
- Software Automation Pipeline
- Tooling for Continuous Delivery / Deployment
- CI/CD Pipelines
- Deployment Strategies / Pattern
- Zero Downtime State-Transition patterns

# The IT Value Stream




# Customer Expectations

- Consumer Expectation
  - constant drumbeat of feature and function delivery
  - quality
  - zero installation headaches and upgrades

→ How to continuously deliver and deploy in quality?

Continuous means, to deploy frequently and in smaller increments

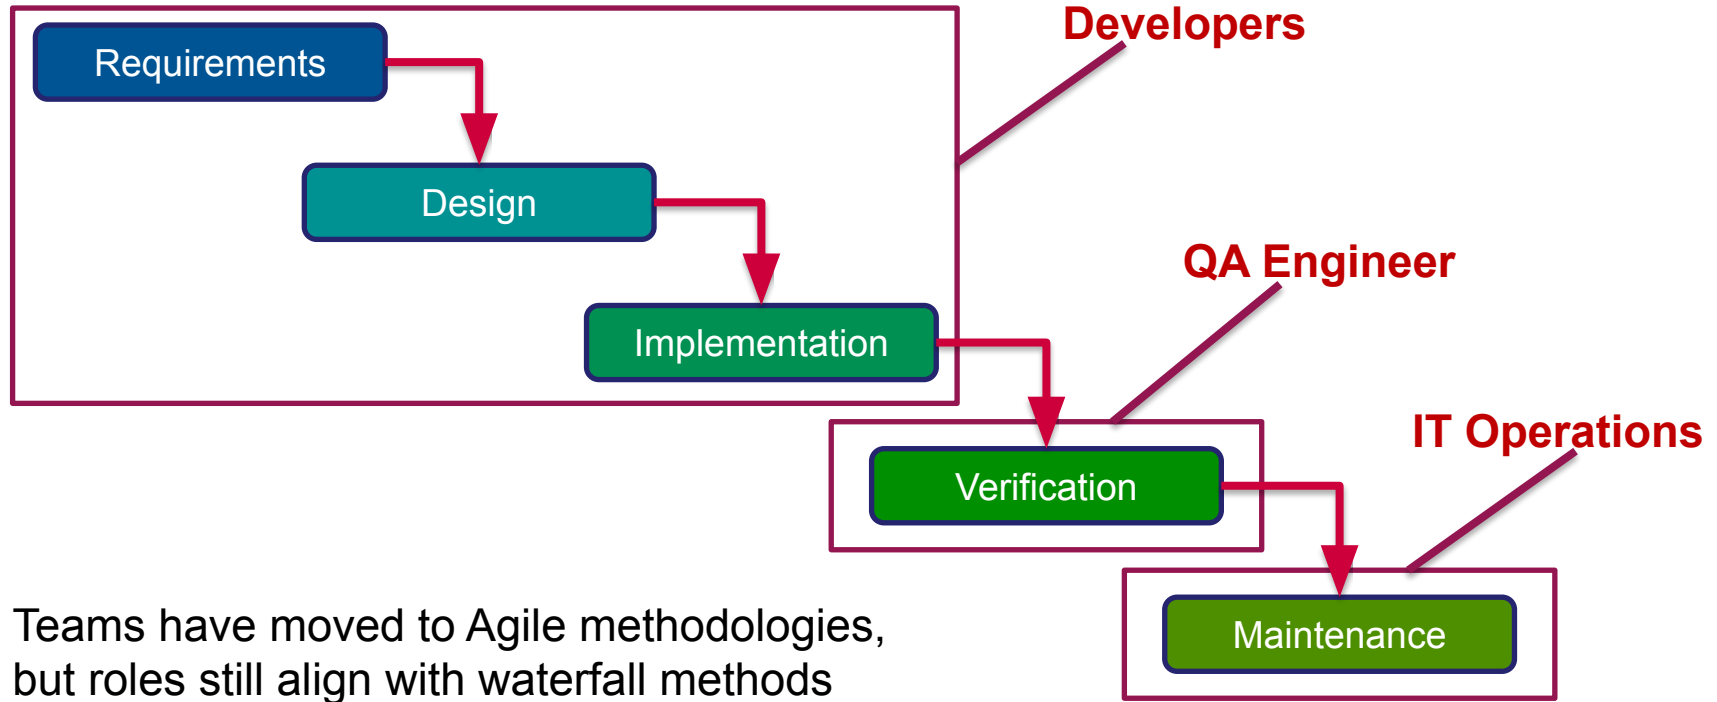
- Faster Time To Market
- Minimize Risks
- Improve Product Quality



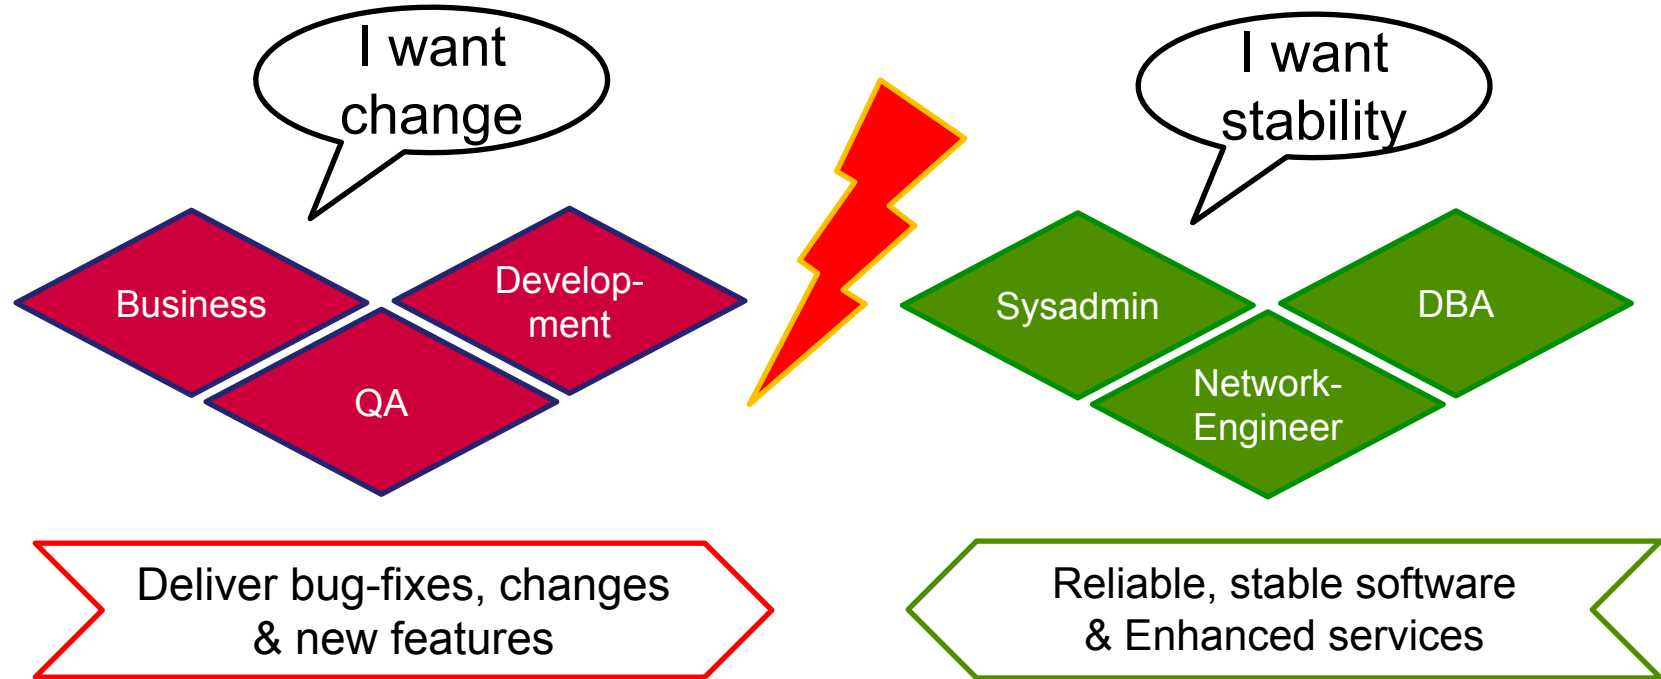
# Main Goals of Continuous Delivery / Deployment

- **Faster Time To Market**
  - Accelerate the process from idea to deployment
  - Immediate feedback
  - Shorter innovation cycle
- **Minimize risks**
  - only small changes
  - prove that software is building
  - find broken build fast and early
  - awareness of current software status
  - eliminate dependencies on key personnel
- **Improve Product Quality**
  - automated testing & automated code auditing
  - documented history of builds to verify issues


# Many Organizational Silos are still Waterfall



# Conflicting Interests & Incentives




# Wall of confusion





# Continuous Life-Cycle between Development and Operations



To evolve your application during its lifetime the Software Automation Pipeline is **executed continuously** for each new feature, bug fix or configuration change.

# Approach - Tearing down walls of confusion

ah ha!



Business



Agile Development  
fixes this



Development  
& QA



DevOps  
fixes this




Operations

ka ching!




# Main driver: Automation

- Machines are really good at doing the same task over and over
- Consistent and Known State
- Fast and Efficient
- 5 mins/day → 2.6 days/year
- What can be automated?
  - Builds
  - Deployments
  - Tests
  - Monitoring
  - Self-healing
  - System rollouts
  - System configuration
  - ...



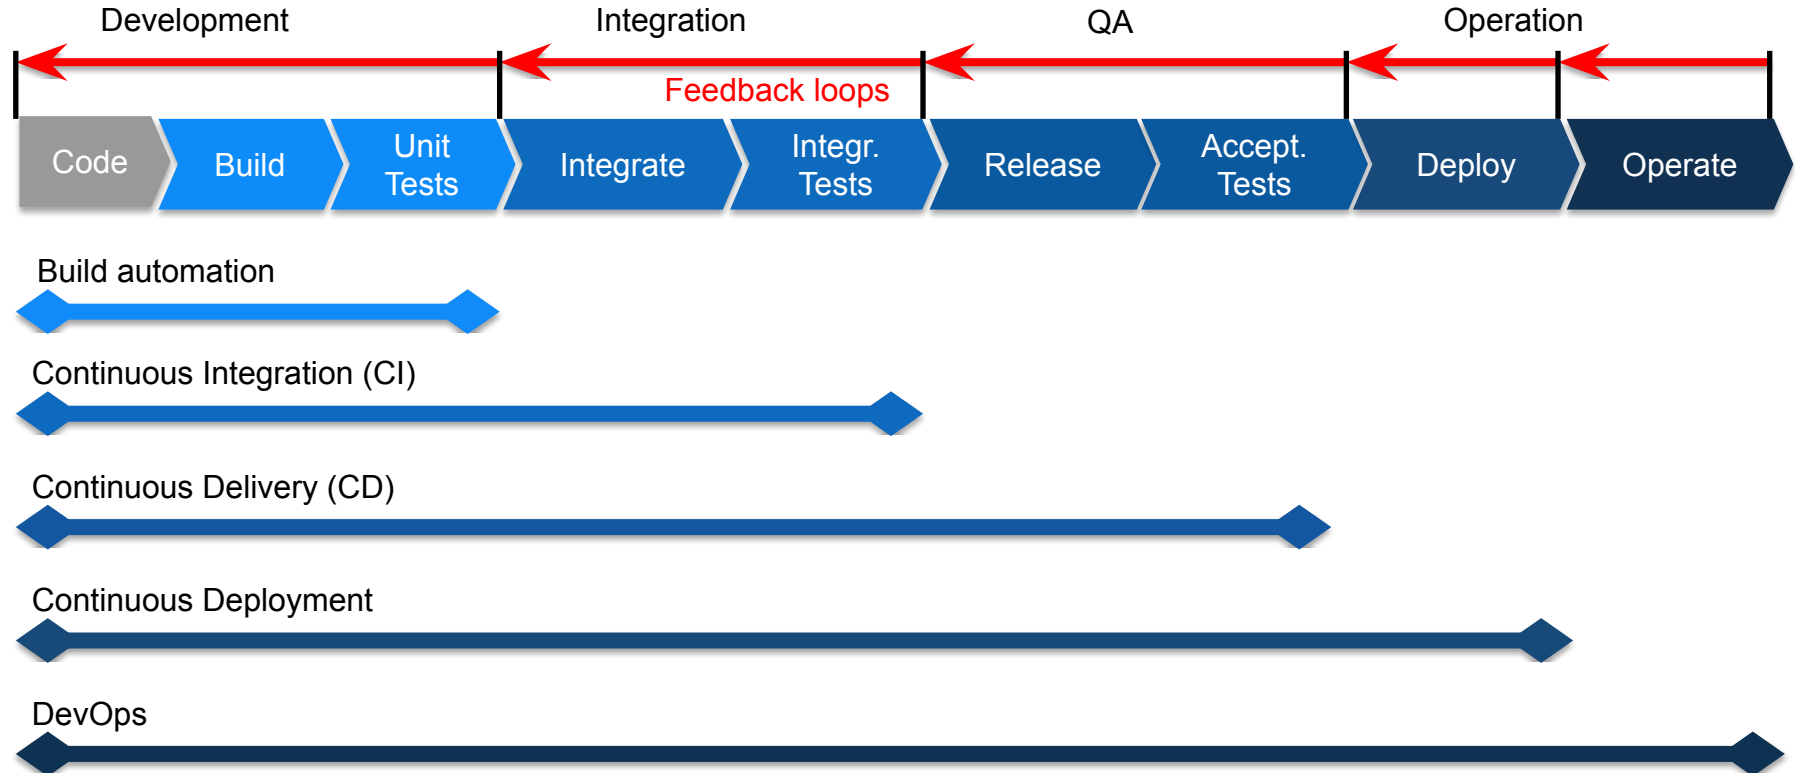
# Software Automation Pipeline



- **Automation** can be implemented in **every stage** of the software process from Development to Integration, QA to Operation.
- Every step can **only continue (automatically)** if it **passes** several tests successfully.
- **Otherwise** the responsible people are **notified** about the failure and the process is halted until fixed → **Feedback loop**

Increase confidence in production readiness

# Phases of the Software Automation Pipeline



# Phases of the Software Automation Pipeline

- **Build automation**
  - **Building** individual components and run **unit tests**
  - Typically run by the developer on his local machine
- **Continuous Integration**
  - Automatically **build, test and integrate** components and run Integration Tests (Code auditing, Security tests, Database tests, UI tests, ...).
  - Typically run on Continuous Integration Server
- **Continuous Delivery**
  - Also **create releases**, deploy to staging environment and run automatic **acceptance** tests (Stress test, Load Tests, Compliance tests,...).
  - Ready for production, but deployment still requires a manual step.
- **Continuous Deployment**
  - Automatically **deploy** to production after successful passing acceptance tests.
- **DevOps**
  - Automatically run the operation of the production system (configuration management, infrastructure provisioning, backup, monitoring, automatic health management, scaling, ...)

# Build Automation

## Typical tasks

Build automation covers the build phase of the SW lifecycle

- **Resolve dependencies**
  - Obtain required versions of 3<sup>rd</sup> party libraries
  - Identify build order
- **Compile** source code
- Run unit **test**
- **Package** Software
- Deploy to Artifact repositories
- Create **documentation**
- **Clean up** temporary files

## Tools

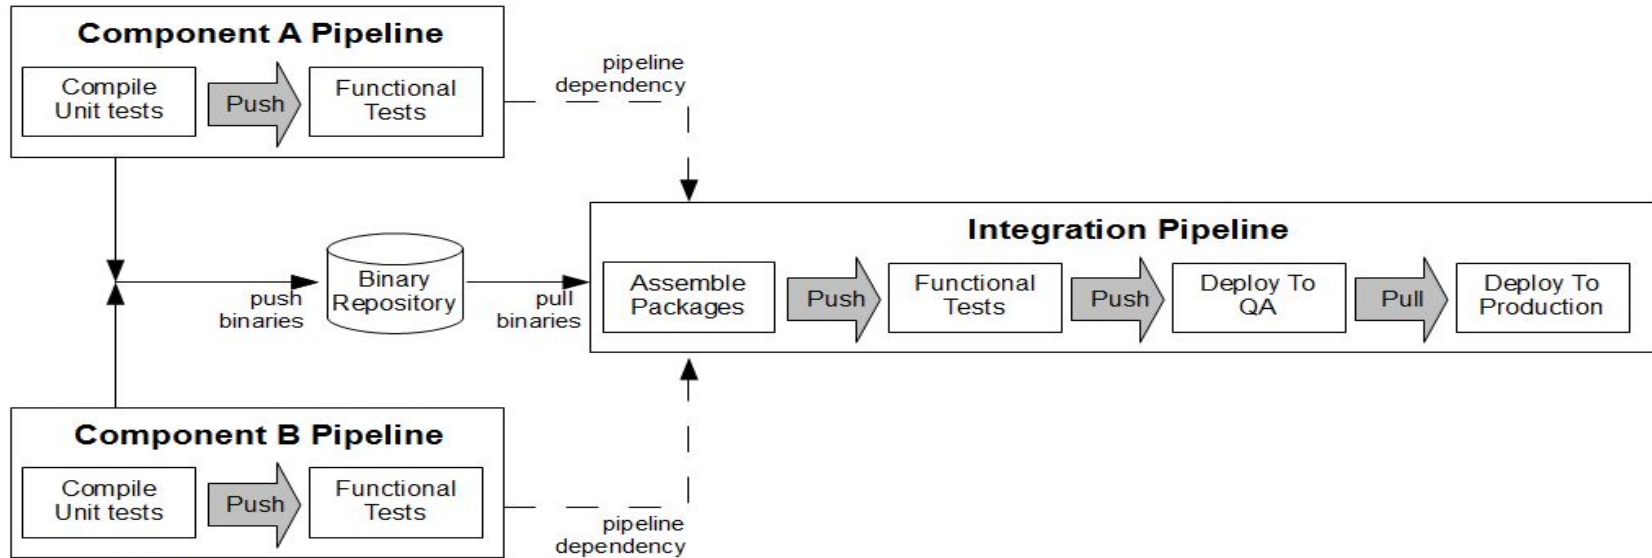
Depending on your environment different tools are very common

- C, C++, ... → make, gmake, SCons, distcc,...
- Python → Waf, Snakemake, ...
- Ruby → Rake,...
- JavaScript/Node.js → Grunt, Gulp, Bower,...
- Microsoft Environment → MS Build, Visual Build, Psake,...
- Java, JVM languages → Apache Ant, Apache Maven, Gradle, sdt, ...

Most tools work in / for multiple environments

# Continuous Integration

Continuously integrate pushed individually developed components (from different developers/teams), assemble them to a complete application and run functional integration tests against it.



"Integration Pipeline"

© Adapted from Humble and Farley 2011



# Continuous Delivery vs Continuous Deployment

- In **Continuous Delivery**, **a human manually deploys** the release into the production environment.



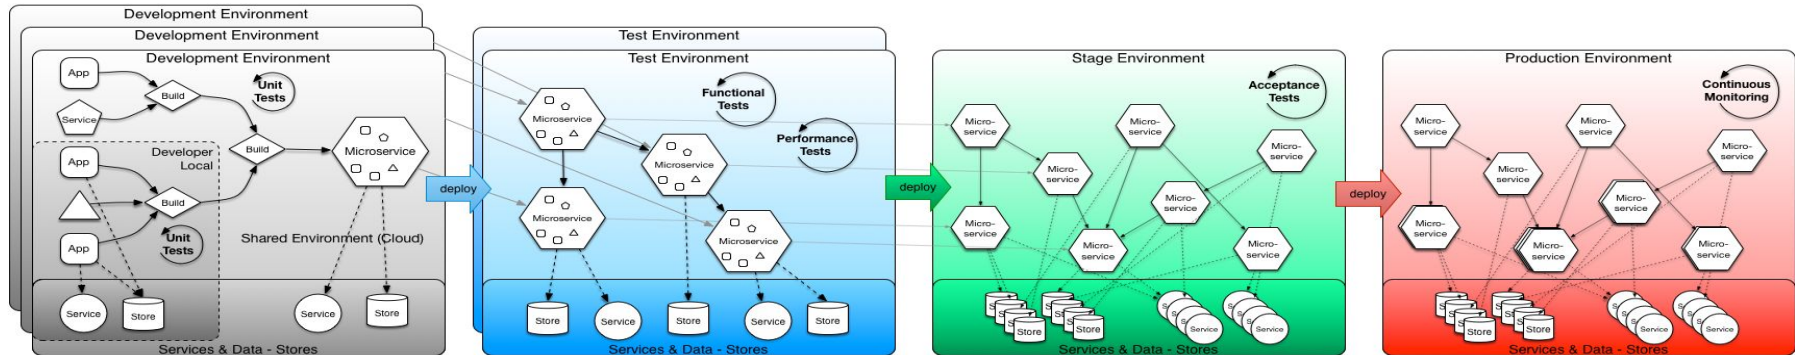
- **Continuous Deployment** evolves the continuous delivery process one step further: **every change that happens to pass automated tests is deployed to production automatically.**



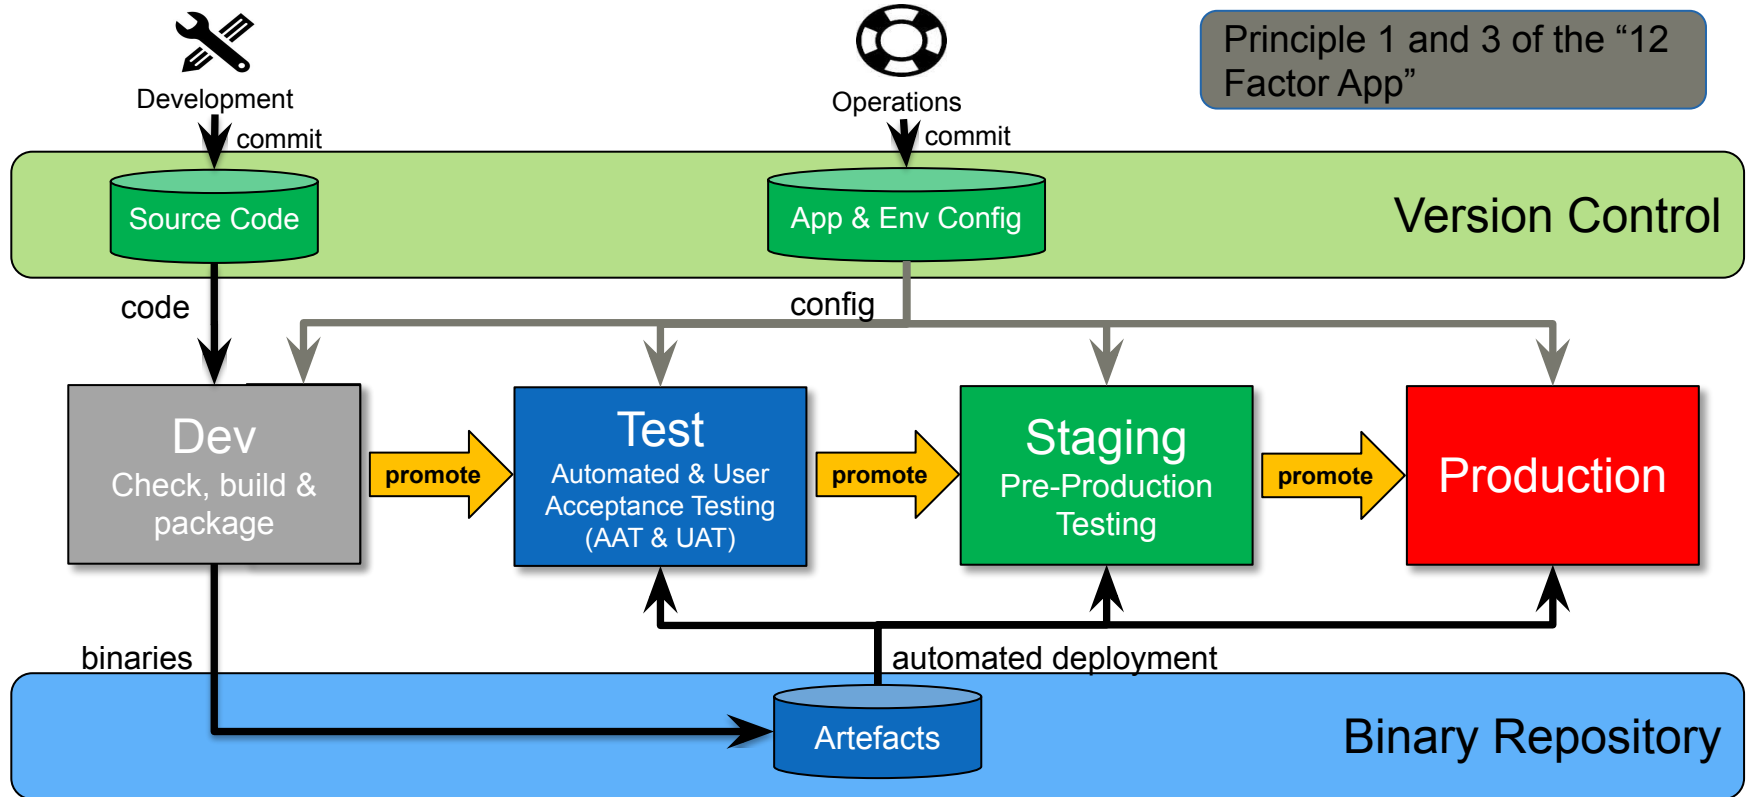
# Multi-Stage Delivery and Environments

Provide different environments for each phase / stage of the process:

- **Development** (1+): Run the application per developer / team to develop specific features
- **Test** (1+): Run Integration, Functional and Performance Test in a dedicated test environment, close to the production environment.
- **Staging**: Exact copy of the production environment to run Acceptance and Operational Tests. Specifically also test the deployment process & scaling
- **Production**: Environment accessible to the end user containing real production data



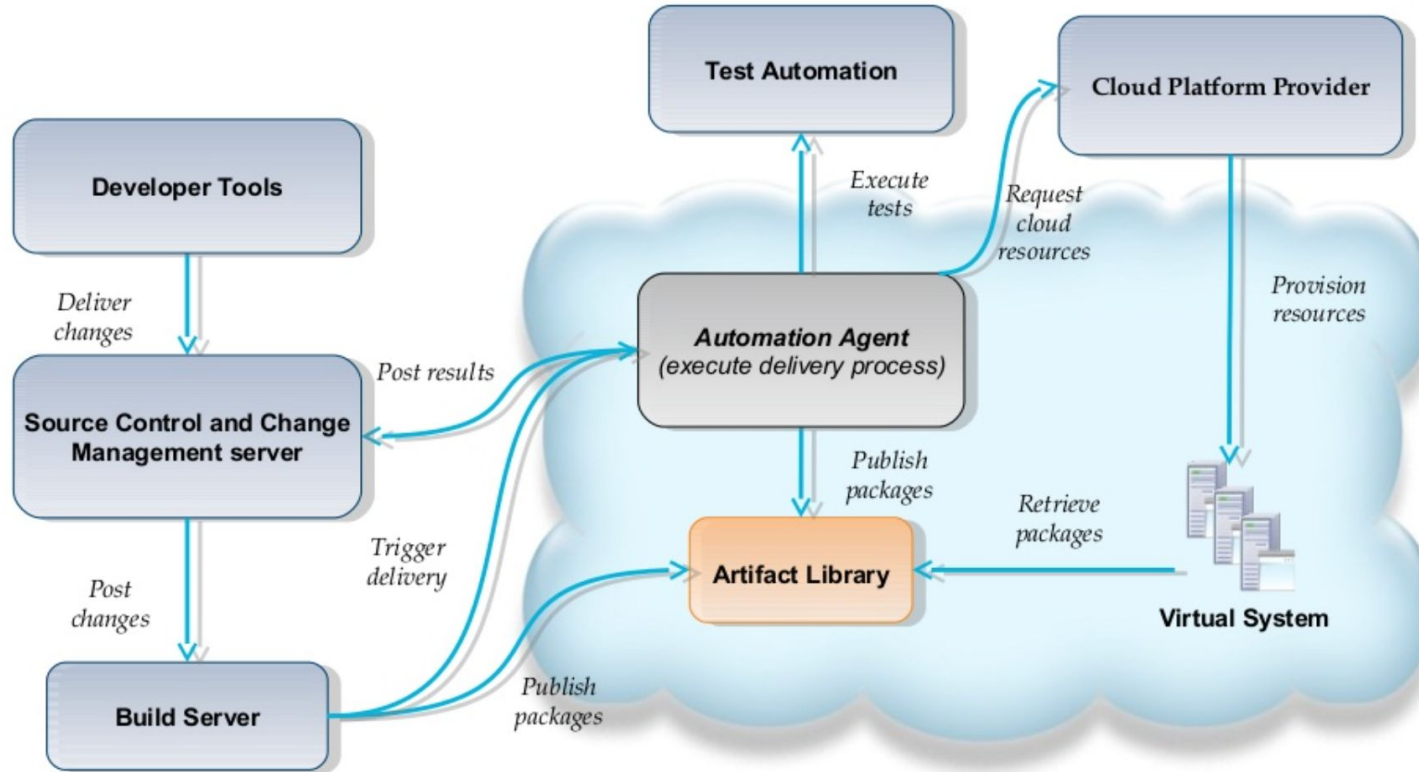
# Flow through the Multi-Stage Pipeline



# Best Practices Multi-Stage Pipelines

- Code (Dev) and Config (Ops) changes go **always to version control**
- Binary artifacts are only built once. **Same artifacts** are used in all environments  
→ *no environment specific builds*
- **Different Configs** to support environment specific requirements provided by environment variables
- **Same tooling** is used to deploy to all environments
- **Typical Actions per Stage**
  - **Development:** Syntax check, code metrics, compile, unit tests, package
  - **Test:** Can be split into multiple stages. Is using stubbed or mocked data
    - AAT (Automated Acceptance) → Component/Integration-T, Feature/Story-level-T
    - UAT (User Acceptance) → UI-T, Usability-T, Showcase-T, Client-T
  - **Staging / Pre-Production:** Network-T, Capacity-T, Performance-T Smoke-T
  - **Production:** Post-Deployment-T, Smoke-T, Cont. Monitoring, Rollback & Re-Deploy

# Continuous Delivery / Deployment Tooling and Systems



# Tooling - Components

- **Version Control Server (VCS)** for Code and Configuration  
e.g. **Git**, Mercurial, SVN,...
- **Artifact Repository** management to store binary artifacts  
e.g. generic repositories like Jfrog Artifactory, Sonatype Nexus, GitHub Packages, GitLab, or type specific like (Docker) Image Registry, NPM
- **Build Server** to trigger and control the build, deploy and test tasks  
e.g. Jenkins-X, Travis-CI, GitHub-Actions, GitLab CI, **Tekton**, **Argo Workflows**,...
- **Automation Agent** to orchestrate the deployment process  
e.g. Shell-Scripts, Puppet, Chef, Ansible, Heat, CloudFormation, **ArgoCD**, **Flux**, ...
- **Monitoring infrastructure** to get feedback from the runtime environment  
e.g. ELK (ElasticSearch-Logstash-Kibana), Graylog, Splunk
- **Secure-Store** to manage and provide security sensitive data like credentials and certificates e.g. HashiCorp Vault, Square Keywhiz, ...

# Tooling - Build-Server


- Traditional CI-Servers (Jenkins, Travis, ...)
  - Easy to configure individual tasks (e.g. check-out and test)
  - Very challenging to configure workflows (sequence of tasks) or pipelines
- CI-Servers with Workflow support
  - Jenkins Pipeline aka Blue Ocean (<https://jenkins.io/projects/blueocean/>)
  - Travis CI, GitHub Actions, GitLab CI, ...
- Modern pipeline based build servers
  - Jenkins-X (<https://jenkins-x.io/>)
  - Tekton CD (<https://tekton.dev/>)
  - Argo Workflow (<https://argoproj.github.io/argo-workflows/> )
  - Concourse CI (<http://concourse.ci>)

# Blue Ocean Pipeline UI and Editor

CCP2-DVOP2-Lab-Solution / development

Cancel

Save



← Deploy development branch ...

Steps

**Print Message**

Deploying development branch....

**Print Message**

"GIT\_COMMIT is \${env.GIT\_COMMIT}"

**Print Message**

"GIT\_BRANCH is \${env.GIT\_BRANCH}"

**Run arbitrary Pipeline script**

```
def server = Artifactory.server 'ccp2-arti...'
...
```

**Shell Script**

```
tar xvzf artifacts.tar.gz
```

**Push to Cloud Foundry**

```
https://api.lyra-836.appcloud.swisscom....
```

✓ CCP2-DVOP2-Lab-Solution 11

Pipeline

Änderungen

Tests

Artefakte



Ausloggen



Branch: development

2m 13s

Änderungen von mach

Commit: 7251cd1

2 hours ago

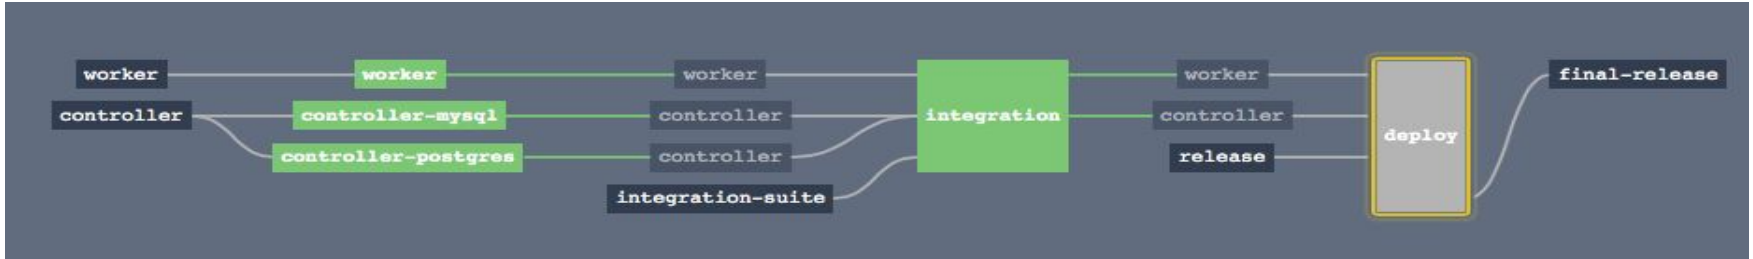
Replayed #10





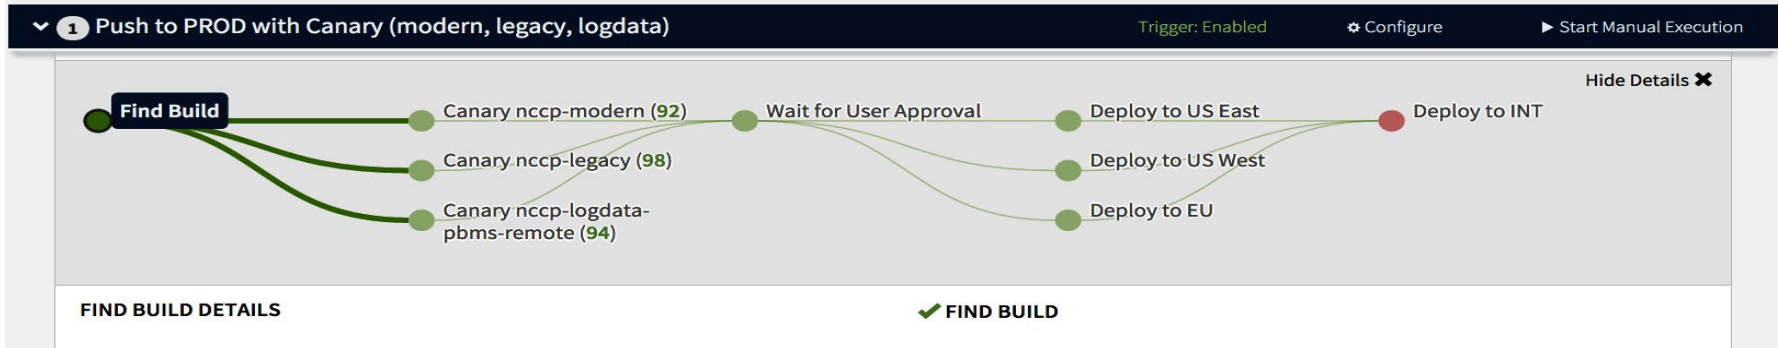
# Tooling - Build-Server Workflow Examples

- Concourse CI (<http://concourse.ci>)



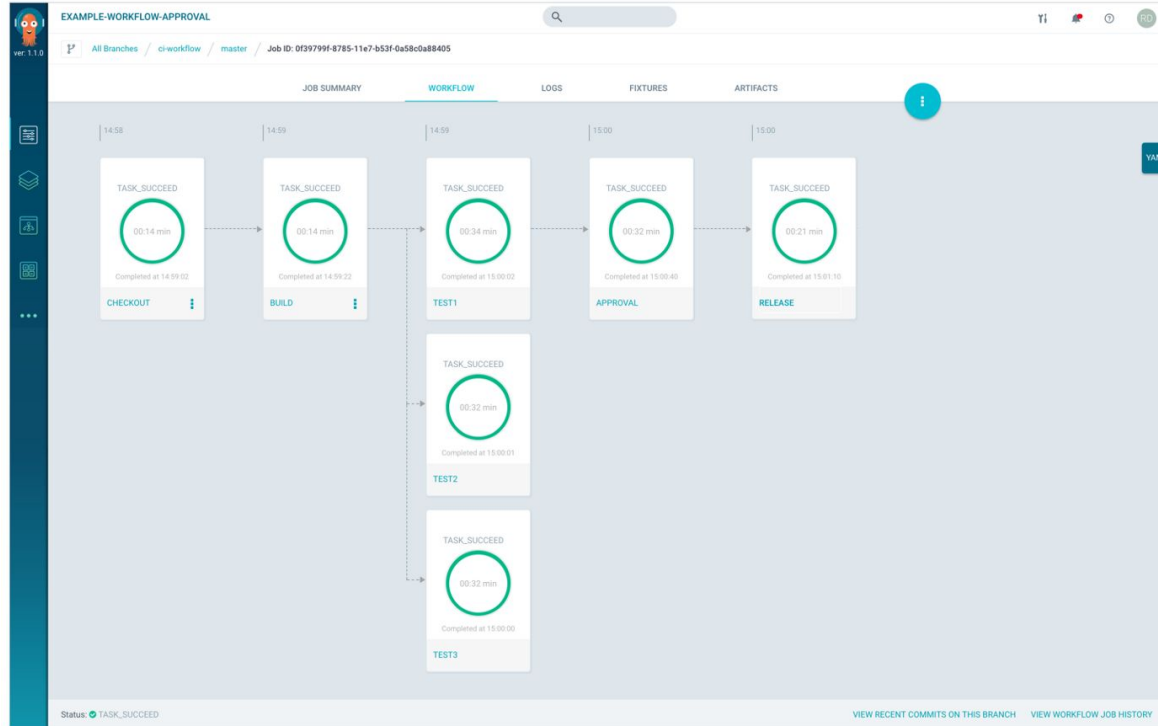
<https://ci.concourse.ci/teams/main/pipelines/main?groups=develop>

- Spinnaker CI (<http://www.spinnaker.io>)



# Tooling - Build-Server Workflow Examples

- Argo Workflow (<https://argoproj.github.io/argo-workflows/>)



## Traditional CI-Server

- Proven and well known in community
- Scalable using master nodes and agents
- Many plugins and integrations in dev tools
- New: Support for docker and kubernetes
- Newest releases (Jenkins 2+) support pipelines
- Modernized Blue Ocean UI for pipelines



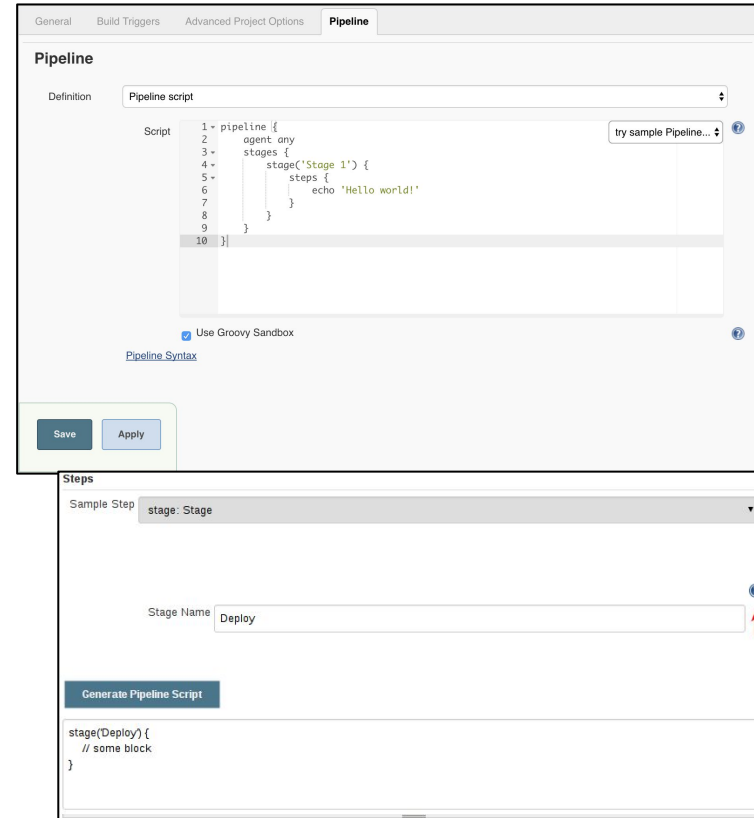
# Jenkins

# Jenkins Pipeline Concepts

- **Pipeline** → user-defined model of a CD pipeline
  - Defines the entire build process, which typically includes stages for building an application, testing it and then delivering it.
- **Node / Agent** → Worker executing a Pipeline
  - Node is used in the Scripted syntax; Agent in the declarative
  - Stages can be executed on different agents (sequential or parallel)
- **Stage** → Conceptual distinct subset of tasks
  - Used to visualize progress
  - In declarative syntax the can be skipped using a when clause
  - Can run in parallel and on different agents
- **Step** → A single task to execute
  - Typically every command in is its own step (e.g. `sh './gradle assemble'`)
  - Commands can be extended using plugins (e.g. `pushToCloudFoundry()` )

# Creating a pipeline

- Multiple ways to create Pipelines
  - Using the classic UI
    - edit text in browser
    - built in snippet generator
  - Using BlueOcean Editor
    - Automatically reads Jenkinsfile and allows to edit
    - Needs write permission to update Jenkinsfile in repository
    - May reorder your code and loose comments
  - Using Jenkinsfile in SCM (source control mgmt)
    - Jenkinsfile is automatically checked out and read from the repository root
    - Versioned together with project



# Jenkinsfile – Basic structure (declarative)

```
pipeline {  
  agent { label 'Linux' }  
  options {  
    timeout(time: 1, unit: 'HOURS')  
  }  
  environment {  
    CF_CRED = credentials('init-cloudlab')  
  }  
  triggers {  
    cron('H */4 * * 1-5')  
  }  
  stages {  
    stage ('Build') {  
      steps {  
        echo "Building..."  
        sh './gradlew assemble'  
      }  
      post {  
        failure {  
          echo "Build failed"  
        }  
      }  
    }  
    ...  
  }  
}
```

```
// begin of the declarative pipeline  
// declare agents (default: any). (also possible in stage)  
// options declare global behavior (timeout, retry,...)  
// e.g. skipDefaultCheckout → do not checkout git repo on agent  
  
// set global environment (also possible in stage)  
e.g. set CF_CRED_USR and CF_CRED_PSW vars from Jenkins credential  
  
// when to trigger build (cron, pollSCM, ...) e.g. for nightly builds;  
often not needed → triggered using webhook from SCM  
  
// start stages section (list of stages)  
// first stage called 'Build'  
// steps to execute in this stage  
// each command is a single step  
  
// commands to execute after the stage has completed  
// depending on result of stage (success, failure, aborted, always,...)
```

# Jenkinsfile – Basic structure (declarative) cont.

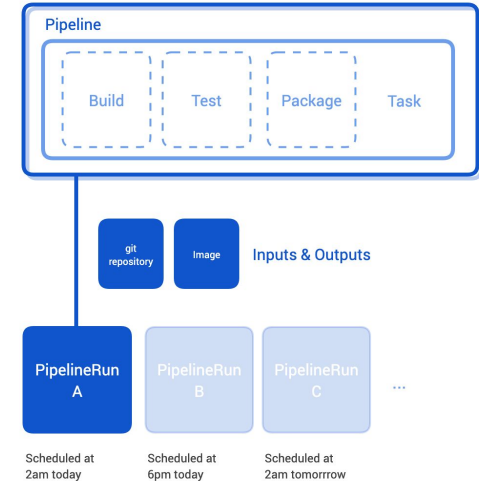
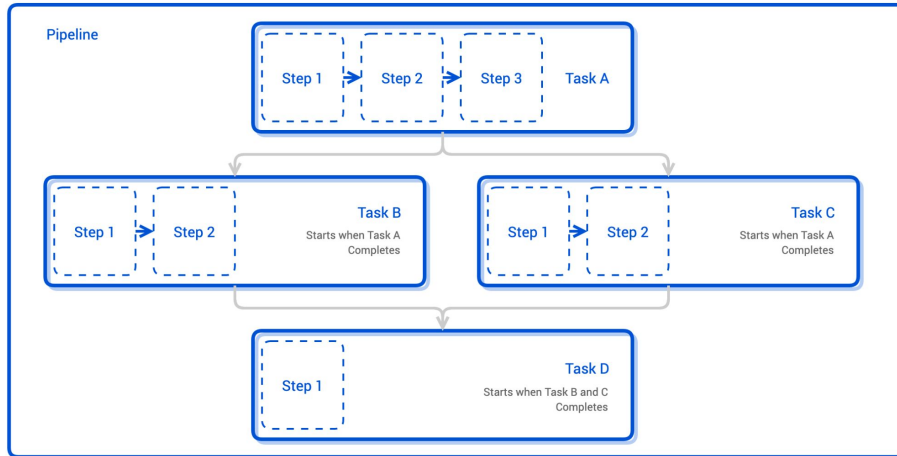
```
...
stage ('Deploy') {
    when {
        branch 'master'
        anyOf {
            environment name: 'DEPLOY_TO', value: 'production'
            environment name: 'DEPLOY_TO', value: 'staging'
        }
    }
    steps {
        echo "Deploy..."
        script {
            def server = Artifactory.server 'ccp2-artifactory'
            def uploadSpec = """{
                "files": [{
                    ...
                }]
            }"""
            server.upload(uploadSpec)
        }
    }
}
}
```

// second stage called 'Deploy'  
// stage is executed, if all conditions are true  
// e.g. only on master branch  
// and if DEPLOY\_TO is set to 'production' or 'staging'

// execute a script snippet  
supports plugins who do not yet support declarative syntax

# Tekton

- Cloud Native CI/CD Solution
- Runs on K8s and uses K8s Concepts
- Uses Custom Resource Descriptions (CRD) to define Resources like Pipeline, Tasks, ...





# Tekton Concepts

- **Step:** Operation on a workflow (compile, run test, package, create image, ...) Each step runs in a specific container image
- **Task:** Sequence of **steps** executed in order.  
A task is run in a Pod, each step becomes a running container in the Pod
- **Pipeline:** Collection of Tasks run in a directed acyclic graph (DAG)  
May run tasks in parallel and sequence (fan-out/fan-in scenarios)
- **Inputs / Outputs:** Sources and Targets to Read / Store artefacts  
e.g. Git-Repositories, Image-Registry, ...
- **PipelineRun & TaskRun:** Specific execution of a Pipeline or a Task  
e.g. Scheduled at specific intervals, triggered by events, ...

All Resources are defined using CRDs

# Tekton Examples

```
// tasks.yaml
apiVersion: tekton.dev/v1beta1
kind: Task
metadata:
  name: hello
spec:
  steps:
    - name: echo
      image: alpine
      script: |
        #!/bin/sh
        echo "Hello World"
----
apiVersion: tekton.dev/v1beta1
kind: Task
metadata:
  name: goodbye
spec:
  steps:
    - name: goodbye
      image: ubuntu
      script: |
        #!/bin/bash
        echo "Goodbye World!"
```

```
// pipeline.yaml
apiVersion: tekton.dev/v1beta1
kind: Pipeline
metadata:
  name: hello-goodbye
spec:
  tasks:
    - name: hello
      taskRef:
        name: hello
    - name: goodbye
      runAfter:
        - hello
      taskRef:
        name: goodbye
```

```
// pipeline-run.yaml
apiVersion: tekton.dev/v1beta1
kind: PipelineRun
metadata:
  name: hello-goodbye-run
spec:
  pipelineRef:
    name: hello-goodbye
```

```
// task-run.yaml
apiVersion: tekton.dev/v1beta1
kind: PipelineRun
metadata:
  name: hello-goodbye-run
spec:
  pipelineRef:
    name: hello-goodbye
```

All resources can be deployed using: `kubectl apply -f`

```
kubectl apply -f task.yaml
kubectl apply -f pipeline.yaml
kubectl apply -f pipeline-run.yaml
kubectl logs pipelinerun hello-goodbye-run -f -n default
[hello : hello] Hello World!
[goodbye : goodbye] Goodbye World!
```

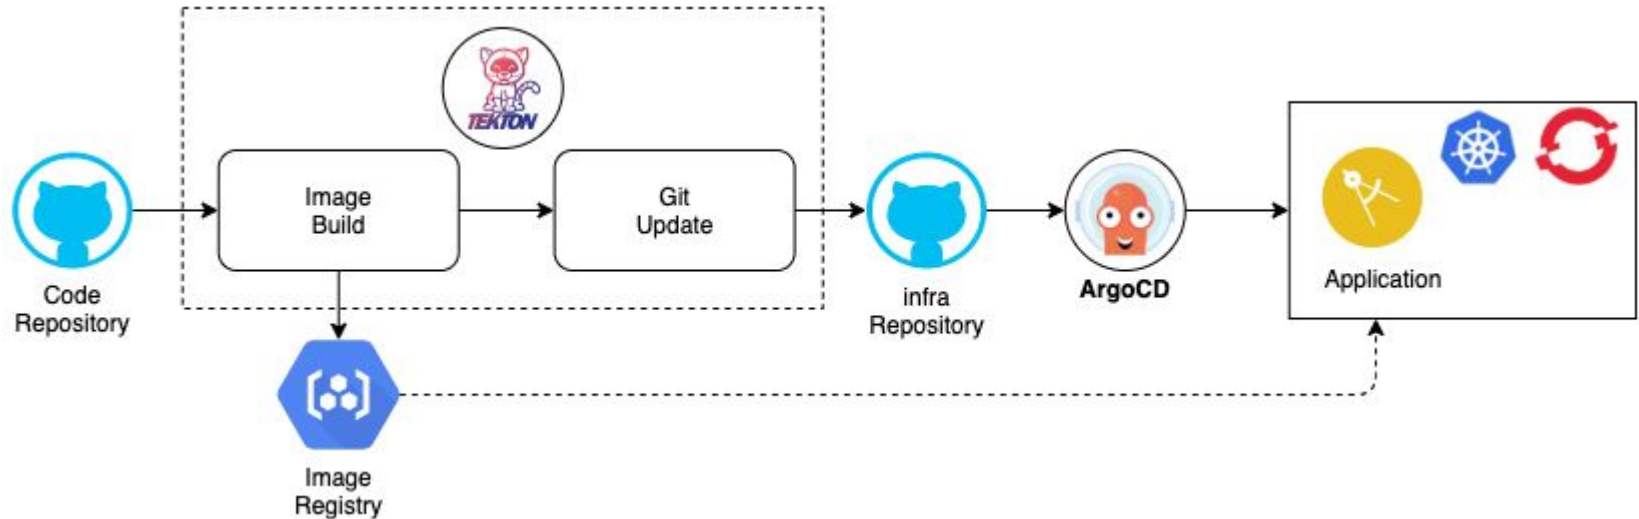
# Tooling - Automation Agent

## Process to run automation logic (on Build-Server or distributed)

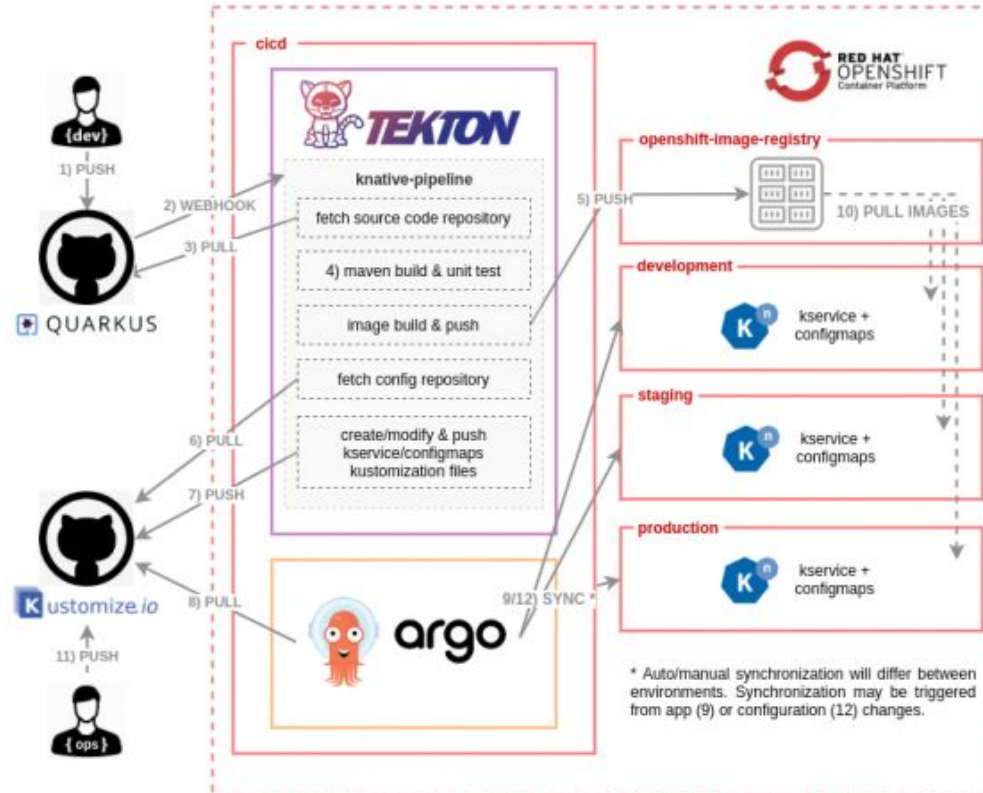
- **Shell-Script frameworks using CLI clients**
  - Interact with Tooling servers: git, mvn, docker, ...
  - Deploy to Cloud
    - PaaS: CF → cf push, OpenShift → git push, GCP → gcloud app deploy
    - IaaS: AWS → aws, Google Cloud Compute → gcloud compute ...
  - Works for small applications, too fragile for complex apps
  - Option to keep scripts simple and model the system structure in the CI/CD system (lock-in)
- **Configuration management and orchestration tools (Puppet, Chef, Ansible, ...)**
  - Infrastructure focused, limited support for application platforms or migration
- **Platform provided automation**
  - Some Deployment services exist, but not there yet for complex application structures
  - Requires to maintain redundant model of system structure → Lock-In

# GitOps

- All "Operations" are triggered using Git-Events (Push, PR, Tag, ...)
- Deployment-Configuration stored in Git-Repository(ies)



# GitOps Workflow in Openshift



# Deployment Strategies

- Expectations on **uptime and responsiveness** are high
  - e.g. **Cloud-Apps (SaaS)** are in use 24x7
- **Downtime is costly** in terms of money, customers and credibility
- Goal: **Zero-downtime deployment**
- Key to zero-downtime deployment is **decoupling** the various **parts of a release** so it can happen as independently as possible
- There exist several proven **deployment patterns**, which can be used individually or in combination:
  - Feature flags / toggles
  - Blue-Green deployment
  - Canary deployment (Rolling update)

# Feature Flags / Toggles

- New software features are deployed to production.  
They can be enabled / disabled at runtime based on configuration
- Depending on strategy features can be enabled for specific
  - users / role → test users, A/B tests
  - clients → IP based, Type web-/mobile-/fat, brand
  - instances → canary deployment
  - geo-locations → coordinated rollout
  - data schema / API versions → seamless migration to new schema & API version
- Decoupling of deployment and enablement of SW functionality
- Allows early (beta) release of new features
- Fast disabling of features in case of problems

# Feature Flags / Toggles - Implementation

Simple Implementation using a  
global generic Config Class.

```
public Spline[] reticulateSplines() {  
    if( toggleConfig.featureIsEnabled("use-new-SR-algorithm") ){  
        return enhancedSplineReticulation();  
    } else {  
        return oldFashionedSplineReticulation();  
    }  
}
```

```
public class FeatureToggleConfig {  
    Map<String, Boolean> featureConfig = new HashMap<>();  
  
    public void setFeature(String featureName, boolean enabled) {  
        featureConfig.set(featureName, enabled);  
    }  
  
    public boolean featureIsEnabled(featureName){  
        return featureConfig.get(featureName);  
    }  
}
```

Or use a framework like Togglz  
(<https://www.togglz.org>)



# Feature Flags / Toggles - Categories

## Release Toggles

- Deploy application with new features and enable them eventually (maybe never)
- Decouples deployment from releasing
- Transitional (should be removed after some time in production)

## Experiment Toggles

- Enable different implementations for different groups of people
- Often used in A/B testing scenarios
- Transitional (should be removed after tests are completed)


## Ops Toggles

- Used to control operational aspects (e.g. reduce functionality during a DB update)
- Usually transitional, some long living

## Permission Toggles

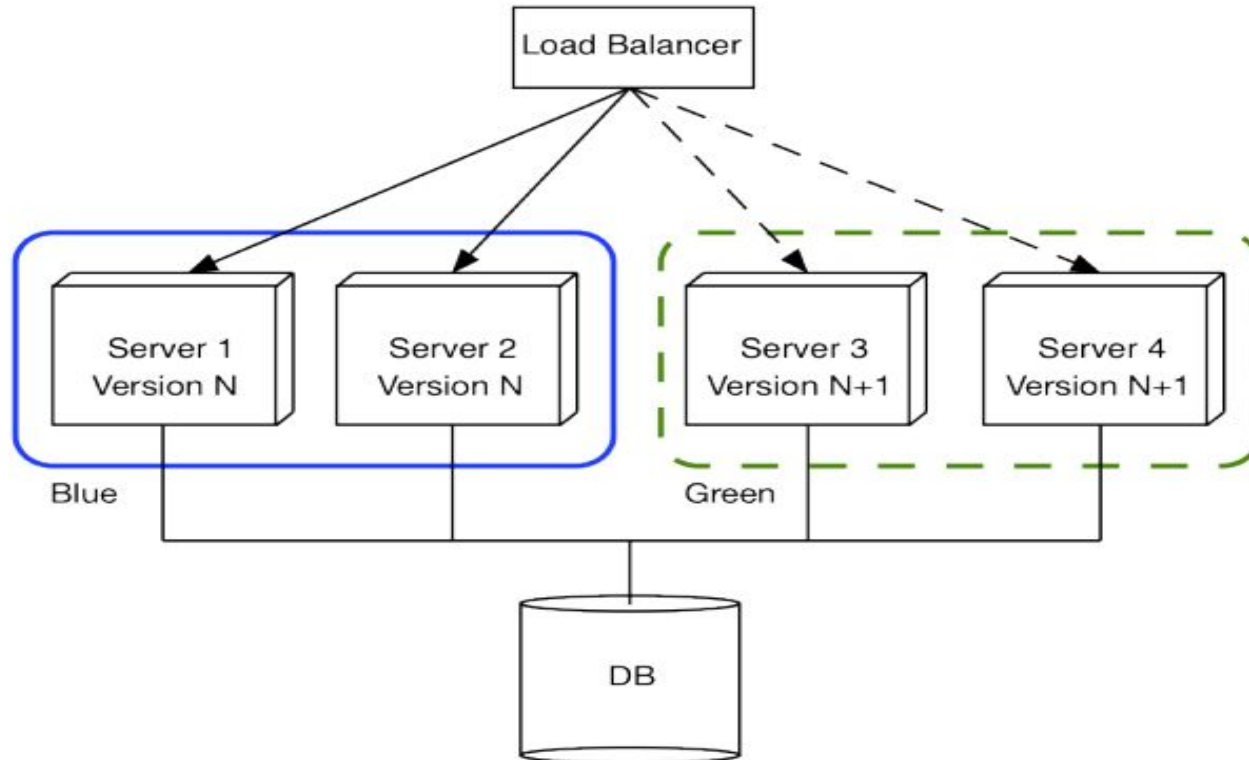
- Enabling some features for a certain group of (internal) people (e.g. alpha / beta testers)
- Very often long living but also dynamic (continuous new features)

# Blue-Green Deployment



- Two identical environments (blue & green)
- Green is running productively
- Deploy new release to blue environment, let it warm up
  - smoke tests to verify it works
  - move to the new version by changing the router configuration
- Blue becomes productive. Green is used for the next release.
- If something goes wrong, change router back to green env.

# Blue-Green Deployment

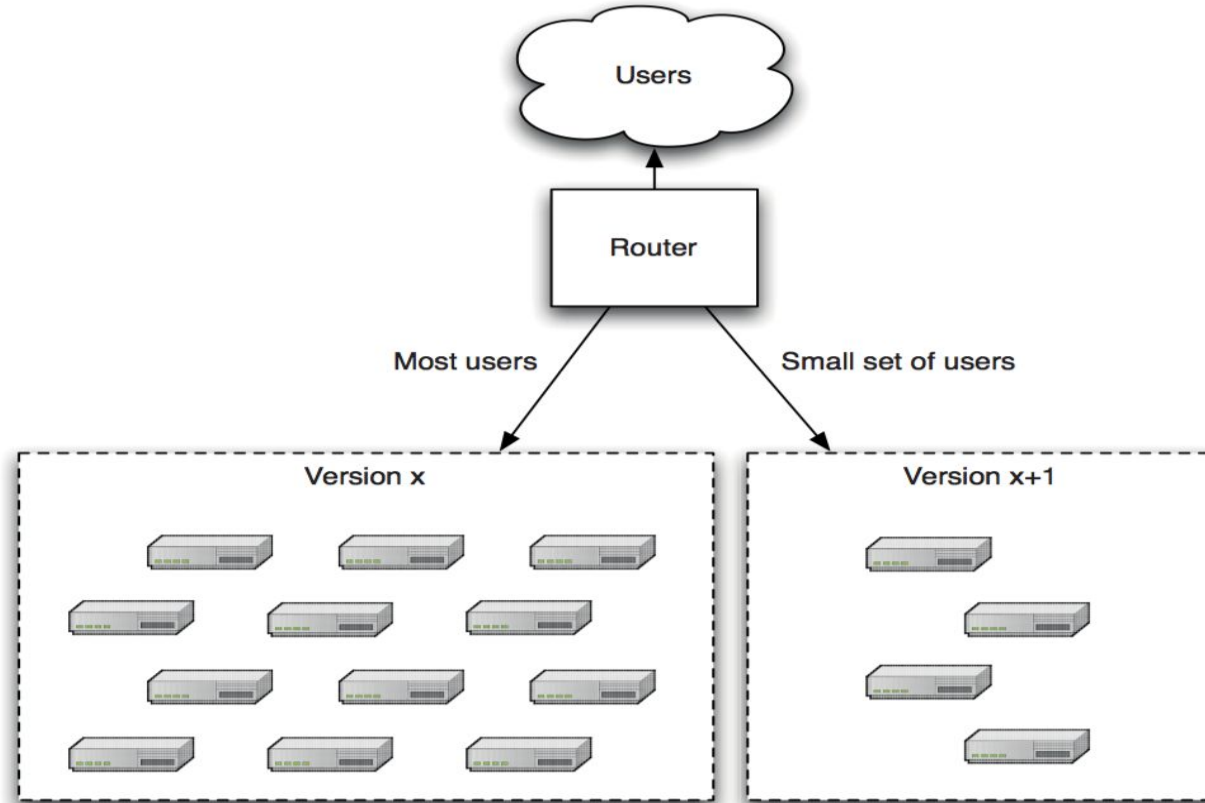


# Canary Deployment (Rolling update)

Instead of instant switching all requests from version  $x$  to the  $x+1$ , users are gradually moved to the new version (e.g. number of new instances are increased and old instances decreased)


- Benefits:
  - Make rolling back easy
  - A/B testing can be done by routing some users to new version and some to the old
  - You can check if the application meets capacity requirements gradually
- Disadvantages
  - Harder to manage in smaller installations
  - Imposes further constraints on DB schema upgrades

# Canary Deployment



# Zero Downtime State transition

- Zero Downtime Migration of cloud-native applications is easy
  - service/code is stateless → state has to be externalized
  - use feature flag, blue-green or canary deployment



- But what if you need to migrate the schema of your database?
    - No maintenance-window to migrate the database
    - No updates may go lost
    - No performance hints
- Zero downtime State transition is hard

# State transition options?

- Stop old app, migrate data, start new app → not an option
- Reduce app functionality → influence on customer
  - Switch app to read only, create a copy of the db, migrate to new schema, switch app to new version
- Synchronization → error prone, performance?
  - Create copy of db, migrate the schema, run a process which keeps both versions in sync, switch app to new version
- Design the application to allow migration → needs planning
  - plan migration in several small steps instead of one large change
  - needs back and forward compatibility in each step
  - use well known building-blocks (patterns) for migration steps

# State transition building-blocks

- Multiple (baby) steps per transition
- Each step is executed in sequence:
  - DB: → DB migration, one-time process
  - Code: → Zero-downtime (B/G, canary) deploy of all instances to new version
- Possible to go back within the block, until it is completed (otherwise start a new reverse transition)

## Add a Field/Column

1. DB: Add new Column
2. DB: Preset value (NULL, default, computed)
3. Code: read from and write to the Column

## Change a Field/Column (name, type, format)

1. DB: add new column (no constraints e.g. NotNull)
2. Code: read from old column, write to both
3. DB: copy data from old to new column  
(for large datasets do it in multiple shards)  
Add required constraints (eg. NotNull) to new column
4. Code: read from new column, write to both
5. DB: delete constraints from old column
6. Code: read from, write to new column only
7. DB: delete the old column

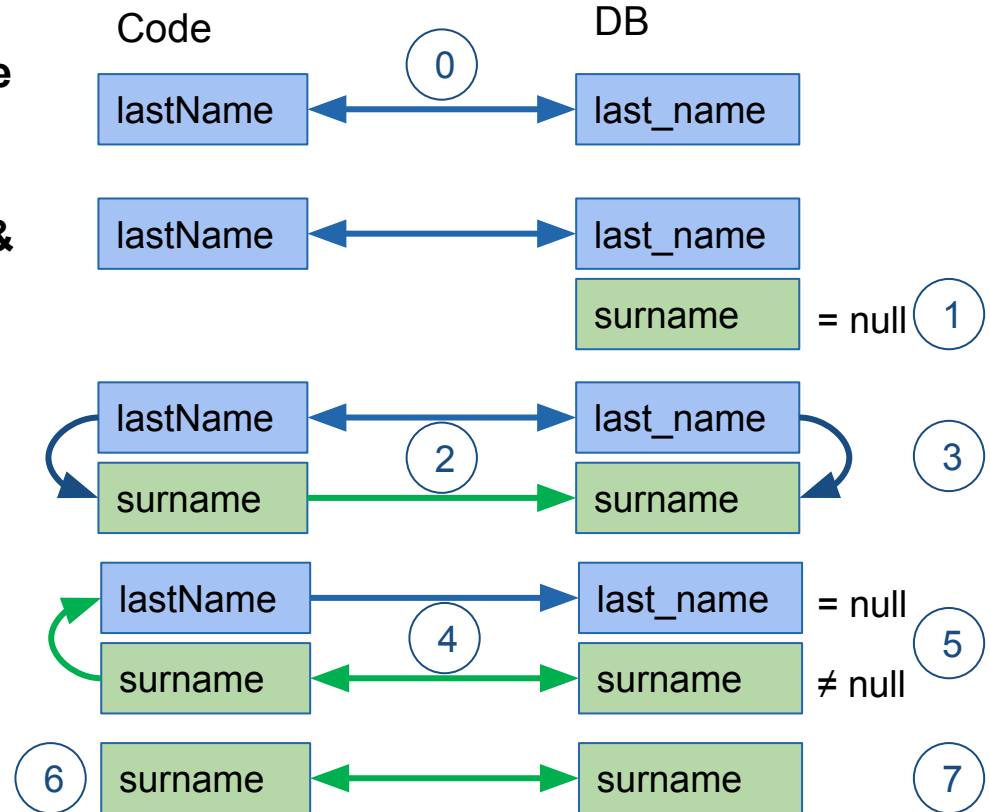
## Delete a Field/Column

1. DON'T! It is a destructive operation  
→ Keep the Column for a retention period
2. Code: stop reading, but keep writing the column
3. (in consolidation phase after retention period)  
Code: stop, writing the column  
DB: Delete the Column



# Example - Rename field/column: last\_name → surname

0. DB field is **last\_name** / Code **lastName**
1. DB: add field **surname**
2. Code: add **surname** field, still **read from last\_name, write to last\_name & surname**
3. DB: copy values from **last\_name** to **surname**
4. Code: **read from surname, write to last\_name & surname**
5. DB: remove **NotNull** constraint from **last\_name**, add to **surname**
6. Code: remove **lastName**, **read & write to surname only**
7. DB: **delete last\_name** column




# More complex operation

- Add/Change/Delete a Relation → Add/Change/Delete a Column
- Add/Split/Merge/Delete Tables → Series of Column operations
- What about Referential Integrity constraints (NotNull, foreign keys,...)
  - Contain no business value
  - (only) a safety-net to avoid corruption→ remove before and recreate after transition

# Best practice


- Use migration tool like flyway or liquibase
  - allows back and forward transition for each step
  - changes in version control
- Decouple the database
  - use simpler database model per microservice
  - less dependencies and side-effects
  - easier to change
- Use migration friendly Architecture
  - Use Event-Sourcing and CQRS
  - Allows to recreate the view model

# Decoupling the data model



monolith - single database

Single database and common data model  
→ lots of dependencies and side-effects  
→ difficult to change





microservices - application databases

Individual database and data model per service  
→ less dependencies and side-effects on other services  
→ flexible to migrate and adopt

# Continuous Delivery – How?

- A lot of good literature is available



From the authors of *The Visible Ops Handbook*



## The Phoenix Project

A Novel About IT, DevOps,  
and Helping Your Business Win

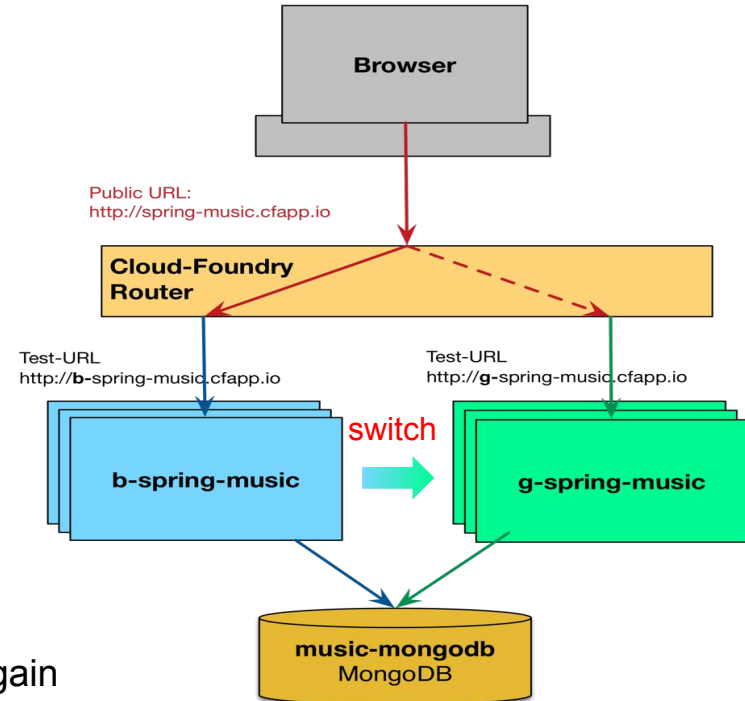
Gene Kim, Kevin Behr, and George Spafford

# Appendix

# Blue-Green Deployment

There are several options to do blue green deployment in CF  
Most complete version is:

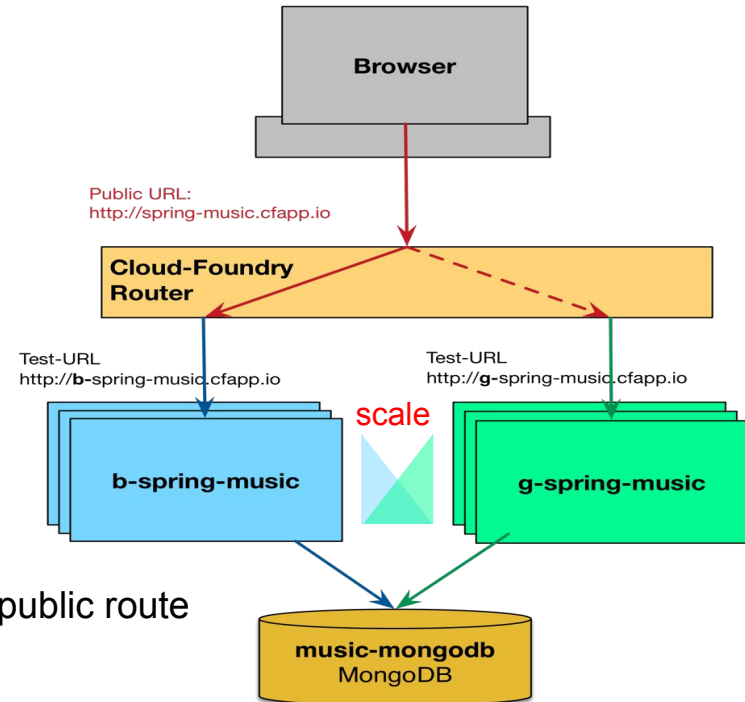
1. Current version is deployed as **b-spring-music**,  
Public route <https://spring-music.cfapp.io> points to it
2. Push new version with name **g-spring-music**,  
scale it up to the same number of instances  
and bind it to the required services (db)
3. Test the new version on the test URL  
<https://g-spring-music.cfapp.io> (smoke tests)
4. **Add** public route to the new instance  
`cf map-route g-spring-music cfapp.io -n spring-music`  
and **remove** it from the old instance  
`cf unmap-route b-spring-music cfapp.io -n spring-music`  
(for a very short moment the requests were routed to  
both version using a round-robin algorithm)
5. After some time, if everything runs correctly  
the old app can be deleted. If not, you can switch back again



# Canary Deployment with CF

Basically an extension to Blue-Green deployment:

1. Current version is deployed as **b-spring-music**,  
Public route <https://spring-music.cfapp.io> points to it
2. Push new version with name **g-spring-music**,  
scale it ~~up~~ to the ~~same~~ **minimal** number of instances (1)  
and bind it to the required services (db)
3. Test the new version on the test URL  
<https://g-spring-music.cfapp.io> (smoke tests)
4. Add public route to the new instance  
`cf map-route g-spring-music cfapp.io -n spring-music`
5. Gradually **scale up** instances of new **g-spring-music**  
and **scale down** instances of old **b-spring-music**  
(distribution of requests is proportional to number of  
instances of old and new version)
6. If b-spring music is down to 1 instance remove it from the public route  
`cf unmap-route b-spring-music cfapp.io -n spring-music`  
After some time, if everything runs correctly  
the old app can be deleted. If not, you can switch back again





# DevOps Practices

- Version Control For All
- Automated Testing
- Proactive Monitoring and Metrics
- Kanban/Scrum
- Visible Ops/Change Management
- Configuration Management
- Incident Command System
- **Continuous Integration / Deployment / Delivery**
- “Put Developers On Call”
- Virtualization/Cloud/Containers
- Toolchain Approach
- Transparent Uptime/Incident Retrospectives

# Things Not To Do

- Only Token Gestures
  - “Ops team, change your name to DevOps team!”
  - “Put DevOps in those job titles!”
- Only Implement Tools
  - Changing tools without changing tactics leaves the battlefield strewn with bodies
- Create more silos (Dev, DevOps, Ops)
- Devalue operations Or development knowledge
- Anything you're not measuring the impact of