

Cortex-M Architecture

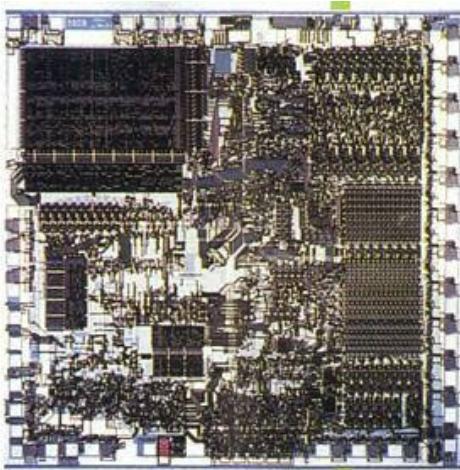
Computer Engineering 1

CT Team: A. Gieriet, J. Gruber, B. Koch, M. Loeser, M. Meli,
M. Rosenthal, M. Ostertag, A. Rüst, J. Scheier

Motivation

■ Intel

1978

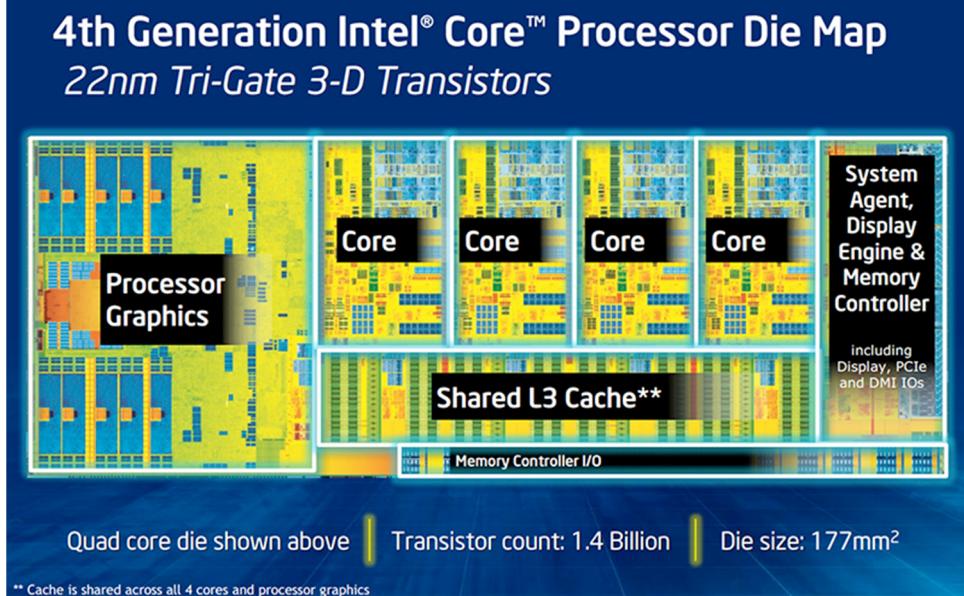


8086
33 mm²
29k transistors
5 MHz



50'000 x
transistors
720 x
clock speed

2013



core i7 Haswell
177 mm²
1400M transistors
3.6 GHz



Motivation

■ ARM (Cortex-M)

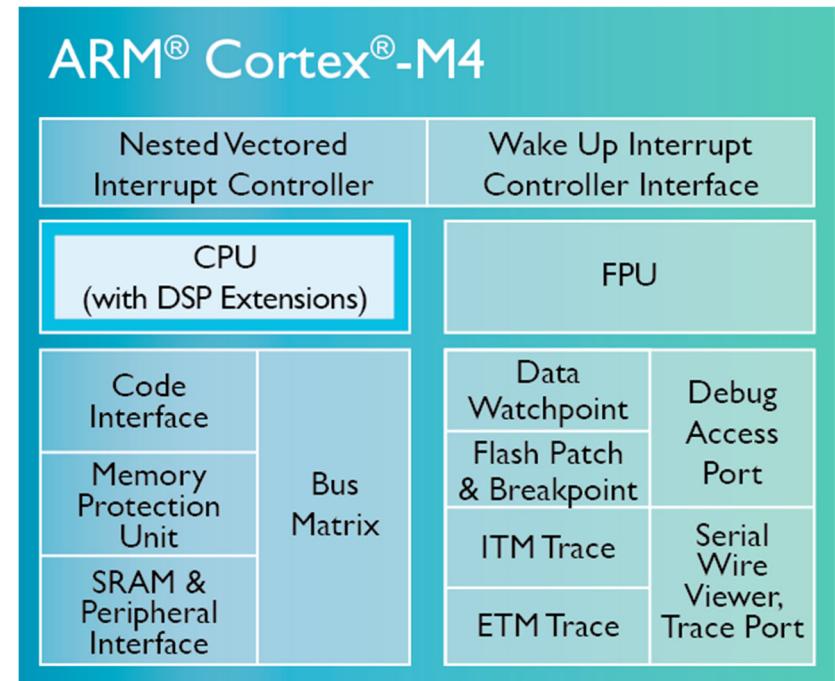
1985



1'400 x
transistors

36 x
clock speed

2014



ARM1
25k Transistors
5 MHz

ARMv7 - Cortex-M4
35 Mio Transistors
180 MHz
4 mm²

Agenda

- **Hardware Platform**
 - STM32F4 and evaluation board, ARM Processor Portfolio
- **CPU Model**
 - Register, ALU, Flags, Control Unit
- **Instruction Set**
 - Assembly, Instruction Types, Cortex-M0
- **Program Execution**
 - Fetch and Execute
- **Memory Map**
 - ARM, ST, CT-Board
- **Integer Types**
 - Sizes, Little Endian vs. Big Endian, Alignment
- **Object File Sections**
 - Code, Data, Stack

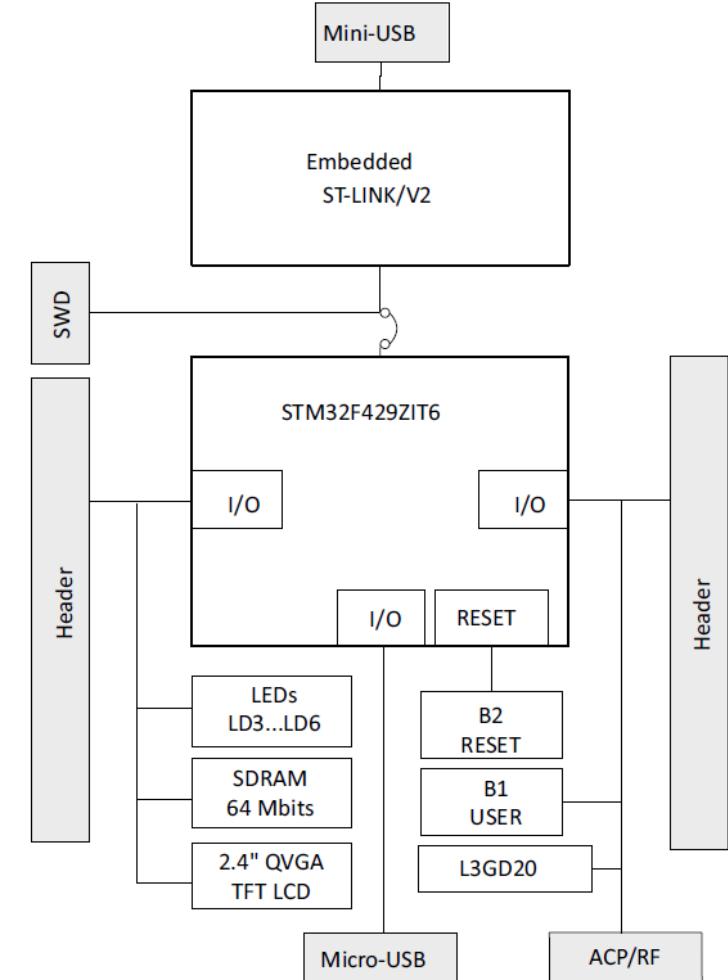
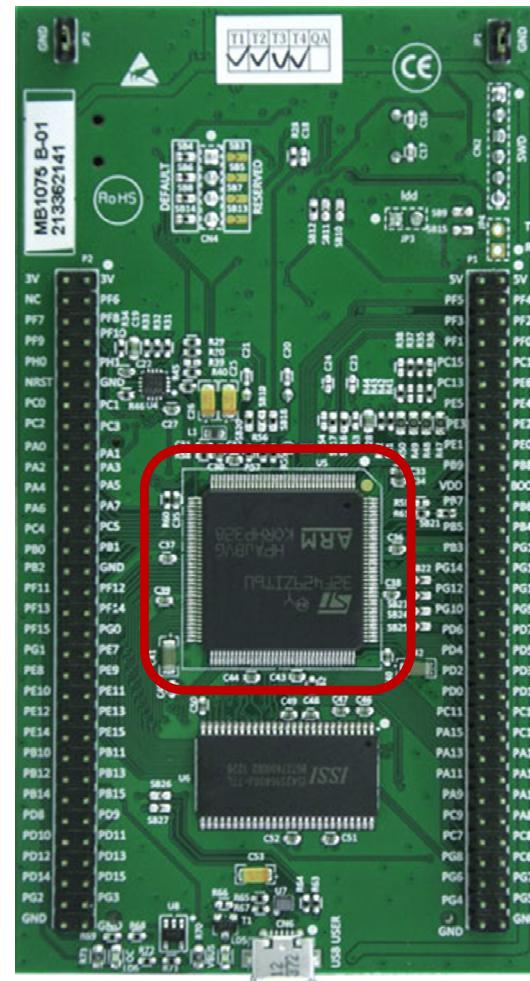
Learning Objectives

At the end of this lesson you will be able

- to describe what an 'Instruction Set Architecture' is
- to outline the Cortex-M architecture and enumerate the main components and their functions
- to enumerate the instruction type categories
- to understand the structure of the Cortex-M instruction set
- to explain how a processor executes a program
- to recall the registers of the Cortex-M, their layout and their functions
- to explain and draw a memory map
- to calculate the size in bytes for a memory block given by its start and end address
- to determine the end address of a memory block given by the start address and the number of bytes
- to understand that sizes of integer types in C depend on the architecture and that portability can be enhanced by using the C99 types in stdint.h
- to explain the difference between 'little endian' and 'big endian' and to show how multi-byte integer values are mapped to individual bytes
- to list the three typical memory sections of an object file and to explain their content

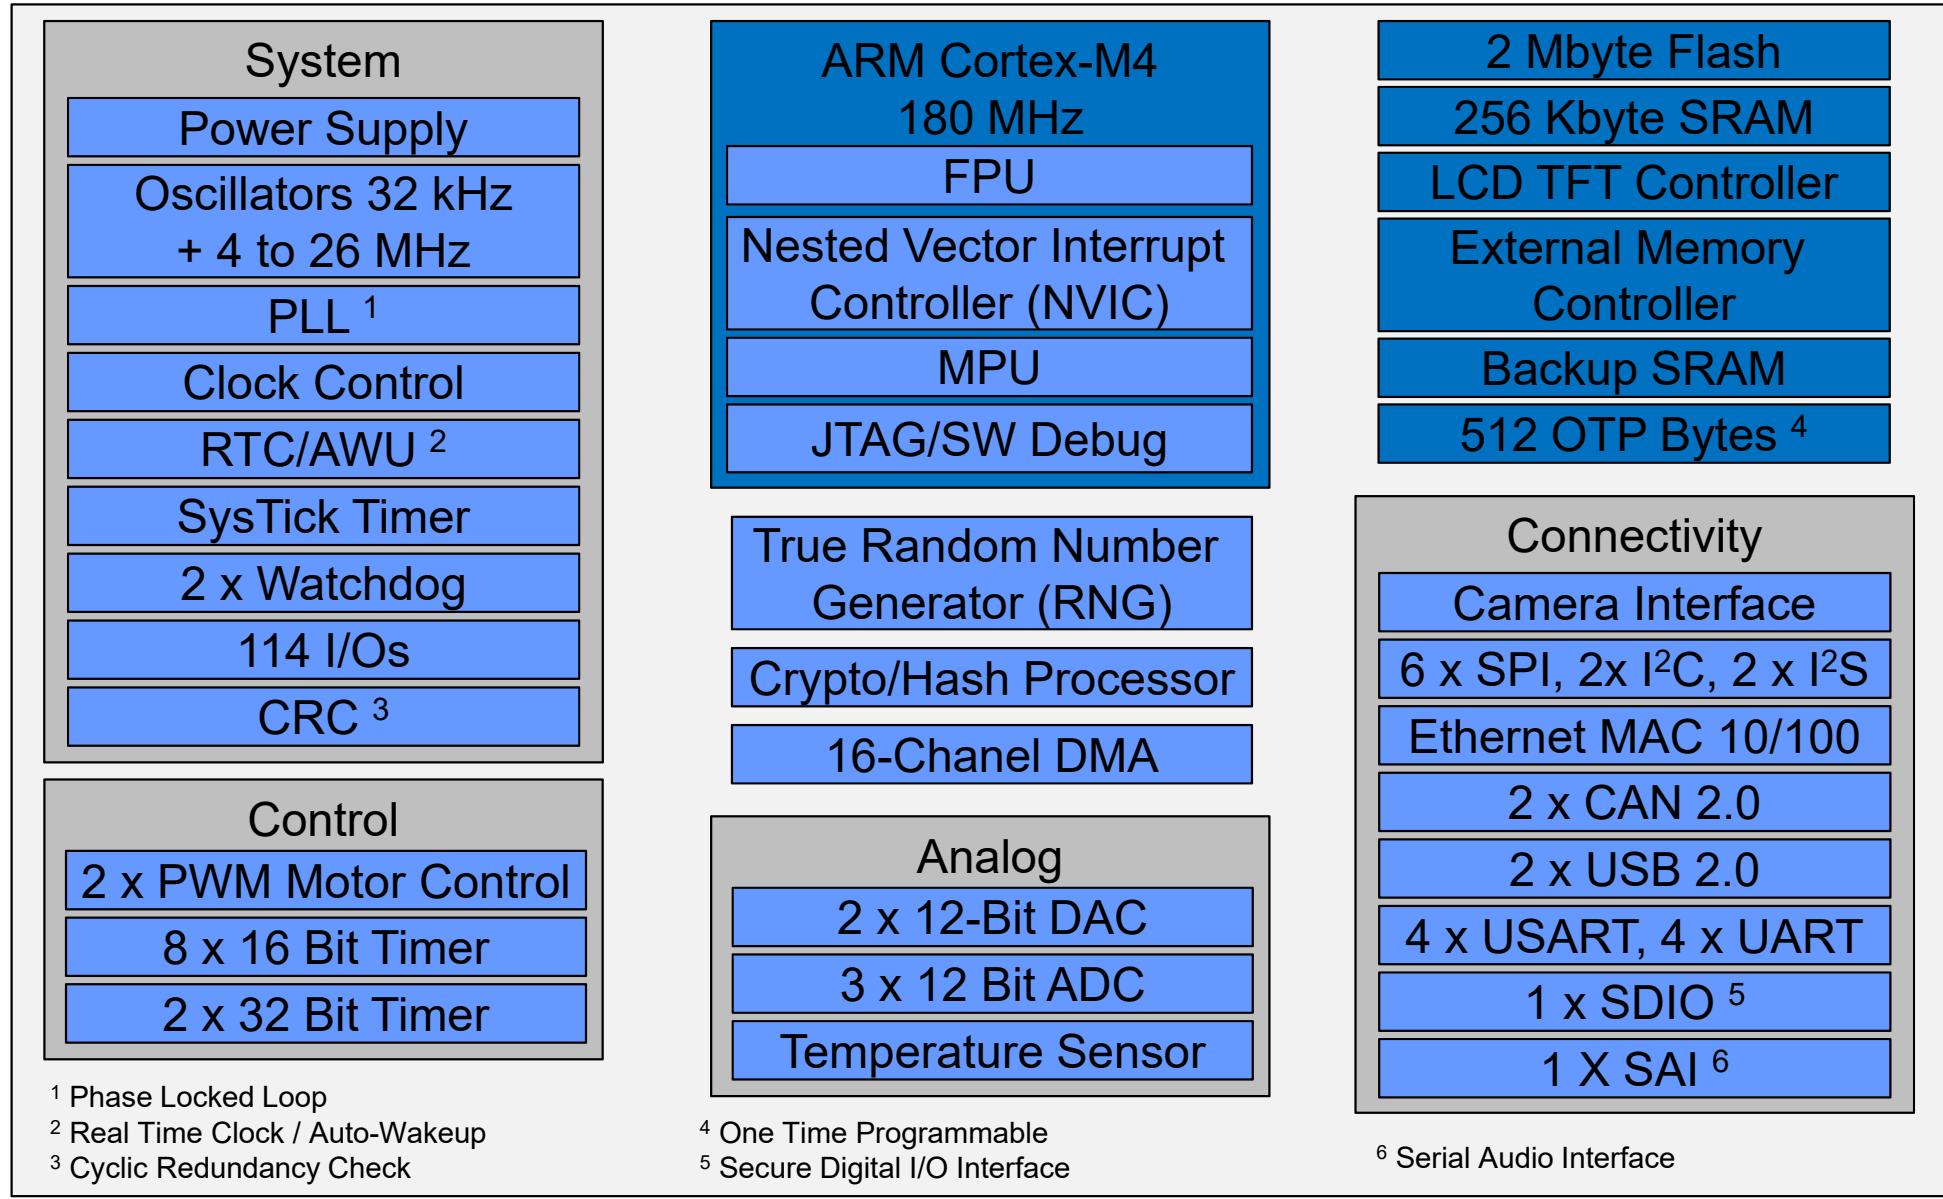
Hardware Platform: STM32F4-Discovery

■ Evaluation board STM32F4-DISCO



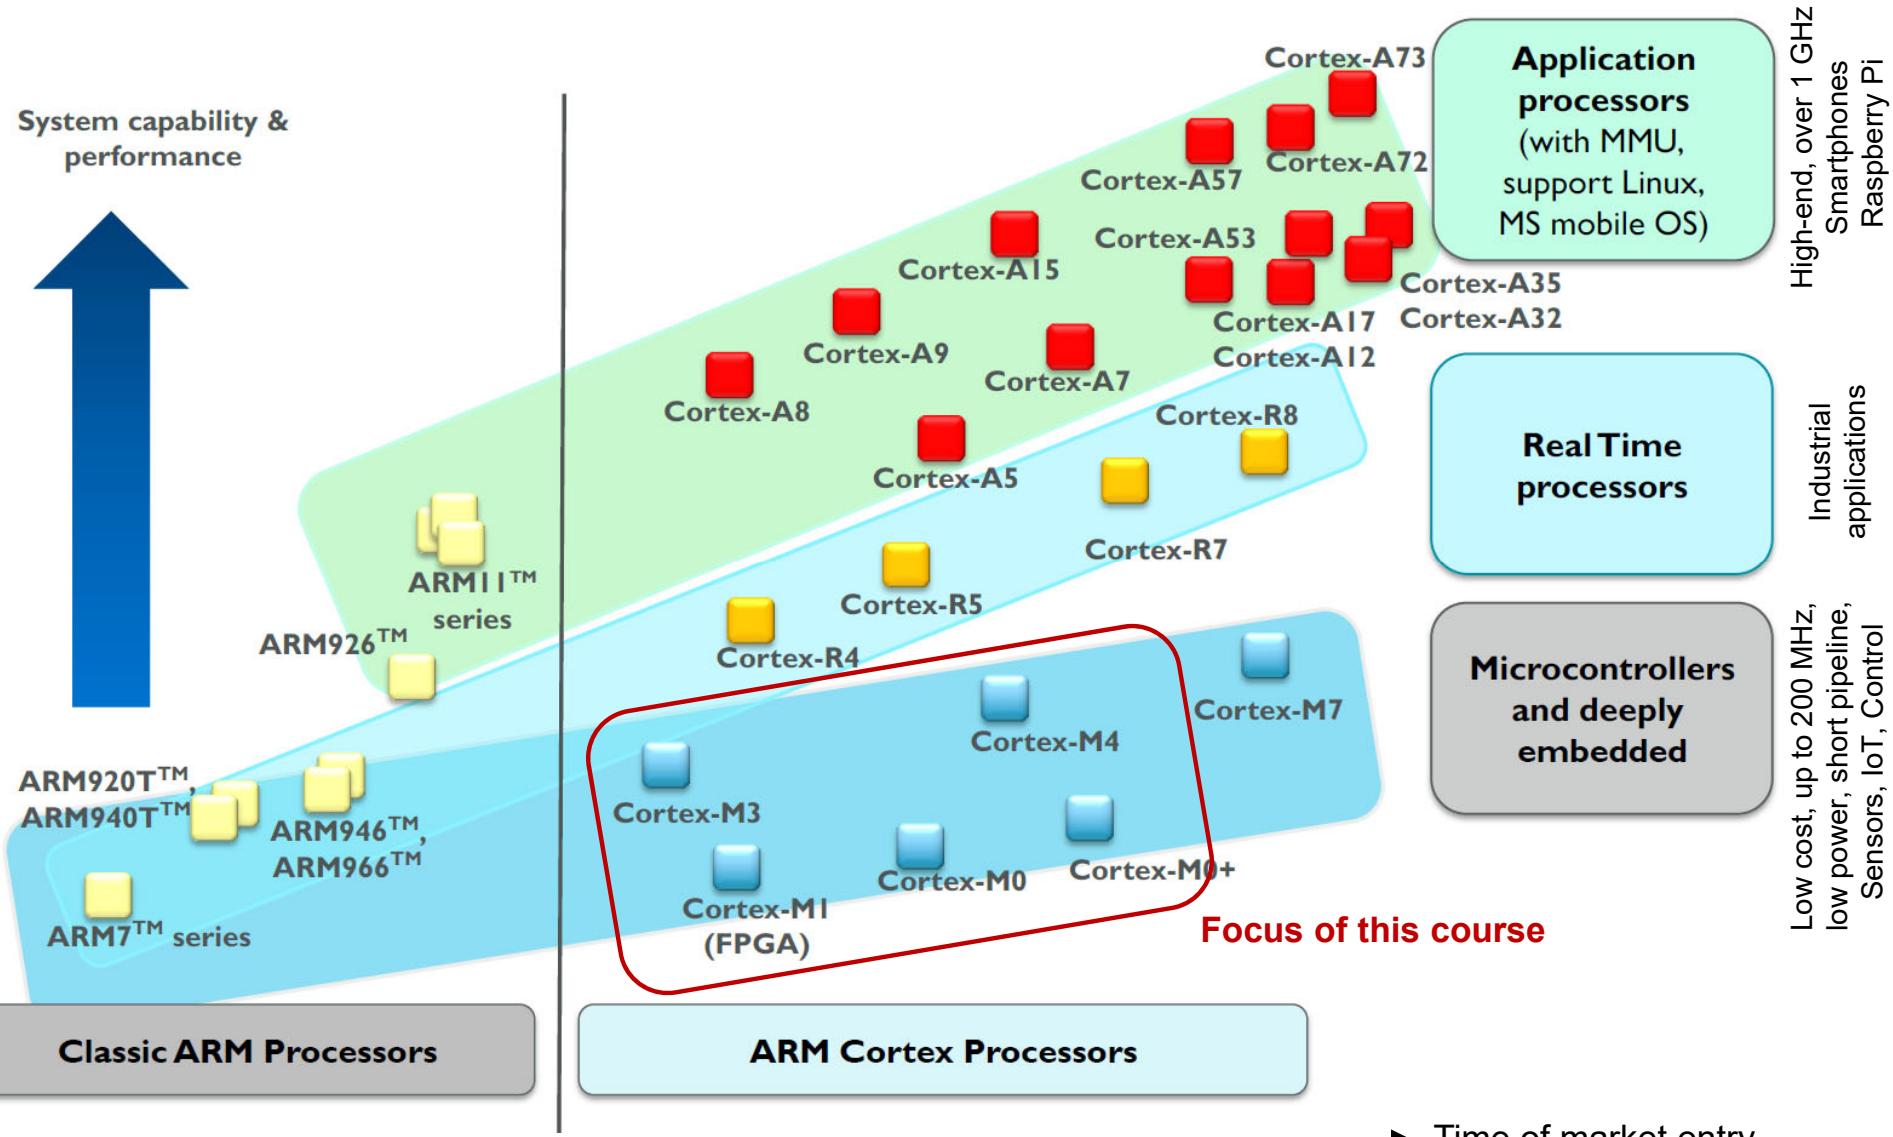
source: STMicroelectronics

Hardware Platform: STM32F429I



based on source of STMicroelectronics

Hardware Platform: ARM Processor Portfolio

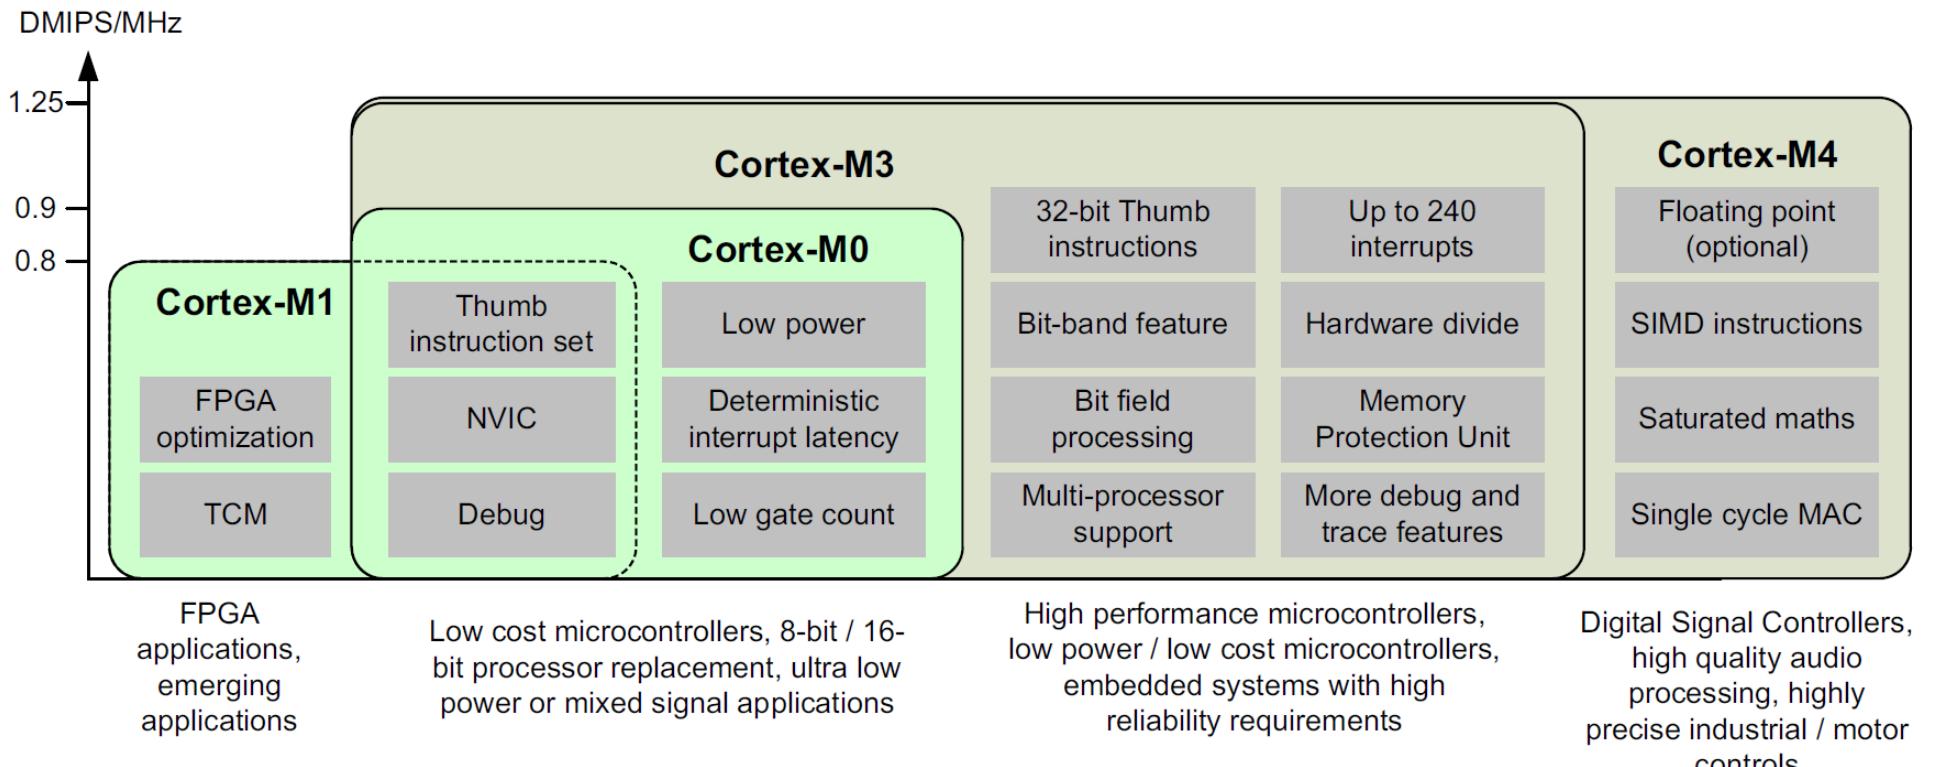


Source: ARM Limited. White paper: ARM® Cortex®-M for Beginners

Hardware Platform

■ Cortex-M Processor Family

- Hardware used in this course is a Cortex-M4
- But most of the time we only use the simpler Cortex-M0 subset



source: Joseph Yiu, *The definitive Guide to the Cortex-M0*

■ Instruction Set Architecture (ISA) → CT1

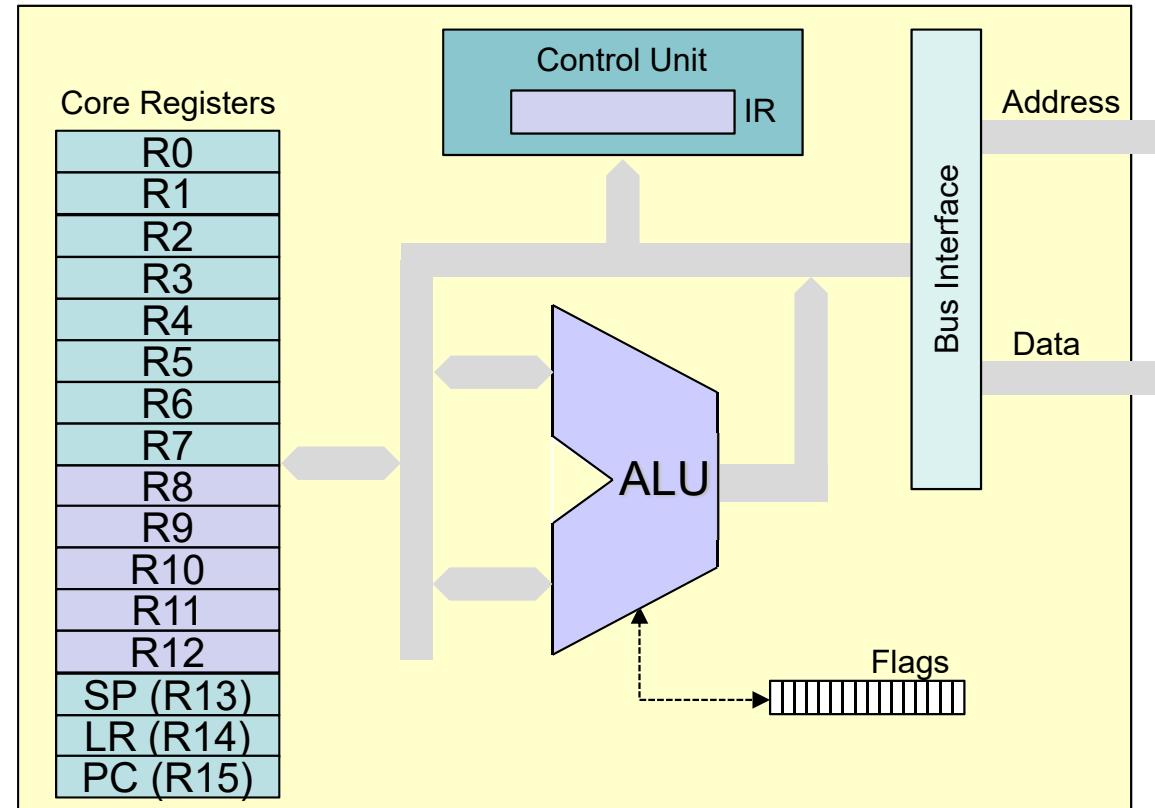
What the programmer sees of a computer

- Instruction Set
 - Available instructions?
- Processing width
 - 8-bit/16-bit/32-bit?
- Register set
 - How many registers? Which size?
- Addressing modes
 - How can memory and IO be accessed?
- ARM Cortex-M
 - ARMv6-M → Cortex-M0
 - ARMv7-M → Cortex-M3/M4 (Superset of ARMv6-M)

CPU Model

■ CPU Components

- Core Registers
- 32-bit ALU
- Flags (APSR)
- Control Unit with IR (Instruction Register)
- Bus Interface

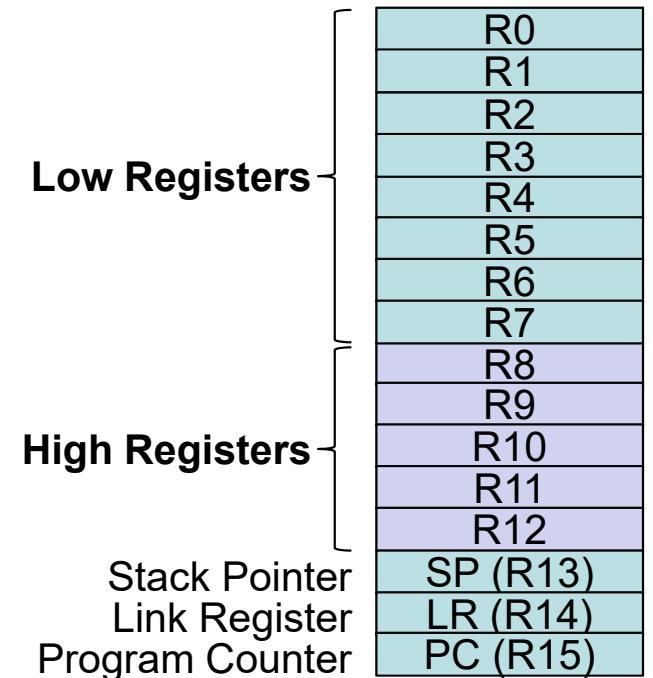


simple CPU model based on the Programmers'
Model of the ARM Cortex-Mx CPUs¹⁾

¹⁾ without pipelining

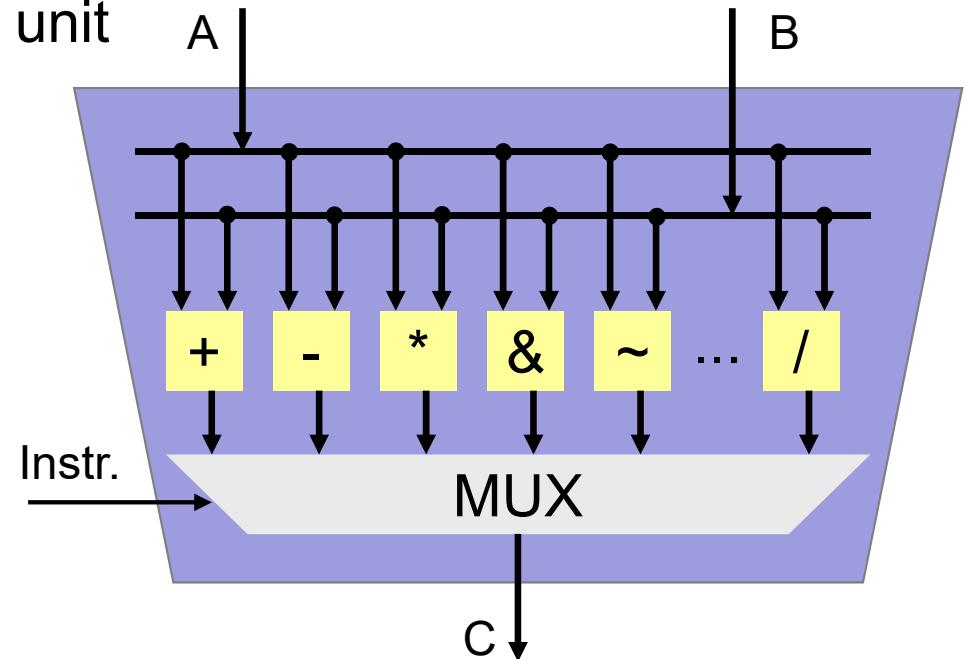
■ 16 Core Registers

- Each 32-bit wide
- 13 General-Purpose Registers
 - Low Registers R0 – R7
 - High Registers R8 – R12
 - Used for temporary storage of data and addresses
- Program Counter (R15)
 - Address of next instruction
- Stack Pointer (R13)
 - Last-In First-Out for temporary data storage
- Link Register (R14)
 - Return from procedures



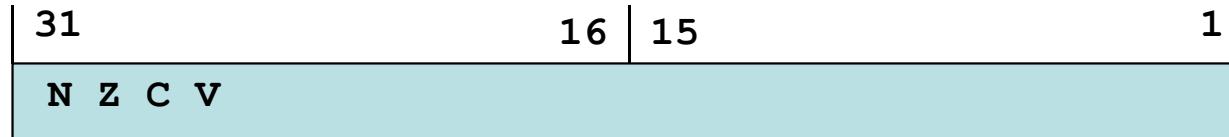
■ ALU – Arithmetic Logic Unit

- 32-bit wide data processing unit
 - inputs A and B
 - result C
- integer arithmetic
 - addition / subtraction
 - multiplication / division
 - sign extension
- logic operations
 - AND, NOT, OR, XOR
- shift/rotate
 - left / right



■ APSR¹⁾ or Flag-Register

- Bits set by CPU based on results in ALU



N Negative
Z Zero
C Carry
V Overflow

As seen in the IDE

Register	Value
R0	0x20000078
R1	0x20000278
R2	0x20000278
R3	0x20000278
R4	0x08000860
R5	0x00000000
R6	0x00000000
R7	0x00000000
R8	0x00000000
R9	0x00000000
R10	0x08000860
R11	0x08000860
R12	0x00000000
R13 (SP)	0x20000678
R14 (LR)	0x08000243
R15 (PC)	0x08000270
xPSR	0x41000000
N	0
Z	1
C	0
V	0

¹⁾APSR: Application Processor Status Register

■ Control Unit

- Instruction Register (IR)
 - Machine code (opcode) of instruction that is currently being executed
- Controls execution flow based on instruction in IR
- Generates control signals for all other CPU components

■ Bus Interface

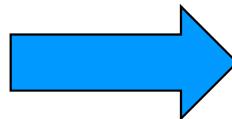
- Interface between internal CPU bus and external system-bus
 - contains registers to store addresses



■ Processors interpret binary coded instructions

- But binary is hard for programming
- Therefore instructions in human readable text form
 - → assembly
- Assembler (tool) does the translation
 - assembly → binary

ADDS R0 ,R0 ,R1



0001'1000'0100'0000
=
0x1840

■ Assembly Program

- Label (optional)
- Operands
- Instruction (Mnemonic)
- Comment (optional)

Label	Instr.	Operands	Comments
demoprg	MOVS	R0 , #0xA5	; copy 0xA5 into register R0
	MOVS	R1 , #0x11	; copy 0x11 into register R1
	ADDS	R0 ,R0 ,R1	; add contents of R0 and R1 ; store result in R0
	LDR	R2 ,=0x2000	; load 0x2000 into R2
	STR	R0 , [R2]	; store content of R0 at ; the address given by R2

■ Instruction Types

- **Data transfer**

- Copy content of one register to another register
- Load registers with data from memory
- Store register contents into memory

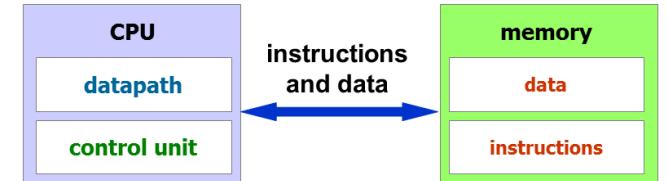
- **Data processing**

- Arithmetic operations → + - * / ...
- Logic operations → AND, OR, ...
- Shift / rotate operations

- **Control flow**

- Branches
- Function calls

- **Miscellaneous (various)**

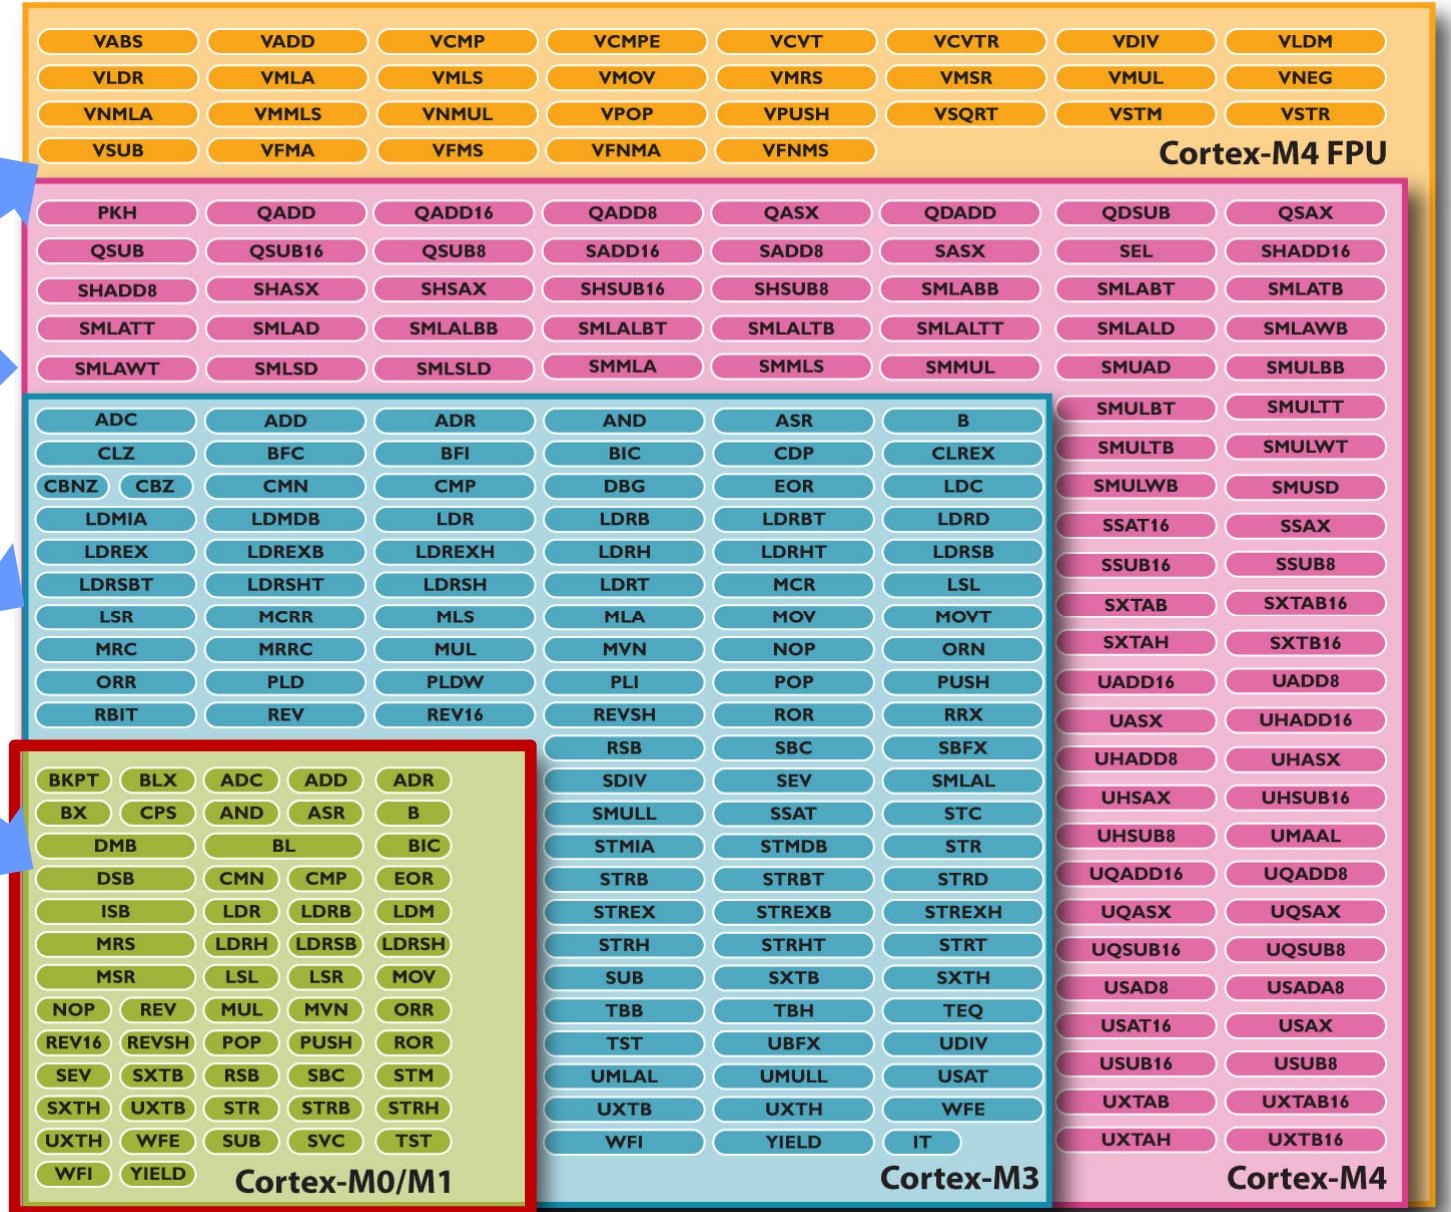


Type	Frequency
Data transfer	43%
Control flow	23%
Arithmetic	15%
Compare	13%
Logical	5%
Miscellaneous	1%

Data processing

Instruction Set

- Instructions Cortex-M
- ARMv7-M Architecture
- ARMv6-M Architecture
- CT1
Instruction Set Cortex-M0
=
Subset Cortex-M3/4



■ Overview Cortex-M0

Instruction Type	Instructions
Move	MOV, MOVS
Load/Store	LDR, LDRB, LDRH, LDRSB, LDRSH, LDM, STR, STRB, STRH, STM
Add, Subtract, Multiply	ADD, ADDS, ADCS, ADR, SUB, SUBS, SBCS, RSBS, MULS
Compare	CMP, CMN
Logical	ANDS, EORS, ORRS, BICS, MVNS, TST
Shift and Rotate	LSLS, LSRS, ASRS, RORS
Extend	SXTH, SXTB, UXTH, UXTB
Reverse	REV, REV16, REVSH
Branch	B, BL, B{cond}, BX, BLX
Stack	POP, PUSH
Processor State	BKPT, CPS, MRS, MSR, SVC
No Operation	NOP
Hint / Synchronization	DMB, DSB, ISB, SEV, WFE, WFI, YIELD

Instruction Set

■ Cortex-M0: 16-bit Thumb instruction encoding

15	14	13	12	11	10	9	8	7	6	5	4	3	2	1	0
0	0	0	opcode		imm5			Rm		Rd		shift by immediate, move register			
0	0	0	1	1	0	opc		Rm		Rn		Rd		add/subtract register	
0	0	0	1	1	1	opc		imm3		Rn		Rd		add/subtract immediate	
0	0	1	opcode		Rdn				imm8					add/sub./comp./move immediate	
0	1	0	0	0	0	0	opcode			Rm		Rd		data-processing register	
0	1	0	0	0	1	opcode	DN		Rm		Rd			special data processing	
0	1	0	0	0	1	1	1	L	Rm	0	0	0		branch/exchange	
0	1	0	0	1	Rt			PC-relative imm8						load from literal pool	
0	1	0	1	opcode		Rm			Rn		Rt			load/store register offset	
0	1	1	B	L		imm5			Rn		Rt			load/store word/byte imm. offset	
1	0	0	0	L		imm5			Rn		Rt			load/store halfword imm. offset	
1	0	0	1	L	Rt			SP-relative imm8						load from or store to stack	
1	0	1	0	SP	Rd			imm8						add to SP or PC	
1	0	1	1	x	x	x	x	x	x	x	x	x	x	x	miscellaneous
1	1	0	0	L	Rn			register list						load/store multiple	
1	1	0	1	cond			imm8							conditional branch	
1	1	0	1	1	1	1	0	x	x	x	x	x	x	x	undefined instruction
1	1	0	1	1	1	1	1		imm8						service (system) call
1	1	1	0	0		imm11									unconditional branch
1	1	1	0	1	x	x	x	x	x	x	x	x	x	x	32-bit instruction
1	1	1	1	x	x	x	x	x	x	x	x	x	x	x	32-bit instruction

Note: There are some inconsistencies in ARMs use of the notation Rd (destination) and Rt (target). The slides use the ARMv6-M Architecture Reference Manual (ARM DDI 0419C (ID092410)) as a reference.

Instruction Set

■ Cortex-M0 Example Program

- Assembler converts each Assembly instruction to 16-bit opcode
- Memory addresses in steps of 2 (Zweierschritte)
 - Reason: opcodes are 16-bit (two bytes) long, e.g. **0x20A5**

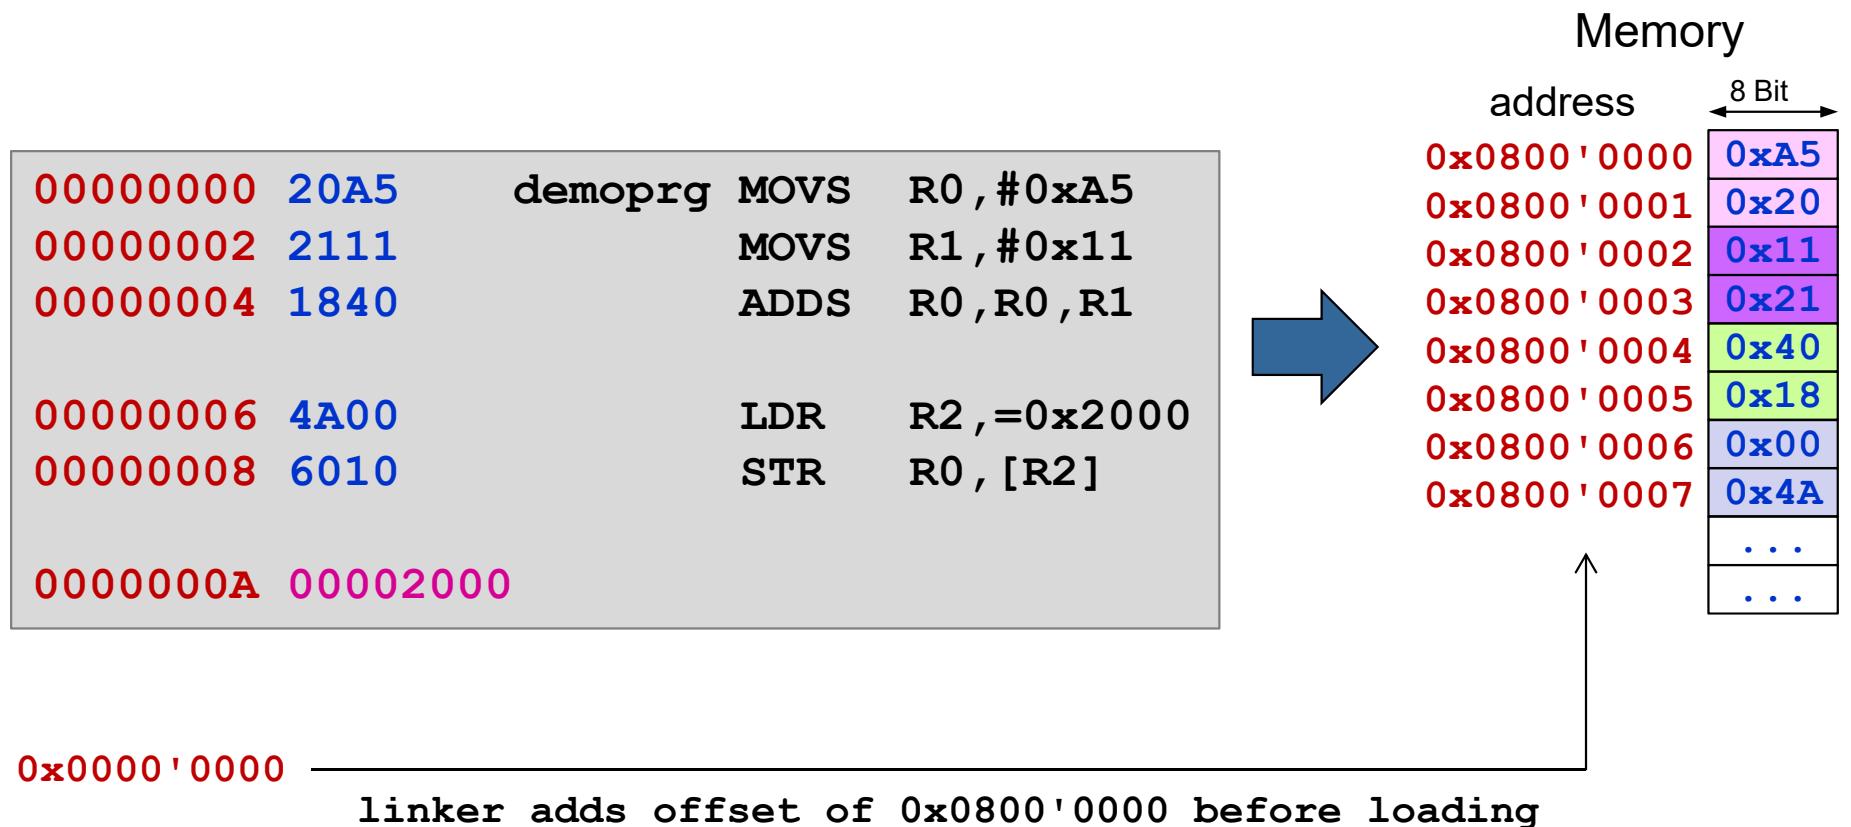
<i>Memory address</i>	<i>Opcode</i>	<i>Listfile</i>		
00000000	20A5	demoprg	MOVS R0, #0xA5	; copy 0xA5 into R0
00000002	2111		MOVS R1, #0x11	; copy 0x11 into R1
00000004	1840		ADDS R0, R0, R1	; add contents of R0 and R1 ; store result in R0
00000006	4A00		LDR R2, =0x2000	; load address into R2
00000008	6010		STR R0, [R2]	; store content of R0 at ; the address given by R2
1)	0000000A	00002000		

1) Binary ← generated by Assembler (tool) → Source code

Instruction Set

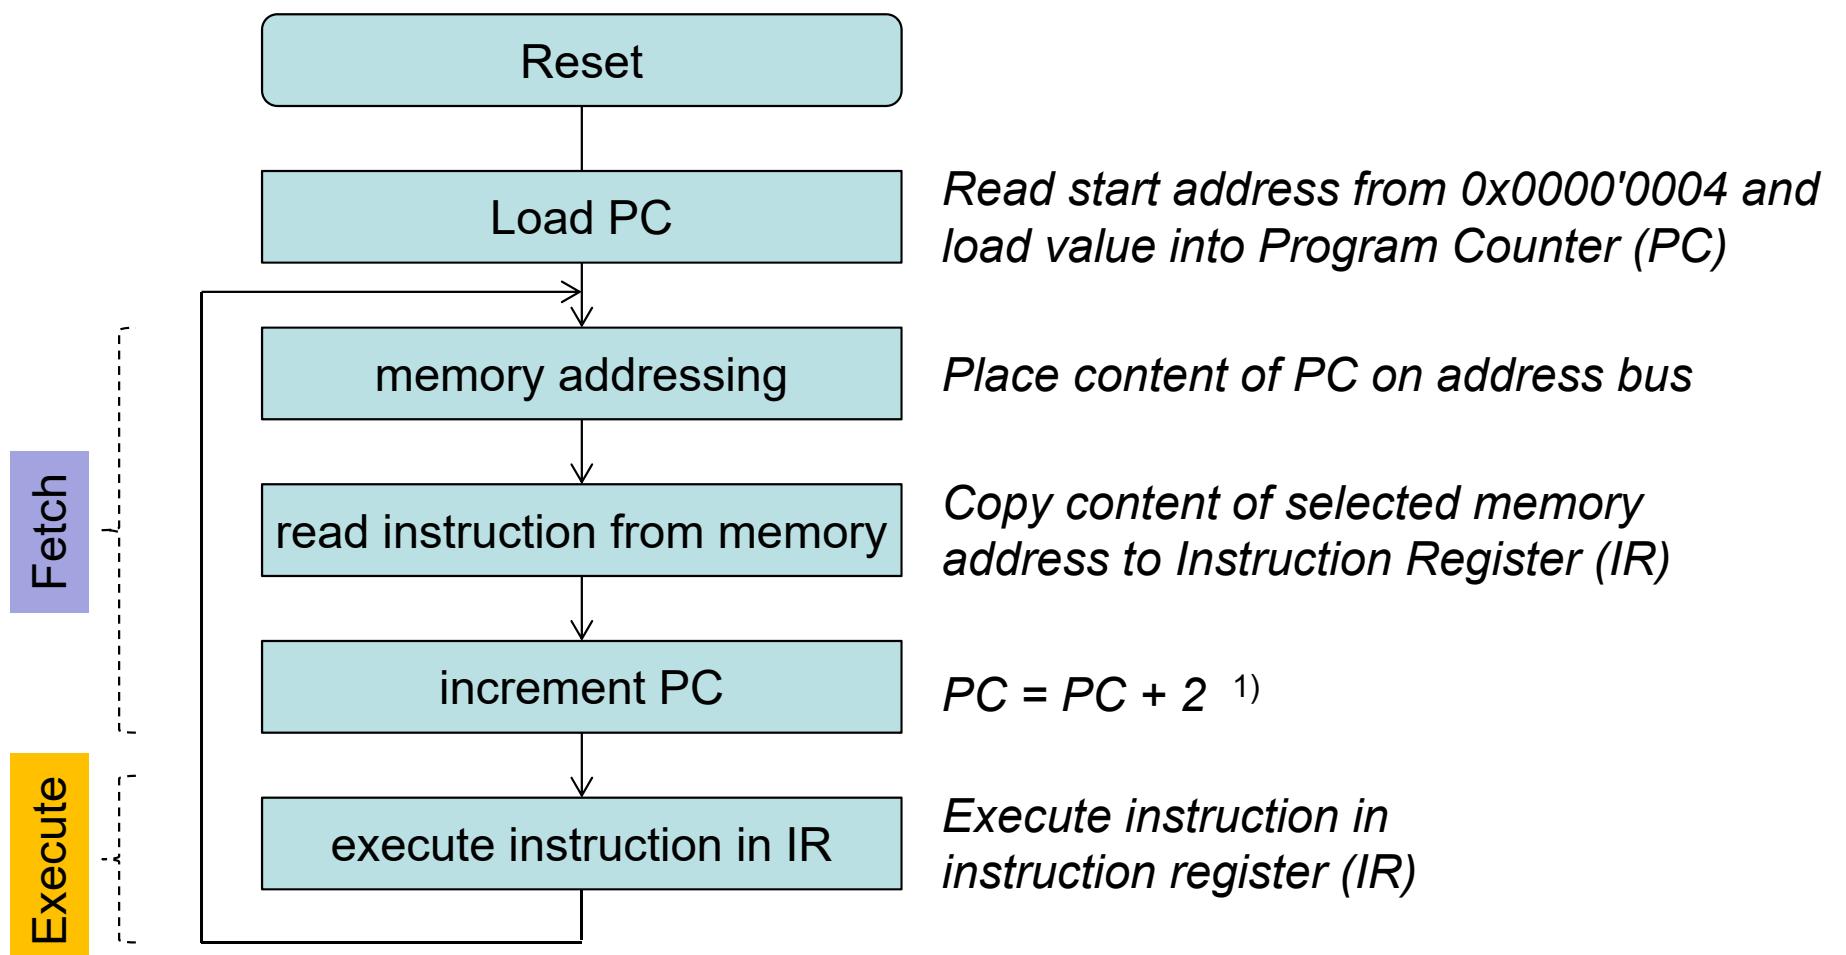
■ Load program into memory

- Example assumes code area in memory at address 0x0800'0000¹⁾



Program Execution

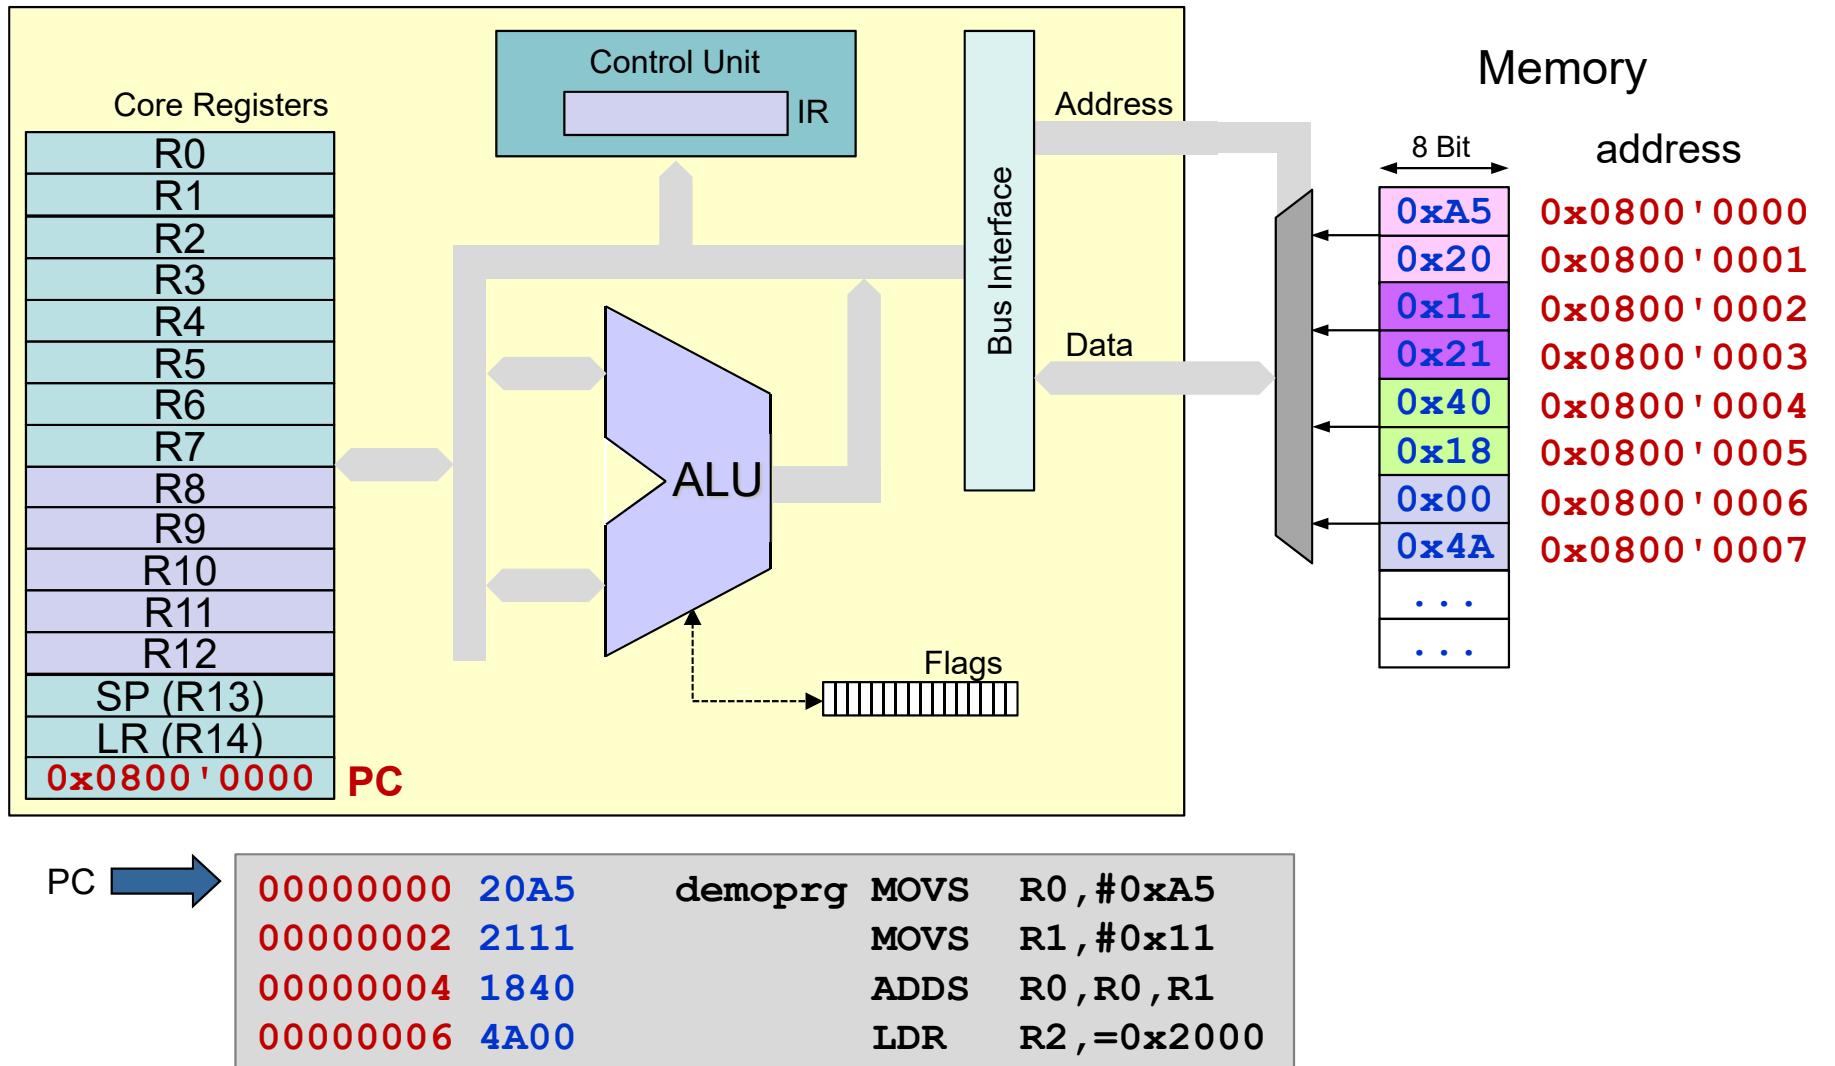
■ Sequence



1) '2' is used here for simplicity. In fact the PC is incremented by the length of the executed instruction.

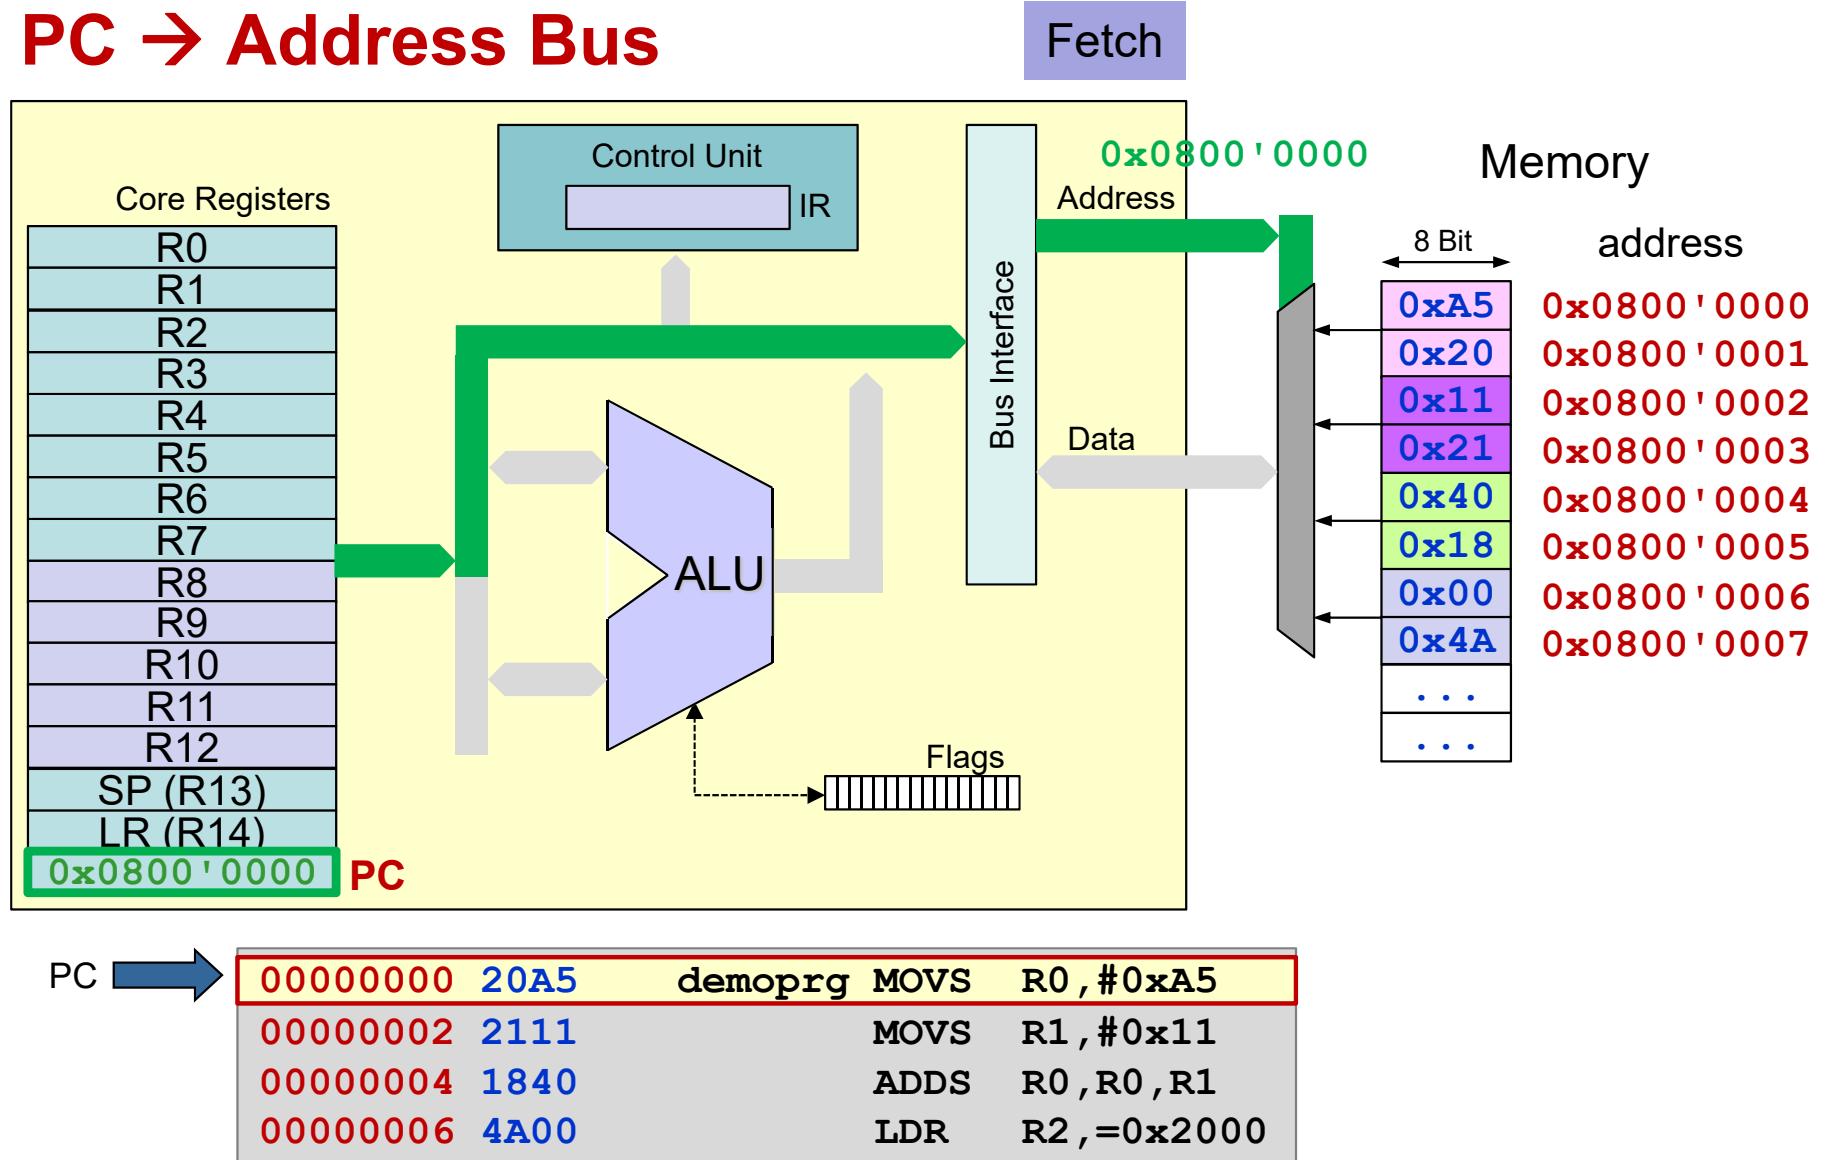
Program Execution

- "Reset": Read $0x0000'0004 \rightarrow PC$



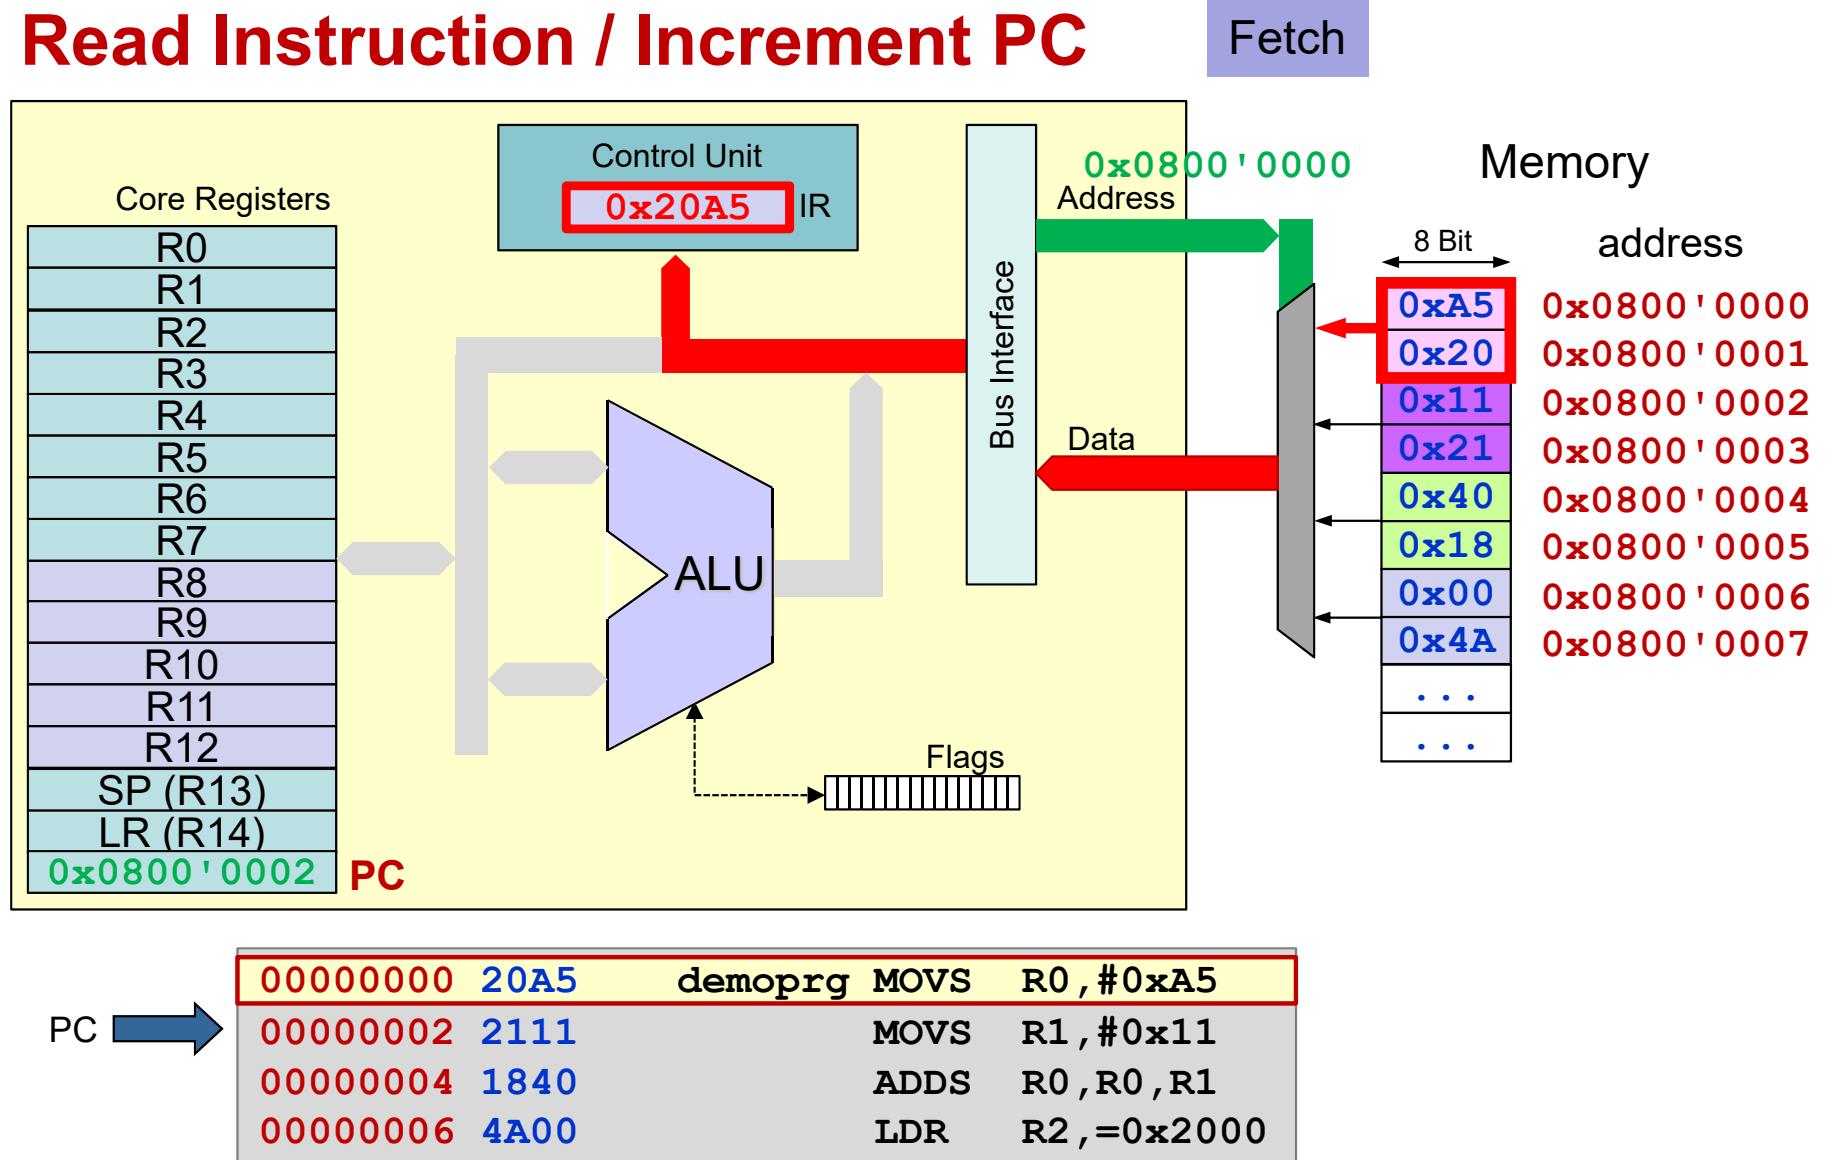
Program Execution

■ PC → Address Bus



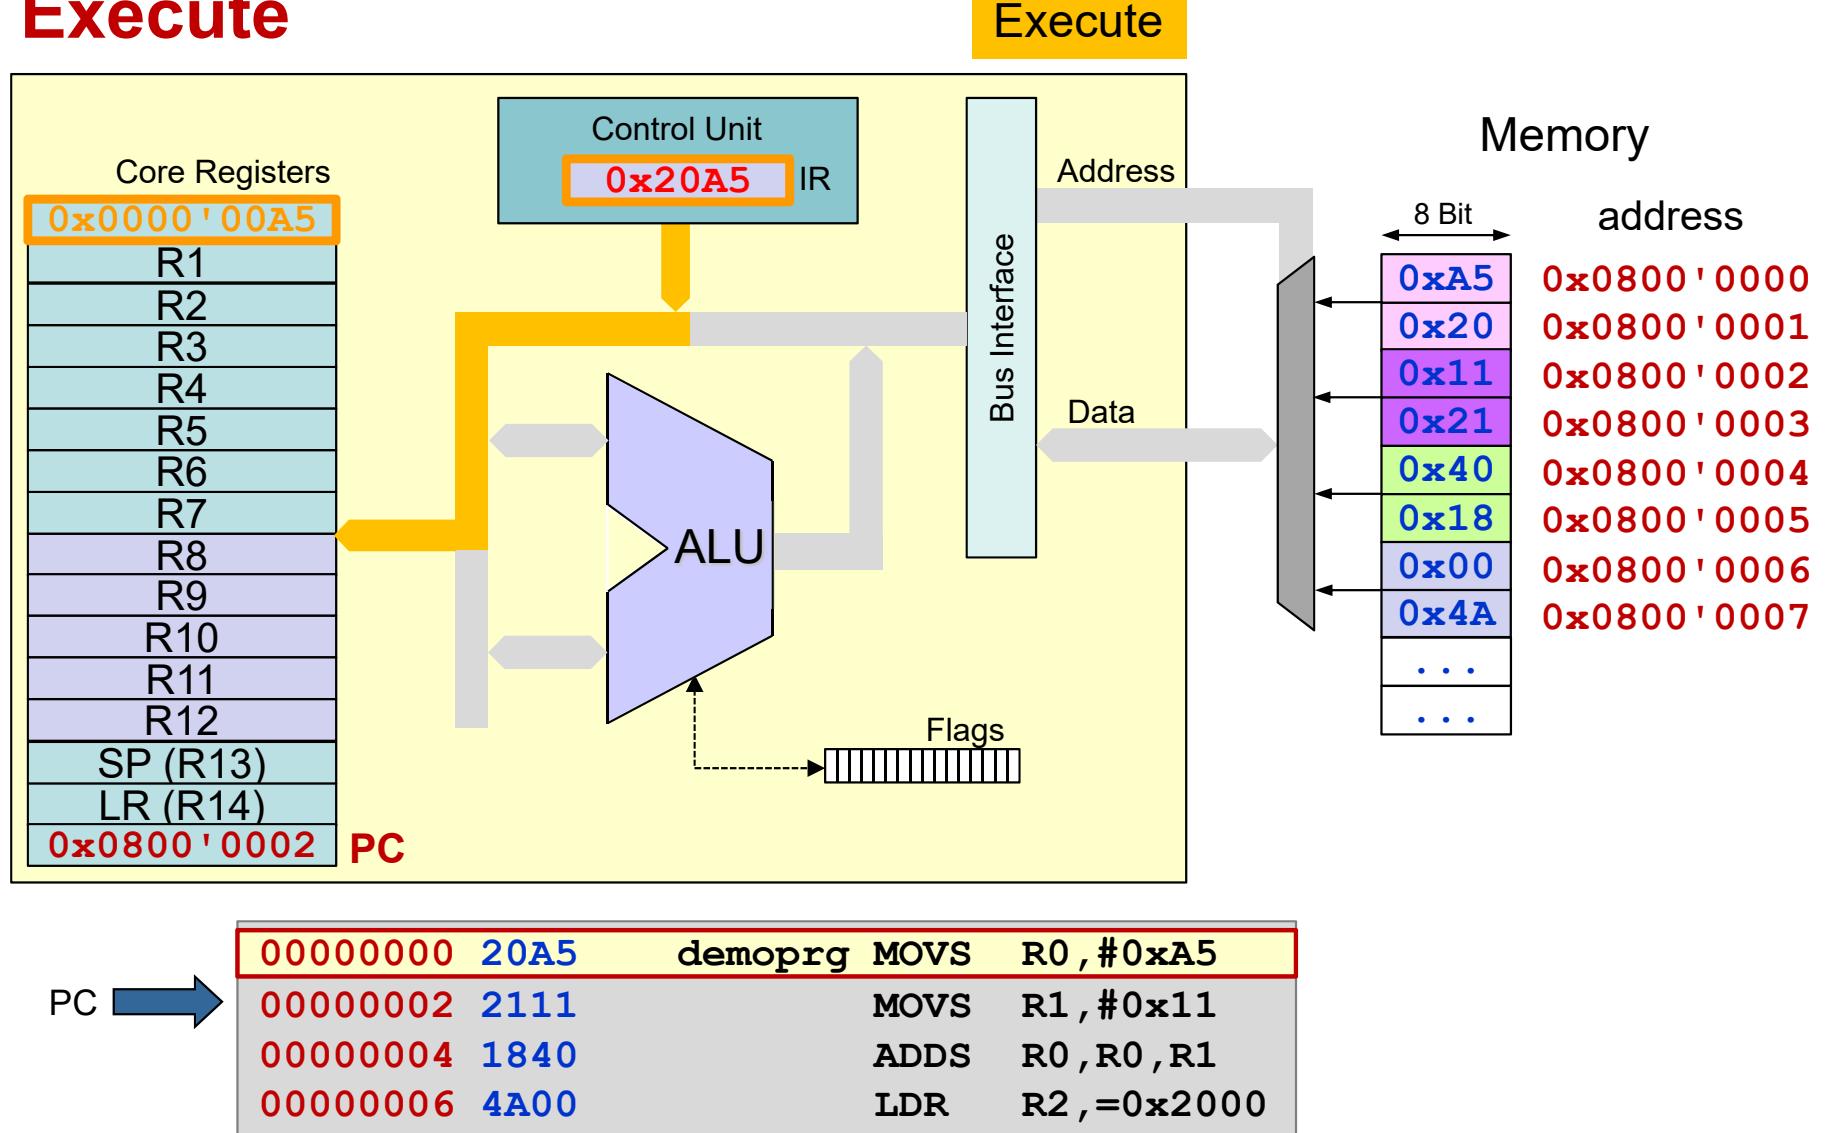
Program Execution

■ Read Instruction / Increment PC



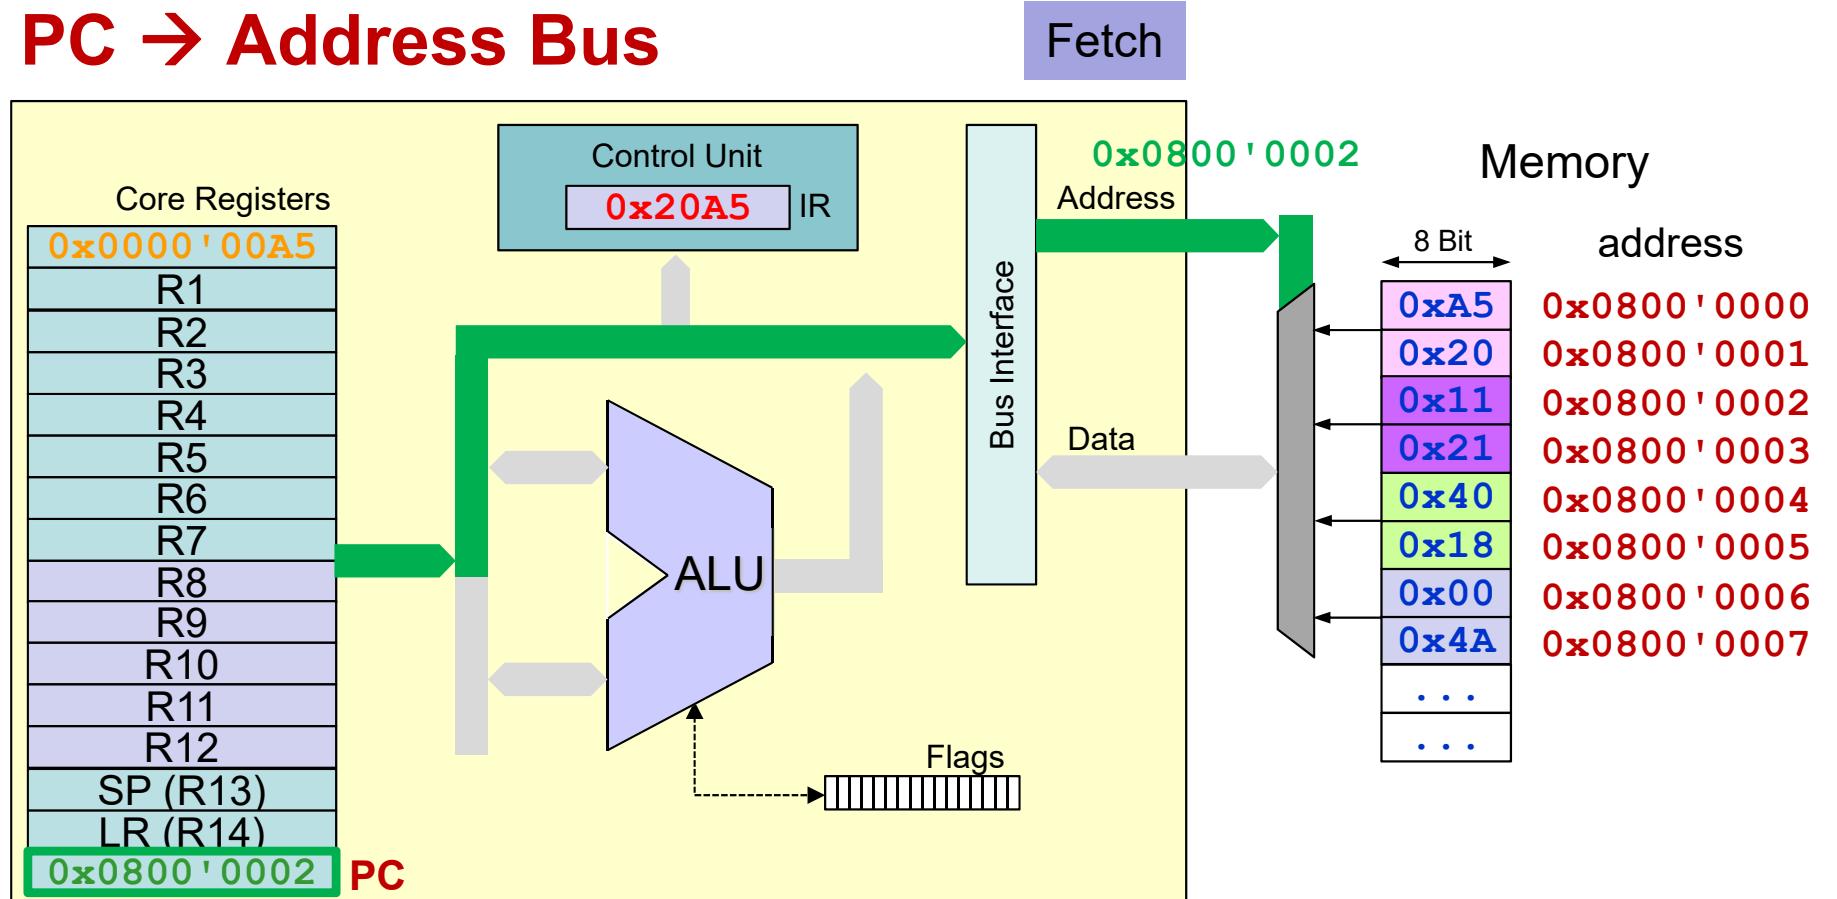
Program Execution

Execute



Program Execution

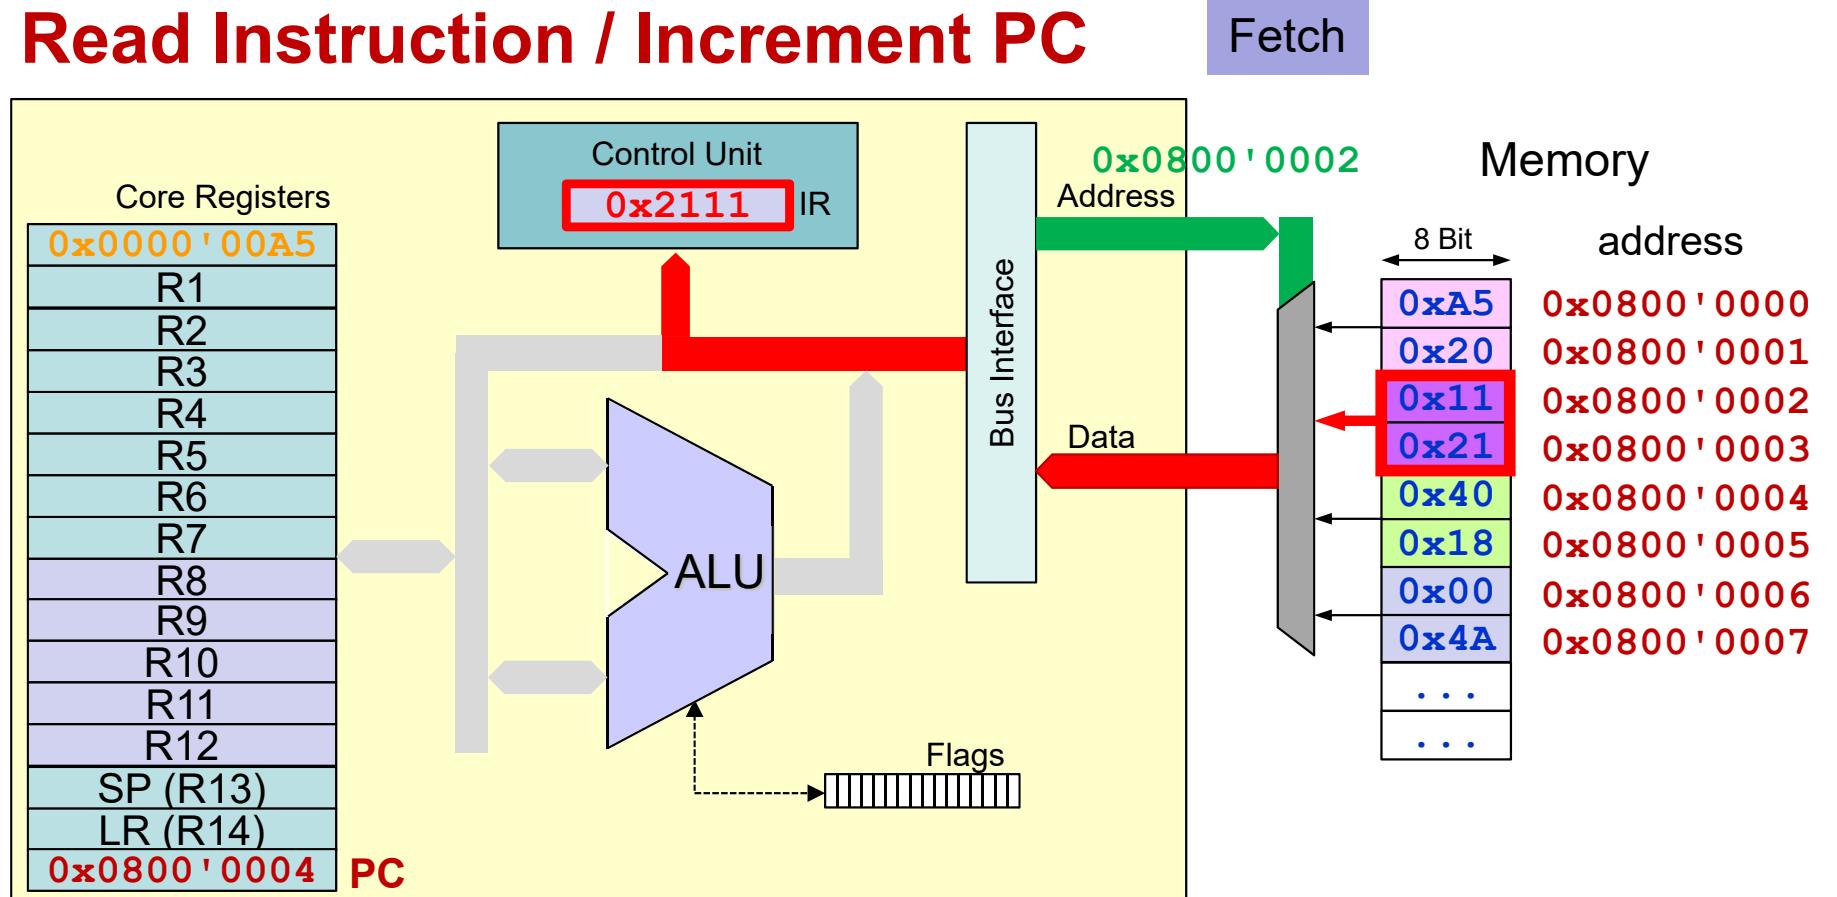
■ PC → Address Bus



PC →	00000000 20A5	demoprg	MOVS R0 , #0xA5
	00000002 2111		MOVS R1 , #0x11
	00000004 1840		ADDS R0 , R0 , R1
	00000006 4A00		LDR R2 , =0x2000

Program Execution

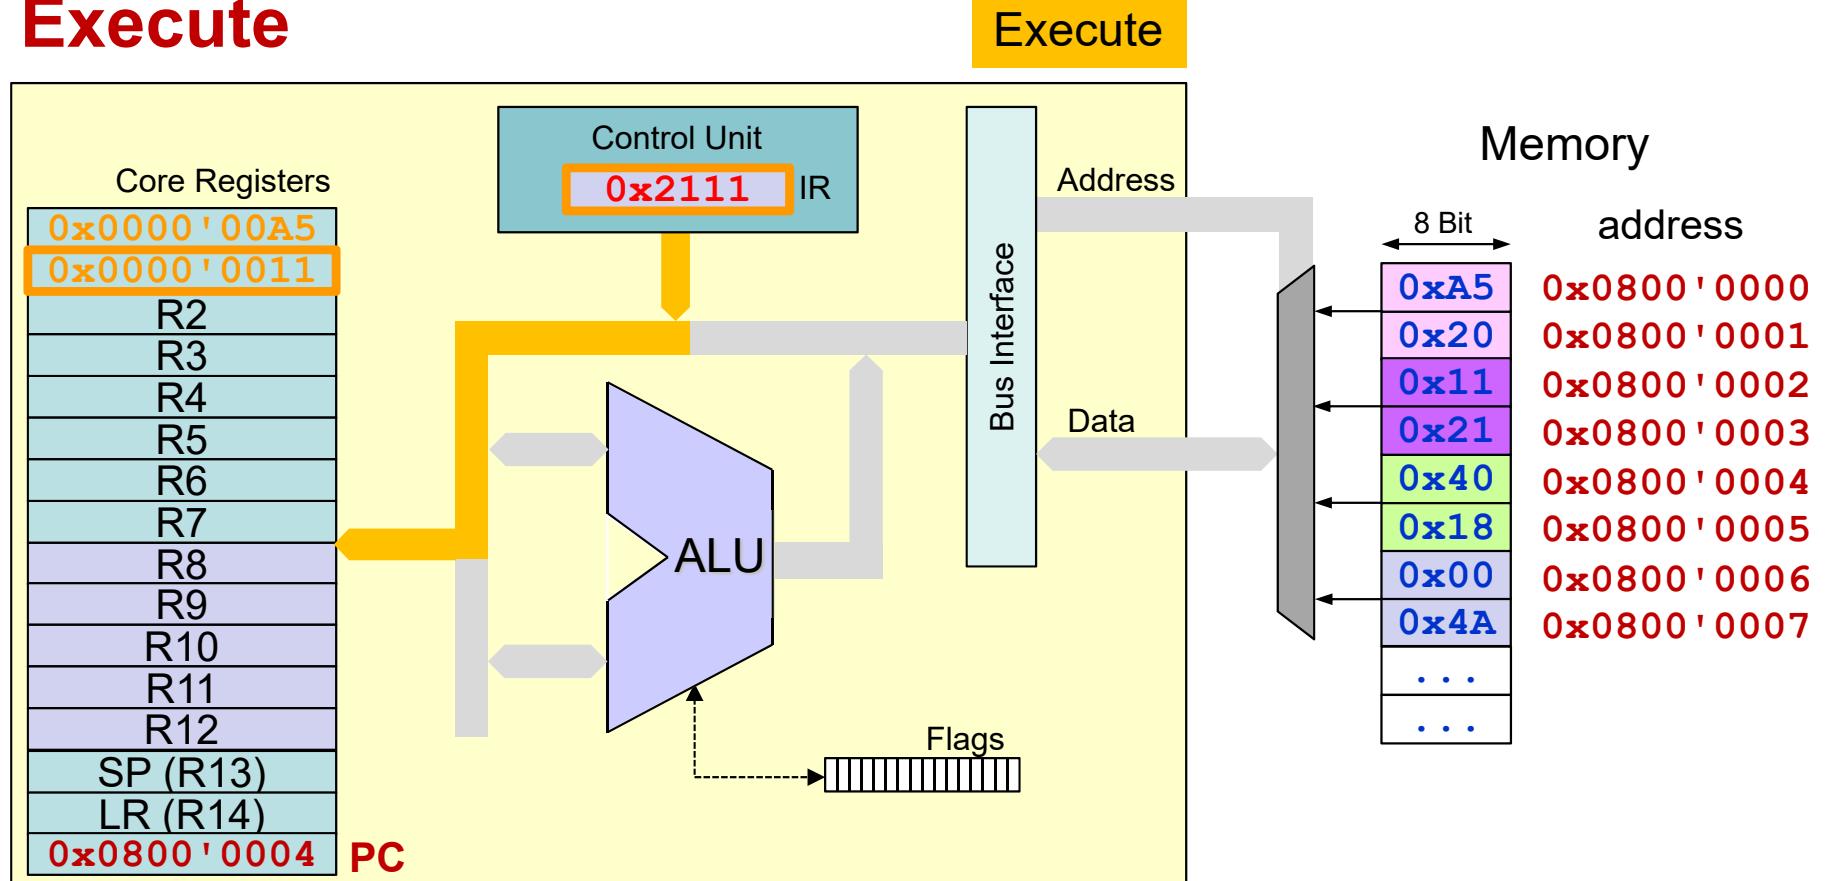
■ Read Instruction / Increment PC



00000000	20A5	demoprg	MOVS	R0 , #0xA5
00000002	2111		MOVS	R1 , #0x11
00000004	1840		ADDS	R0 , R0 , R1
00000006	4A00		LDR	R2 , =0x2000

Program Execution

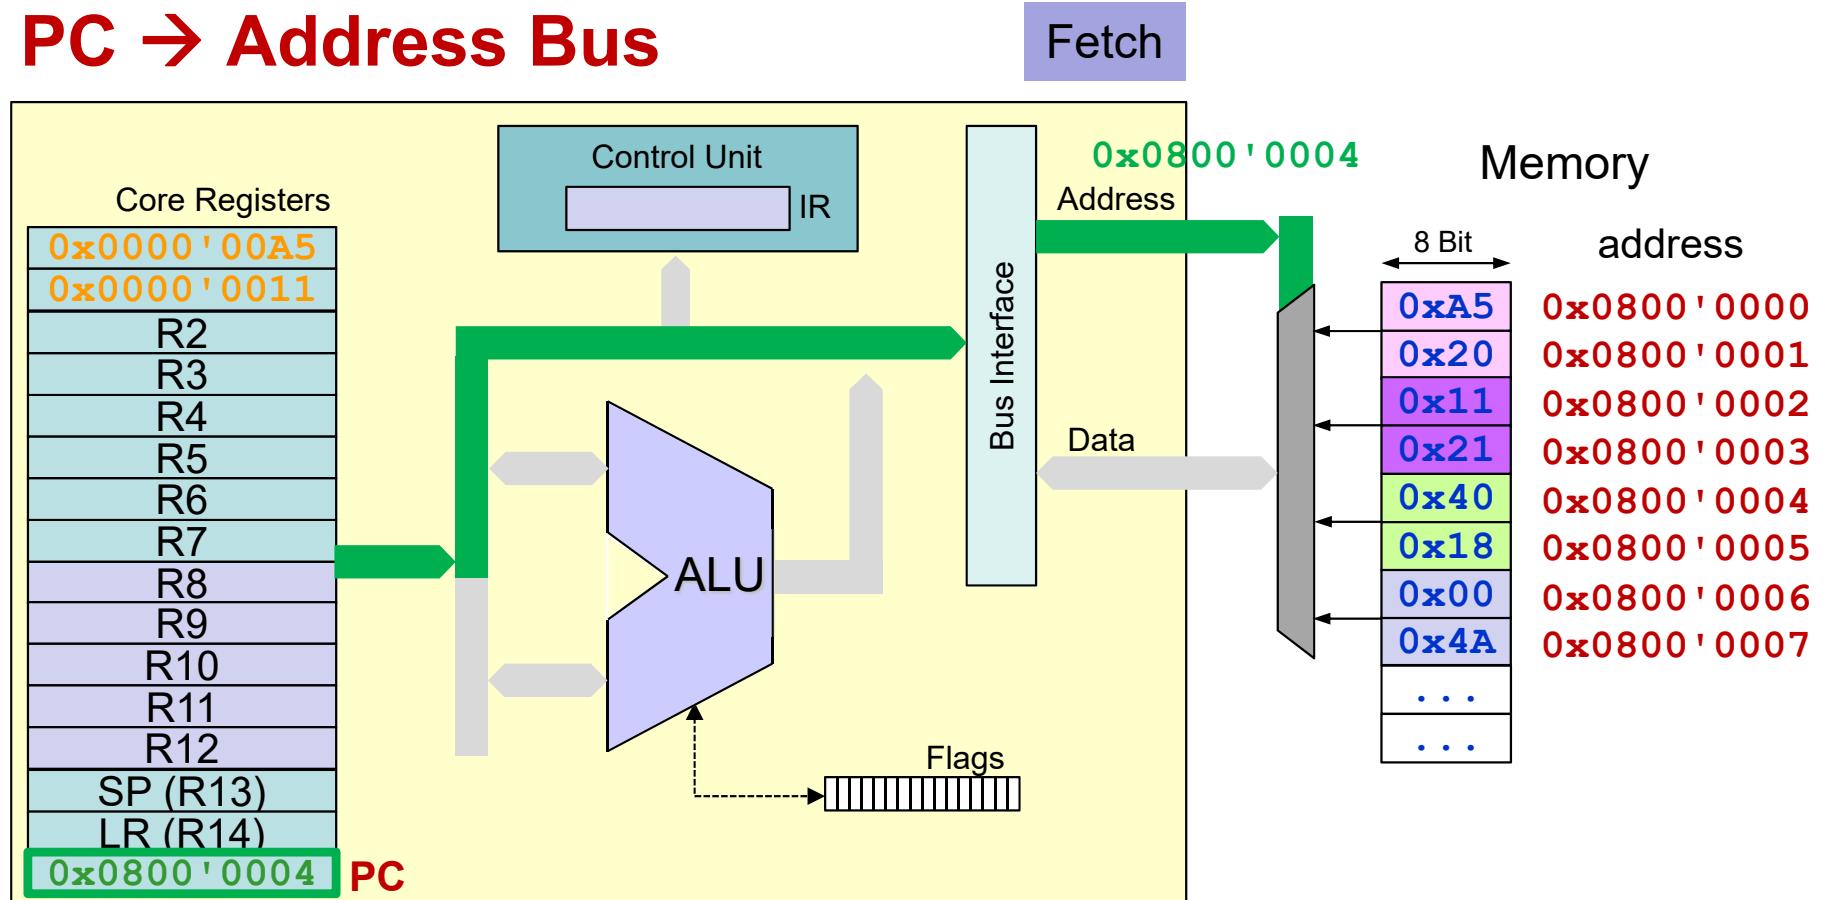
Execute



00000000	20A5	demoprg	MOVS	R0 , #0xA5
00000002	2111		MOVS	R1 , #0x11
00000004	1840		ADDS	R0 , R0 , R1
00000006	4A00		LDR	R2 , =0x2000

Program Execution

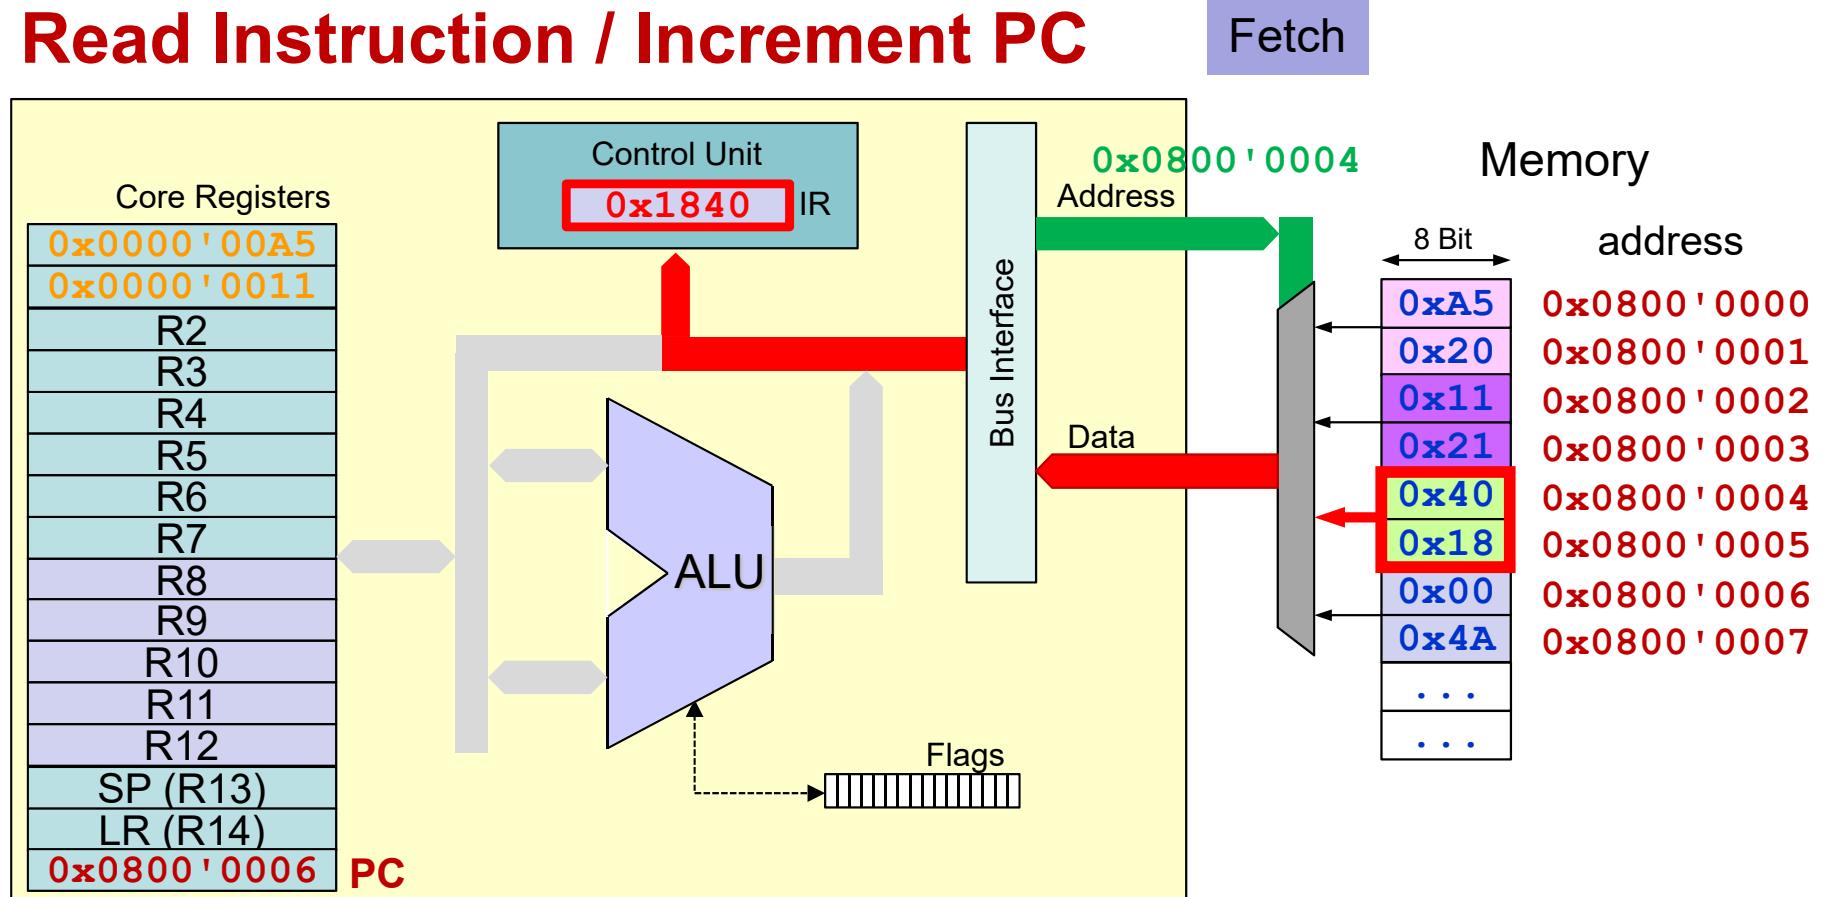
■ PC → Address Bus



00000000	20A5	demoprg	MOVS	R0 , #0xA5
00000002	2111		MOVS	R1 , #0x11
00000004	1840		ADDS	R0 , R0 , R1
00000006	4A00		LDR	R2 , =0x2000

Program Execution

■ Read Instruction / Increment PC

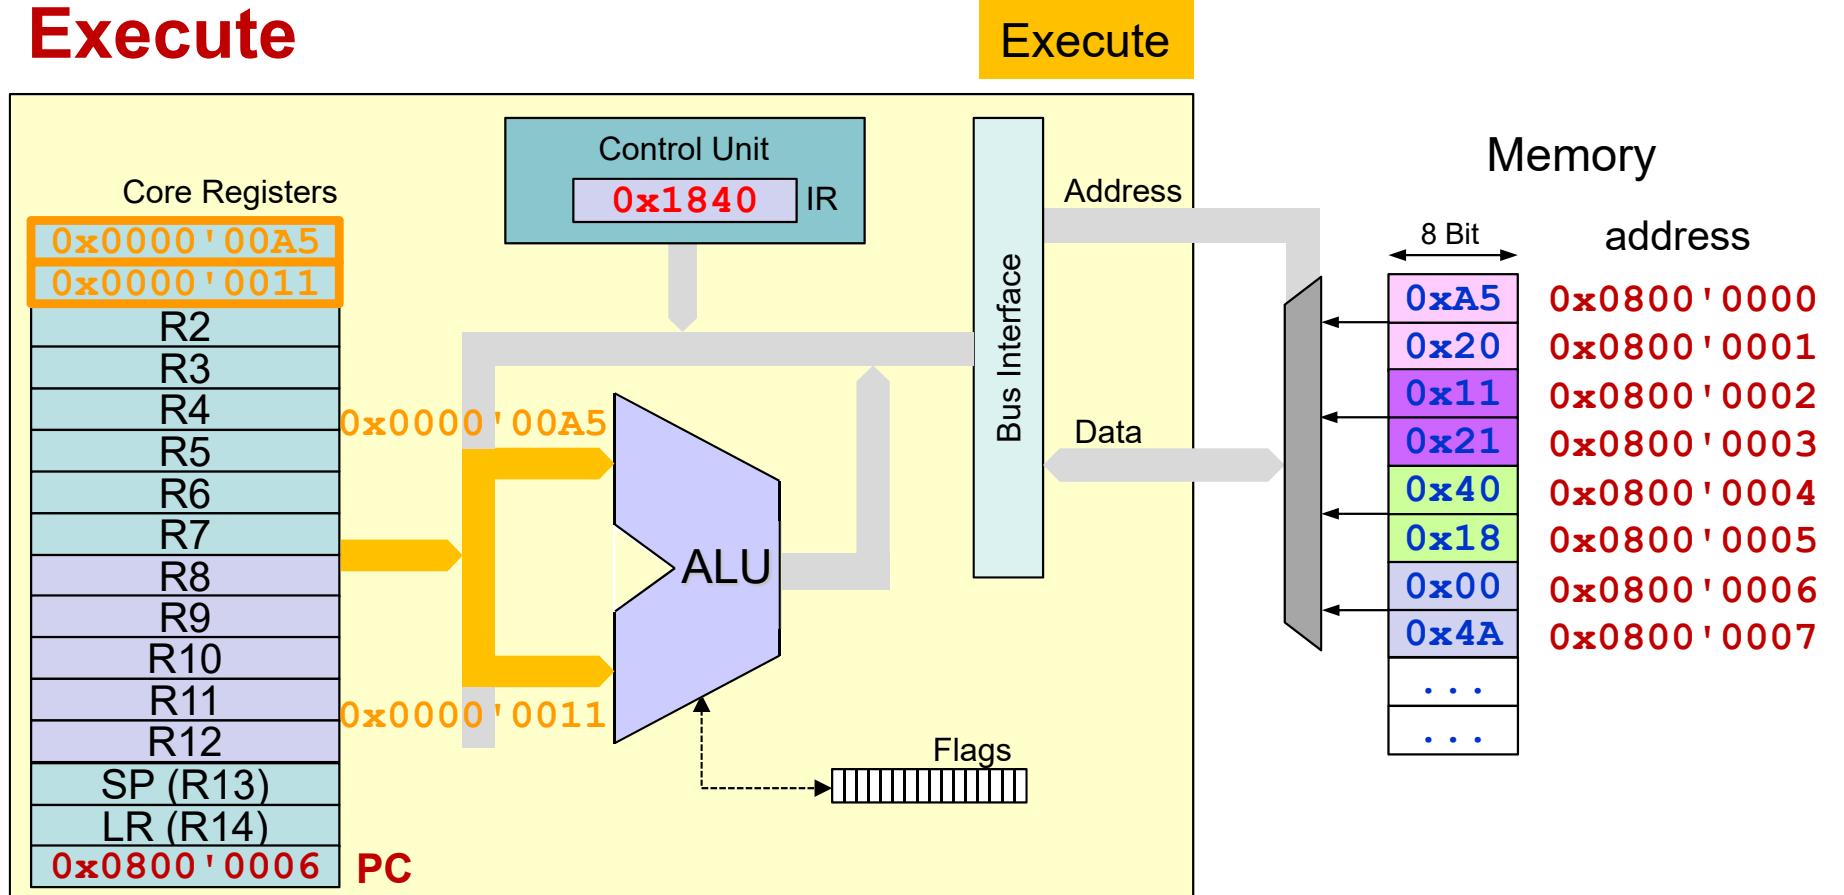


00000000	20A5	demoprg	MOVS	R0 ,#0xA5
00000002	2111		MOVS	R1 ,#0x11
00000004	1840		ADDS	R0 ,R0 ,R1
00000006	4A00		LDR	R2 ,=0x2000

PC →

Program Execution

Execute

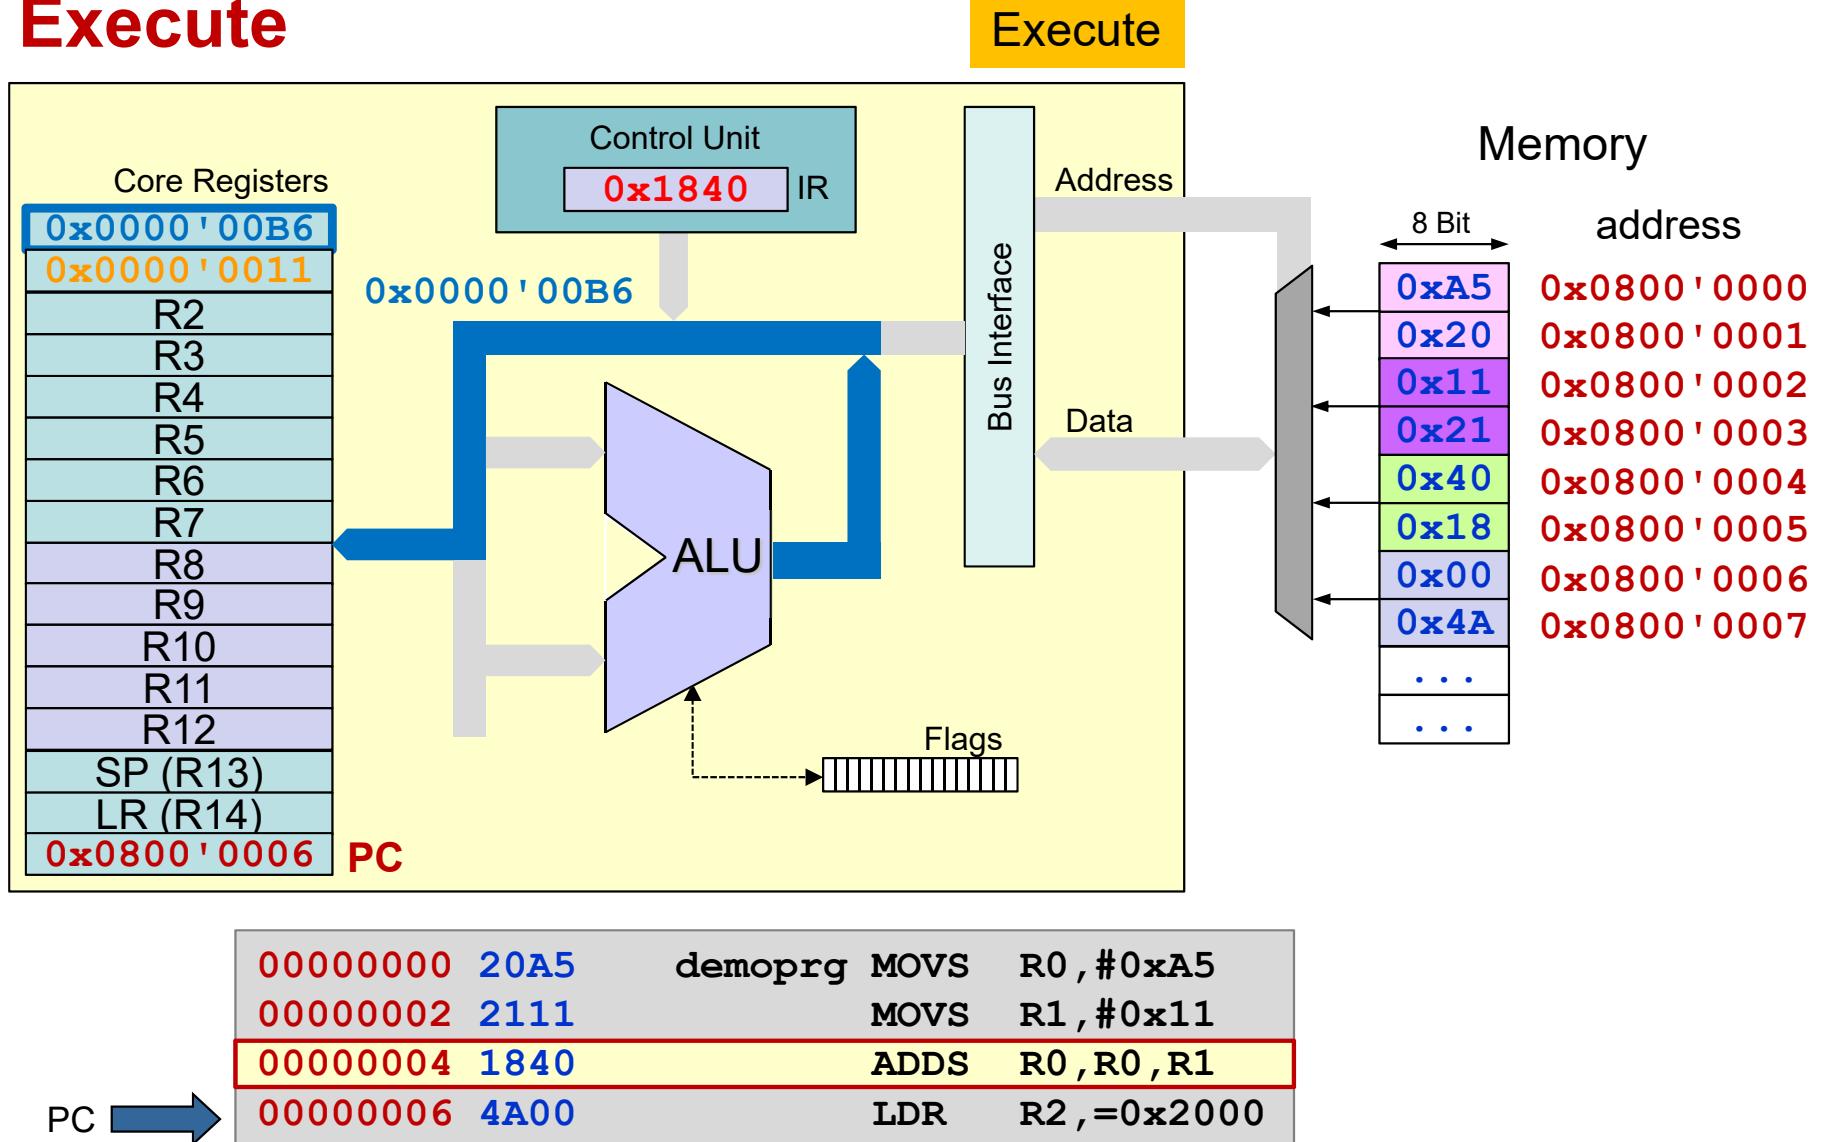


00000000	20A5	demoprg	MOVS R0 , #0xA5
00000002	2111		MOVS R1 , #0x11
00000004	1840		ADDS R0 , R0 , R1
00000006	4A00		LDR R2 , =0x2000

PC →

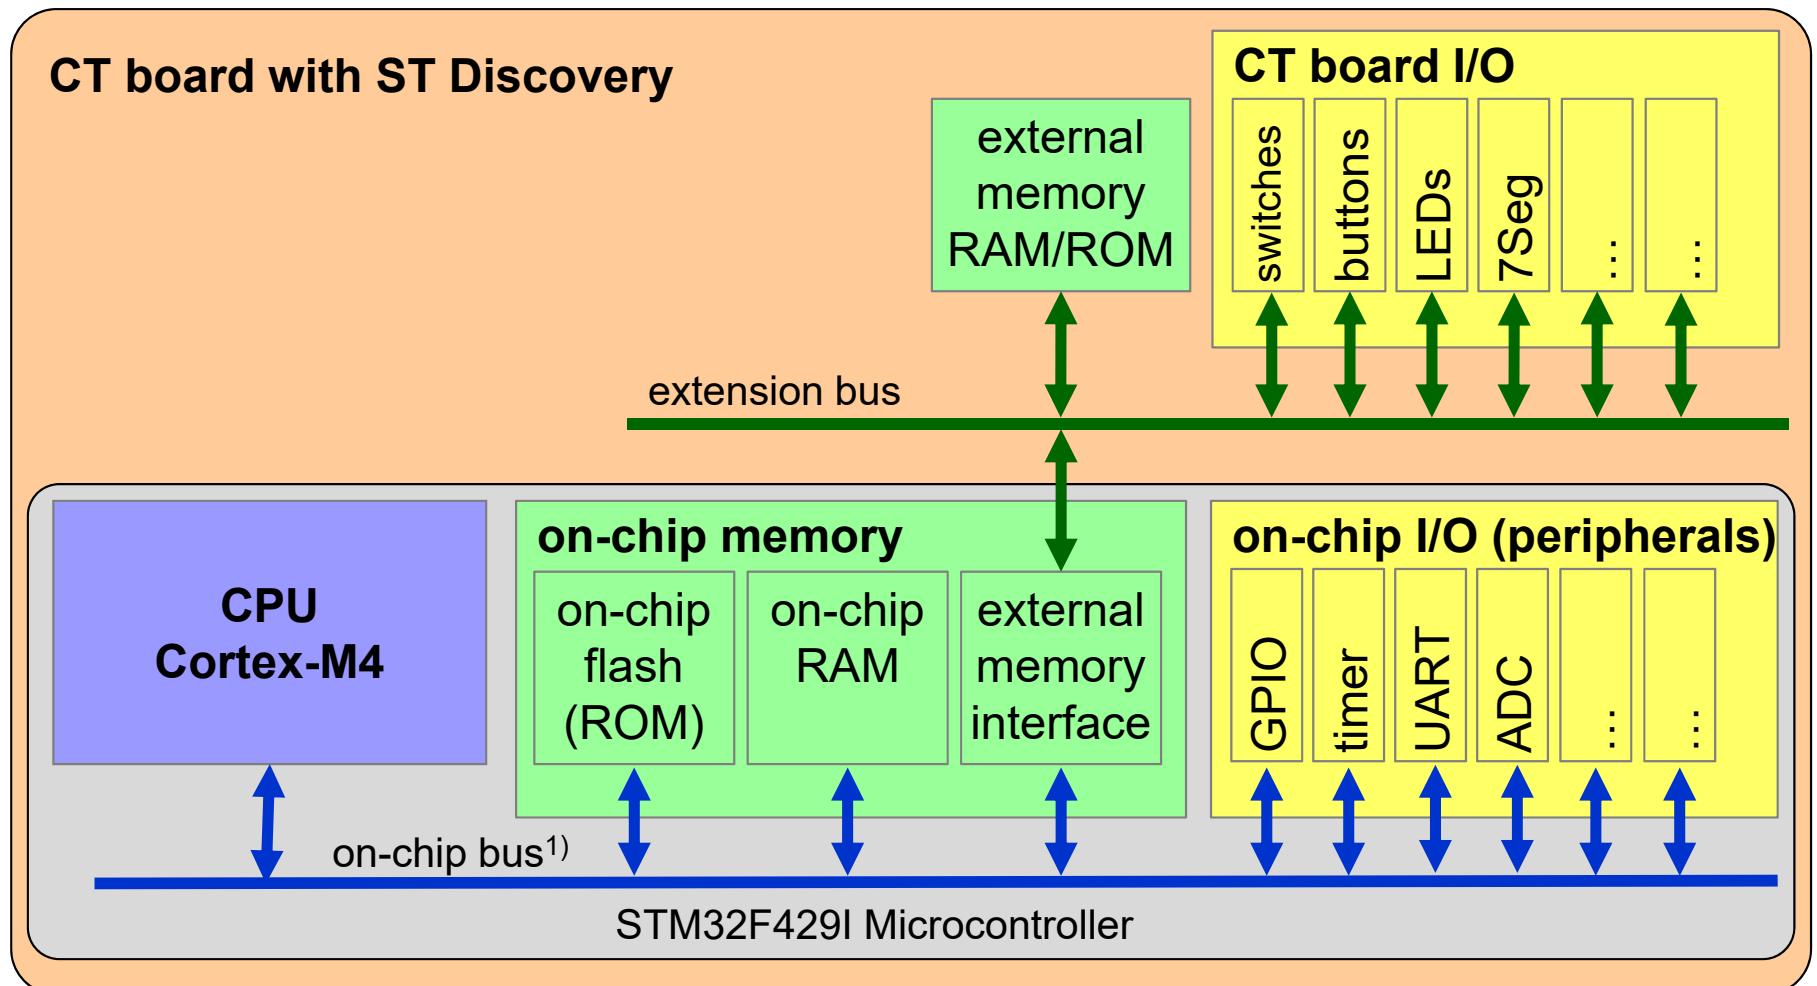
Program Execution

Execute



Memory Map

■ Hardware View



¹⁾ The implementation partitions the on-chip bus into several busses called AHBx and APBx

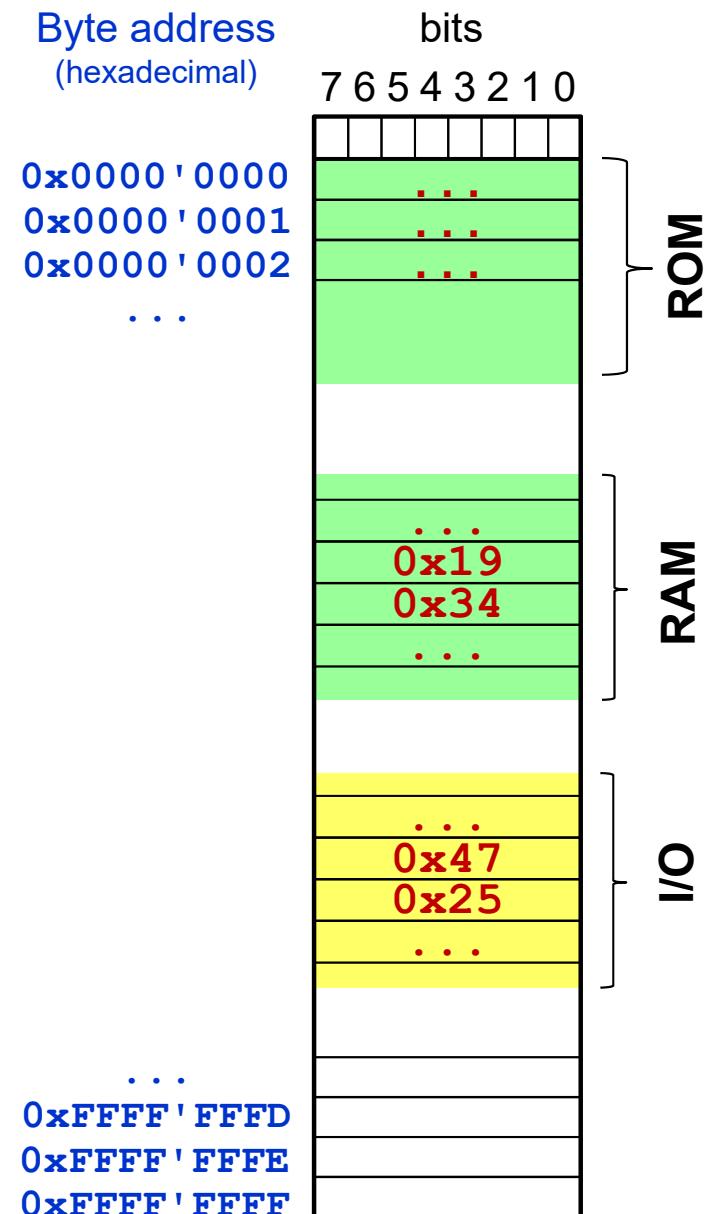
Memory Map

■ Memory Layout

- System Address Map
- Graphical layout of main memory
- Linear array of bytes
- What is located where (at which address) in memory?
 - Location of RAM (readable and writable)
 - Location of ROM (only readable)
 - Location of I/O registers

Memory maps in CT1/CT2 will be drawn with lowest address at the top and highest address at the bottom. This simplifies work with assembly listing and tables.

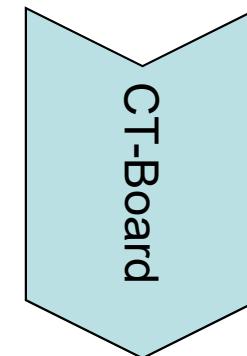
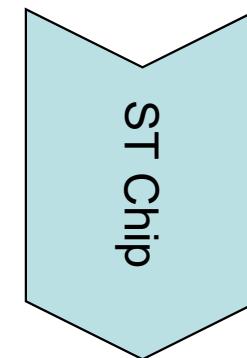
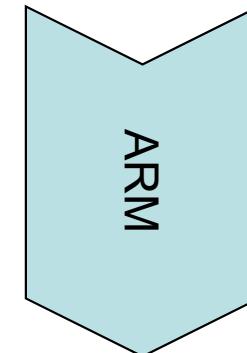
Caution: ARM and ST documentation are the other way round.
Lowest address at the bottom and highest address at top.



Memory Map

■ Address Allocation

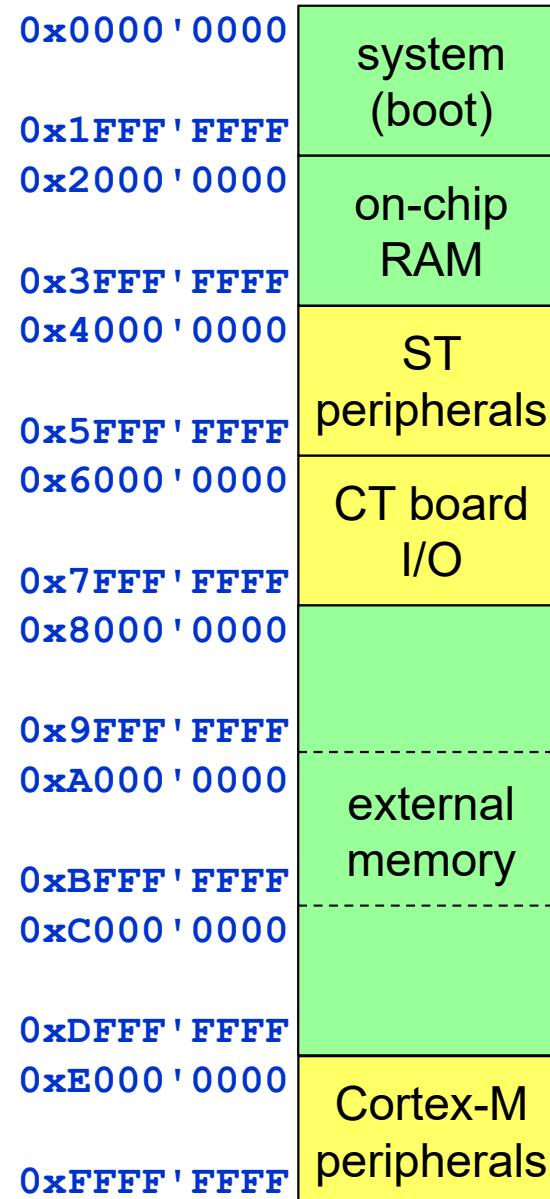
- ARM policies
 - Cortex-M specific
 - guide lines for chip manufacturer
- ST design decisions
 - chip specific
 - number and size of on-chip RAMs
 - size of flash
 - control register for peripherals
- CT board design decisions
 - board specific
 - LEDs, switches, etc



Memory Map

■ CT-Board

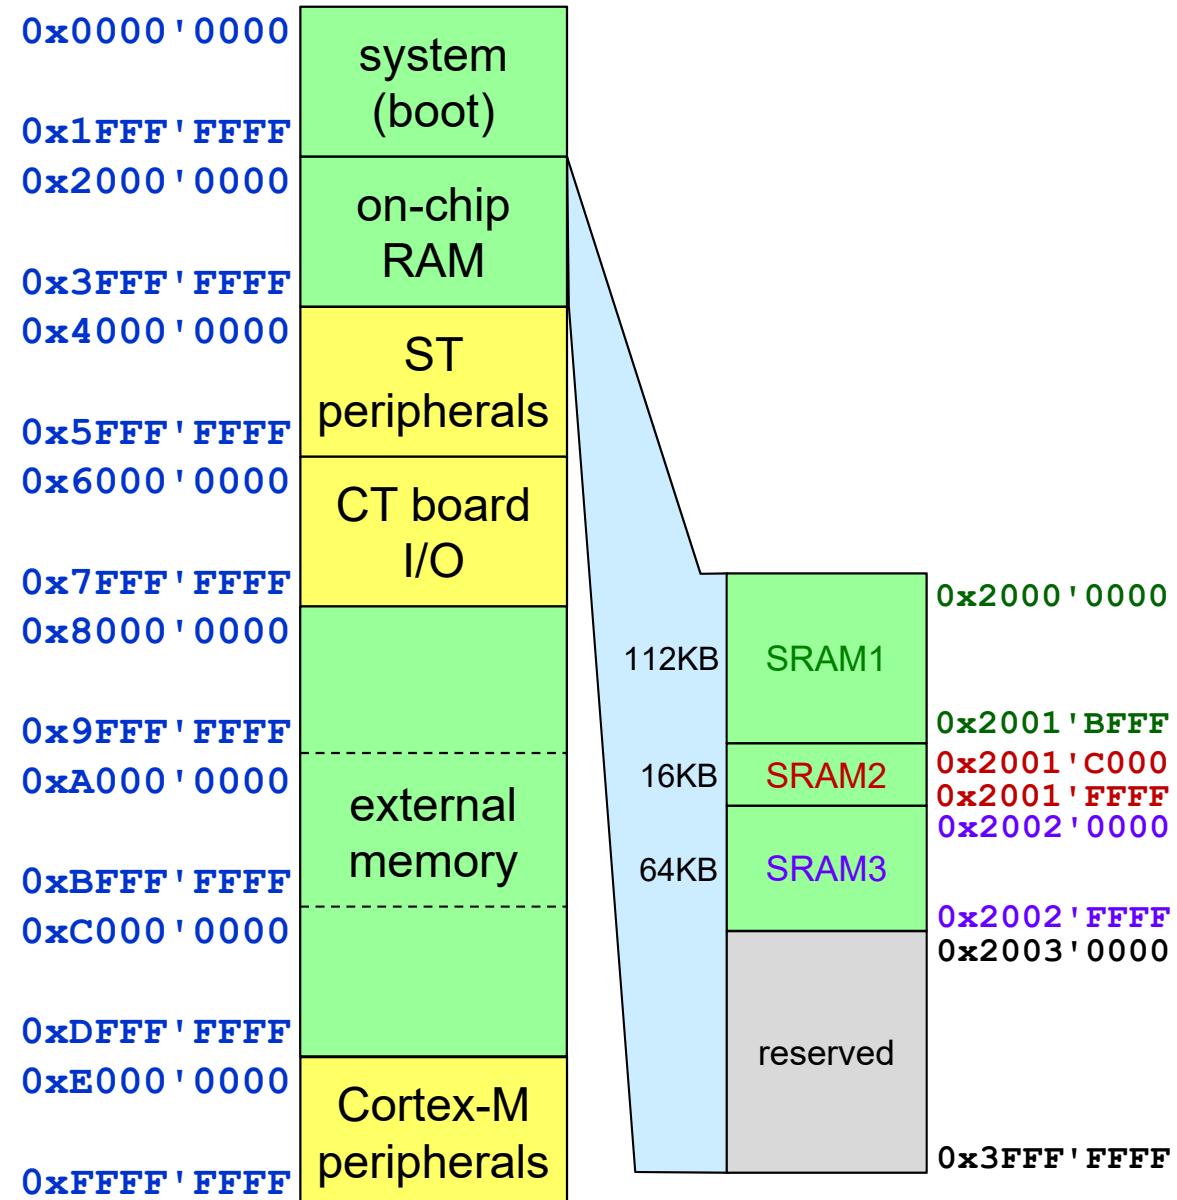
- Address space
4 GByte = 2^{32}
- From **0x0000'0000** to **0xFFFF'FFFF**
- Partitioned into 8 blocks of 512-MByte each



Memory Map

■ ST chip specific

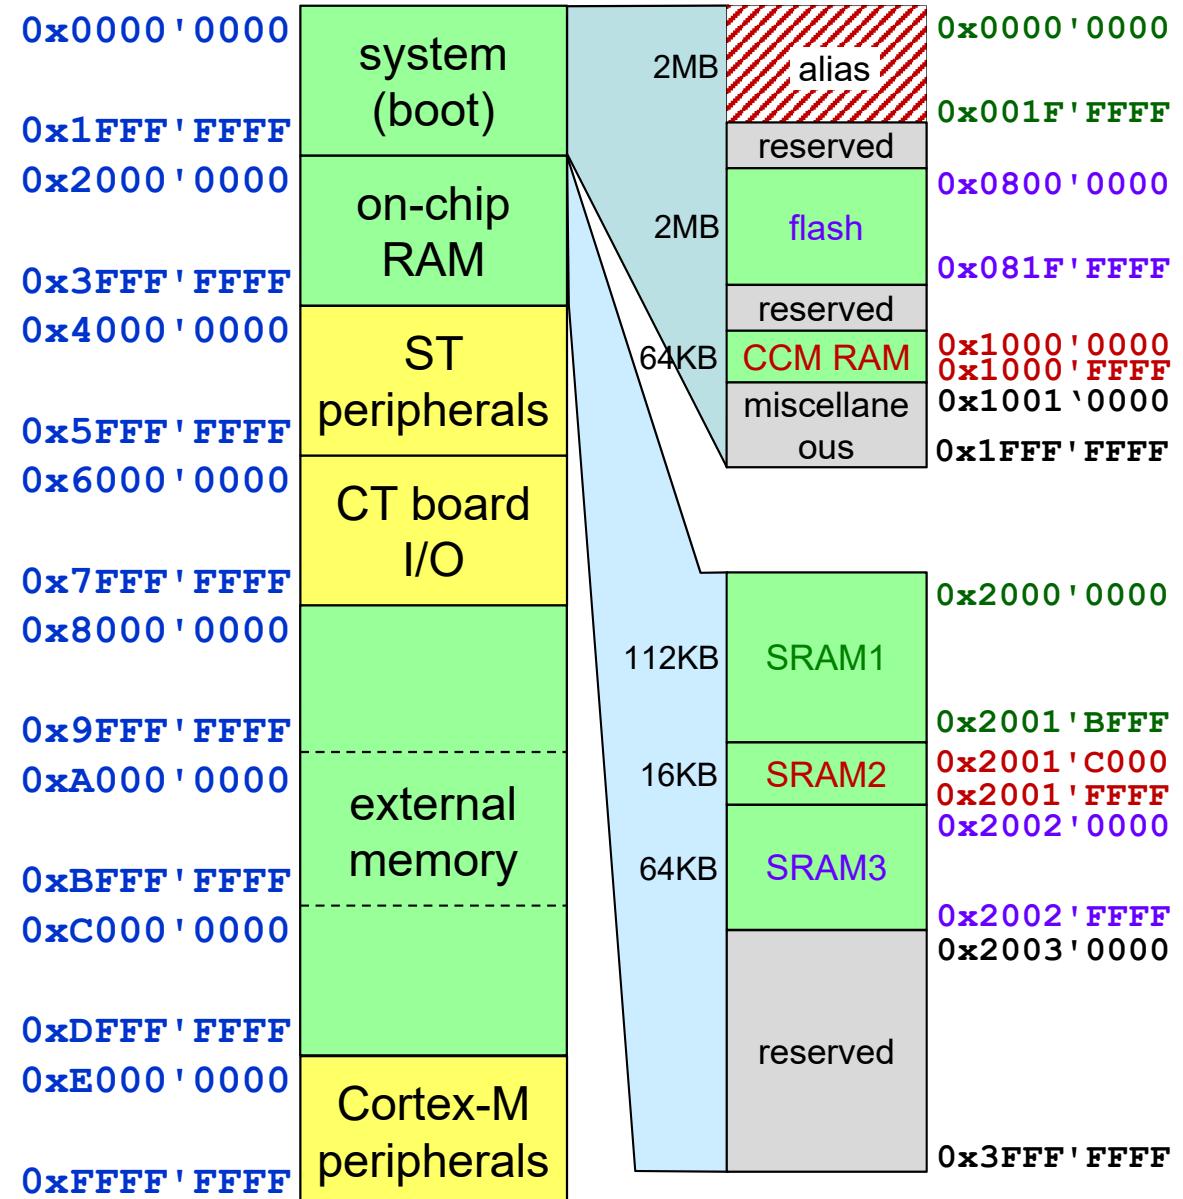
- SRAM1
 - 112 KByte
- SRAM2
 - 16 KByte
- SRAM3
 - 64 KByte



Memory Map

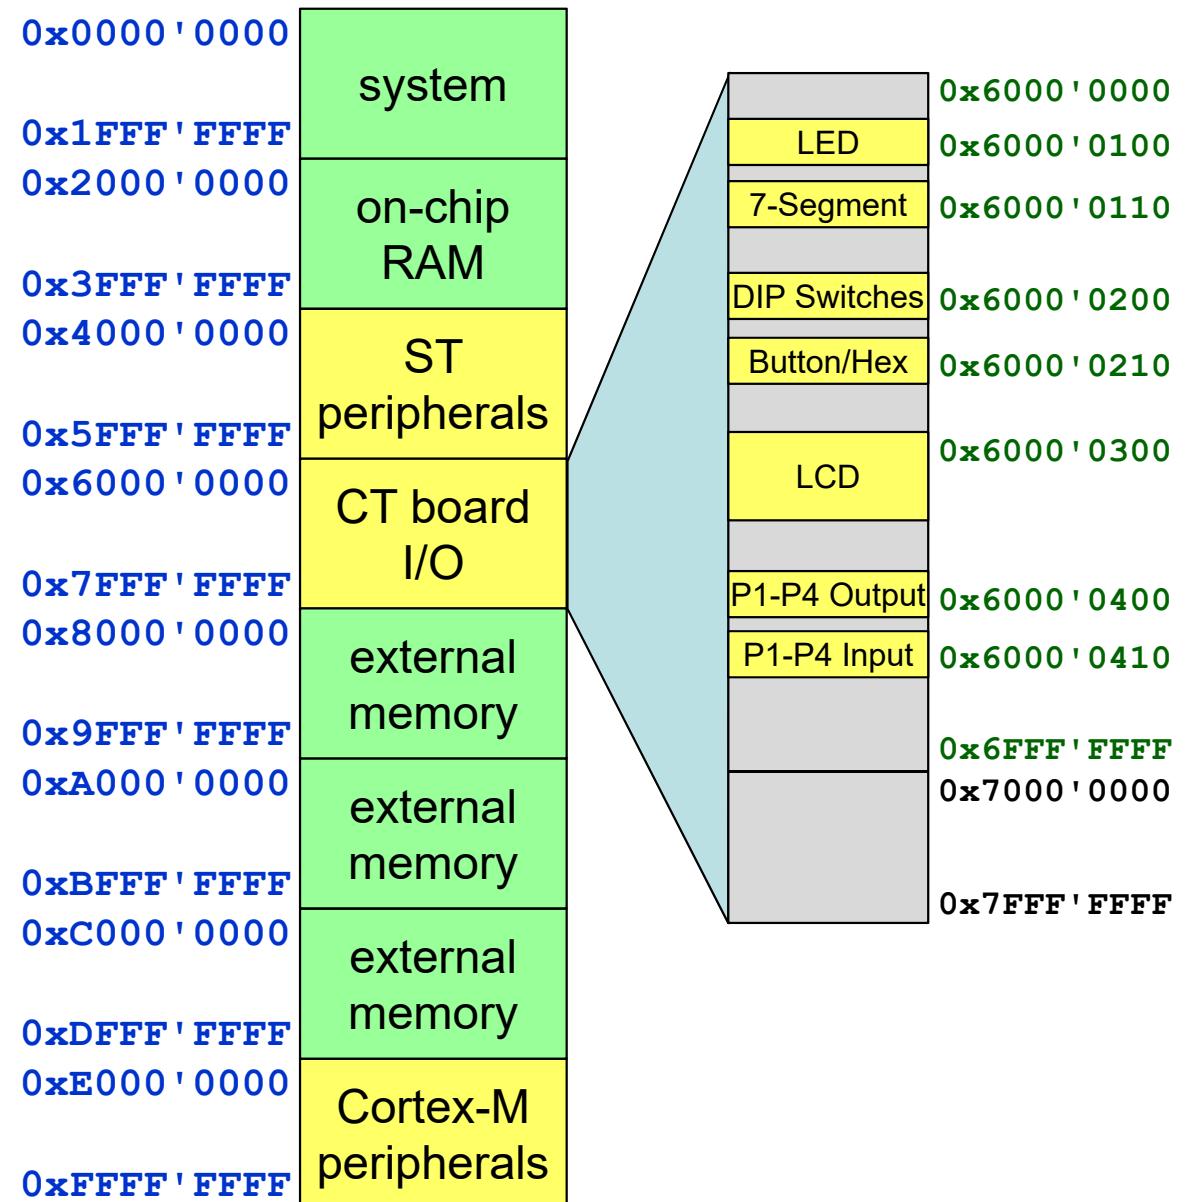
■ ST chip specific

- Flash
 - non-volatile memory
- CCM RAM
 - core coupled memory
 - very fast RAM
- Alias
 - user configurable mirror
 - physical memory can appear at two locations



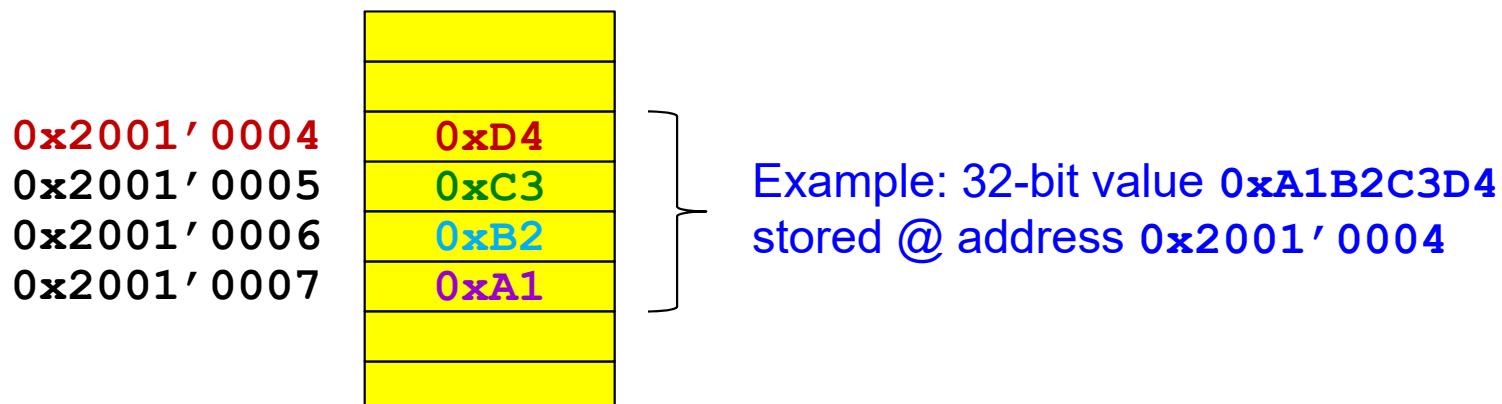
Memory Map

■ CT-Board I/O



■ Multi-byte Integer Types

- Usually memory is organized in bytes
 - one address per byte → reasons: space and history
- An integer type often requires several bytes
 - e.g. 4 byte addresses are required to store a 32-bit integer



Integer Types

■ Integer Types in C

- Sizes of integer types are platform dependent

Sizes in byte

8051			Cortex-Mx: Keil (ARM)			x86-64 (i7): gcc		
char	1		char	1		char	1	
short	2		short	2		short	2	
int	2		int	4		int	4	
long int	4		long int	4		long int	8	
			long long int	8		long long int	8	
char *	2		void *	4		void *	8	

Integer Types

■ ARM Cortex-M

C99 / specified width



C-Type – unsigned integers	Size	Term	inttypes.h / stdint.h
<code>unsigned char</code>	8 Bit	Byte	<code>uint8_t</code>
<code>unsigned short</code>	16 Bit	Half-word	<code>uint16_t</code>
<code>unsigned int</code>	32 Bit	Word	<code>uint32_t</code>
<code>unsigned long</code>	32 Bit	Word	<code>uint32_t</code>
<code>unsigned long long</code>	64 Bit	Double-word	<code>uint64_t</code>

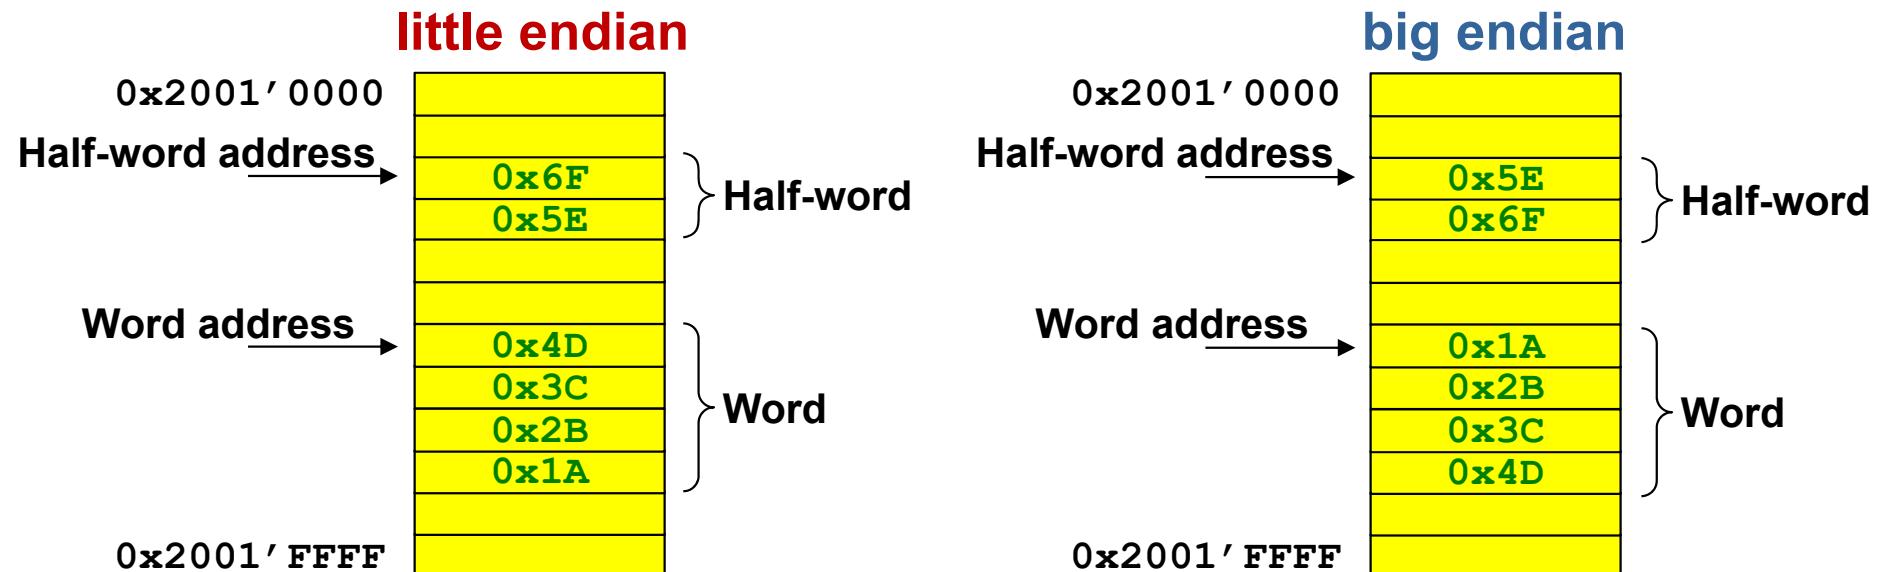
C-Type – signed integers	Size	Term	inttypes.h / stdint.h
<code>signed char</code>	8 Bit	Byte	<code>int8_t</code>
<code>short</code>	16 Bit	Half-word	<code>int16_t</code>
<code>int</code>	32 Bit	Word	<code>int32_t</code>
<code>long</code>	32 Bit	Word	<code>int32_t</code>
<code>long long</code>	64 Bit	Double-word	<code>int64_t</code>

Integer Types

■ How are groups of bytes arranged in memory?

- little endian → *least significant byte at lower address*
e.g. Intel x86, Altera Nios, ST ARM (STM32)
- big endian → *most significant byte at lower address*
e.g. Freescale (Motorola), PowerPC

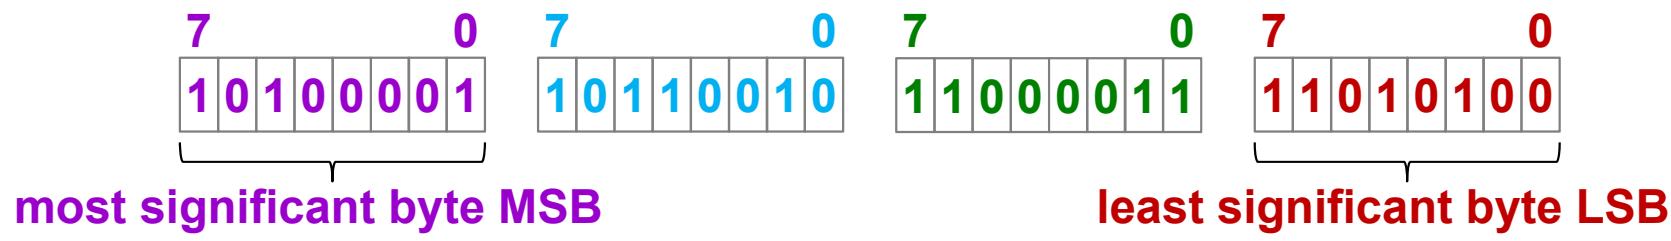
Examples: **0x1A2B'3C4D** for Word and **0x5E6F** for Half-word



Integer Types

■ Example

- Store Word **0xA1B2' C3D4** at Address **0x2001' 0004**



little endian

0x2001' 0004
0x2001' 0005
0x2001' 0006
0x2001' 0007



big endian

0x2001' 0004
0x2001' 0005
0x2001' 0006
0x2001' 0007



■ Alignment

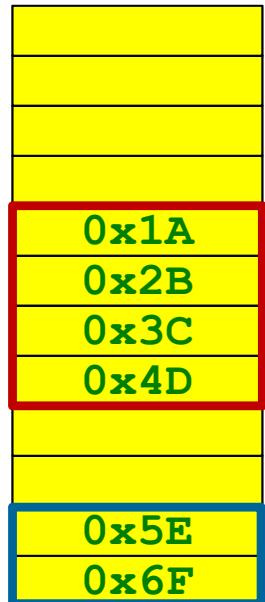
- Half-word aligned Variables aligned on even addresses
- Word aligned Variables aligned on addresses that are divisible by four

Word Aligned

0x2001' 0000

0x2001' 0004

0x2001' 0008



Half-word Aligned

0x2001' 0000

0x2001' 0002

0x2001' 0004

0x2001' 0006

0x2001' 0008

0x2001' 000A

Object File Sections

■ CODE

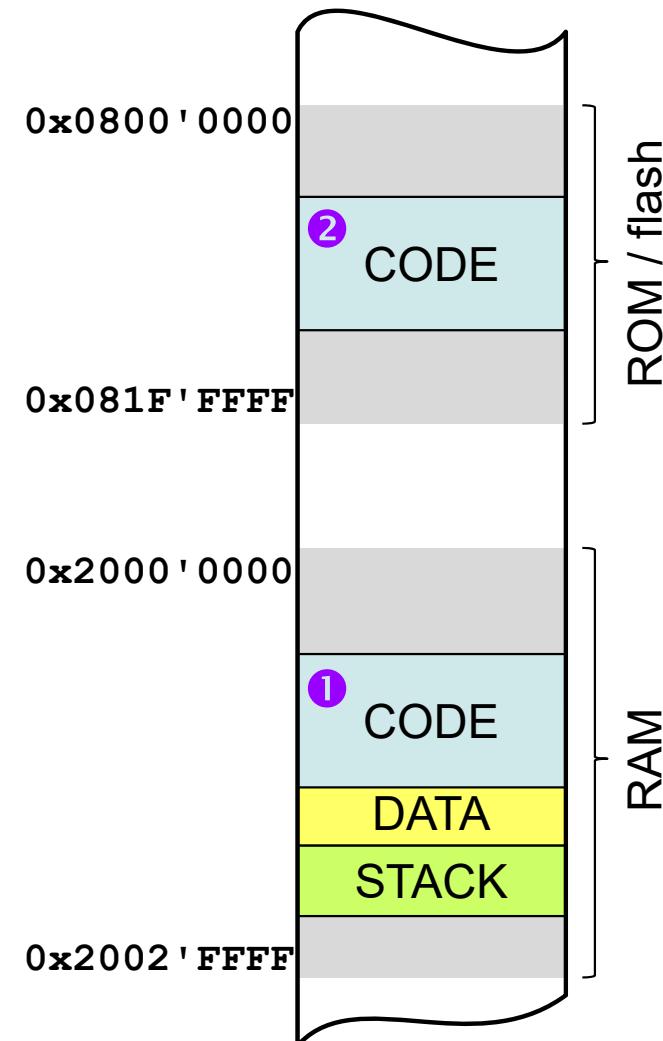
- Read-only → RAM or ROM
- Instructions (opcodes)
- Literals ¹⁾

■ DATA ²⁾

- Read-write → RAM
- Global variables
- *static* variables in C
- Heap in C → `malloc()`

■ STACK

- Read-write → RAM
- Function calls / parameter passing
- Local variables and local constants

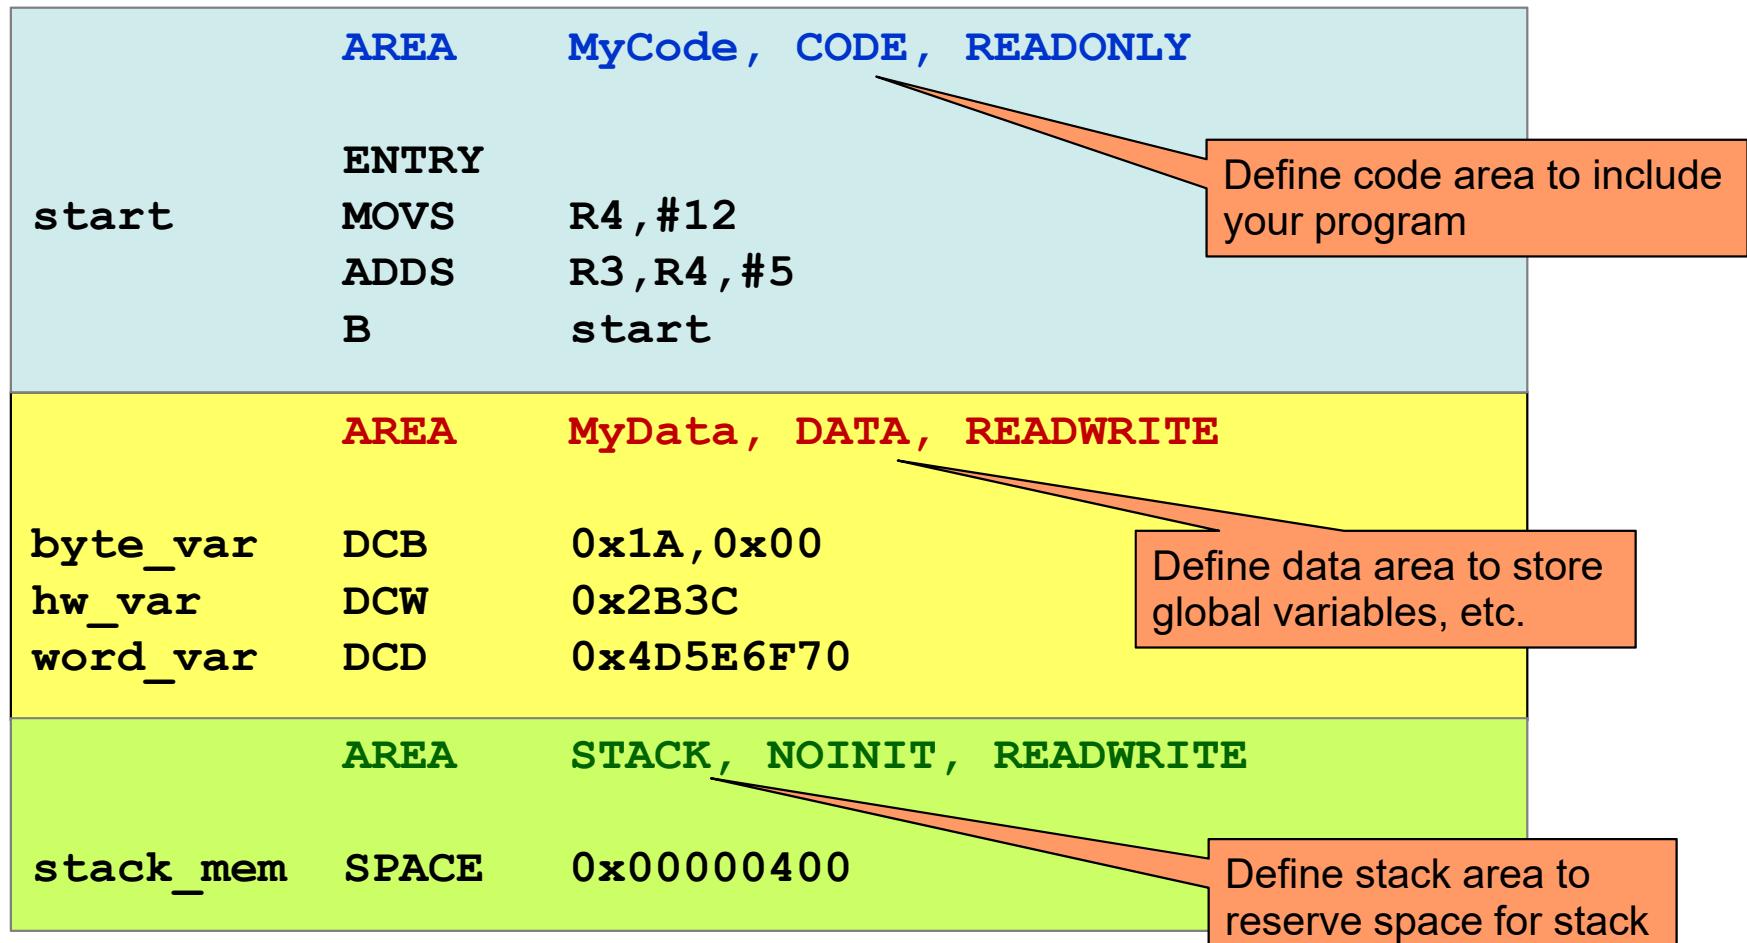


¹⁾ Literal: a fixed/constant value in source code

Object File Sections

■ Assembly Program Structure

- AREA directive



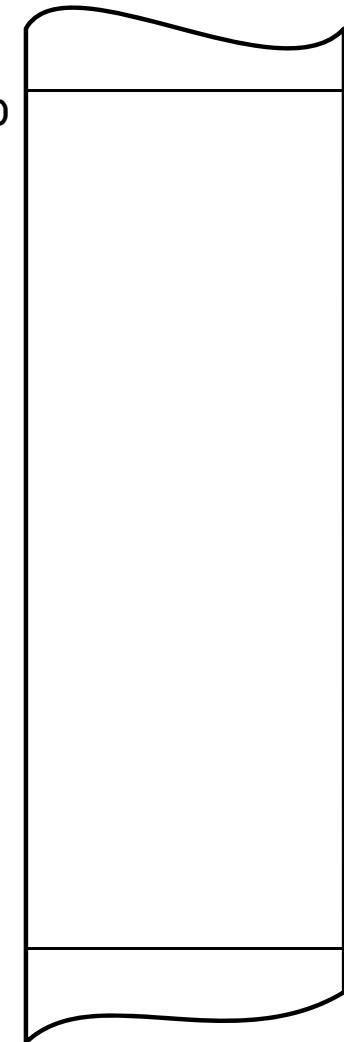
Memory Map / Object File Sections

Assume

- A program uses the following memory segments during execution

- Code 0x0800'1000 to 0x0800'17FF
- Data 0x2001'0000 to 0x2001'01FF
- Stack 0x2001'0200 to 0x2001'05FF

0x0800'0000



- Exercise

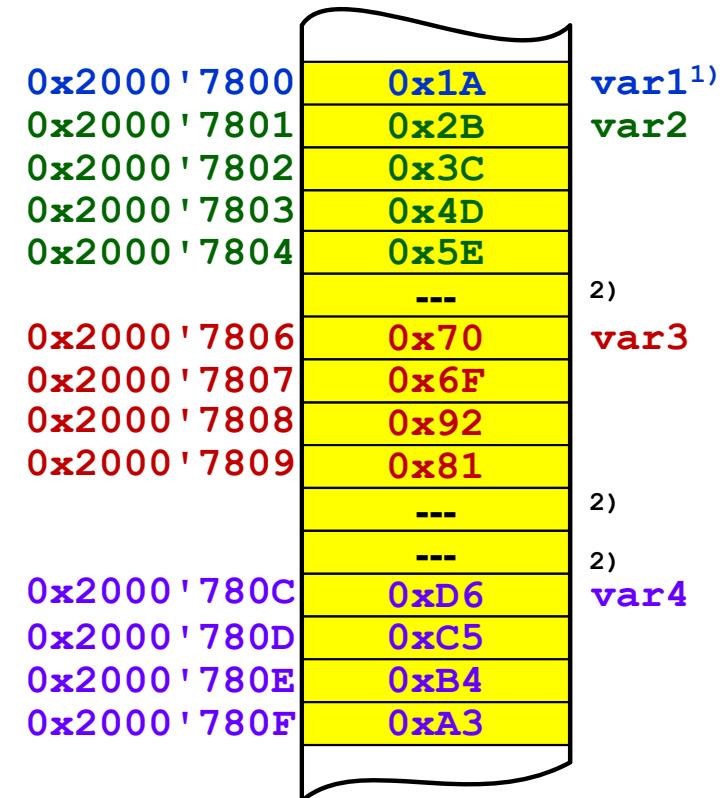
- Draw the memory map with the three sections
- For each section mark the first and the last address with its value
- How many memory cells (bytes) does each section contain?
- For each section: In which STM32F4 memory is it located?

Object File Sections

■ Memory Allocation in Assembly

- Directives for initialized data
 - DCB bytes
 - DCW half-words (half-word aligned)
 - DCD words (word aligned)
 - Can be located in DATA or CODE area

```
AREA example1, DATA, READWRITE
var1    DCB   0x1A
var2    DCB   0x2B, 0x3C, 0x4D, 0x5E
var3    DCW   0x6F70, 0x8192
var4    DCD   0xA3B4C5D6
```

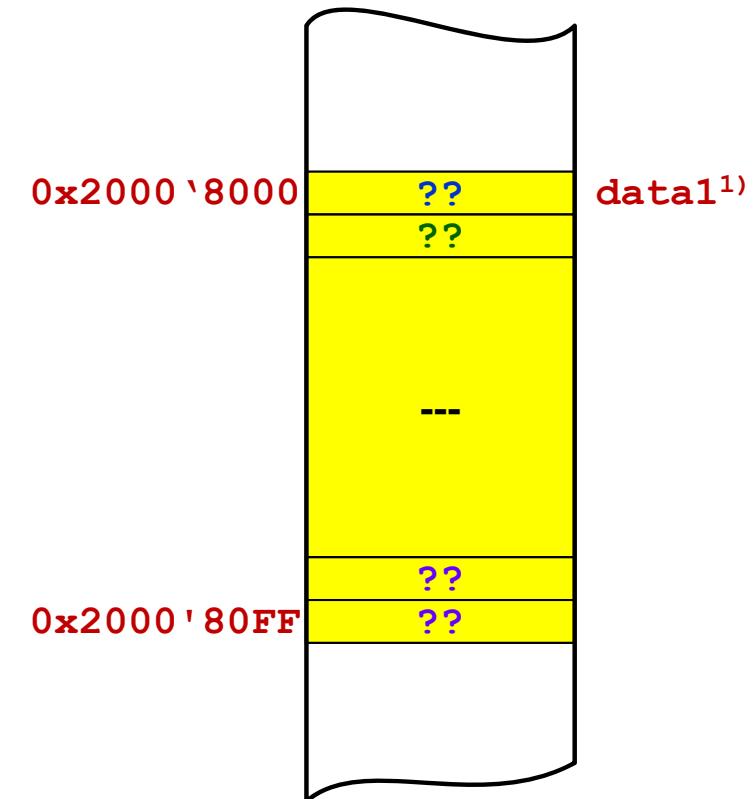


Object File Sections

■ Memory Allocation in Assembly

- Directives for uninitialized data
 - **SPACE** or % with number of bytes to be reserved
 - Reserves number of bytes without initializing them

```
AREA example2, DATA, READWRITE
data1 SPACE 256
```



¹⁾ if we assume that example2 starts at 0x2000'8000

Object File Sections

■ Code Example

show_variables.c

```
uint32_t g_init_var = 0x4D5E6F70;

uint32_t g_noinit_var;

const uint32_t g_const = 0xDDDEFF00;

void show_variables(void)
{
    uint32_t local_var;

    const uint32_t local_const =
                    0x7261504F;

    static uint32_t static_local_var;

    local_var = 0xD4E5F607;
    ...
}
```

Into which sections will the colored objects be allocated?

Object File Sections

Code Example

```
show_variables.c
uint32_t g_init_var = 0x4D5E6F70;           global variable with initialization
uint32_t g_noinit_var;                      global variable with no init value
                                              specified in C source code
const uint32_t g_const = 0xDDDEFF00;         global constant
void show_variables(void)
{
    uint32_t local_var;                     local variable
    const uint32_t local_const =
        0x7261504F;                      local constant
    static uint32_t static_local_var;        local variable with qualifier static
    local_var = 0xD4E5F607;                literal
    ...
}
```

The diagram illustrates the mapping of C code elements to object file sections. A vertical grey bar on the left separates the global scope (above) from the local scope (below). Annotations are color-coded: green for global variables, orange for local variables, pink for global constants, and blue for literals.

- global variable with initialization: `g_init_var`
- global variable with no init value specified in C source code: `g_noinit_var`
- global constant: `g_const`
- local variable: `local_var`
- local constant: `local_const`
- literal: `0x7261504F`
- local variable with qualifier static: `static_local_var`
- literal: `0xD4E5F607`

Literal A fixed value in source code **without** an associated symbol name
Constant “variable” that cannot be changed after first assignment

Object File Sections

Code Example

show_variables.c

```
uint32_t g_init_var = 0x4D5E6F70;  
  
uint32_t g_noinit_var;  
  
const uint32_t g_const = 0xDDDEFF00;  
  
void show_variables(void)  
{  
    uint32_t local_var;  
  
    const uint32_t local_const =  
        0x7261504F;  
  
    static uint32_t static_local_var;  
  
    local_var = 0xD4E5F607;  
    ...  
}
```

assembly

```
AREA GlobalVars, DATA, READWRITE  
  
g_init_var      DCD  0x4d5e6f70  
g_noinit_var    DCD  0x00000000  
static_local_var DCD  0x00000000
```

```
AREA MyConsts, DATA, READONLY  
  
g_const          DCD  0xddeeff00
```

```
AREA MyCode, CODE, READONLY,  
  
show_variables    ...  
                  ...  
                  BX   lr  
  
lit1              DCD  0x7261504f  
lit2              DCD  0xd4e5f607
```

stored in
registers¹⁾

literal

¹⁾ if possible. E.g. if you use the address operator (&) on a local variable it will be allocated on the stack.

Conclusion

■ Components Cortex-M CPU

- Core Registers: R0-R12, SP, LR, PC
- 32-bit ALU
- Flags (APSR)
- Control Unit with IR (Instruction Register)
- Bus Interface

■ Instruction Types

- Data transfer, data processing, control flow

■ Program Execution

- Fetch – Execute

■ Memory Map

■ Integer Types

- Size depends on architecture → use C99 types for portability
- 'Little Endian' vs. 'Big Endian', alignment

■ Object File Sections

- CODE, DATA, STACK