

Cache

Computer Engineering 2

**CT Team: A. Gieriet, J. Gruber, R. Gübeli, M. Meli, M. Rosenthal,
A. Rüst, J. Scheier, M. Thaler**

- **Principle of locality**
- **Cache mechanics**
- **Cache organization**
 - Fully associative
 - Direct mapped
 - N-way set associative
- **Performance metrics**
- **Cache misses**
- **Replacement strategies**
- **The programmers perspective**

■ Situation

- Processor
 - Fast cycle time
- Fast DRAM¹
 - Slow cycle time (up to 100x)
 - Efficiently reads only in bursts

→ Bridging the gap such that pipelining is effective!

■ Goal

- Access “slower” memory in bursts and maintain a fast cache for fast single accesses
- But: Data integrity must be carefully managed, such that both, cache and main memory have the same data

¹ *Dynamic RAM*

■ At the end of this lesson you will be able

- to understand the principles of cache memory
- to explain the principle of locality
- to enumerate advantages and disadvantages of different cache models
 - Fully associative
 - Direct mapped
 - N-way set associative
- to enumerate types of cache misses
- to understand how cache size and cache hit rate are related
- to name different replacement strategies

■ Principle of locality

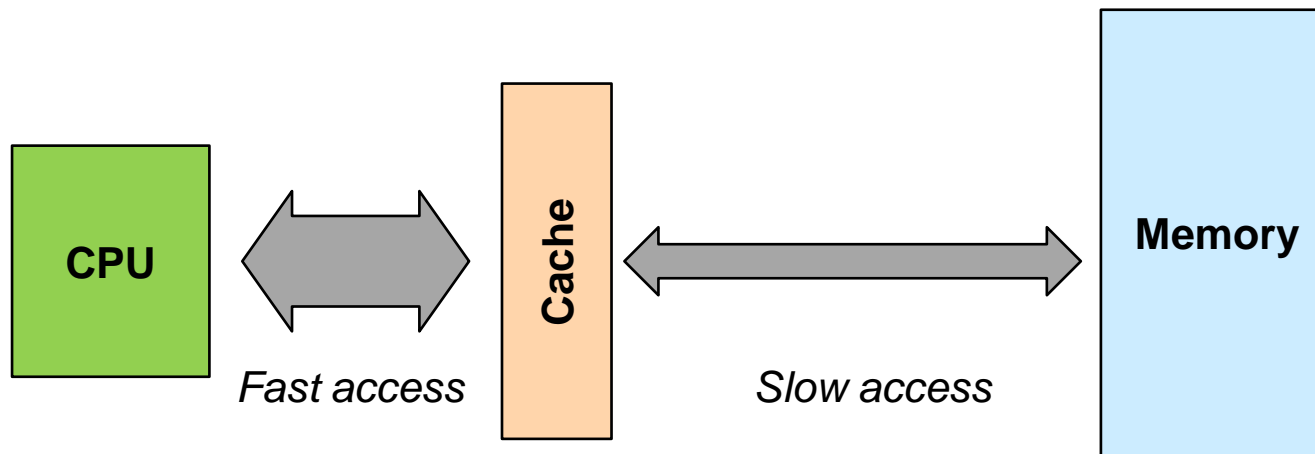
Programs usually access small regions of memory in a given interval of time

- Temporal locality
 - Current data location is likely being accessed again in near future
- Spatial locality
 - Current data location is likely being close to next accessed location

```
for(i = 0; i < 100000; i++) {           // incremental access
    a[i] = b[i];                         // → spatial locality
}
if (a[1234] == a[4321]) {               // → temporal locality
    a[1234] = 0;
}
```

■ Definition cache

- Computer memory with short access time
- Storage of frequently or recently used instructions or data



CPU

Very small
Registers (Word)

Cache

Small, fast, expensive
Caches subset of memory blocks

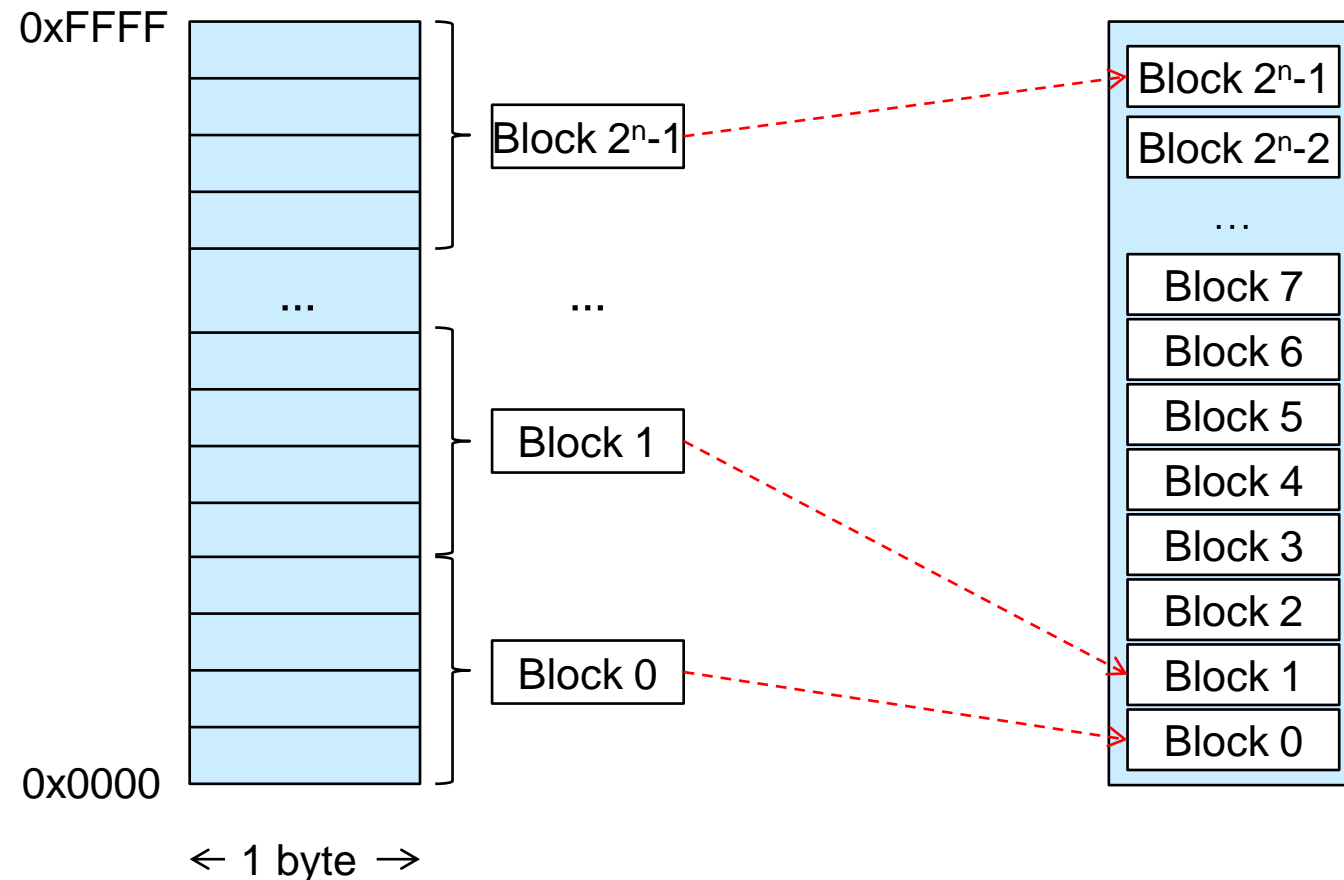
Memory

Large, slow, cheap
Partitioned into memory blocks

Cache Mechanics


■ Memory blocks

- All examples given in this lecture are using*
- hypothetical 16 bit addressing*
 - blocks of 4 bytes*



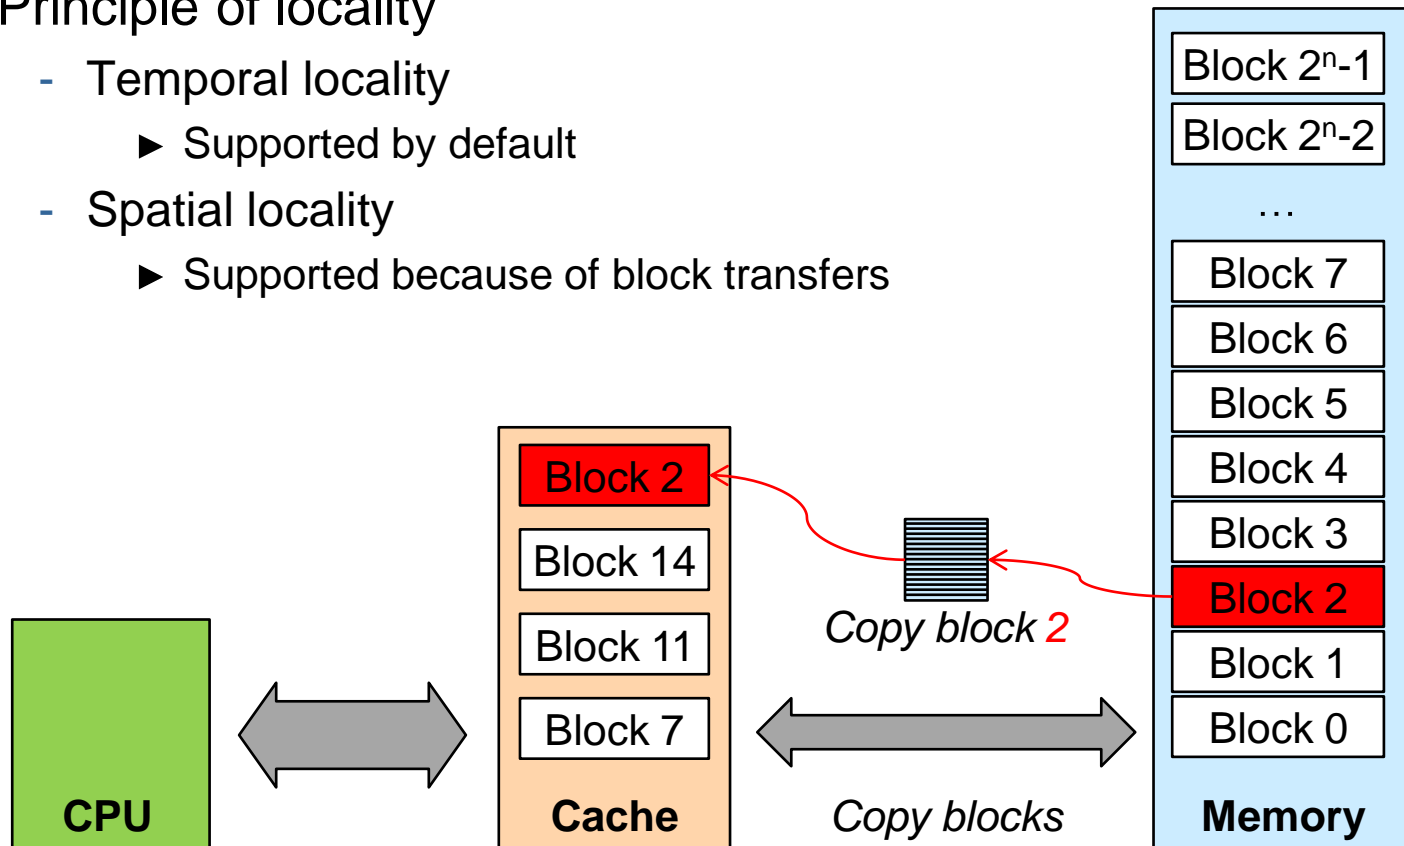
■ Memory blocks

- Selected blocks of main memory copied to faster cache memory



■ Memory blocks

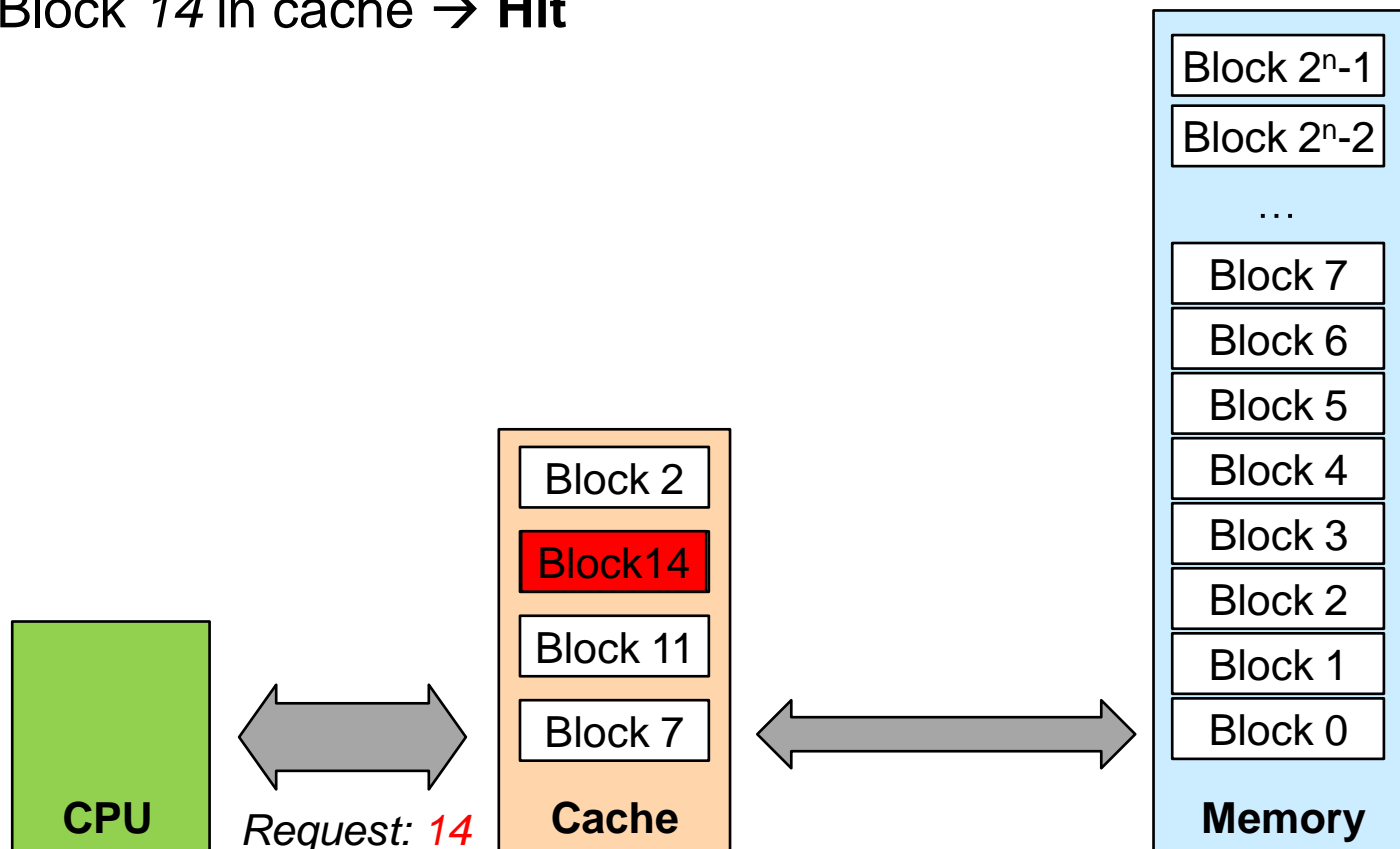
- Blocks of main memory copied to faster cache memory
- Principle of locality
 - Temporal locality
 - ▶ Supported by default
 - Spatial locality
 - ▶ Supported because of block transfers



Cache Mechanics

■ Cache hit

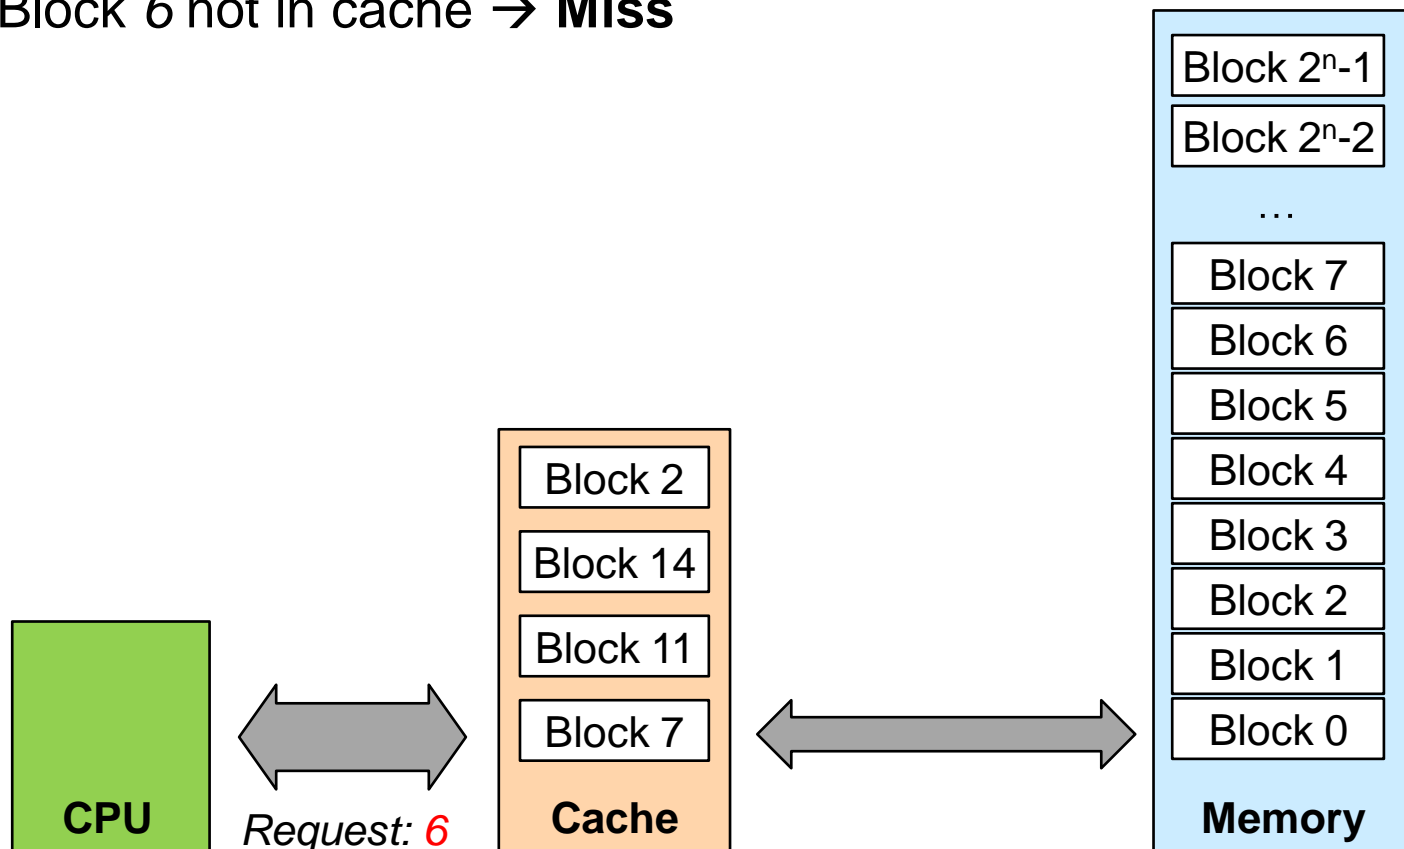
- Data in block 14 is needed
- Block 14 in cache → **Hit**



Cache Mechanics

■ Cache miss (I)

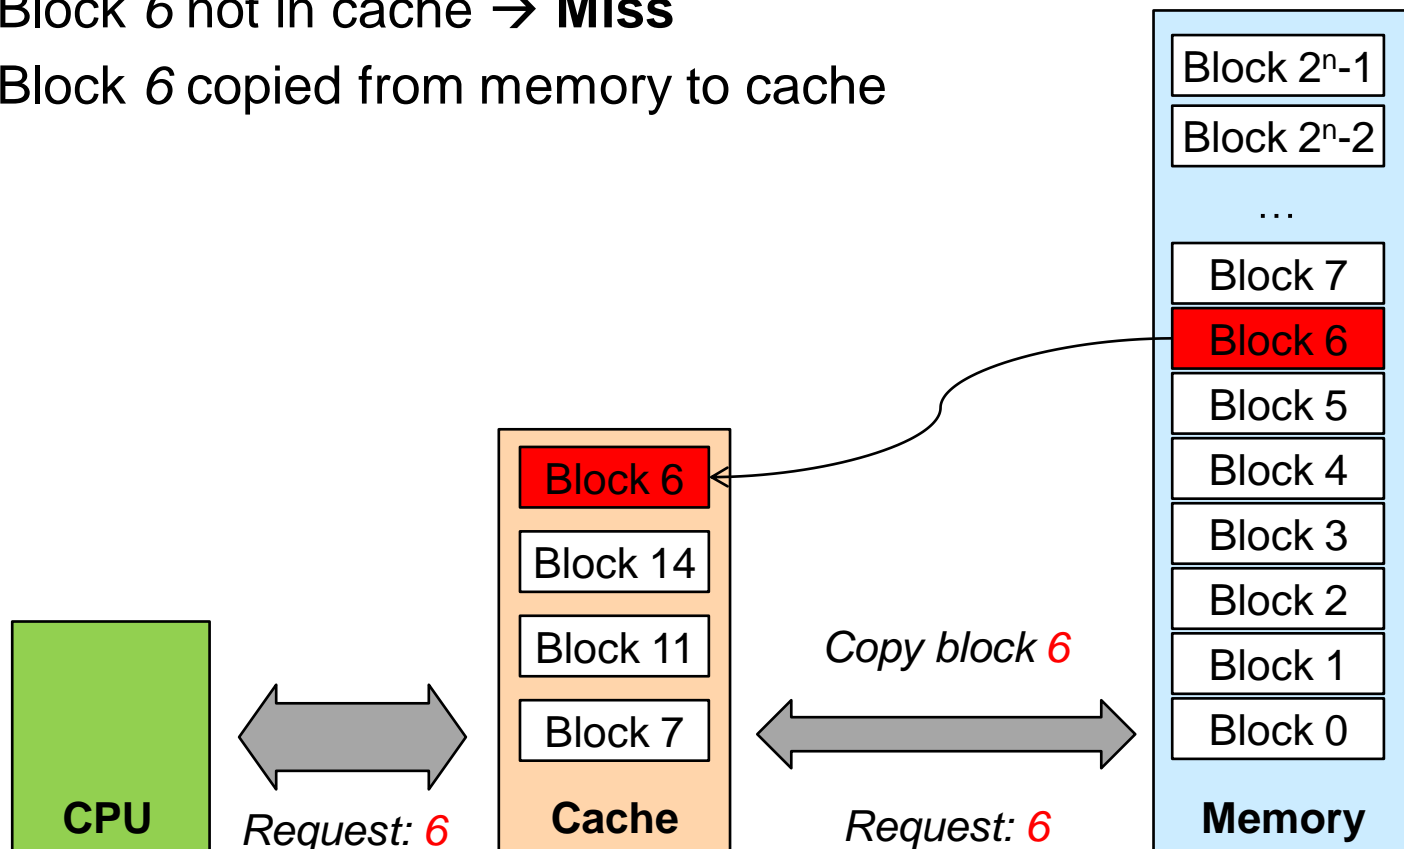
- Data in block 6 is needed
- Block 6 not in cache → **Miss**



Cache Mechanics

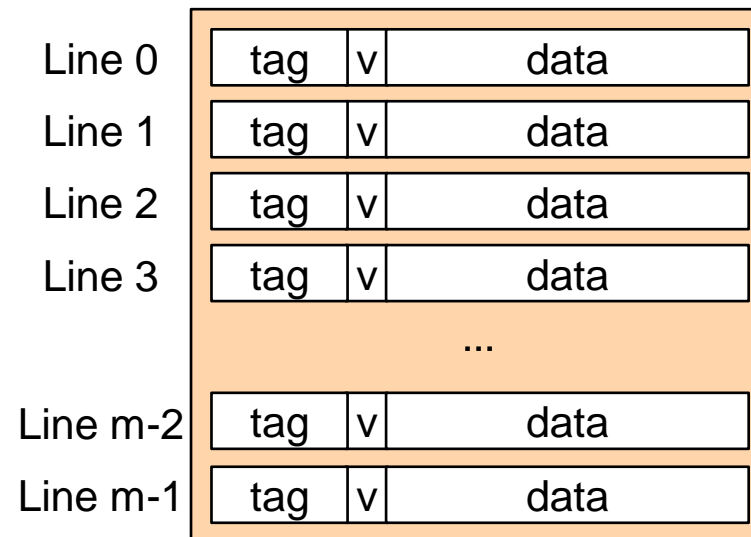
■ Cache miss (II)

- Data in block 6 is needed
- Block 6 not in cache → **Miss**
- Block 6 copied from memory to cache



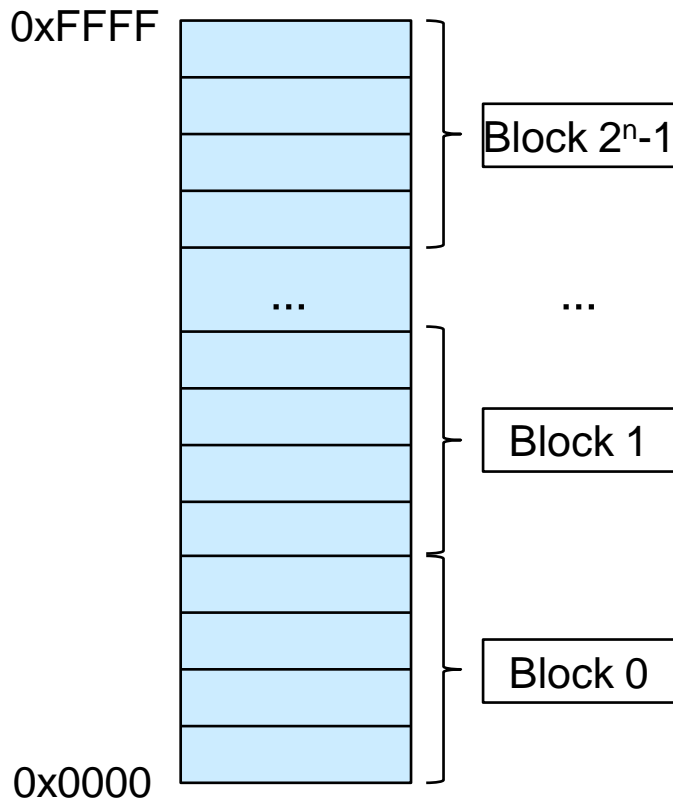
■ Organized in lines

- Valid bit v → indicates that line contains valid data
- Tag → unique identifier for memory location
- Data → data of exactly one memory block
- m = overall number of cache lines



Cache Organization

■ Addressing

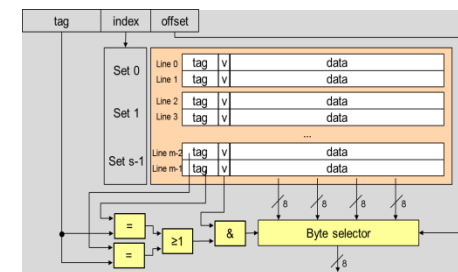
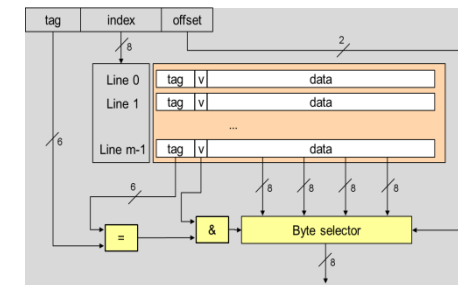
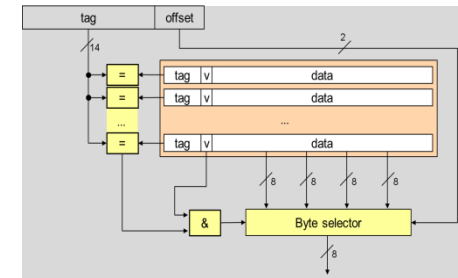


Block	Address	Block identification bits	offset
Block 2 ⁿ -1	0xFFFF	1111 1111 1111 11	11
	0xFFFE	1111 1111 1111 11	10
	0xFFFD	1111 1111 1111 11	01
	0xFFFC	1111 1111 1111 11	00
...			
Block 1	0x0007	0000 0000 0000 01	11
	0x0006	0000 0000 0000 01	10
	0x0005	0000 0000 0000 01	01
	0x0004	0000 0000 0000 01	00
Block 0	0x0003	0000 0000 0000 00	11
	0x0002	0000 0000 0000 00	10
	0x0001	0000 0000 0000 00	01
	0x0000	0000 0000 0000 00	00

Cache Organization

■ Three different cache models

- Fully associative
- Direct mapped
- N-way set associative




■ Addressing

- Tag contains complete block identification

Block	Address	Block identification bits	offset
Block 1	0x0007	0000 0000 0000 01	11
	0x0006	0000 0000 0000 01	10
	0x0005	0000 0000 0000 01	01
	0x0004	0000 0000 0000 01	00

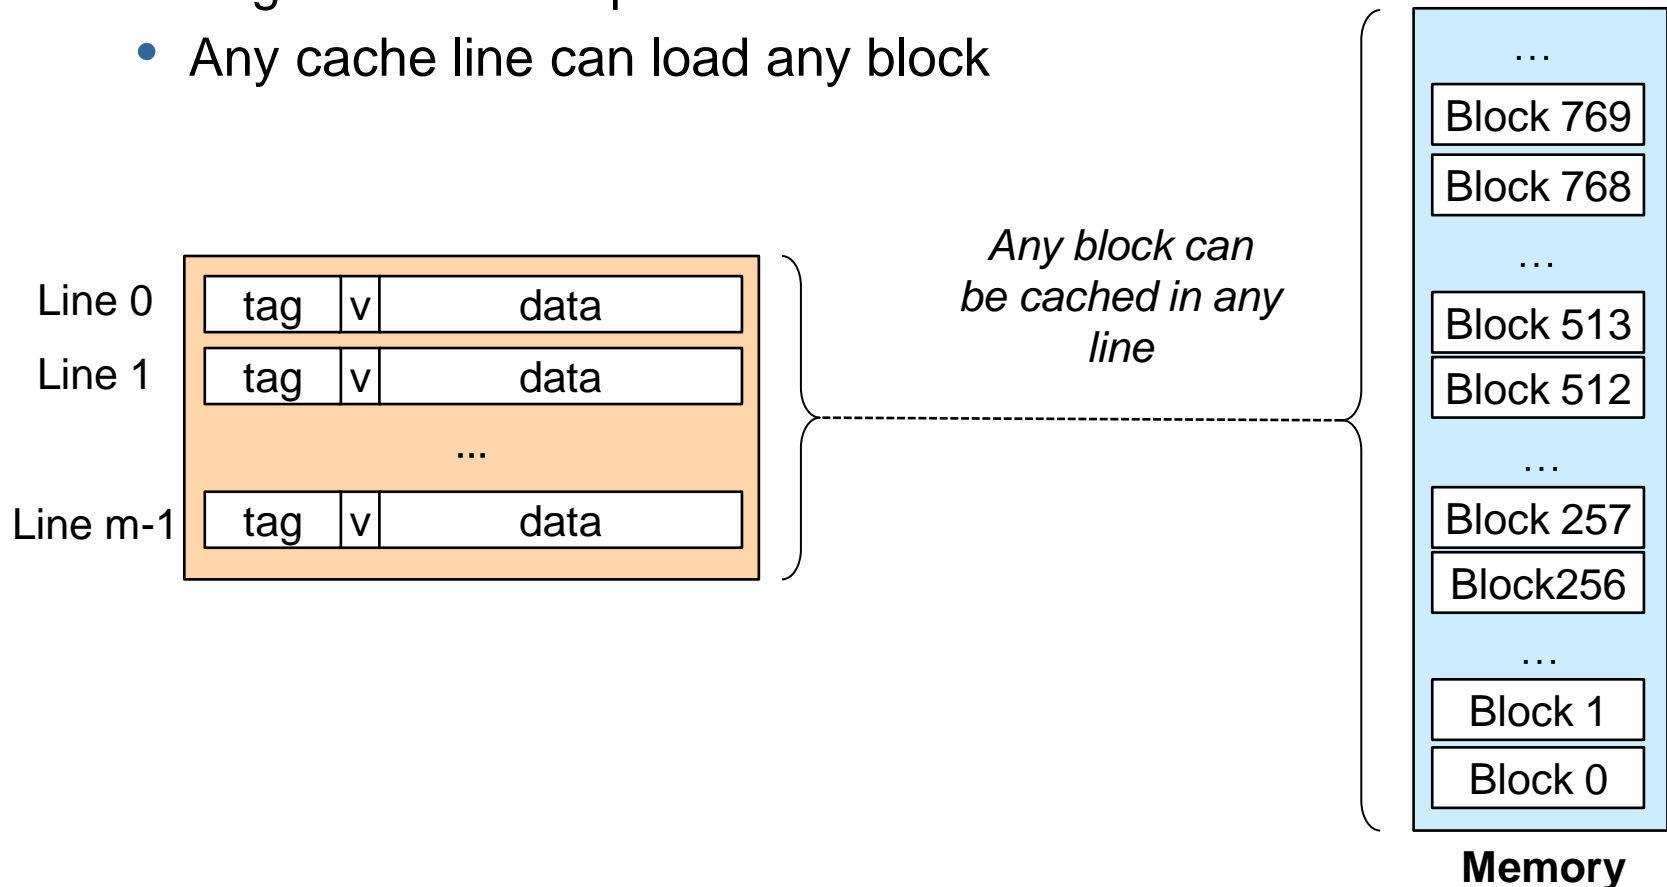
16 bit address
from processor



Fully Associative

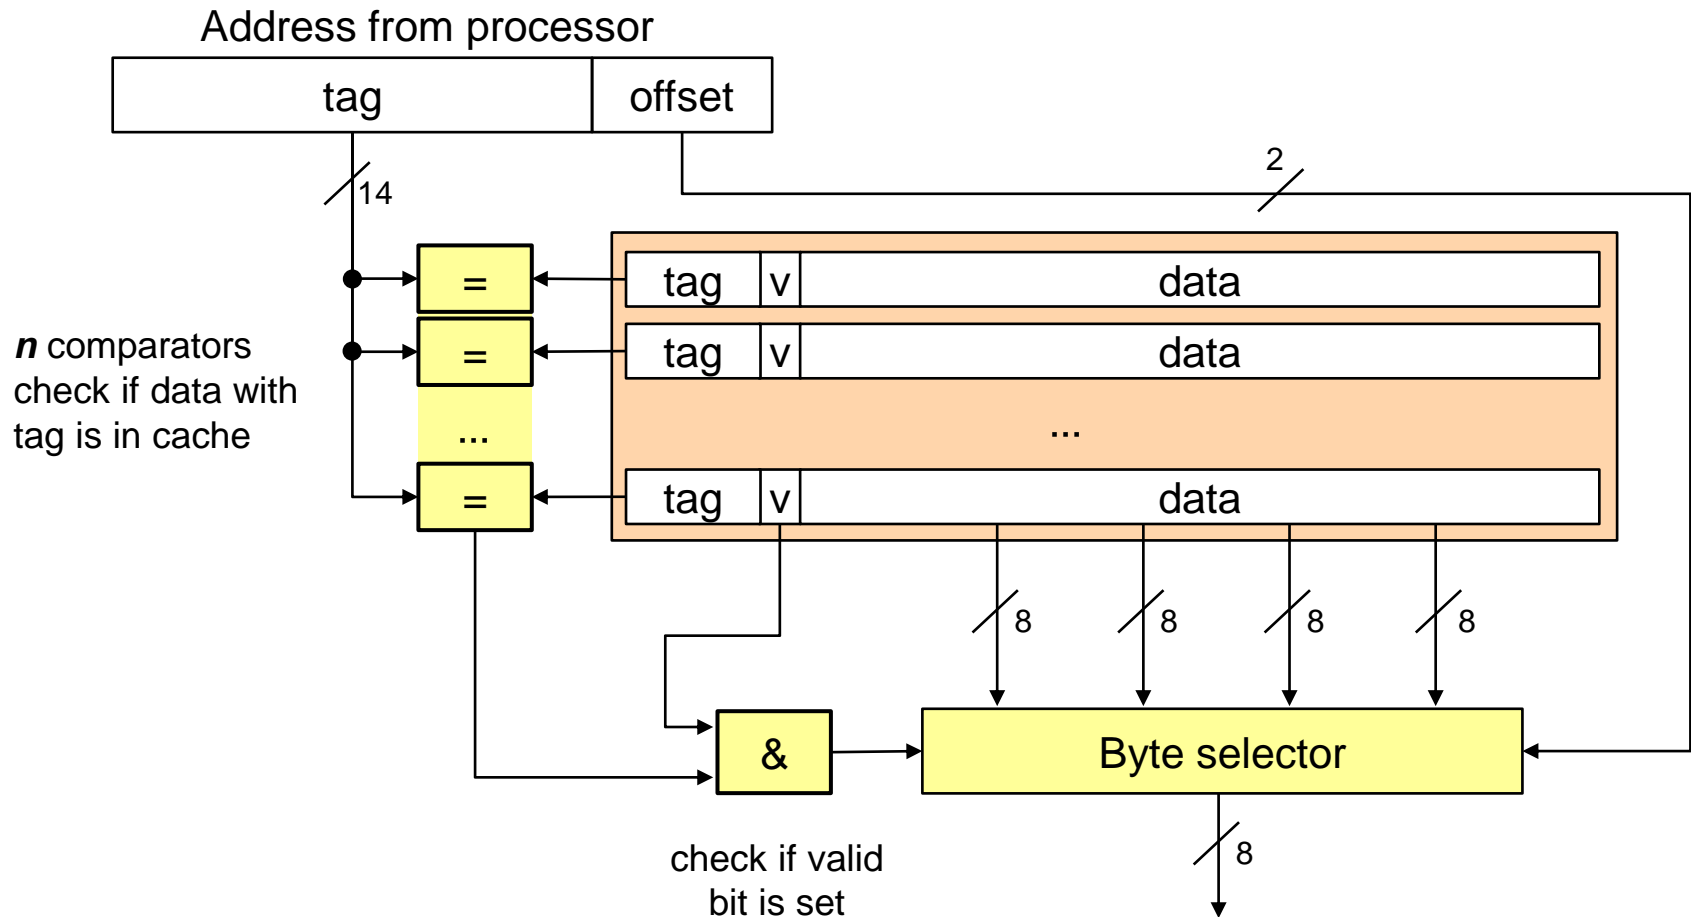
■ Organization

- Tag contains complete block identification
- Any cache line can load any block




Fully Associative

■ Architecture



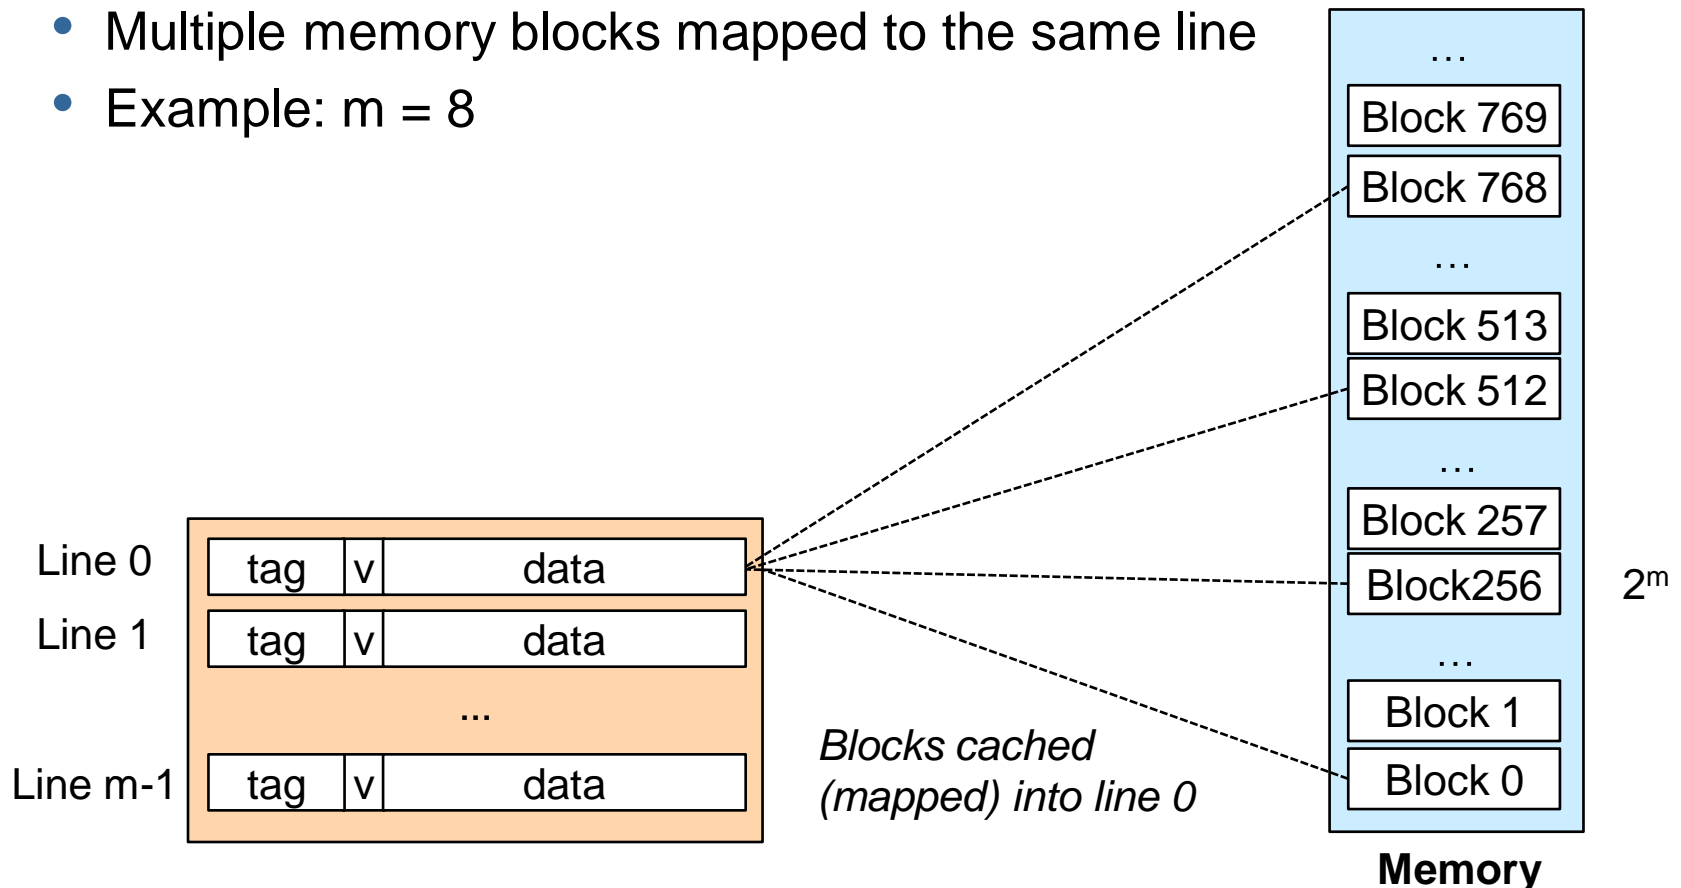
■ Addressing

- Block identification spitted into tag and index

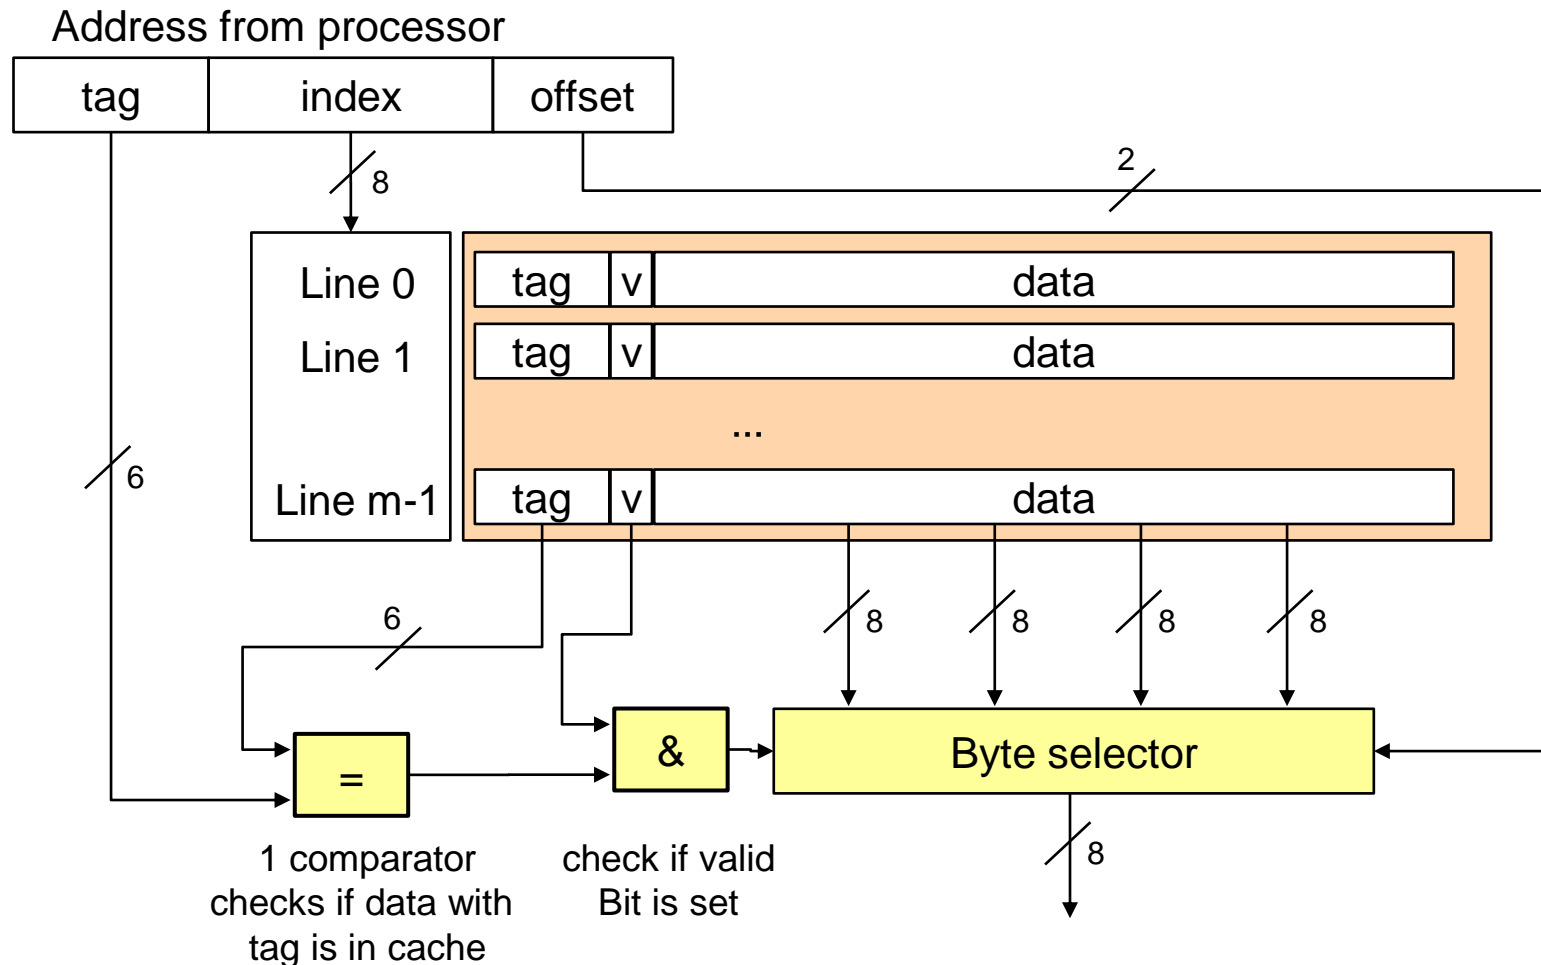


■ Organization

- Each memory block is mapped to exactly one cache line
- Multiple memory blocks mapped to the same line
- Example: $m = 8$



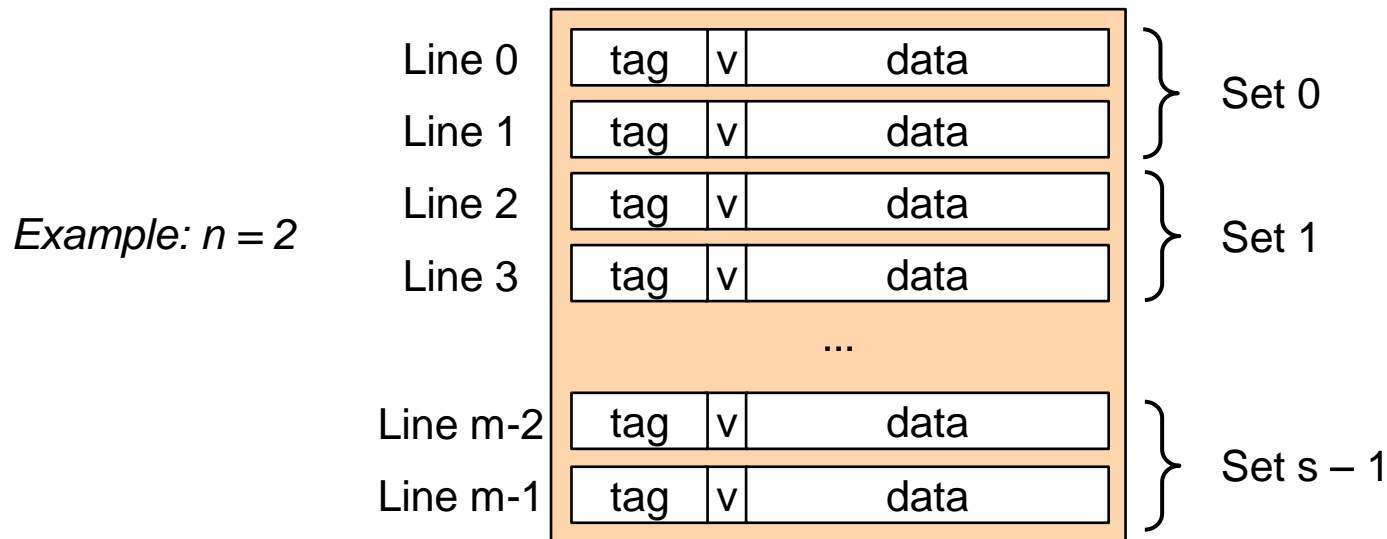
■ Architecture



N-Way Set Associative

■ Organization

- Partition into sets
 - $s = m/n$ number of sets
 - n lines per set („N-way“)
 - b bytes per line
- $s \times n \times b$ data bytes




N-Way Set Associative

■ Addressing

- Example: $n = 2$
 - Maximum index corresponds to number of sets ($s = m/n$)

Block	Address	Block identification bits	offset
Block 1	0x0007	0000 0000 0000 01	11
	0x0006	0000 0000 0000 01	10
	0x0005	0000 0000 0000 01	01
	0x0004	0000 0000 0000 01	00

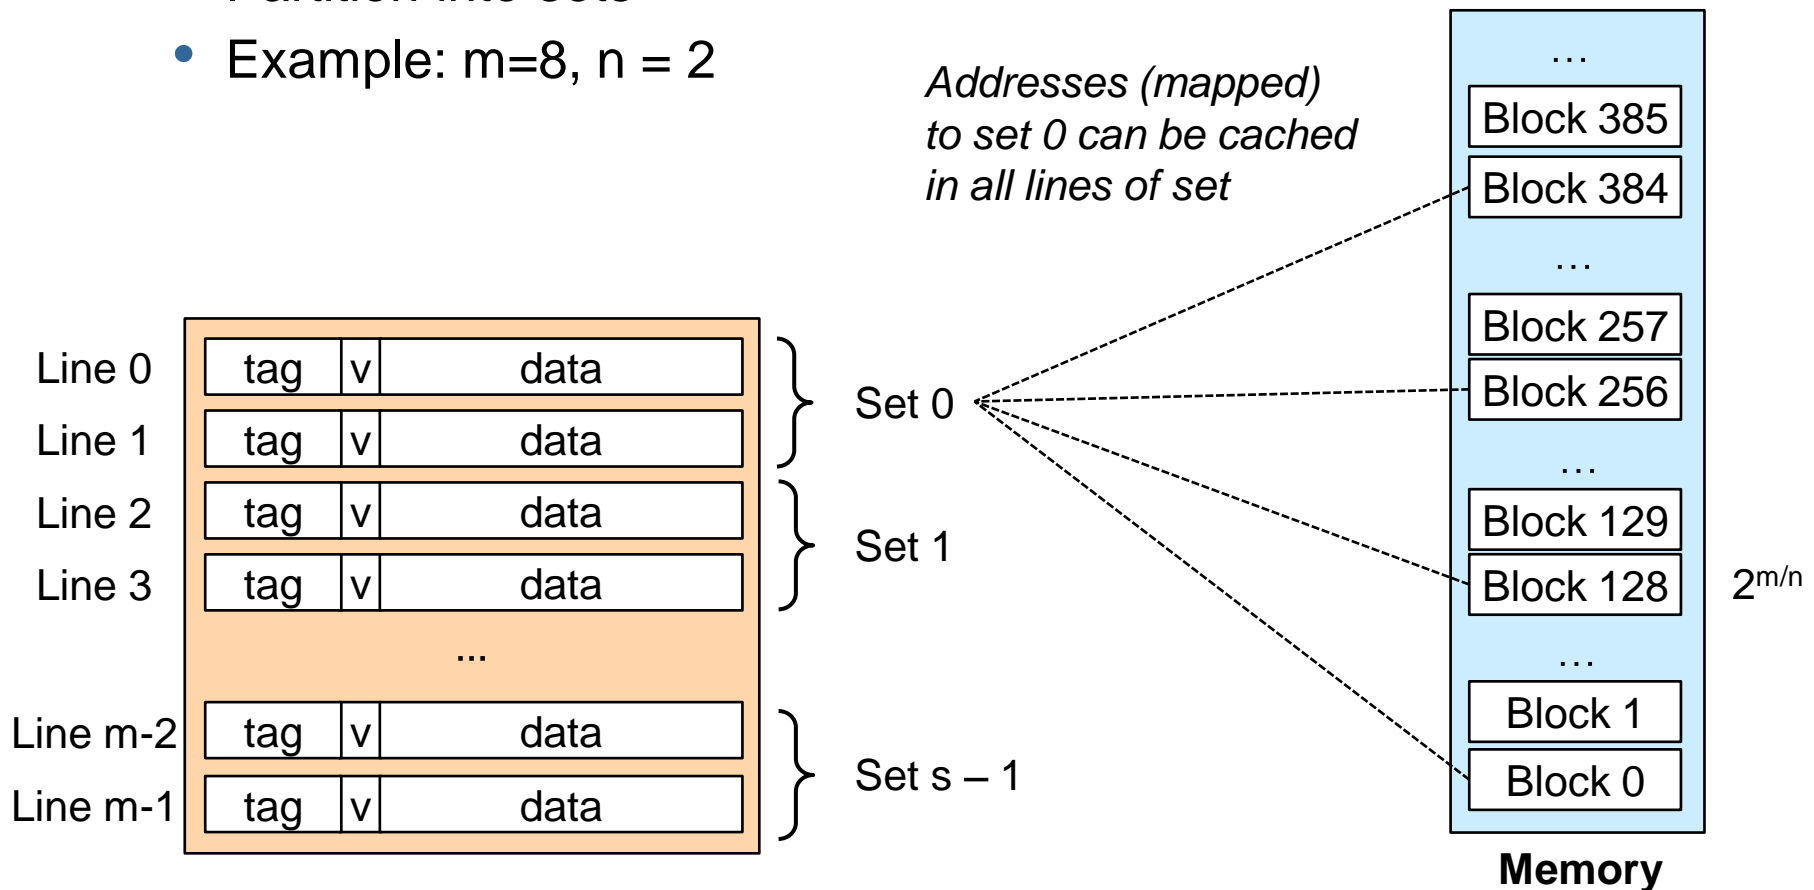
16 bit address
from processor



N-Way Set Associative

■ Organization

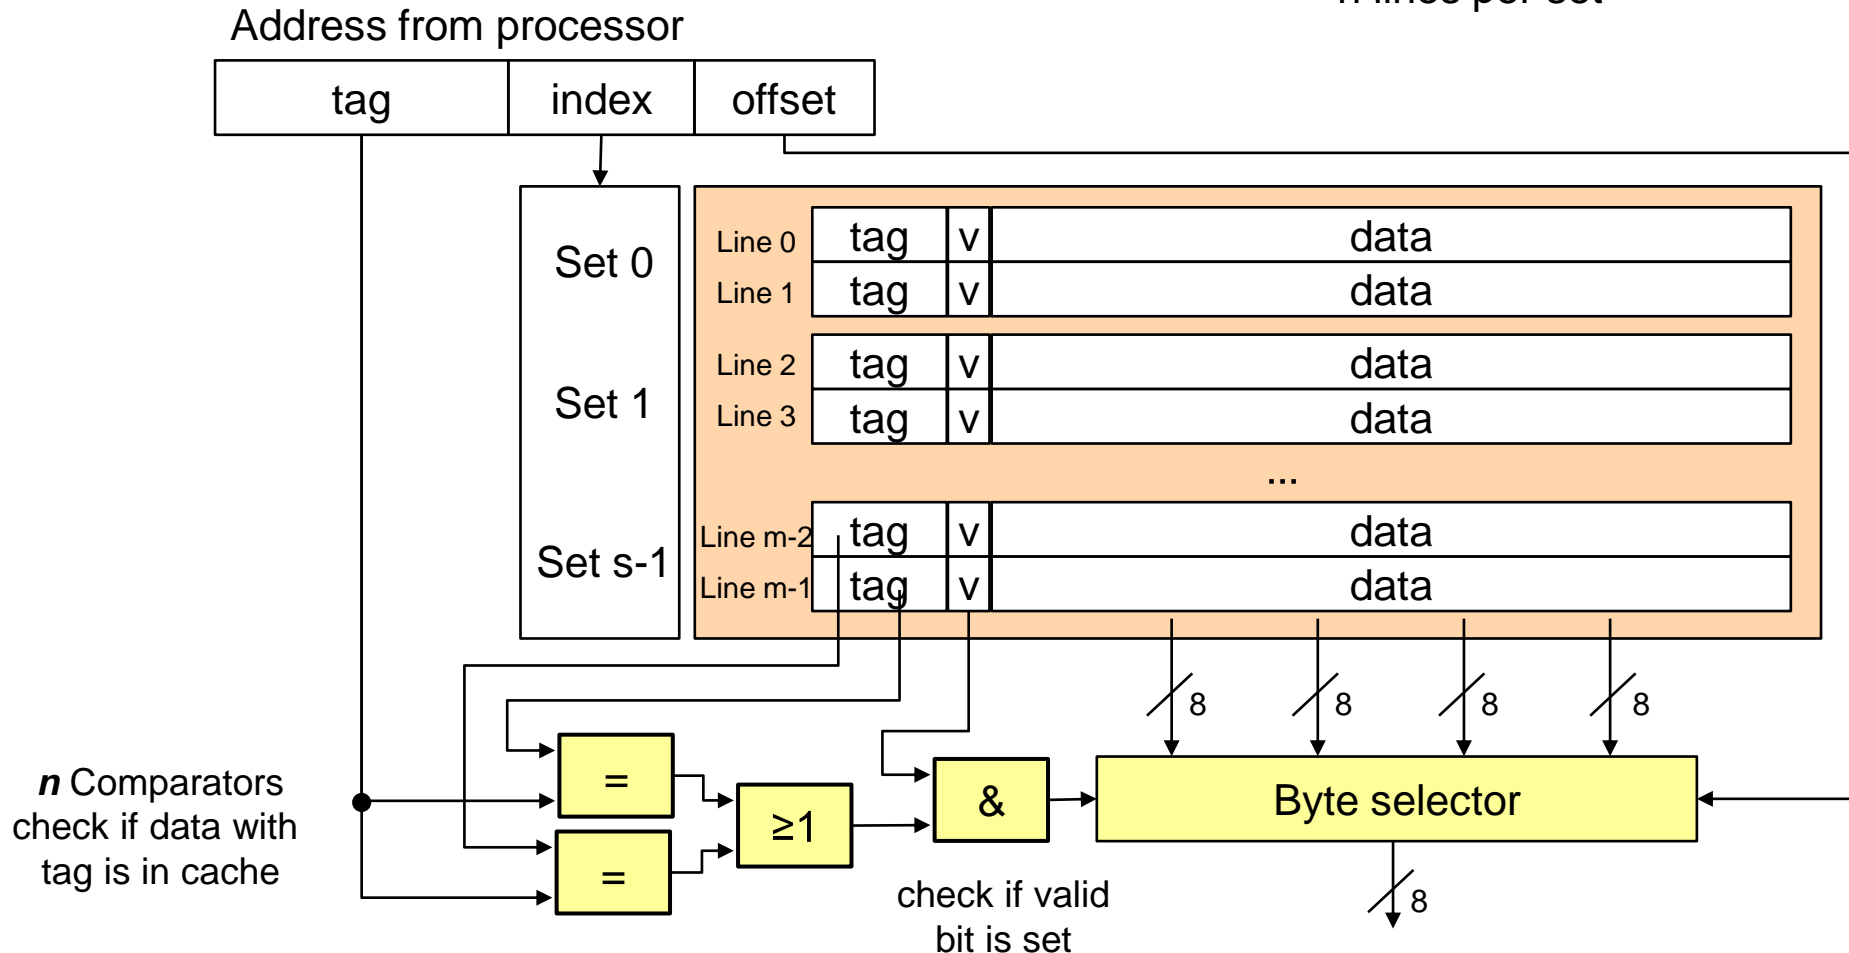
- Partition into sets
- Example: $m=8$, $n = 2$



N-Way Set Associative

■ Architecture

Example: $n = 2$
 $s = m/n$ number of sets
 n lines per set



■ Comparison

Organization	Fully associative	Direct mapped	N-way set associative
Number of sets	1	m	m/n
Associativity	m(=n)	1	n
Advantages	<ul style="list-style-type: none">• Fast, flexible• Highest hit rates• Advanced replacement strategies	<ul style="list-style-type: none">• Simple logic• Replacement strategy defined by organization	Combination of both other concepts to combine advantages and to compensate disadvantages
Disadvantages	<ul style="list-style-type: none">• Complex logic: one comparator per line• Requires large area on silicon• Replacement can be complex	<ul style="list-style-type: none">• Lower hit rates	

■ Cold miss

- First access to a block

■ Capacity miss

- Working set larger than cache

■ Conflict miss

- Multiple data objects map to same slot

■ Miss rate

- Fraction of memory references not found in cache
$$\text{misses} / \text{accesses} = 1 - \text{hit rate}$$

■ Hit time

- Time to deliver a block in the cache to the processor

■ Miss penalty

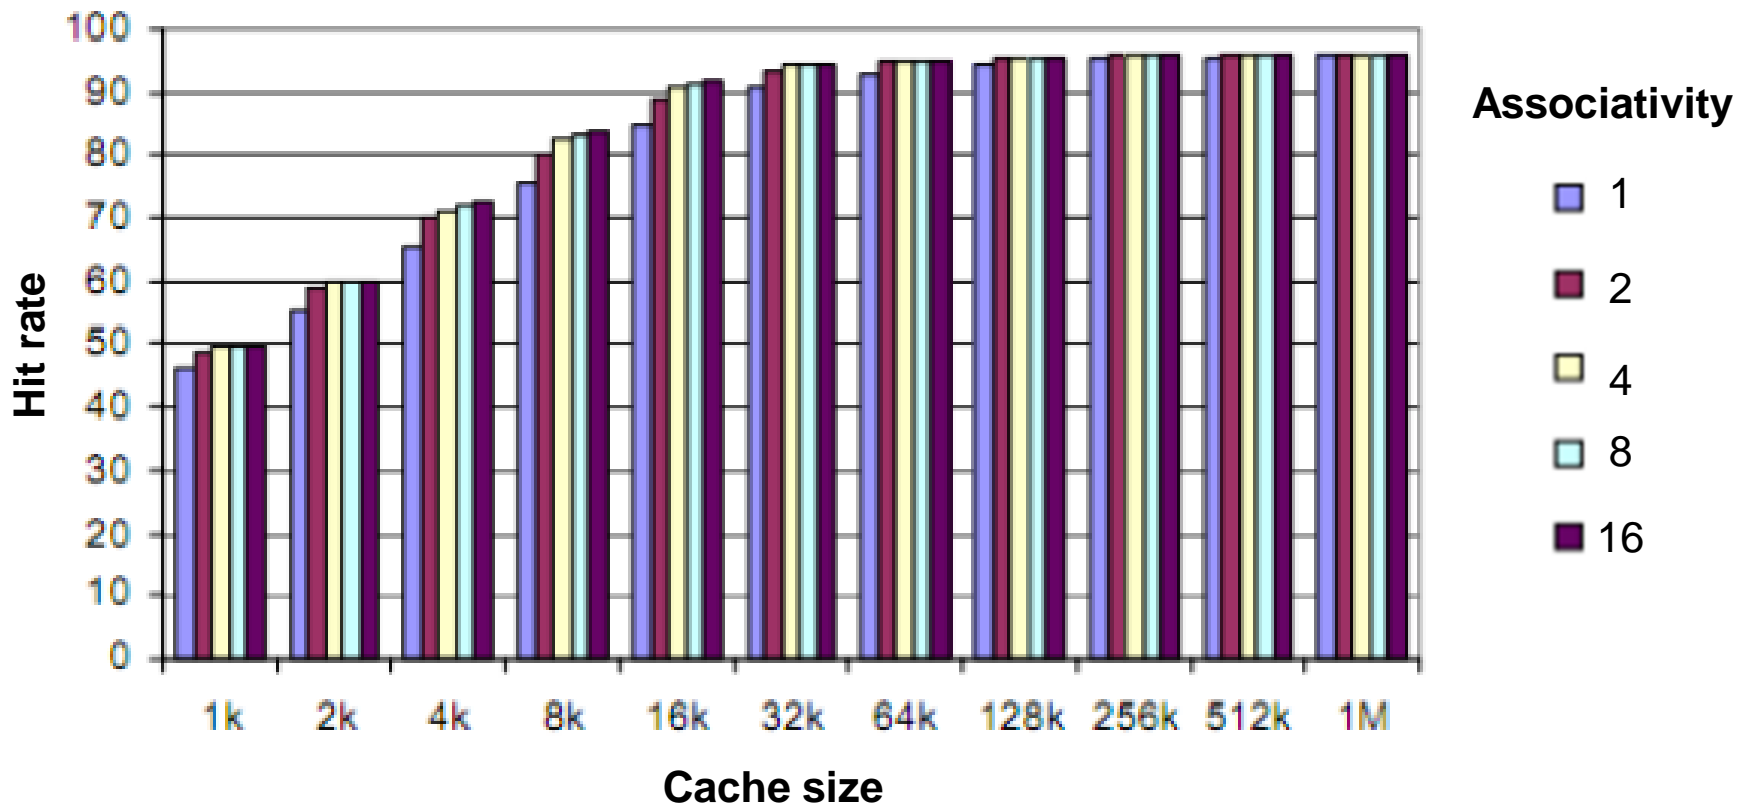
- Additional time required to fetch data from memory because of a cache miss

- **High cache hit rate is important!**
- **99% hit rate can be twice as fast as 97%**

- **Example:**
 - Cache hit time: 1 processor cycle
 - Miss penalty: 100 processor cycles
 - Average access time is:
 - **97% hits:** $1 \text{ cycle} + 0.03 * 100 \text{ cycles} = 4 \text{ cycles average}$
 - **99% hits:** $1 \text{ cycle} + 0.01 * 100 \text{ cycles} = 2 \text{ cycles average}$

■ Cache size versus hit rate

- Typical hit rate for Associativity 1,2,4,8,16 and cache size



■ Selecting cache line to replace by

- LRU: Least recently used
 - LFU: Least frequently used
 - FIFO: First In–First Out oldest
 - Random Replace: randomly chosen
- } *additional information needed in cache*
- } *simple, but still good performance*

→ Relevance only for Fully- / N-way associative cache

→ Hard coded in cache implementation

■ What to do on a write hit*?

- Write-through
 - Write immediately to memory
- Write-back
 - Delay write to memory until replacement of line (needs a valid bit)

■ What to do on a write miss**?

- Write-allocate
 - Load into cache and update line in cache
- No-write-allocate
 - Writes immediately to memory


* *Write hit: data already in cache*

** *Write miss: data not in cache*

The Programmers Perspective

- **For loops over multi-dimensional arrays**
 - Example: matrices (2-dim arrays)
- **Change order of iteration to match layout**
 - Gets better spatial locality
 - Layout in C: **last index changes first!**

```
for(j = 0; j < 10000; j++){  
    for(i = 0; i < 40000; i++){  
        c[i][j]=a[i][j]+b[i][j];  
    }  
}  
// a[i][j] and a[i+1][j]  
// are 10'000 elements apart
```



```
for(i = 0; i < 40000; i++){  
    for(j = 0; j < 10000; j++){  
        c[i][j]=a[i][j]+b[i][j];  
    }  
}  
// a[i][j] and a[i][j+1]  
// are next to each other
```

The Programmers Perspective

■ Change order of iteration to match layout

```
#include<stdio.h>
#include<time.h>

int main( int argc, char** argv )
{
    int iterations = 1000;
    int size = 1024;

    int array[size][size];
    int v = 0;

    if (argc==2) {
        printf("j is inner\n");
        for (int n=0; n<iterations; n++) {
            for (int i=0; i<size; i++) {
                for (int j=0; j<size; j++) {
                    array[i][j] = v;
                    v++;
                }
            }
        }
    } else {
        printf("j is outer\n");
        for (int n=0; n<iterations; n++) {
            for (int j=0; j<size; j++) {
                for (int i=0; i<size; i++) {
                    array[i][j] = v;
                    v++;
                }
            }
        }
    }

    printf("%i-%i done\n", array[0][0], array[size-1][size-1]);
    return 0;
}
```

```
j is inner
1047527424-1048575999 done

real    0m3.460s
user    0m3.452s
sys     0m0.000s


j is outer
1047527424-1048575999 done

real    0m18.509s
user    0m18.472s
sys     0m0.000s
```

■ Cache Levels

- Typical Cache Architecture
- Example: ARM Cortex M3/M4 with external memory

Memory → larger, slower, cheaper



■ Cache

- Cache allows fast data access to a certain part of the main memory which is mirrored in a cache
- Spatial and temporal locality make cache principle
- Different cache models were discussed
 - Fully associative
 - Direct mapped
 - N-way associative
- Replacement / Write strategies
- Reducing cache misses by optimizing memory access
- Example for advanced cache architecture