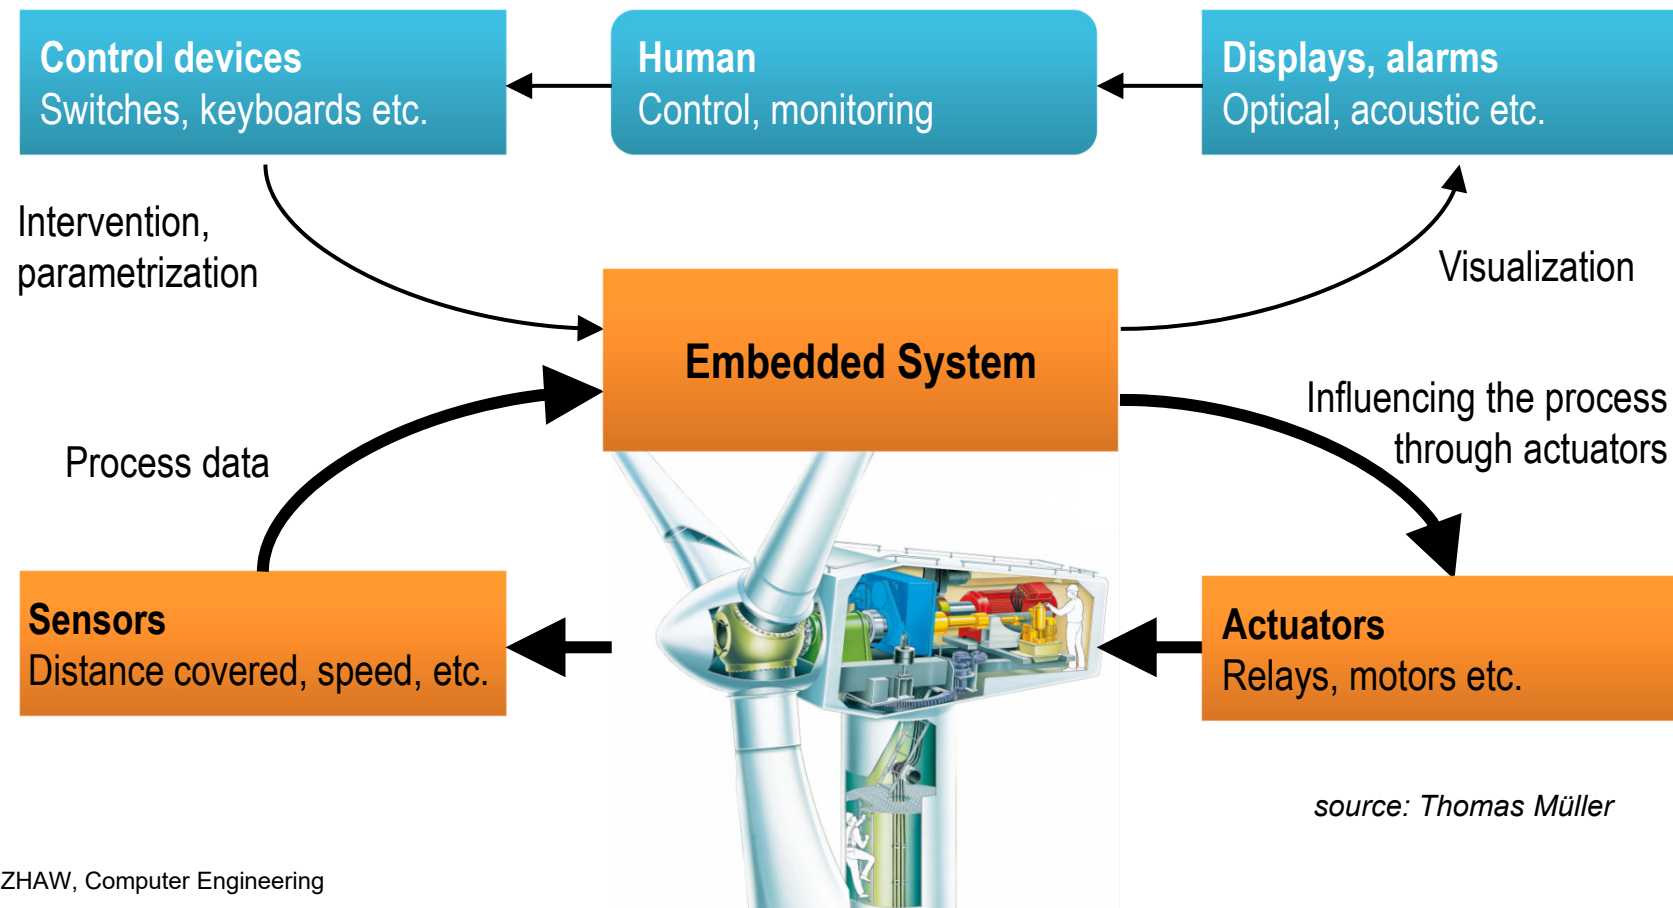


Detecting Events – Interrupt Performance

Computer Engineering 2

- **How do you recognize events?**
 - Cyclic queries or interrupt-driven?



source: Thomas Müller

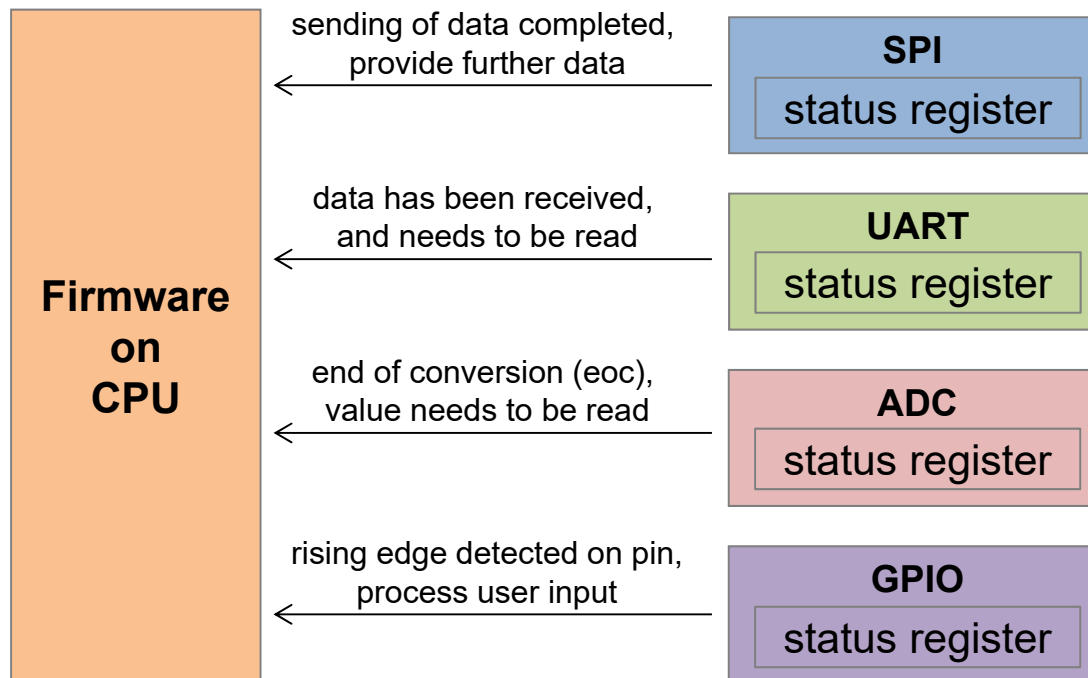
- **Polling**
- **Interrupt Driven I/O**
- **Interrupt Performance**
- **Interrupt Latency**
- **Managing Latency**
- **Interrupt Driven FSM**

At the end of this lesson you will be able


- to explain the methods "Polling" and "Interrupt Driven I/O"
- to name important factors for the two methods
- to enumerate advantages and disadvantages of the methods
- to evaluate for a given situation whether polling or interrupt driven I/O is appropriate
- to quantify timing aspects for interrupt driven I/O
- to comprehend the term "Interrupt Latency" and name potential sources
- to explain the term 'pre-emption'
- to understand approaches to manage interrupt latency
- to name the components and explain the structure of an interrupt driven FSM

■ Peripherals Signal Events to Firmware

- Something happened that requires servicing in firmware



status registers can be addressed and read by firmware



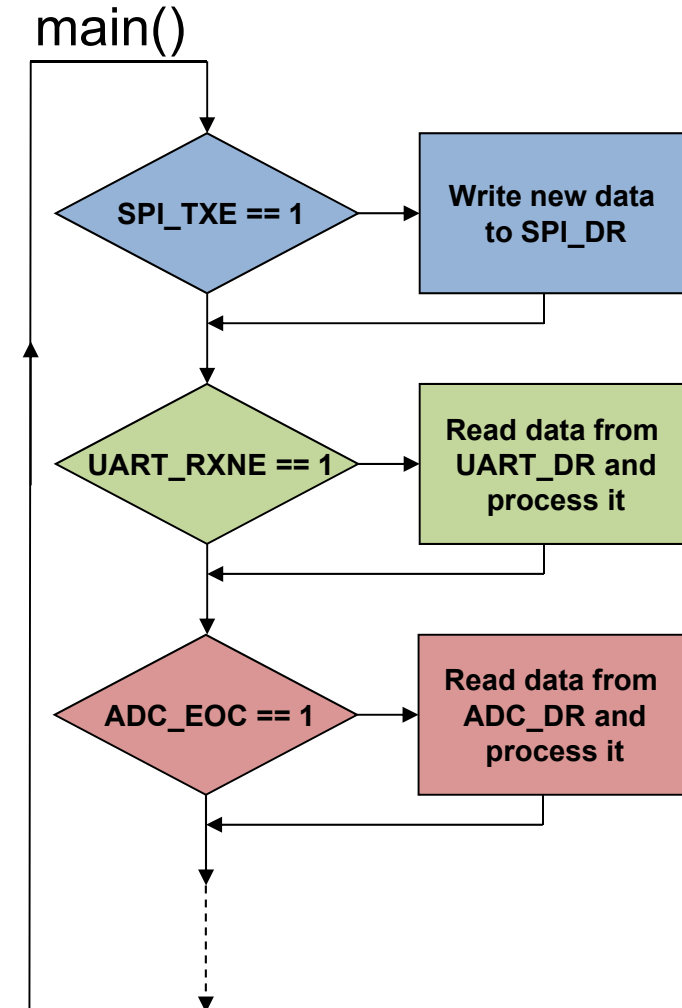
SPI
UART
GPIO
ADC

Serial Peripheral Interface
Universal Asynchronous Receiver Transmitter
General Purpose Input Output
Analog-Digital Converter

■ Periodic Query of Status Information

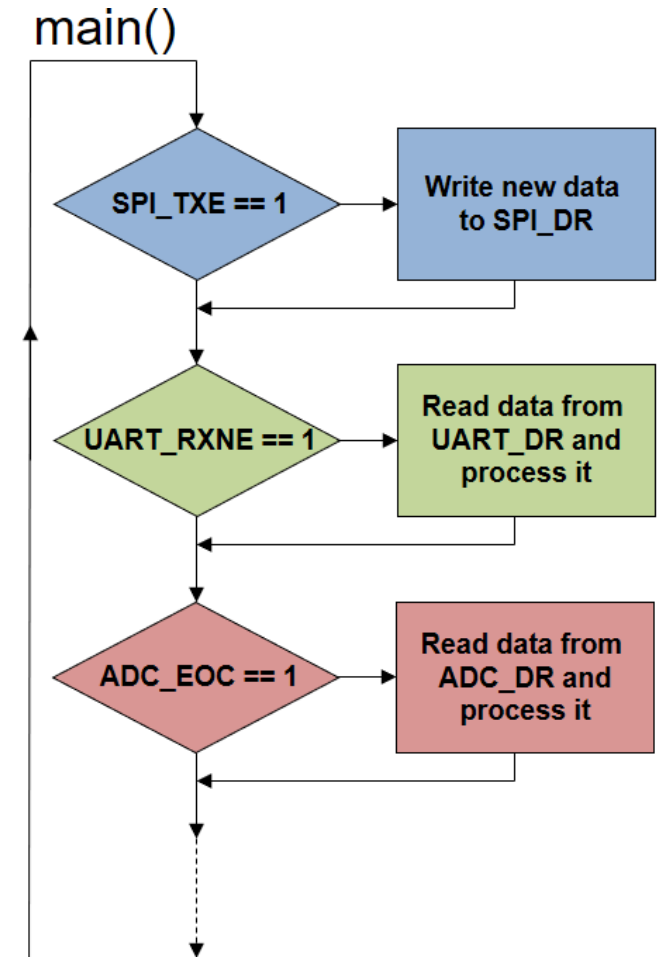
- Synchronous with main program
- Advantages
 - Simple and straightforward
 - Implicit synchronization
 - Deterministic
 - No additional interrupt logic required
- Disadvantages
 - Busy wait → wastes CPU time
 - Reduced throughput
 - Long reaction times
in case of many I/O devices or if the CPU is working on other tasks

1) to poll → abfragen

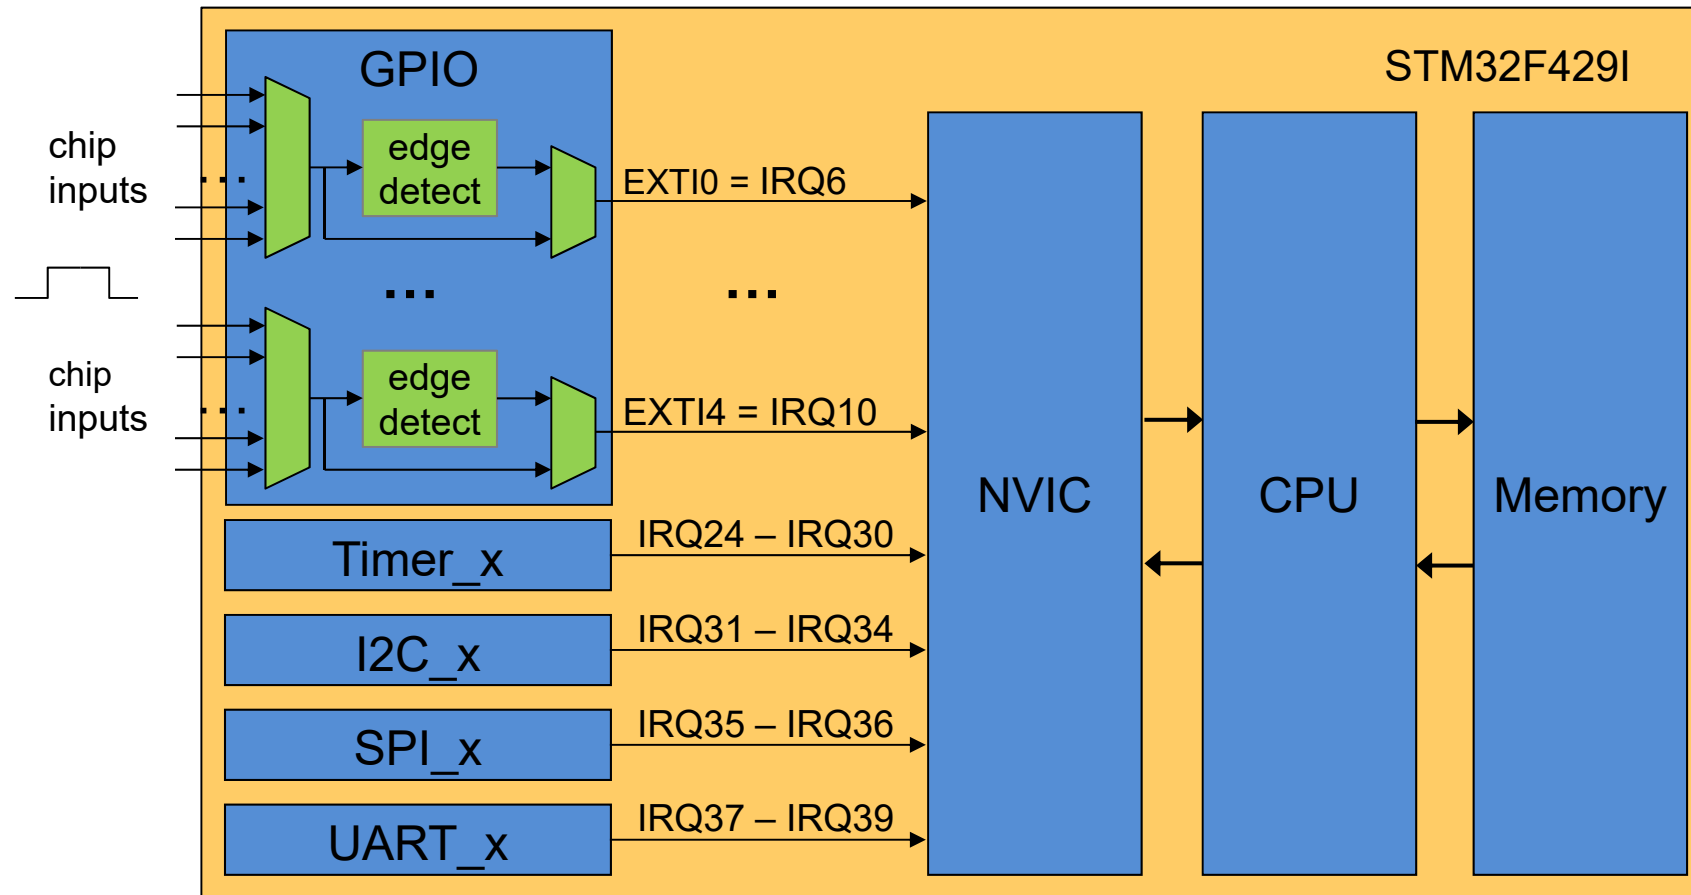


■ Implementation

```
while(1){  
    if (spi_is_txe_set()) {  
        ...  
        spi_write_data(...);  
    }  
  
    if (uart_is_rxne_set()) {  
        uart_data = uart_read_data();  
        ...  
    }  
  
    if (adc_is_eoc()) {  
        adc_data = adc_read_data();  
        ...  
    }  
  
    ...  
}
```



■ Interrupt System STM32F429I



■ Vector Table and ISR

```
; Vector Table Mapped to Address 0 at Reset
AREA RESET, DATA, READONLY
EXPORT __Vectors
EXPORT __Vectors_End
EXPORT __Vectors_Size

__Vectors DCD __initial_sp          ; Top of Stack
          DCD Reset_Handler        ; Reset Handler
          DCD NMI_Handler          ; NMI Handler
          DCD HardFault_Handler    ; Hard Fault Handler
          DCD MemManage_Handler    ; MPU Fault Handler
          DCD BusFault_Handler     ; Bus Fault Handler
          DCD UsageFault_Handler   ; Usage Fault Handler
          DCD 0                    ; Reserved
          DCD 0                    ; Reserved
          DCD 0                    ; Reserved
          DCD 0                    ; Reserved
          DCD SVC_Handler          ; SVCall Handler
          DCD DebugMon_Handler     ; Debug Monitor Handler
          DCD 0                    ; Reserved
          DCD PendSV_Handler       ; PendSV Handler
          DCD SysTick_Handler      ; SysTick Handler

; External Interrupts
DCD WWDG_IRQHandler              ; Window WatchDog
DCD PVD_IRQHandler              ; PVD through EXTI Line detection
DCD TAMP_STAMP_IRQHandler        ; Tamper and TimeStamps through the EXTI line
DCD RTC_WKUP_IRQHandler         ; RTC Wakeup through the EXTI line
DCD FLASH_IRQHandler            ; FLASH
DCD RCC_IRQHandler              ; RCC
DCD EXTI0_IRQHandler            ; EXTI Line0
DCD EXTI1_IRQHandler            ; EXTI Line1
DCD EXTI2_IRQHandler            ; EXTI Line2
DCD EXTI3_IRQHandler            ; EXTI Line3
```

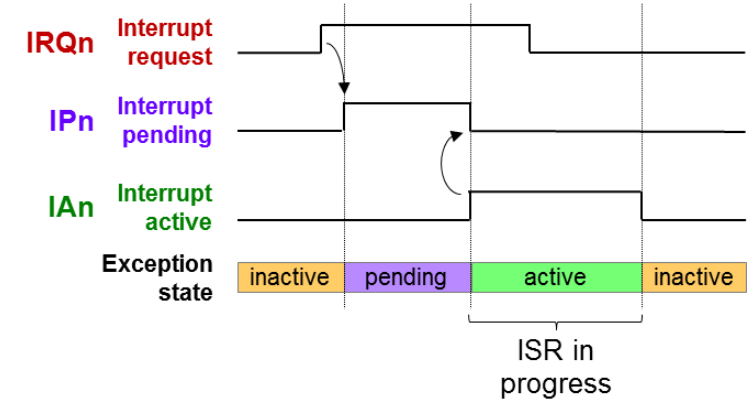
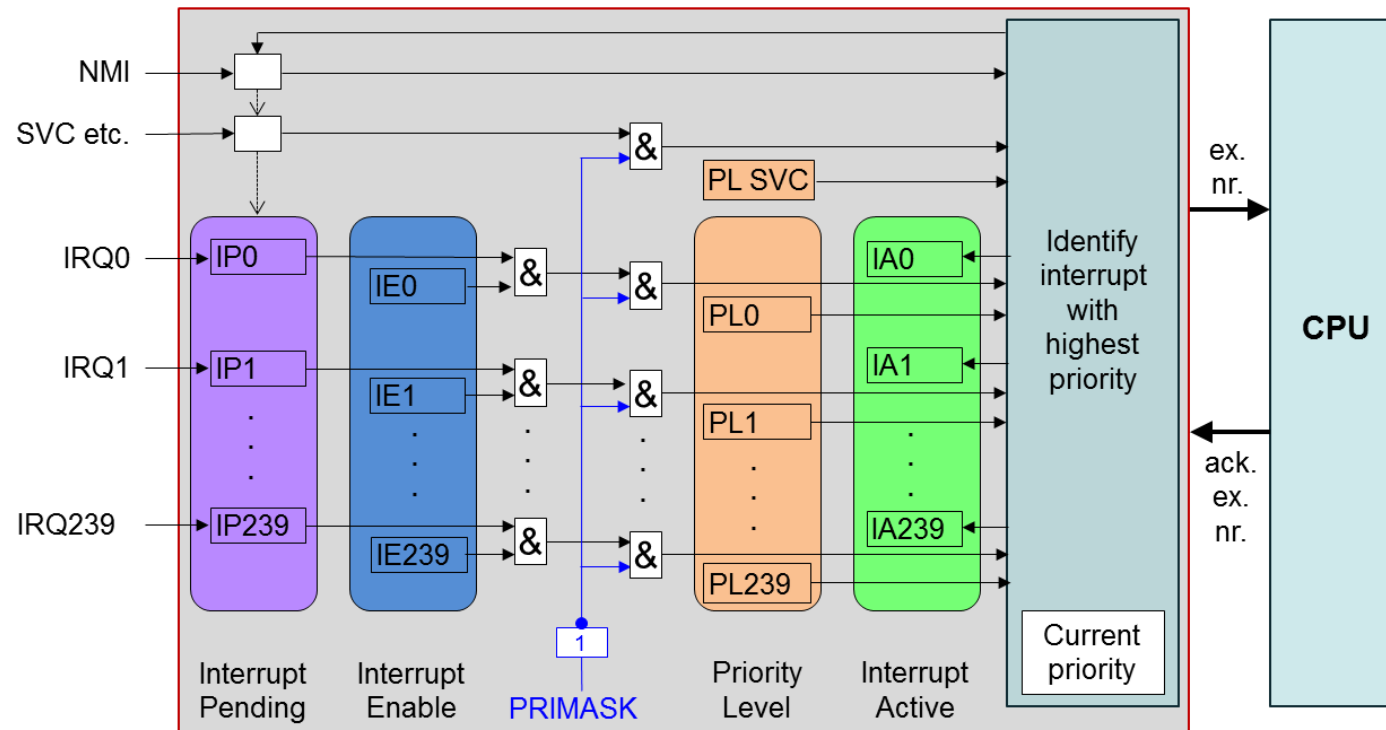
```
/*
 * Interrupt service routines
 * -----
 */

void EXTI0_IRQHandler(void)
{
    // clear flag in status register of peripheral

    // service interrupt
}
```

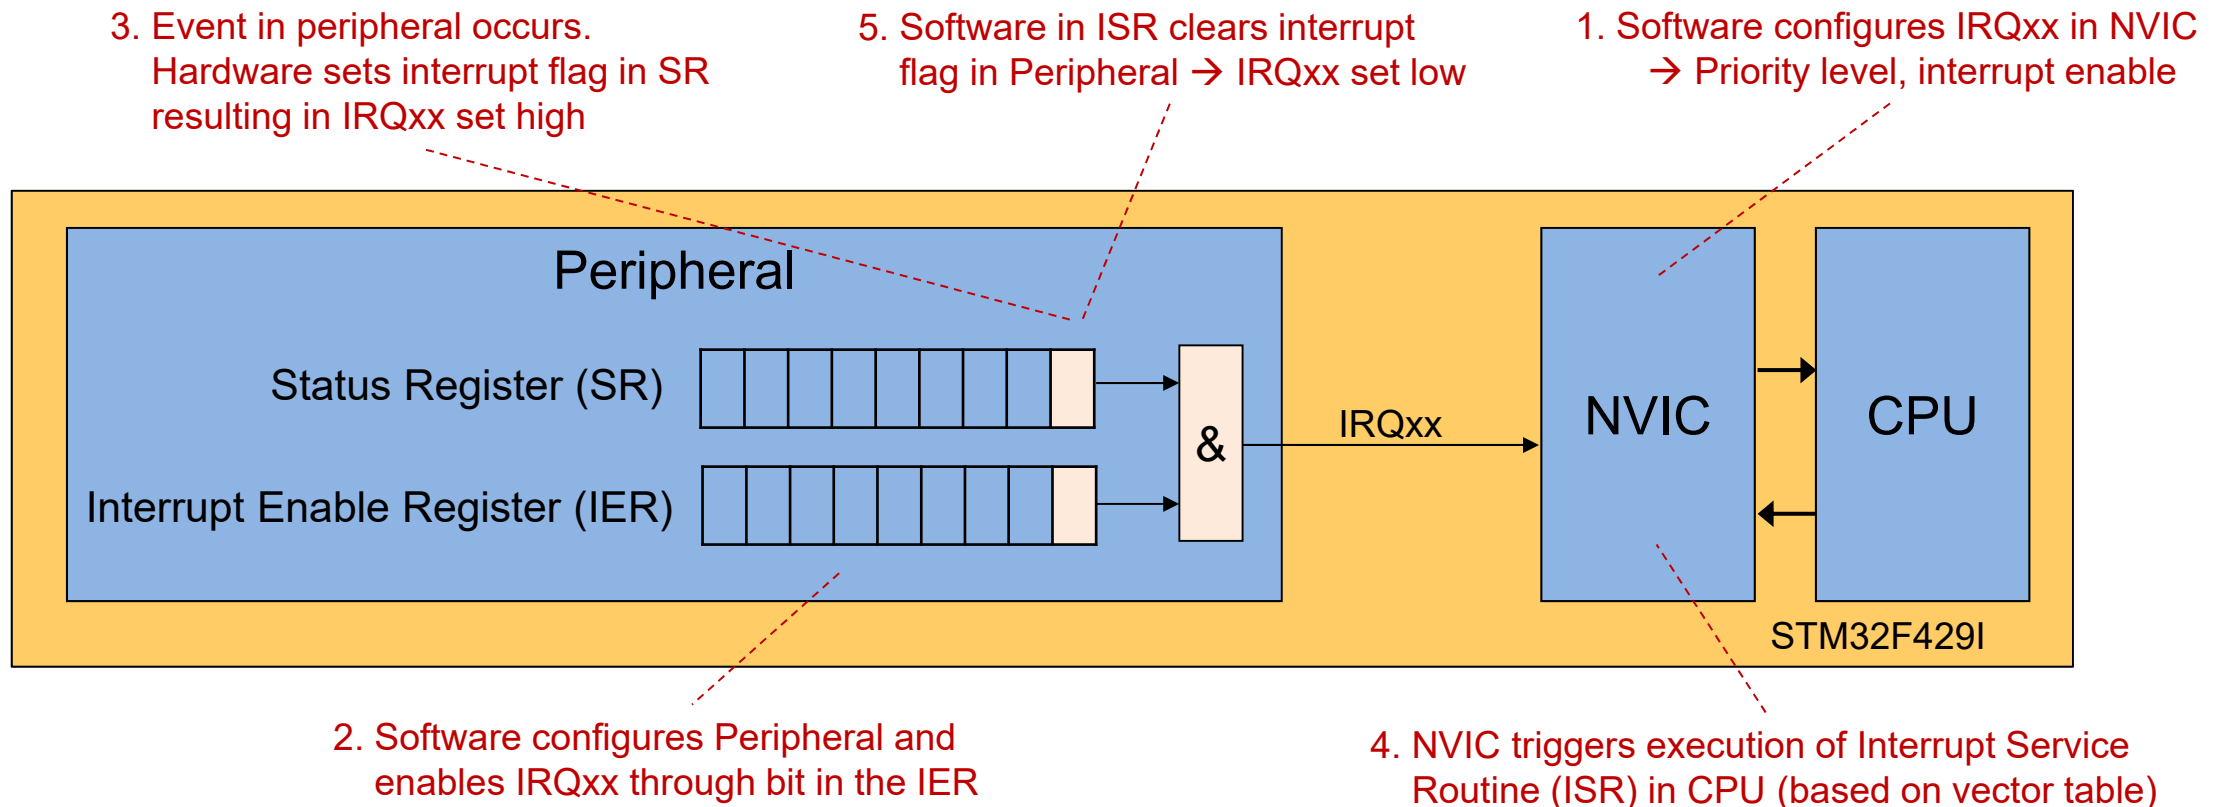
■ NVIC on Cortex-M

- Nested Vectored Interrupt Controller



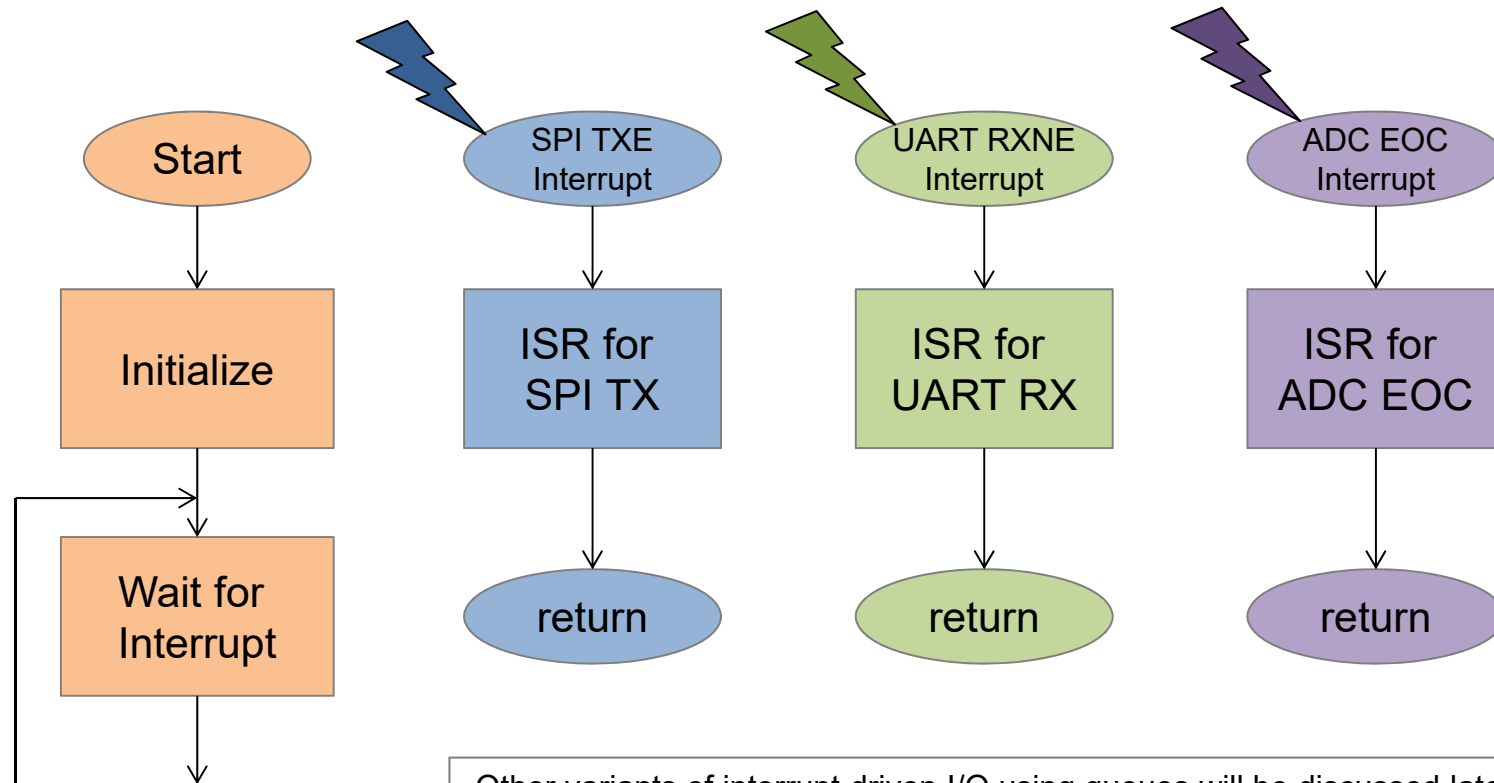
■ Interrupt flags in Status Registers (SR) are hardware set and software reset

- Software has to enable IRQxx line in peripheral



■ Interrupt driven I/O

- Extreme case: All processing done in interrupt service routines





■ Interrupt frequency

- How often does an interrupt occur
- Varies from source to source, e.g.

f_{INT}

Keyboard

→ max. ~ 20 interrupts per second in irregular intervals



Serial interface

e.g. 230'400 baud

→ one interrupt every 8 bit

$230'400 / 8 = 28'800$ interrupts per second

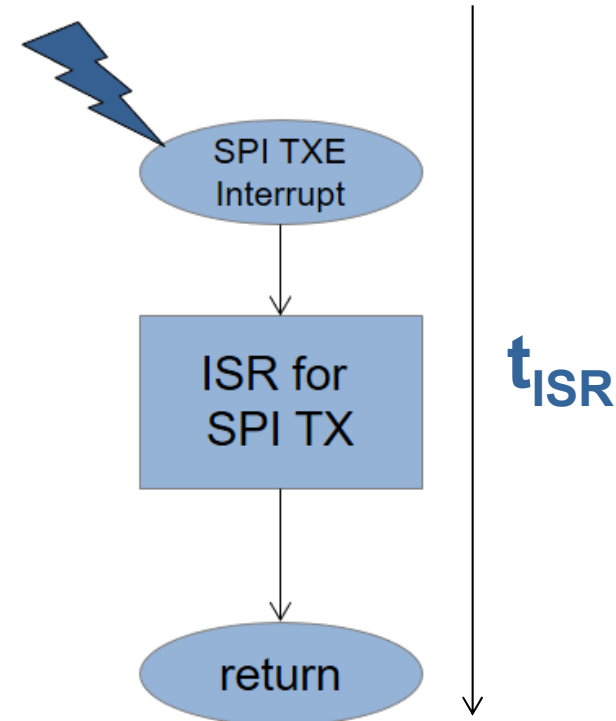




■ Interrupt service time

t_{ISR}

- Required time to process an interrupt
 - I.e. execution time of ISR
- Depends on
 - Number of instructions in ISR
 - Required number of clock cycles per instruction
→ depends on CPU architecture
 - CPU clock frequency
 - Time for switching to and returning from ISR





■ Impact on system performance

- Percentage of CPU time used to service interrupts

$$\text{Impact} = f_{\text{INT}} * t_{\text{ISR}} * 100 \%$$

- Example keyboard

$$\begin{aligned} f_{\text{INT}} &= 20 \text{ Hz} = 20 \frac{1}{s} & t_{\text{ISR}} &= 6 \text{ us } ^{1)} \\ \text{Impact} &= 20 \text{ Hz} * 6 \text{ us} * 100 \% = 0.012 \% \end{aligned}$$

- Example serial interface with 230'400 Baud


$$\begin{aligned} f_{\text{INT}} &= 230'400 / 8 = 28'800 \text{ Hz} & t_{\text{ISR}} &= 6 \text{ us } ^{1)} \\ \text{Impact} &= 28'800 \text{ Hz} * 6 \text{ us} * 100 \% = 17.3 \% \end{aligned}$$

¹⁾ = assumed value: 6 us \approx 1000 cycles @ 168 MHz on CT Board



■ $t_{\text{ISR}} > \text{"Time between two interrupt events"}$

- Some interrupt events will not be serviced (lost)
 - Data will be lost
- f_{INT} as well as t_{ISR} may vary over time
 - Average may be ok, but individual interrupt events may still be lost



Caution in case of several interrupt sources

- Interrupts can occur simultaneously
- Computing power required for both ISRs




■ Strive for short ISRs

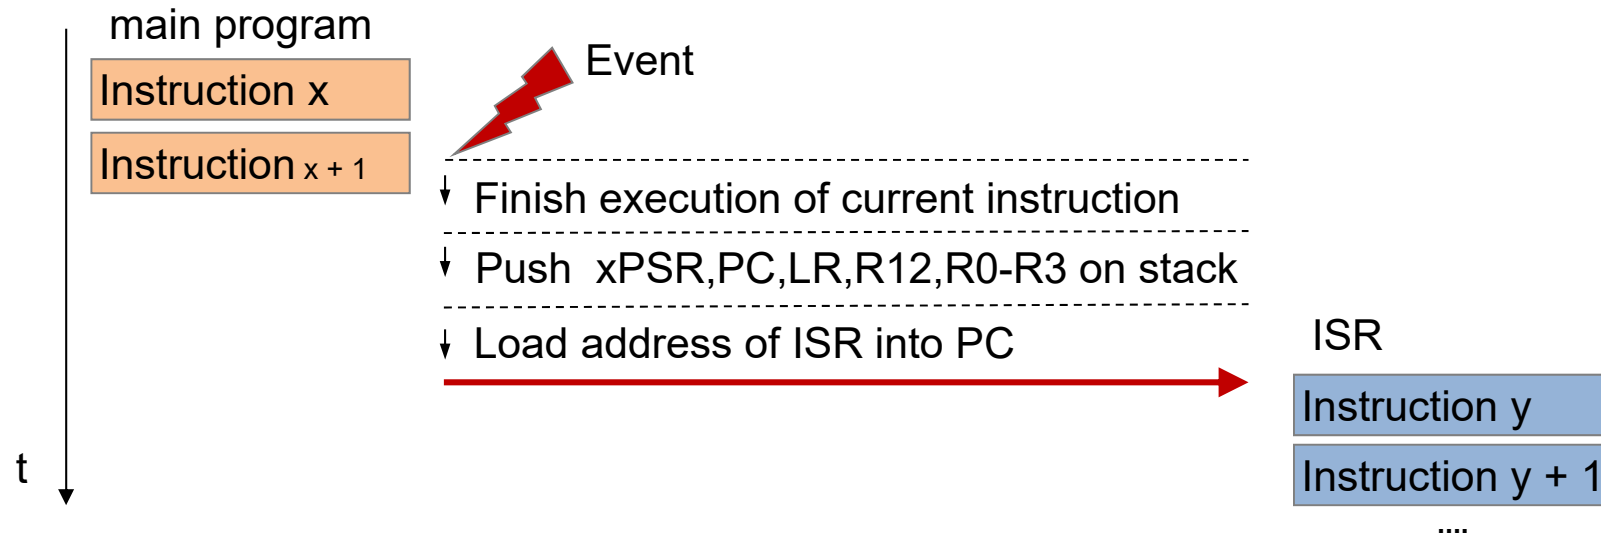
- Simpler debugging
- Execute only time-critical tasks in ISR
 - Move tasks with relaxed time-constraints to main loop
 - Makes ISR available for other time-critical tasks
 - Often simply feed interrupt to queue; dispatch queue in main loop

■ Interrupt latency definition

- Time between interrupt event and start of servicing by ISR
 - How long does it take until first 'useful' instruction in ISR is executed
- Range
 - From about 50 nanoseconds (ns) up to several milliseconds (ms)
- Relevant in cases where guaranteed service times are required
 - E.g. audio/video streaming




■ Latency influenced by hardware (CPU)



- Different instructions may have different execution times
- Multi-cycle instructions on Cortex-M3/M4
 - Some (e.g. SDIV/UDIV) are abandoned and restarted after ISR
 - Some (e.g. LDM/STM) are interrupted and resumed after ISR

■ Latency influenced by software (code)

- Saving additional registers on stack
- Process ongoing or higher prioritized ISRs
- Masked (disabled) interrupts → CPSID i / CPSIE i
- In case several sources are using the same interrupt line
 - SW has to use polling to know which peripheral requires servicing



Interrupt Latency

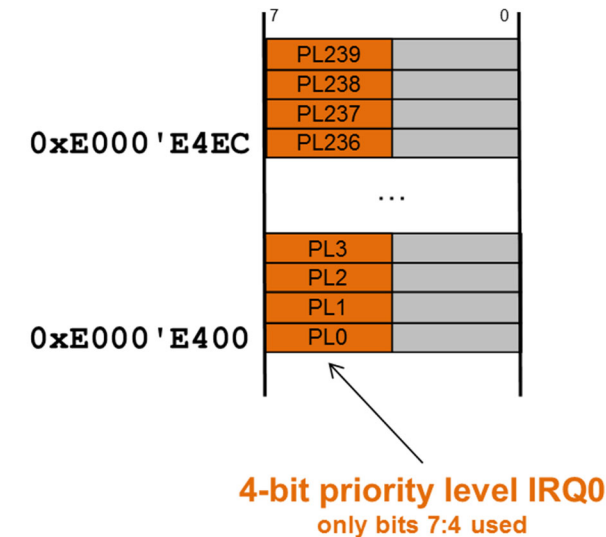
- **Required response time for a specific event**
 - Certain peripherals require fast response ($< 1\mu\text{s}$)
 - Real-time systems, e.g. anti-lock braking system
 - Cannot be guaranteed if maximum interrupt latency is too high

- **f_{INT} too high \rightarrow Too many interrupts**
 - No CPU cycles left for data processing
 - System is busy calling ISRs
 - Neither ISR nor main loop can process any data
 - Performance of CPU is used only for latencies / context switching

■ Fast response times – low latency

- Fast CPU
 - High clock rate
 - Low number of clock cycles per instruction
- Extremely short polling loops can be fast
- Pre-emption with appropriate priorities
 - Priority levels programmed in NVIC
 - Real Time OS

Priority Level Registers for
IRQ0 – IRQ239



source: Techopedia Inc.



Definition - What does *Pre-Emption* mean?


Pre-emption refers to the temporary interruption and suspension of a task, without asking for its cooperation, with the intention to resume that task later. This act is called a context switch and is typically performed by the pre-emptive scheduler, a component in the operating system authorized to pre-empt, or interrupt, and later resume tasks running in the system.

■ Example interrupt priorities

- ISR1 does not pre-empt ISR0
- ISR2 pre-empts ISR0

assuming

IRQ0	PL0 = 0x2	medium priority
IRQ1	PL1 = 0x3	lowest priority
IRQ2	PL2 = 0x1	highest priority

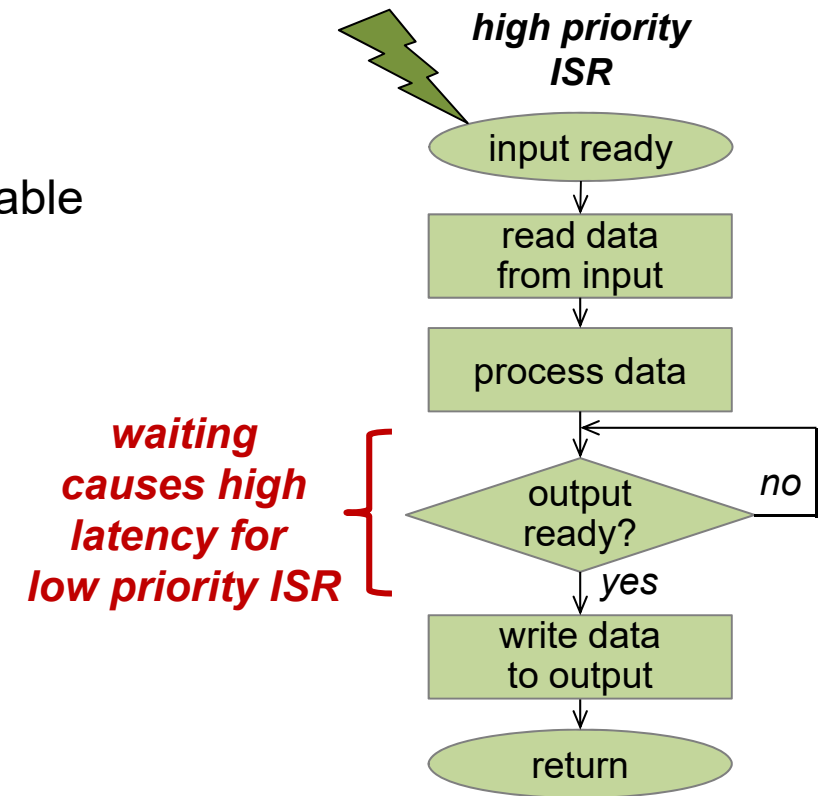


■ Latency consistency

- Some applications can tolerate considerable interrupt latency as long as it is consistent
 - I.e. same amount of latency from one interrupt event to next one
- Example: Measurements with periodic time intervals
 - Path measurements to detect whether an object is being accelerated, e.g. counting of step pulses on an incremental position encoder for a motor
 - Works only with a constant time interval

■ High priority interrupts may cause high latencies for lower priority interrupts

- Example for high priority ISR
 - Interrupt triggers reading of input data
 - Process data
 - Write to output
- Writing may have to wait for output device to become available
 - E.g. because SPI is still transmitting previous data



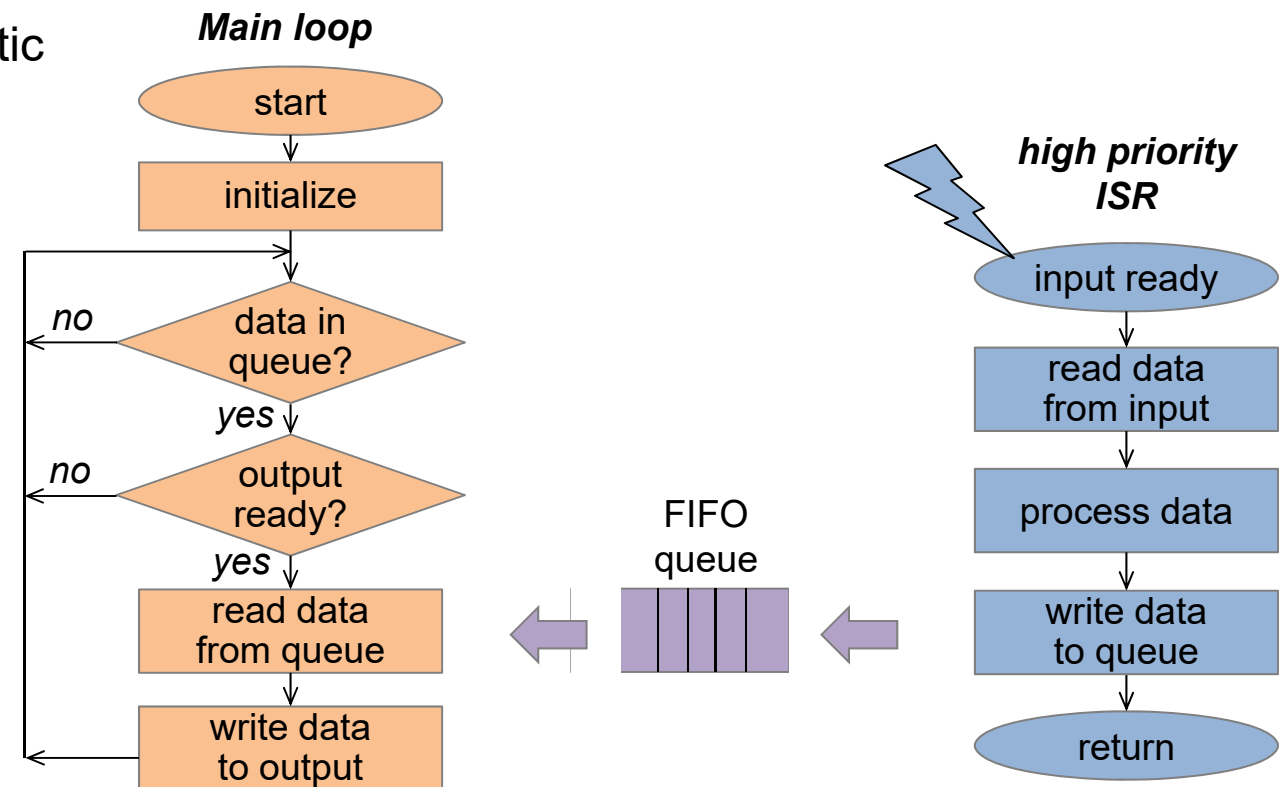
■ Remedy: Move 'waiting loop' to main program

- Use queue (FIFO) for decoupling
- Makes input ISR short and deterministic

Example for functions to enter and retrieve data from a queue

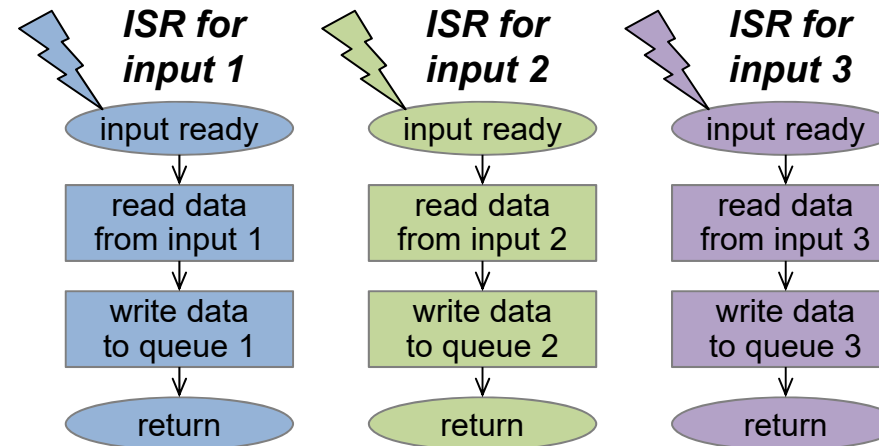

```
/**
 * \brief Enqueues data at the tail of the specified queue
 * \param queue Points to the queue to which the element should be added
 * \param data
 * \return 0 if data NOT added (queue is full),
 *         ~0 if data added (queue not full)
 */
uint32_t queue_enqueue(queue_t *queue, uint32_t data);

/**
 * \brief Grabs and removes the element at the head of the specified queue
 * \param queue Points to the queue from which the element should be polled
 * \return element at the head of queue, returns 0 if no element in queue.
 */
uint32_t queue_dequeue(queue_t *queue);
```



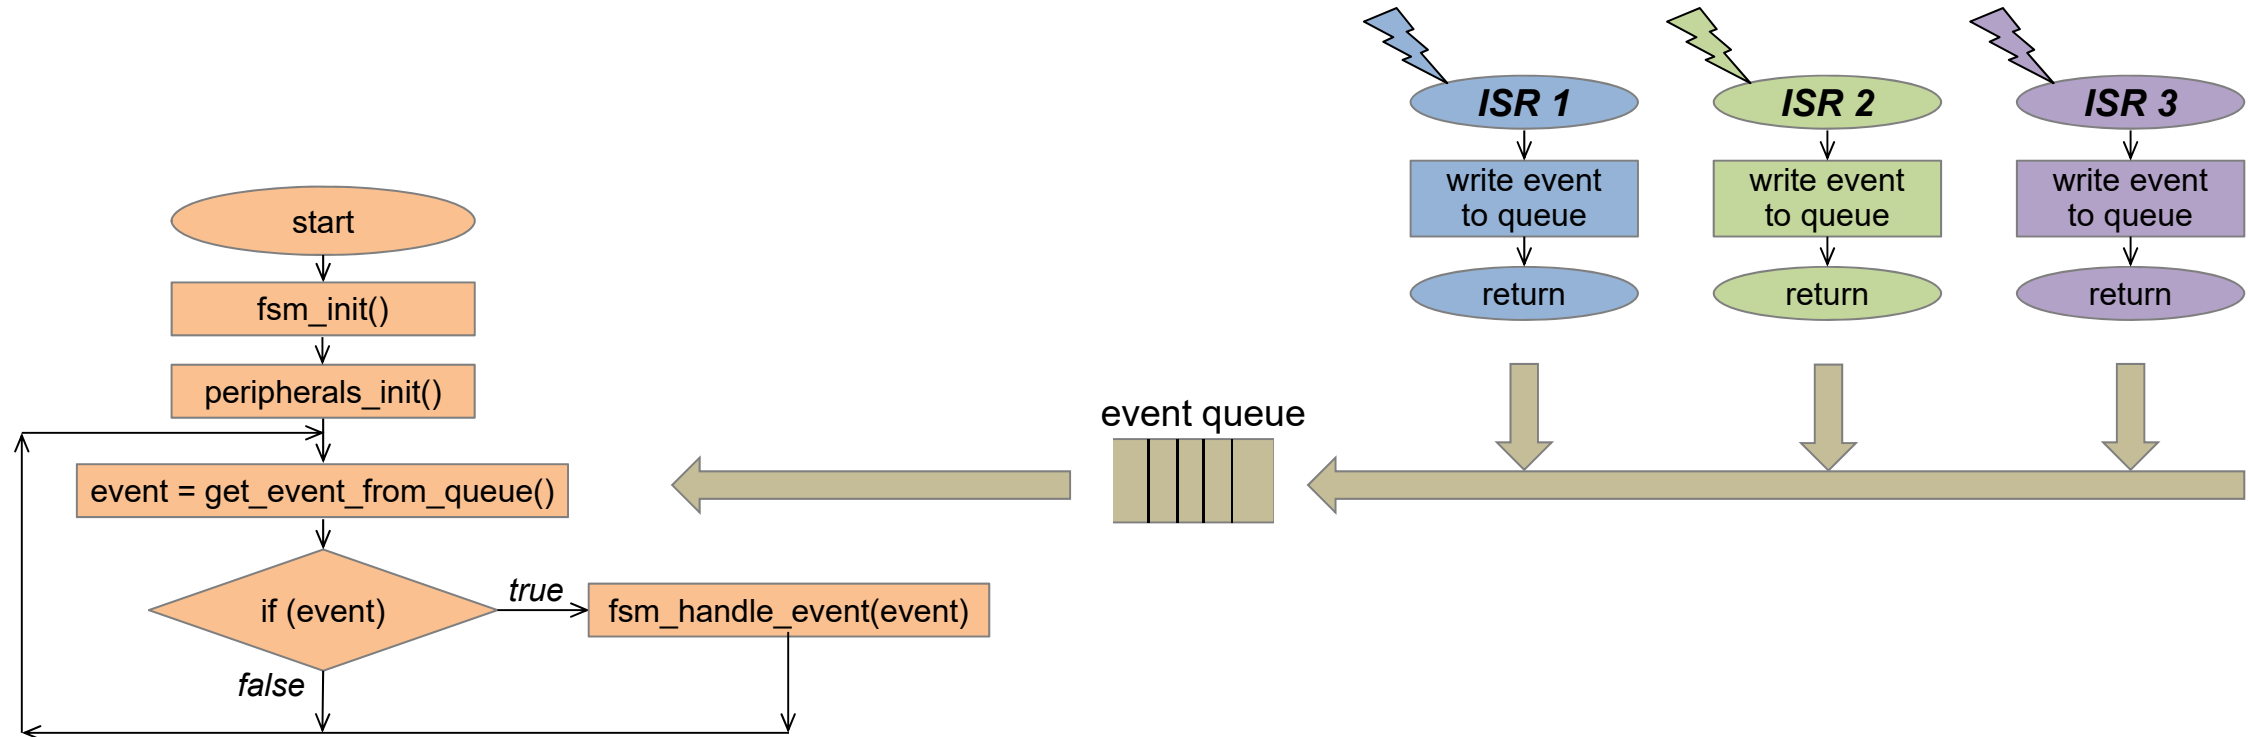
■ Move non-time-critical work from ISRs to main loop

- Several ISRs write input data into their dedicated queues
- Main loop checks all queues and processes their contents
- May result in many tasks in main loop
 - Not all processing tasks have the same priority
 - Requires scheduling of tasks
 - May require real-time OS in case of many ISRs



■ Events Trigger Interrupts

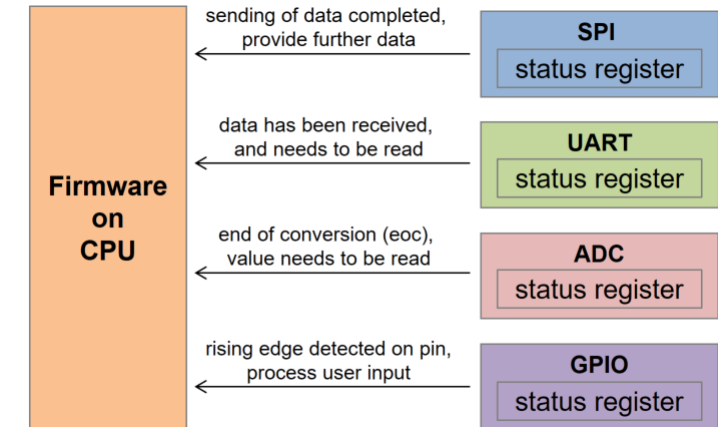
- Several Interrupt Service Routines (ISR) enter events into queue
- FSM in main loop reads events from queue



■ Detection of Events → Polling vs. Interrupt Driven I/O

■ Interrupt Performance

- Interrupt frequency f_{INT}
- Interrupt service time t_{ISR}
- Impact = $f_{INT} * t_{ISR} * 100\%$

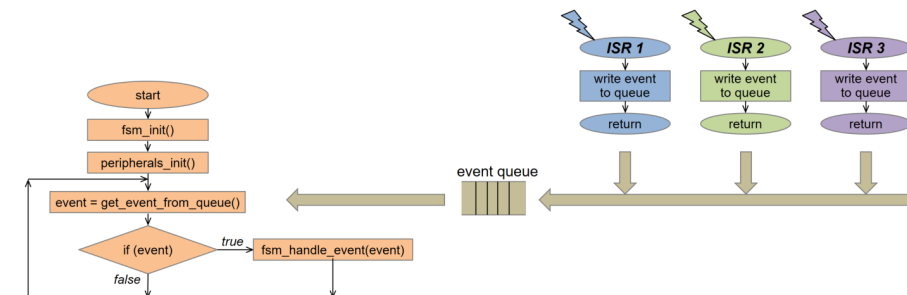


■ Interrupt Latency

- How fast does an application require a response?
- Pre-emption: High priority interrupts may cause high latencies for lower priority interrupt
- Move 'waiting loops' and non-time-critical work from ISR to main loop

■ Interrupt Driven FSM

- Interrupt Service Routines (ISR) enter events into queue
- FSM in main loop reads events from queue



■ Interrupts In General

- "The Art of Assembly Language Programming"
 - by Randall Hyde, Chapter 17
- "Fundamentals of Embedded Software with the ARM Cortex-M3"
 - by Daniel W. Lewis, Chapter 9

■ Interrupt Latency

- <http://www.segger.com/interrupt-latency.html>
- <http://www.ganssle.com/articles/interruptlatency.htm>