

Information Engineering 2

Data Storage: HDFS & Parquet

Prof. Dr. Kurt Stockinger

Semesterplan

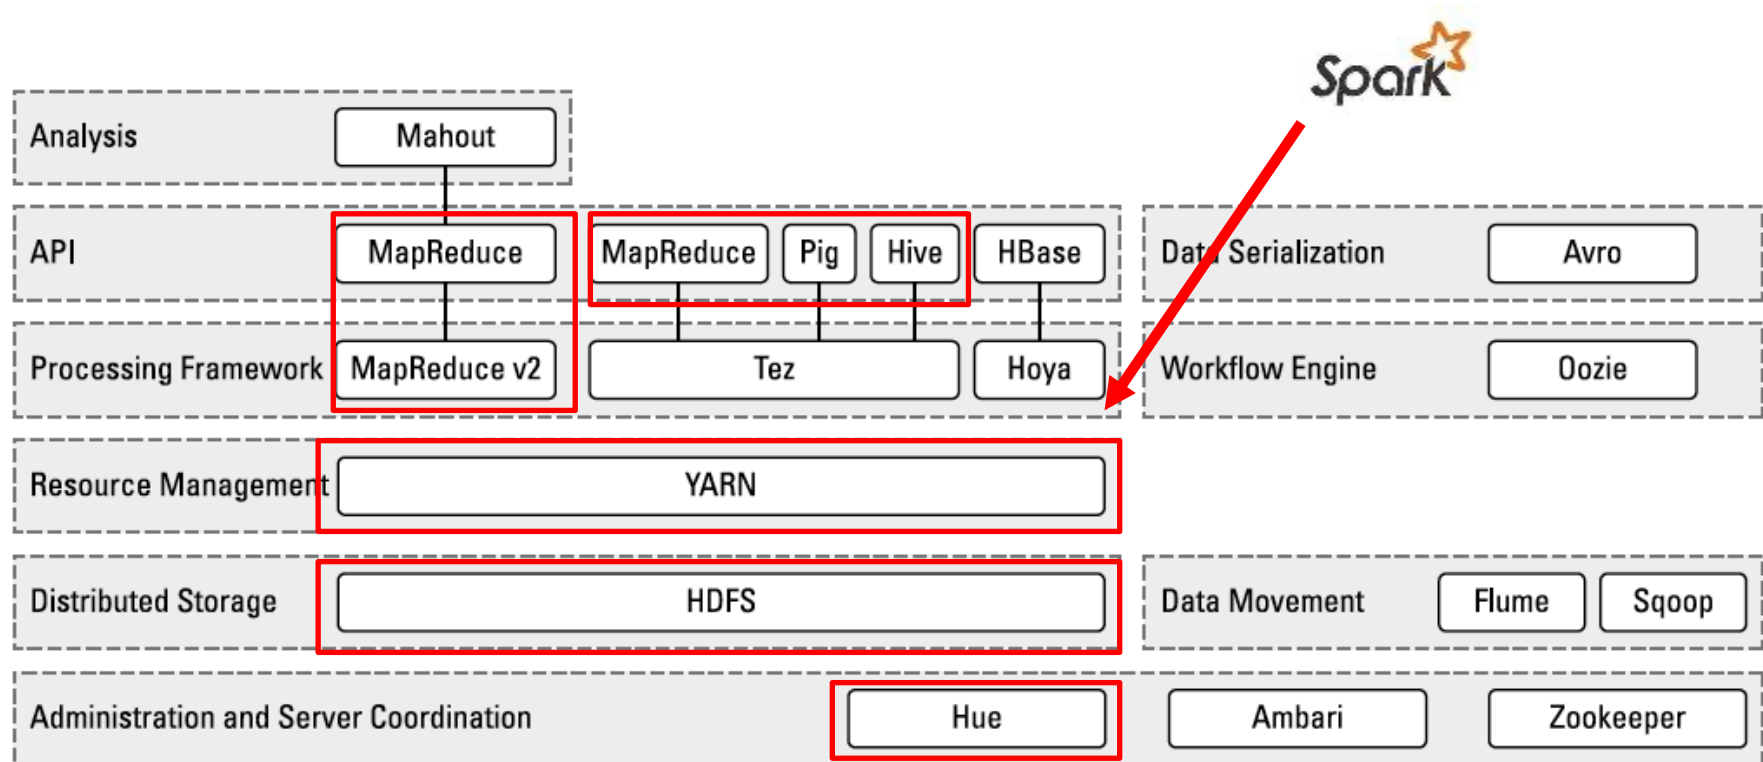
SW	Datum	Vorlesungsthema	Praktikum
1	23.02.2022	Data Warehousing Einführung	Praktikum 1: KNIME Tutorial
2	02.03.2022	Dimensionale Datenmodellierung 1	Praktikum 1: KNIME Tutorial (Vertiefung)
3	09.03.2022	Dimensionale Datenmodellierung 2	Praktikum 2: Datenmodellierung
4	16.03.2022	Datenqualität und Data Matching	Praktikum 3: Star-Schema, Bonus: Praktikum 4: Slowly Changing Dimensions
5	23.03.2022	Big Data Einführung	DWH Projekt - Teil 1
6	30.03.2022	Spark - Data Frames	DWH Projekt - Teil 2 (Abgabe: 4.4.2022 23:59:59)
7	06.04.2022	Data Storage: Hadoop Distributed File System & Parquet	Praktikum 1: Data Frames
8	13.04.2022	Query Optimization	Praktikum 2: Data Storage
9	20.04.2022	Spark Best Practices & Applications	Praktikum 3: Query Optimization & Performance Analysis
10	27.04.2022	Machine Learning mit Spark 1	Praktikum 3: Query Optimization & Performance Analysis (Vertiefung)
11	04.05.2022	Machine Learning mit Spark 2 + Q&A	Praktikum 4: Machine Learning (Regression)
12	11.05.2022	NoSQL Systems	Big Data Projekt - Teil 1
13	18.05.2022	Keine Vorlesung (Arbeit am Projekt)	Big Data Projekt - Teil 2
14	25.05.2022	Keine Vorlesung (Arbeit am Projekt)	Big Data Projekt - Teil 3 (Abgabe: 30.5.2022 23:59:59)

Learning Objectives for Today

- Learn about concepts of **Hadoop Distributed File System (HDFS)**
- Understand data **distribution and replication**
- Learn about **Parquet** file storage
- Understand **performance implications** of using Parquet storage

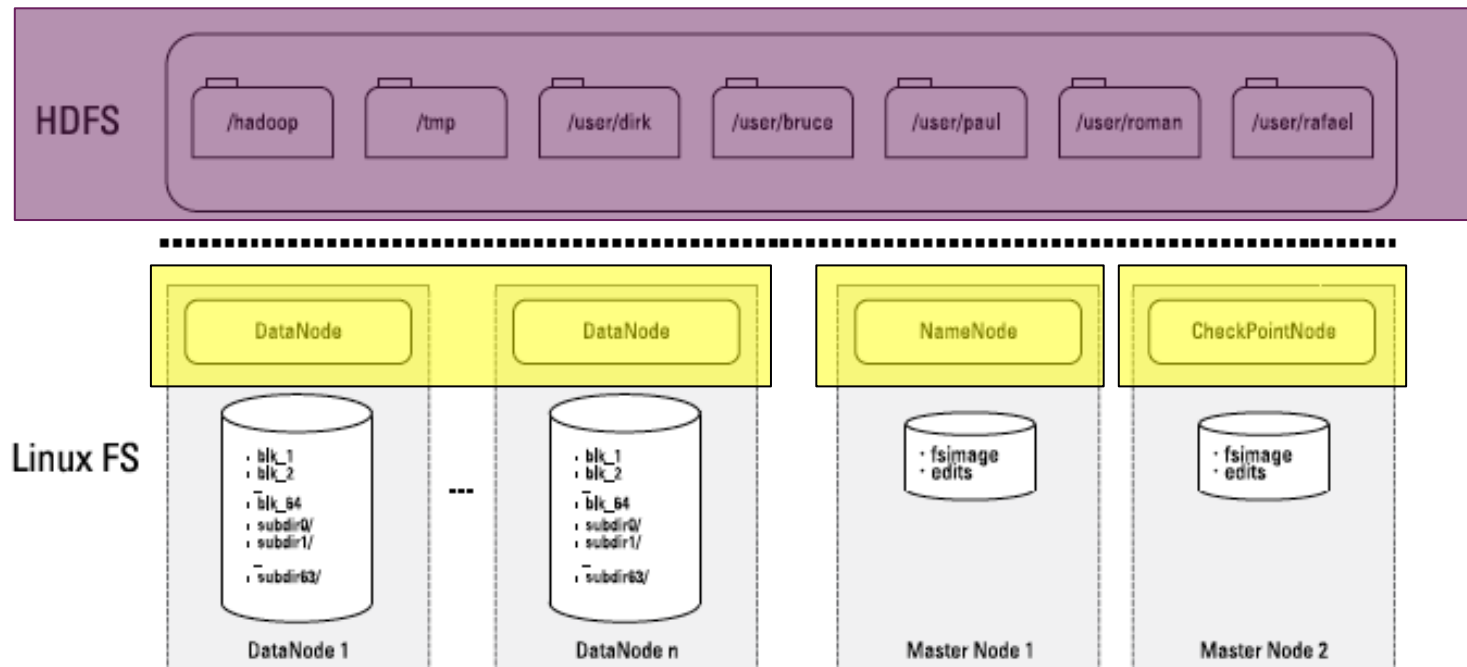
Hadoop

Architecture of Hadoop



Hadoop Distributed File System

- **Parallel file system** for managing large amounts of data
- Runs on top of Linux file system
- Data is distributed onto n machines



Managing Data in HDFS

- Starting situation:
 - Large log file about click streams (access to web pages)
 - HDFS with 4 machines
- How should the file be distributed onto the 4 machines?




Starting Situation


- File of 513 MB



Distribution of the Data Blocks onto Machines



Distribution of the Data Blocks onto Machines





- HDFS block size: typically 128 MB
- Data blocks are distributed irrespective of content
- Goal:
 - Even distribution of blocks to yield highest possible parallelization factor
- Conflict?
 - Isn't it better to consider data content for distribution?



Distribution of Data Blocks to Map-Tasks

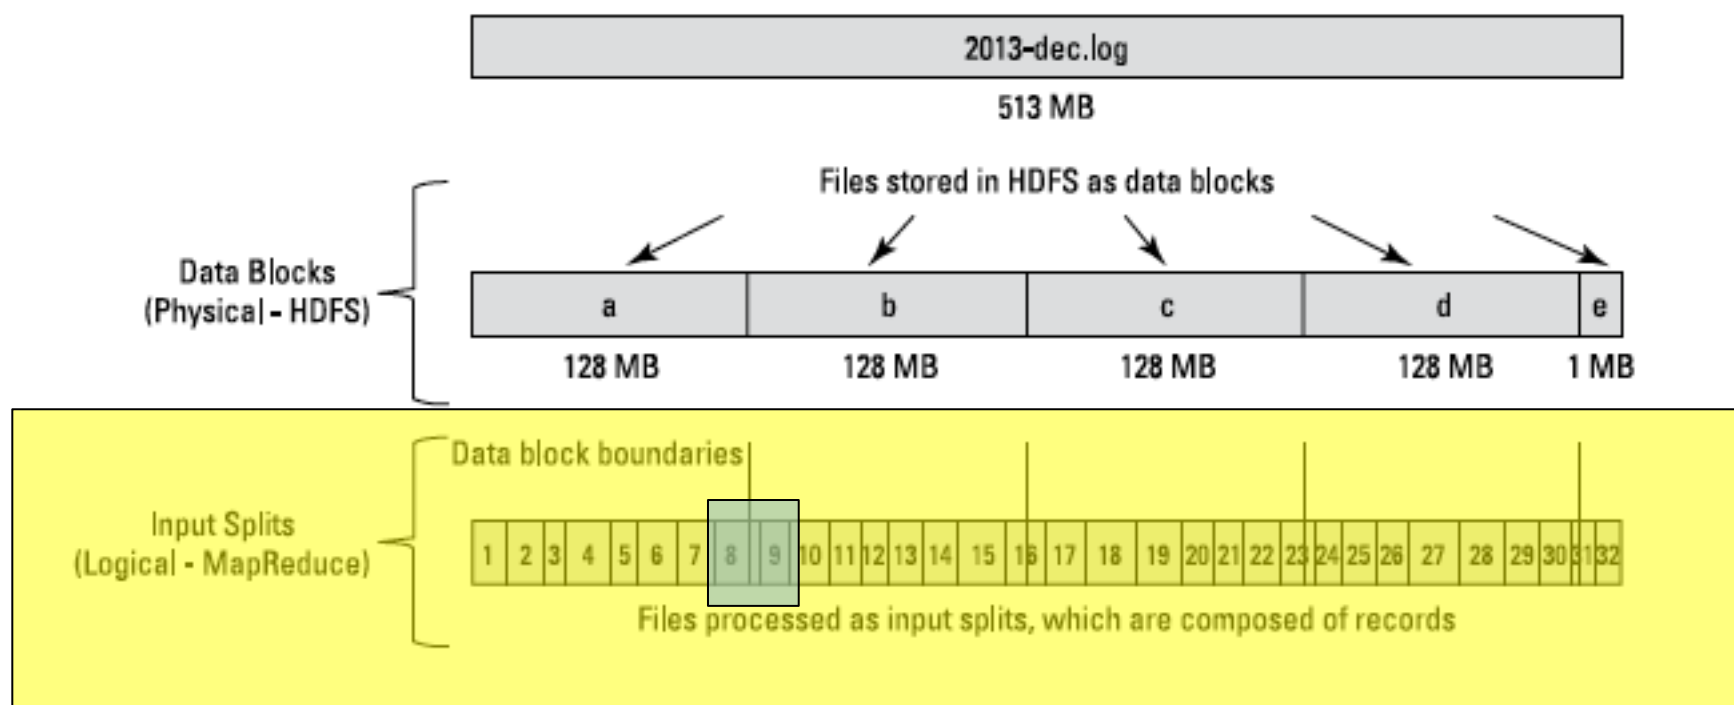
- Input data: "... This is a very long text with a size of 513 MB. ... "



- How are the blocks distributed to mappers?
- Are words split (e.g. "text")?

Distribution of Data Blocks to Map-Tasks

In MapReduce data is read record wise (key/value)



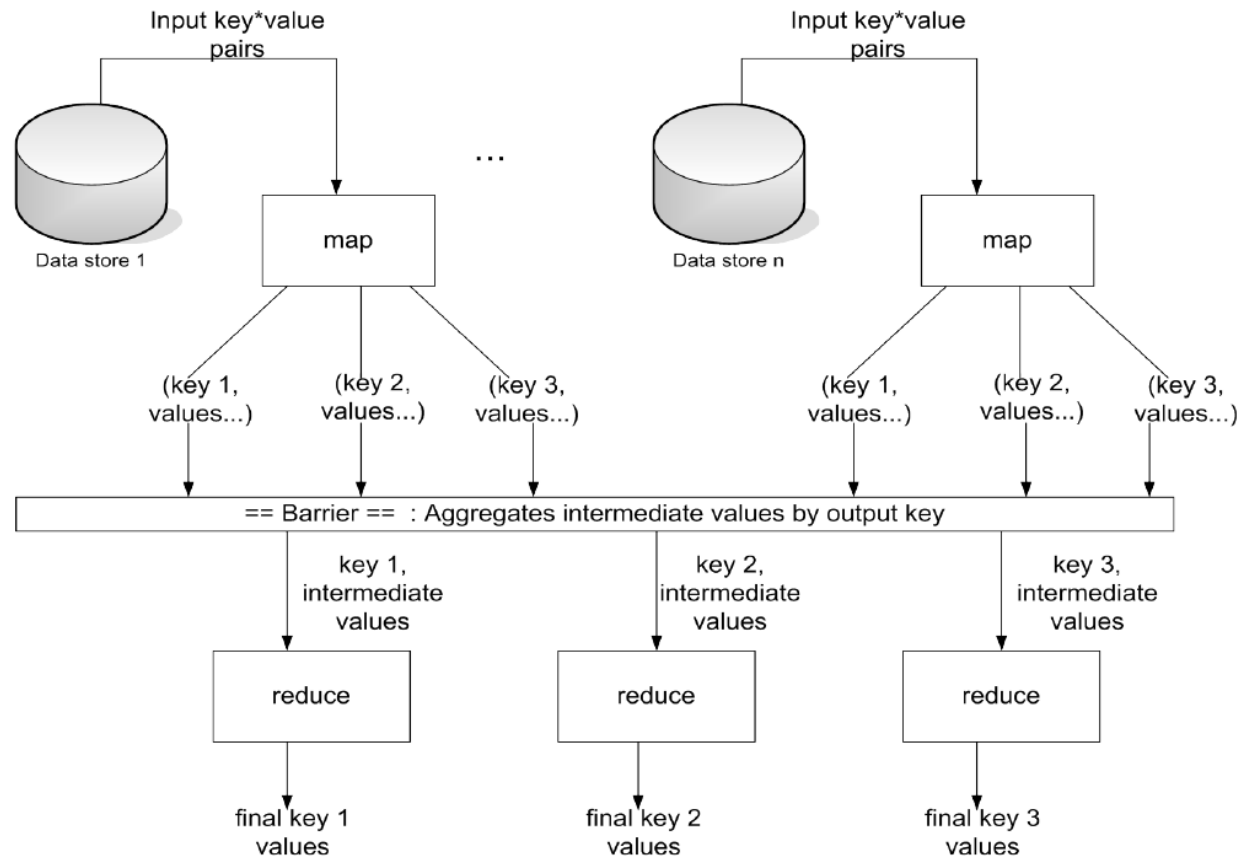
map: (K1, V1) → list(K2, V2)

reduce: (K2, list(V2)) → list(K3, V3)

Mapper - Details

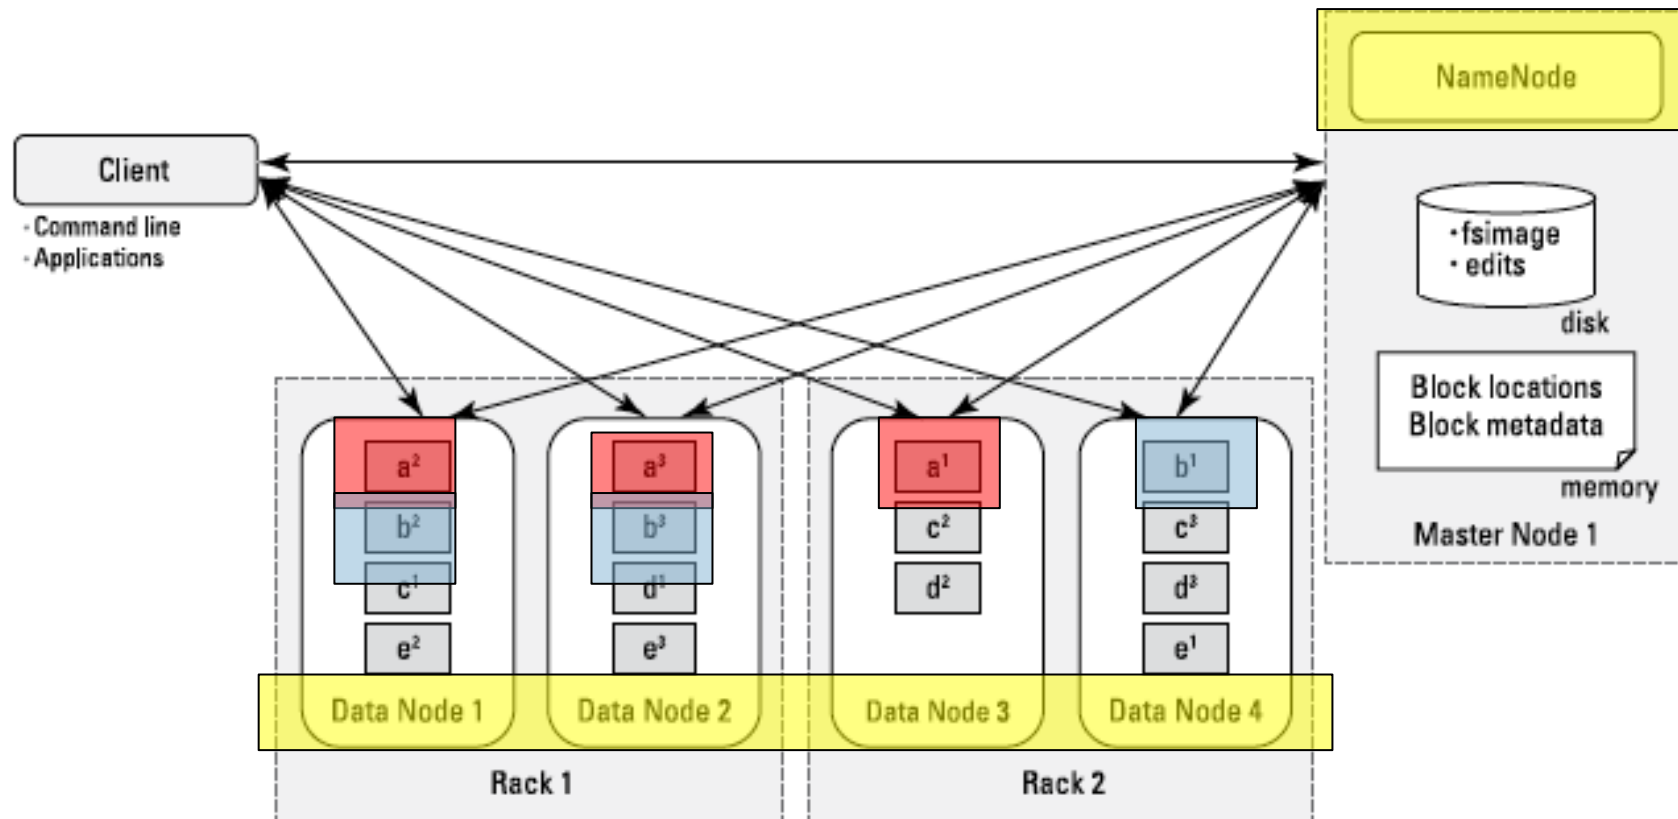
- Data is read **line by line**
- One **mapper task** per input block
- **Logical separation** if data is larger than size of input block:
 - No separation of single fields

Processing Model of MapReduce



Source: Big Data Vorlesung, Kossmann, Tatbul 2012

Data Distribution and Replication

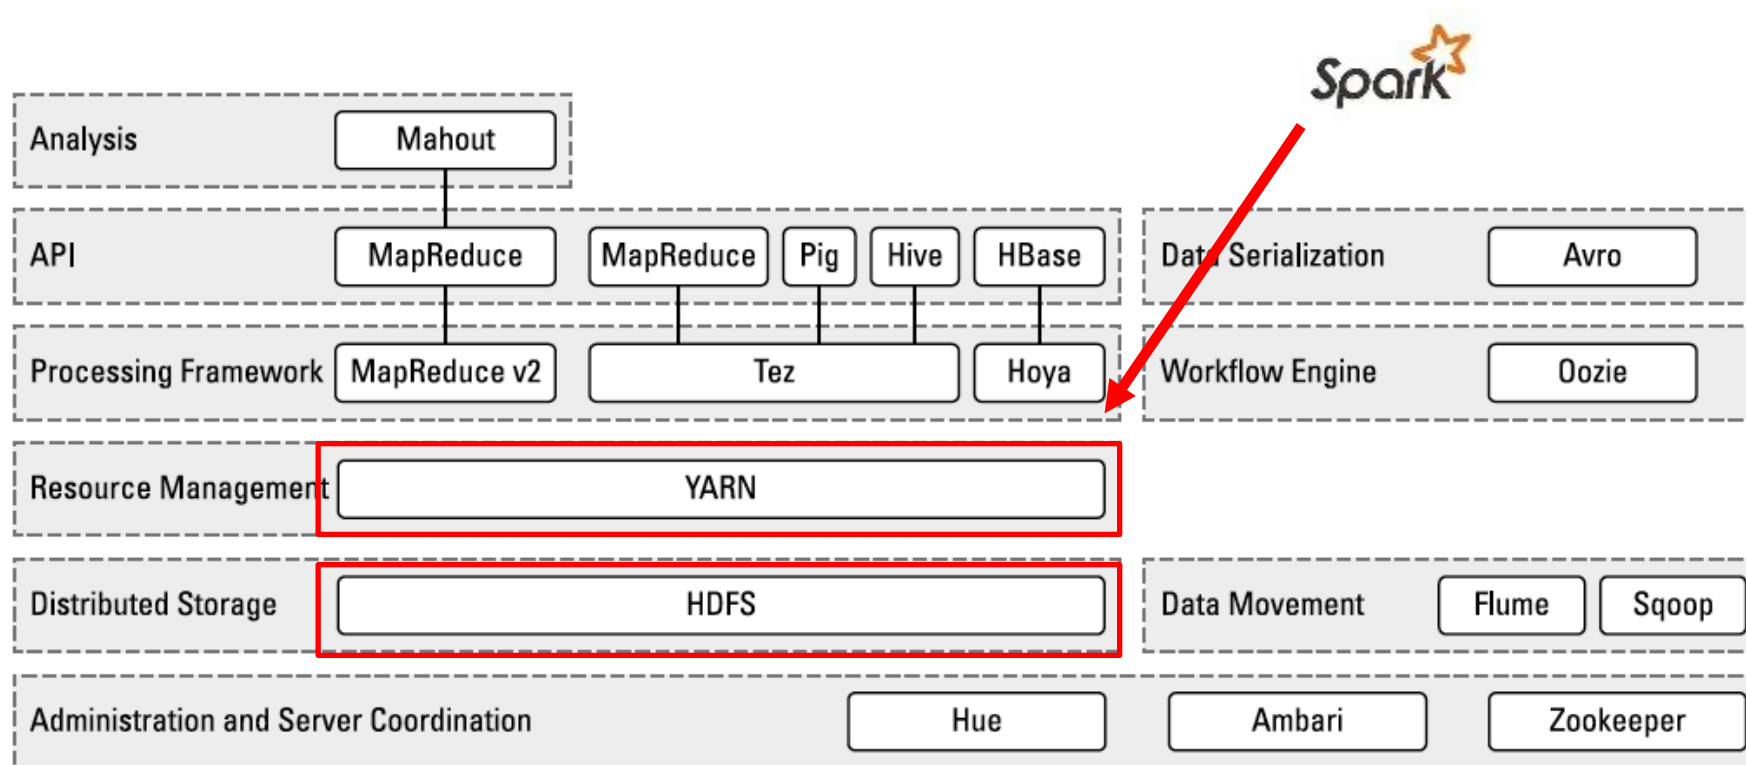


Default replication factor: 3

Relevance for HDFS

- HDFS is the basis for parallel data processing in Hadoop
- Can easily be combined with Spark
- Can also be used as a file system without using MapReduce (e.g. via Spark)

Reflection of Architecture of Hadoop: How does Spark fit in?



Most important HDFS commands

Copy Data from Linux to HDFS

- File: test.txt
- Copy from Linux to HDFS:
 - **hadoop fs -copyFromLocal** test.txt
- Result:
 - Data is copied to HDFS and distributed over all machines

Inspect HDFS Folders

- Show root directory of HDFS:
 - **hadoop fs -ls**
- Explanation:
 - ls ... list directory
- Result:

```
[root@hadoop-master ~]# hadoop fs -ls
```

```
Found 2 items
```

```
drwx----- - root hadoop
```

```
-rw-r--r-- 2 root hadoop
```

```
0 2015-04-01 18:22 .staging
```

```
15 2015-04-01 18:31 test.txt
```

readable and writable

replication factor

file size in bytes

Create Directory

- Command:
 - **hadoop fs -mkdir** output7
- Explanation:
 - mkdir ... make directory

- Result:

```
[root@hadoop-master ~]# hadoop fs -mkdir output7
[root@hadoop-master ~]#
[root@hadoop-master ~]# hadoop fs -ls
Found 3 items
drwx----- - root hadoop          0 2015-04-01 18:22 .staging
drwxr-xr-x - root hadoop          0 2015-04-01 18:36 output7
-rw-r--r-- 2 root hadoop        15 2015-04-01 18:31 test.txt
```

 **d... directory, rwx ... readable, writeable, executable**

Show Content of File

- Command:
 - **hadoop fs -cat** test.txt
- Explanation:
 - cat ... concatenate file and print on the standard output

```
[root@hadoop-master ~]# hadoop fs -cat test.txt  
Test Kurt 123.
```

Copy Data from HDFS to Linux

- Command:
 - **hadoop fs -copyToLocal** output/simulation1.txt simulation2015.txt

Help for Hadoop Commands

- General help:

- **hadoop**

```
[root@hadoop-master ~]# hadoop
Usage: hadoop [--config confdir] COMMAND
    where COMMAND is one of:
    fs                run a generic filesystem user client
    version           print the version
    jar <jar>         run a jar file
    checknative [-a|-h] check native hadoop and compression libraries
    distcp <srcurl> <desturl> copy file or directories recursively
    archive -archiveName NAME -p <parent path> <src>* <dest> create a
    classpath         prints the class path needed to get the
    credential        interact with credential providers
                     Hadoop jar and the required libraries
    daemonlog        get/set the log level for each daemon
or
    CLASSNAME        run the class named CLASSNAME
```

Most commands print help when invoked w/o parameters.

More Details

- Command:
 - **hadoop fs**

```
[root@hadoop-master ~]# hadoop fs
Usage: hadoop fs [generic options]
    [-appendToFile <localsrc> ... <dst>]
    [-cat [-ignoreCrc] <src> ...]
    [-checksum <src> ...]
    [-chgrp [-R] GROUP PATH...]
    [-chmod [-R] <MODE[,MODE]... | OCTALMODE> PATH...]
    [-chown [-R] [OWNER][:[GROUP]] PATH...]
    [-copyFromLocal [-f] [-p] <localsrc> ... <dst>]
    [-copyToLocal [-p] [-ignoreCrc] [-crc] <src> ... <localdst>]
```

Overview of Major Hadoop Commands #1

Command	Usage	Description
ls	Usage: <code>hadoop fs -ls <args></code> Example: <code>hadoop fs -ls /user/hive/warehouse</code>	For a file returns information on the file For a directory it returns list of its direct children as in unix
lsr	Usage: <code>hadoop fs -lsr <args></code> Example: <code>hadoop fs -ls /user</code>	Recursive version of ls. Similar to Unix <code>ls -R</code>
mkdir	Usage: <code>hadoop fs -mkdir <paths></code> Example: <code>hadoop fs -mkdir /user/ dir1</code>	Takes path uri's as argument and creates directory
moveFromLocal	Usage: <code>dfs -moveFromLocal <src> <dst></code>	Displays a "not implemented" message
mv	Usage: <code>hadoop fs -mv URI [URI ...] <dest></code> Example:	Moves files from source to destination. This command allows multiple sources as well in which case the destination needs to be a directory. Moving files across filesystems is not permitted.
put	Usage: <code>hadoop fs -put <localsrc> ... <dst></code> Example: <code>hadoop fs -put localfile1 localfile2 /user/hadoop/hadoopdir</code>	Copy single src, or multiple srcs from local file system to the destination filesystem. Also reads input from stdin and writes to destination filesystem.

Overview of Major Hadoop Commands #2

copyToLocal	Usage: <code>hadoop fs -copyToLocal [-ignorecrc] [-crc] URI <localdst></code> Example: <code>hadoop fs -copyToLocal /temp/file.txt /tmp</code>	Similar to get command, except that the destination is restricted to a local file reference
cp	Usage: <code>hadoop fs -cp URI [URI ...] <dest></code> Example: <code>hadoop fs -cp /user/hadoop/file1 /user/hadoop/file2 /user/hadoop/dir</code>	Copy files from source to destination. This command allows multiple sources as well in which case the destination must be a directory.
cat	Usage: <code>hadoop fs -cat URI [URI ...]</code> Example: <code>hadoop fs -cat hdfs://nn1.example.com/file1 hdfs://nn2.example.com/file2</code>	Copies source paths to stdout.
rm	Usage: <code>hadoop fs -rm URI [URI ...]</code> Example: <code>hadoop fs -rm hdfs://nn.example.com/file /user/hadoop/emptydir</code>	Delete files specified as args. Only deletes non empty directory and files. Refer to rmr for recursive deletes.
rmr	Usage: <code>hadoop fs -rmr URI [URI ...]</code> Example: <code>hadoop fs -rmr /user/hadoop/dir</code>	Recursive version of delete
du	Usage: <code>hadoop fs -du URI [URI ...]</code> Example:	Disk usage: Displays aggregate length of files contained in the directory or the length of a file in case its just a file.

Parquet Storage Format & Spark

How Shall We Store Data in a File/Table?

Title	Released	Label	PeakChart.UK	Certification.BVMI	Certification.RIAA	(omitted for space...)
Led Zeppelin	01/12/1969	Atlantic	6		8x Platinum	...
Led Zeppelin II	10/22/1969	Atlantic	1	Platinum	Diamond	...
Led Zeppelin III	10/05/1970	Atlantic	1	Gold	6x Platinum	...
Led Zeppelin IV	11/08/1971	Atlantic	1	3x Gold	Diamond	...
Houses of the Holy	03/28/1973	Atlantic	1	Gold	Diamond	...
Physical Graffiti	02/24/1975	Swan Song	1	Gold	Diamond	...
Presence	03/31/1976	Swan Song	1		3x Platinum	...
In Through The Out Door	08/15/1979	Swan Song	1		6x Platinum	...
Coda	11/19/1982	Swan Song	4		Platinum	...

Data in Columns on Disk

Title	Date	Chart

Row-Oriented data on disk

Led Zeppelin IV	11/08/1971	1	Houses of the Holy	03/28/1973	1	Physical Graffiti	02/24/1975	1
-----------------	------------	---	--------------------	------------	---	-------------------	------------	---

Column-Oriented data on disk

Led Zeppelin IV	Houses of the Holy	Physical Graffiti	11/08/1971	03/28/1973	02/24/1975	1	1	1
-----------------	--------------------	-------------------	------------	------------	------------	---	---	---

Goals for Data Lake Storage

- What is the requirement for a good storage system?
 - Easy to backup
 - Minimal learning curve
 - Easy integration with existing tools (Spark)
 - Resource efficient:
 - Disk space
 - Disk I/O
 - Network I/O
- Overall goal: **fast queries**

Options for Multi-PB Data Lake Storage


	Files	Compressed Files	Databases
Usability	Great!	Great!	OK to BAD (not as easy as a file!)
Administration	None!	None!	LOTS
Spark Integration	Great!	Great!	Varies
Resource Efficiency	BAD (Big storage, heavy I/O)	OK... (Less storage)	BAD (Requires storage AND CPU)
Scalability	Good-ish	Good-ish	BAD (For multi-petabyte!)
CO\$\$\$\$T	OK...	OK...	TERRIBLE
QUERY TIME	TERRIBLE	BAD	Good!

Parquet Format

“Apache Parquet is a columnar storage format available to any project in the Hadoop ecosystem, regardless of the choice of data processing framework, data model or programming language.”

- Binary Format
- API for JVM/Hadoop & C++
- Columnar
- Encoded
- Compressed
- Machine-Friendly

CSV vs. Parquet: SQL Query



1 master, 3 worker nodes with data on HDFS, file size: 420 GB

CSV vs. Parquet: Cost Consideration

Dataset	Size on Amazon S3	Query Run time	Data Scanned	Cost
Data stored as CSV files	1 TB	236 seconds	1.15 TB	\$5.75
Data stored in Apache Parquet format*	130 GB	6.78 seconds	2.51 GB	\$0.01
Savings / Speedup	87% less with Parquet	34x faster	99% less data scanned	99.7% savings

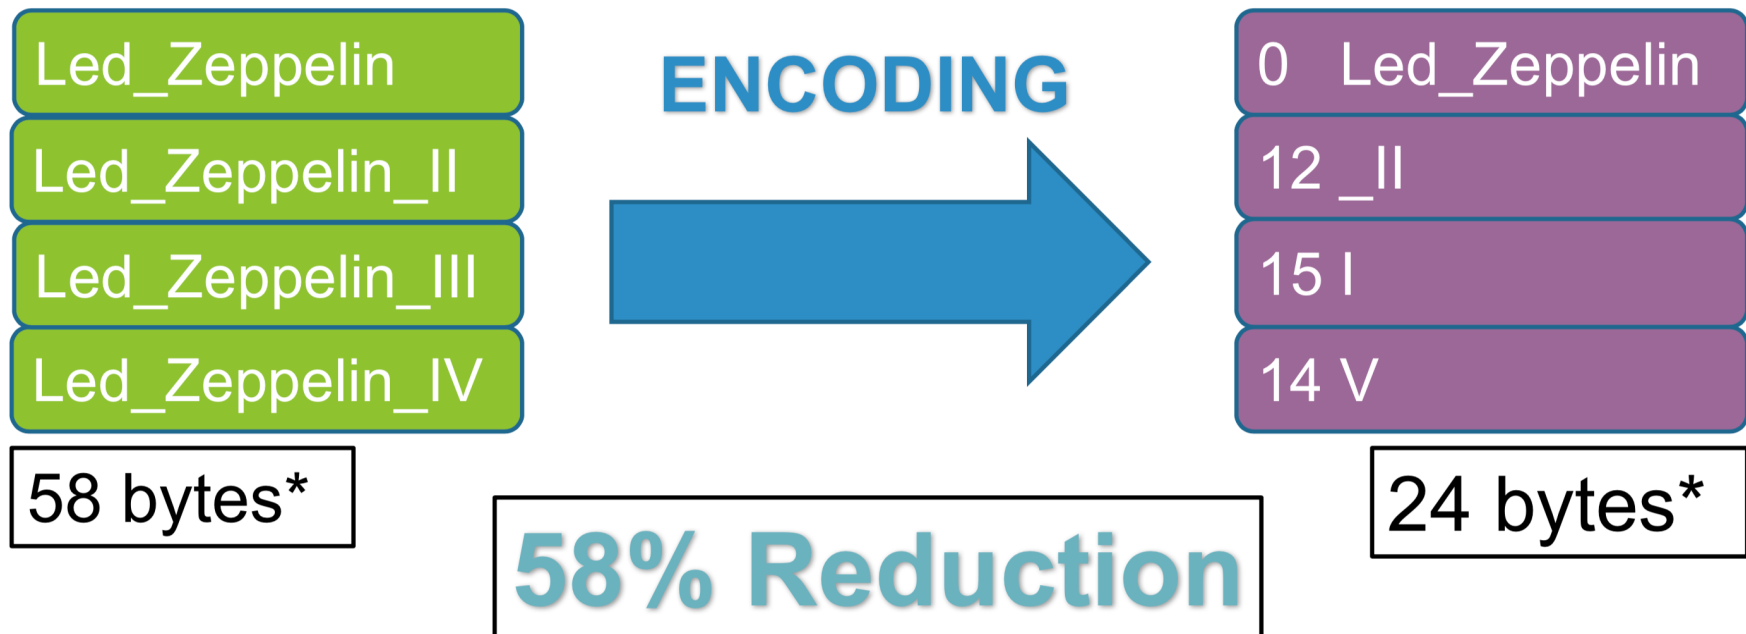
- **Costs for just one query every day within a year**
 - \$2000 for the csv
 - \$3.50 for the parquet file
- **“unproductive waiting time” for an analyst for just one query every day within a year**
 - 30 hours for the csv
 - 42 mins for the parquet file

Source: Andreas Weiler

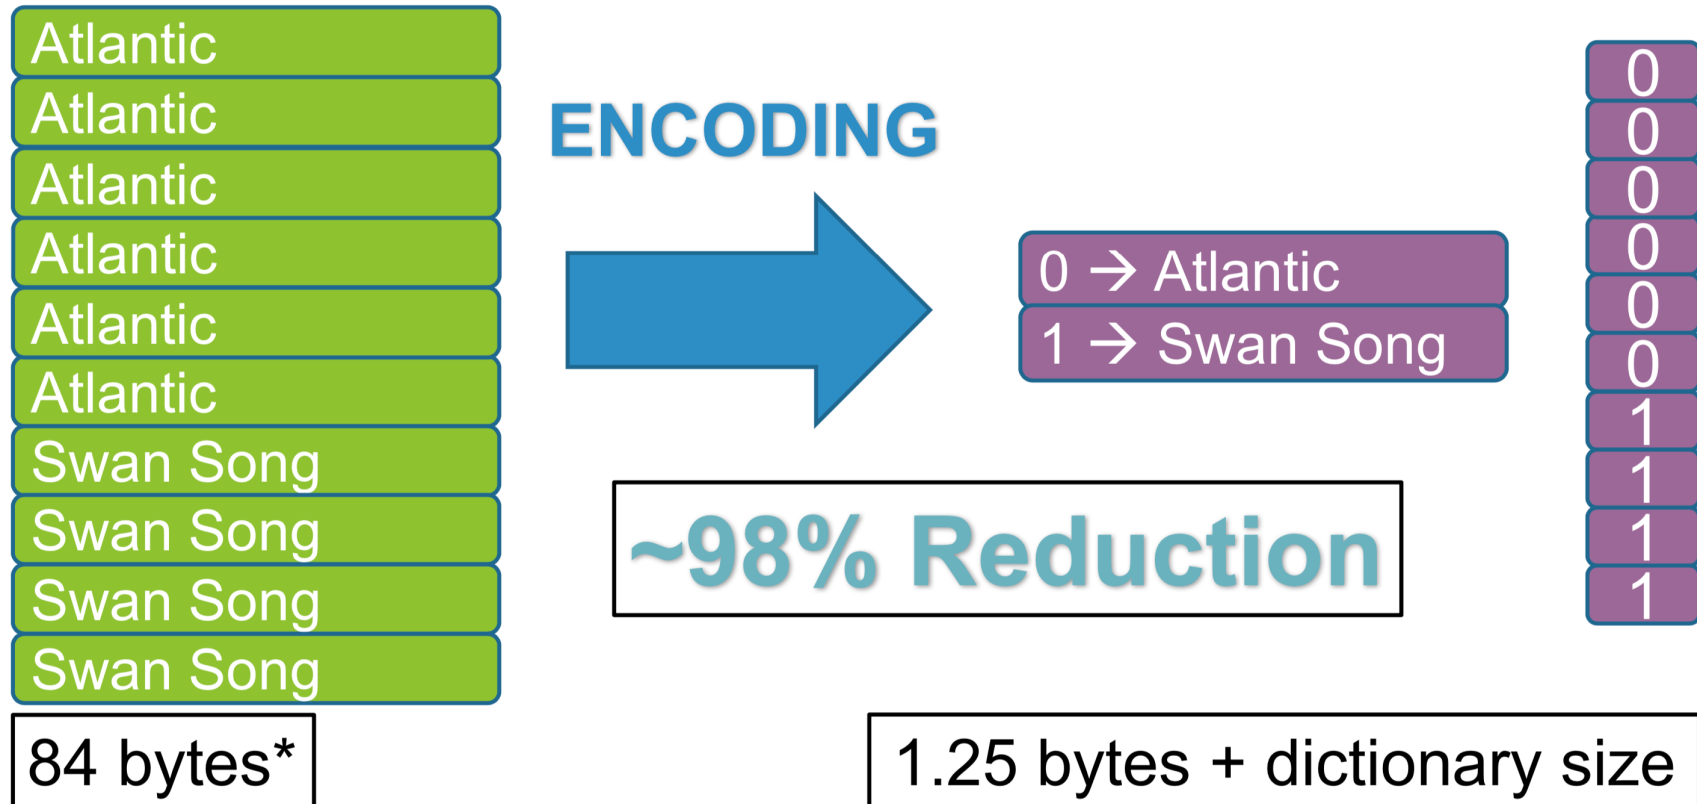
How can we Encode/Compress the Data?

Title	Released	Label	PeakChart.UK	Certification.BVMI	Certification.RIAA	(omitted for space...)
Led Zeppelin	01/12/1969	Atlantic	6		8x Platinum	...
Led Zeppelin II	10/22/1969	Atlantic	1	Platinum	Diamond	...
Led Zeppelin III	10/05/1970	Atlantic	1	Gold	6x Platinum	...
Led Zeppelin IV	11/08/1971	Atlantic	1	3x Gold	Diamond	...
Houses of the Holy	03/28/1973	Atlantic	1	Gold	Diamond	...
Physical Graffiti	02/24/1975	Swan Song	1	Gold	Diamond	...
Presence	03/31/1976	Swan Song	1		3x Platinum	...
In Through The Out Door	08/15/1979	Swan Song	1		6x Platinum	...
Coda	11/19/1982	Swan Song	4		Platinum	...

Encoding: Incremental Encoding



Encoding: Dictionary Encoding



Parquet File: Writing and Reading

```
# Read data from file (tab-separated)
customer = sqlContext.read.format("com.databricks.spark.csv")\
    .option("header","true")\
    .option("delimiter", "|")\
    .option("inferSchema", "true")\
    .load("/FileStore/tables/uqrziax61502358733901/customer.tbl")

# Write as parquet file
customer.write.parquet("customerPq")

# Read from parquet file
customerDF = spark.read.parquet("customerPq")
```

More Encoding Schemes

- Plain (bit-packed, little endian)
- Dictionary-encoding
- Run length encoding
- Etc.

Reading vs. Writing

- Writing a Parquet file is typically significantly slower than writing a CSV file:
 - More operations required (encoding)
- Reading a Parquet file is typically significantly faster than reading a CSV file:
 - Only specific columns need to be read
 - Columns are compressed
- However, data is often written once and read many times
- Hence, benefits of fast reads often outweigh performance penalty of slow writes

Conclusions

- HDFS **distributes blocks** onto machines in a round robin fashion
- Replication guarantees **fault tolerance**
- Using **Parquet can be more performant** than CSV due to encoding and compression
- However, Parquet format **might also be slower**:
 - e.g. if data is not “compressible” (i.e. no repetition in data)