

Information Engineering 2

Bonus Material: Deep Learning Pipelines

Prof. Dr. Kurt Stockinger

Deep Learning Pipelines

- Use existing deep learning libraries ([Keras](#), [TensorFlow](#), [PyTorch](#), [XGBoost](#)) with [Spark](#)

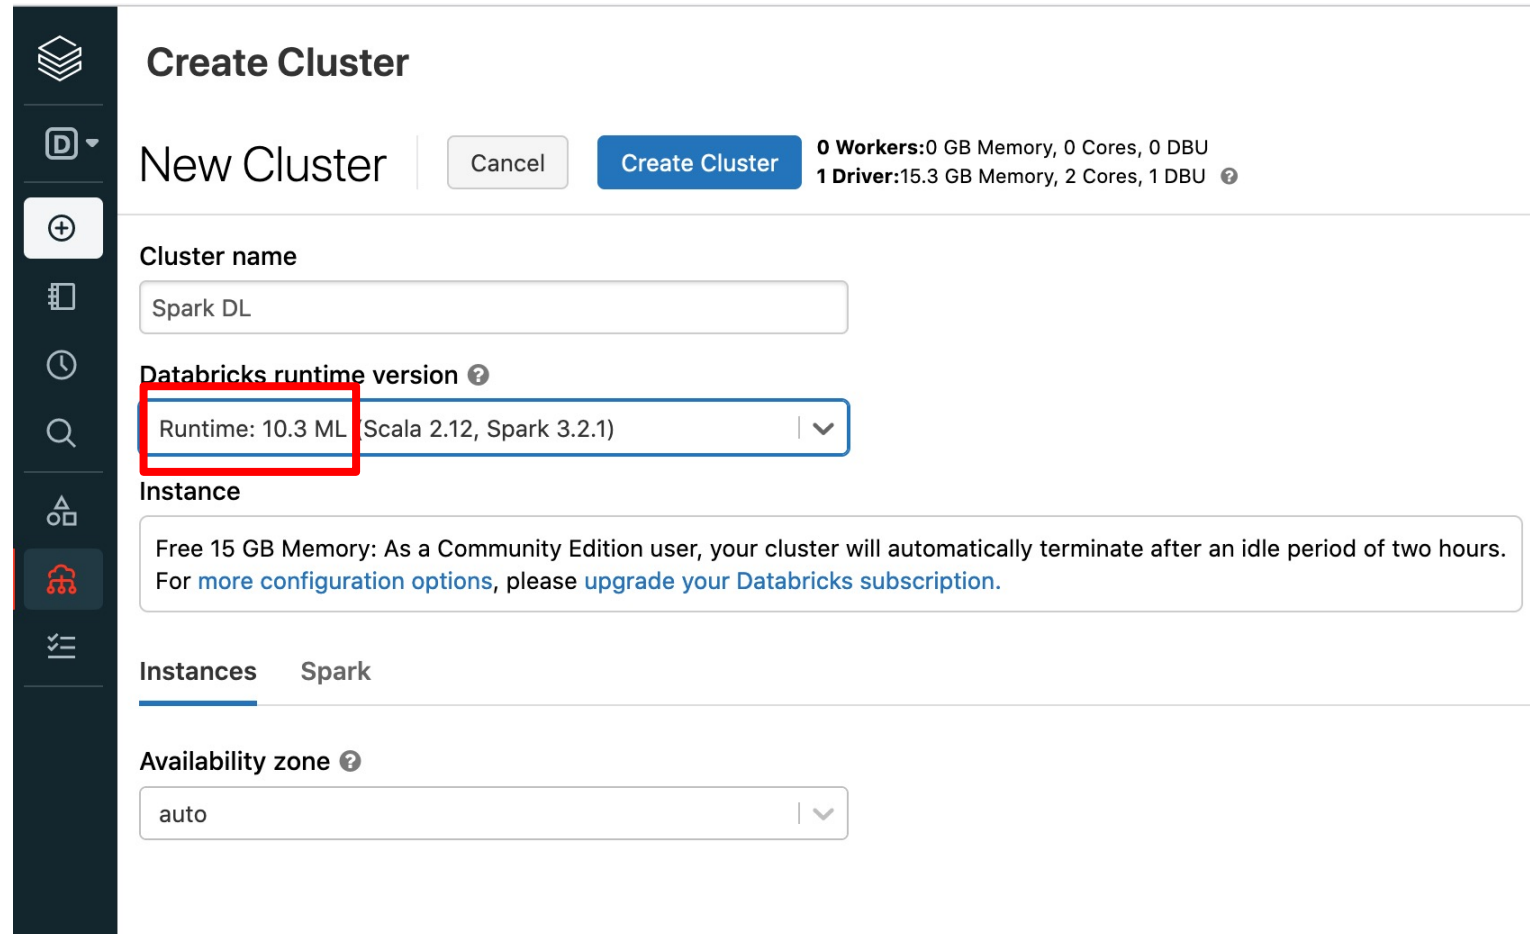
Deep Learning Guide

Databricks provides an environment that makes it easy to build, train, and deploy deep learning (DL) models at scale. Many deep learning libraries are available in [Databricks Runtime ML](#), a machine learning runtime that provides a ready-to-go environment for machine learning and data science. For deep learning libraries not included in Databricks Runtime ML, you can either install libraries as a [Databricks library](#) or use [init scripts](#) to install libraries on clusters upon creation.

Graphics processing units (GPUs) can accelerate deep learning tasks. For information about creating GPU-enabled Databricks clusters, see [GPU-enabled Clusters](#). Databricks Runtime includes pre-installed GPU hardware drivers and NVIDIA libraries such as CUDA.

Documentation and example notebooks: <https://docs.databricks.com/applications/deep-learning/index.html>

Databricks Community Edition: Runtime for Machine Learning



Create Cluster

New Cluster Cancel Create Cluster **0 Workers:**0 GB Memory, 0 Cores, 0 DBU
1 Driver:15.3 GB Memory, 2 Cores, 1 DBU ?

Cluster name
Spark DL

Databricks runtime version ?
Runtime: 10.3 ML (Scala 2.12, Spark 3.2.1) | v

Instance
Free 15 GB Memory: As a Community Edition user, your cluster will automatically terminate after an idle period of two hours. For [more configuration options](#), please [upgrade your Databricks subscription](#).

Instances Spark

Availability zone ?
auto | v

Deep Learning with Keras+Spark: Example Code Snippet

```
import tensorflow as tf
from tensorflow.keras.layers import Dense
from tensorflow.keras.models import Sequential
import mlflow
import mlflow.keras
import mlflow.tensorflow
```

Create the neural network

```
def create_model():
    model = Sequential()
    model.add(Dense(20, input_dim=8, activation="relu"))
    model.add(Dense(20, activation="relu"))
    model.add(Dense(1, activation="linear"))
    return model
```

Compile the model

```
model = create_model()

model.compile(loss="mse",
              optimizer="Adam",
              metrics=["mse"])
```

Detailed example:

https://docs.databricks.com/_static/notebooks/getting-started/get-started-keras-dbr7ml.html